# Use AI to acquire customers: find them, get shortlisted, close them, and let them bring the next one

Jweokk · 2026-10-25 · v0.3.6

Buyers have moved: they ask an AI, let an agent shortlist for them, and read what platforms and other people say.
This book follows the path a buyer actually walks — find, shortlist, close, bring the next one — with 135 methods you can copy, every number traced to its original source.

Copyright © Jweokk. Free for reading and non-commercial sharing; commercial use requires prior written permission.
Contact: weokk2025@gmail.com

# Contents


**01** Preface · How Buyers Find Vendors Has Changed

**02** Chapter 1: See the Battlefield First — How Buyers Pick Vendors Now

**03** Chapter 2: How to Get Your Product Mentioned in AI Answers

**04** Chapter 3: Get Onto the Buyer's First Shortlist

**05** Chapter 4: Let Third Parties Speak for You

**06** Chapter 5: Lead with Your Own Pitfalls

**07** Chapter 6: Be Readable by Machines

**08** Chapter 7: Win the Buyer's AI Agent

**09** Chapter 8: Three Extra Moves for Chinese Markets

**10** Chapter 9: Advertising: From "Selecting an Audience" to "Giving a Target"

**11** Chapter 10: Acquiring Customers on Other People's Platforms: App Stores, In-Platform E-Commerce, and Industry Directories

**12** Chapter 11: Letting Others Sell for You: Channels, Resellers, and Ecosystem Directories

**13** Chapter 12: Turning One Piece of Content into Thirty: The AI Content Factory and Matrix Distribution

**14** Chapter 13 Outreach: From Volume to Precision

**15** Chapter 14 Content Others Cannot Copy: First-Party Data and Real Cases

**16** Chapter 15 Let the Product Acquire: Free Tools, Trials, and Open Interfaces

**17** Chapter 16 Let AI Answer Around the Clock

**18** Chapter 17 Find Customers, Manage Leads, and Run the Sale with AI

**19** Chapter 18 Phone and Private Domains: What Still Works, and What Blows Up the Moment You Touch It

**20** Chapter 19 Let Customers Bring Customers: Word of Mouth, Referrals, and Virality

**21** Chapter 20 How to Know If It's Working: Three Sets of Books and a Holdout Test

**22** Chapter 21 Don't Bet on One Gateway: Dependency Check and Hedging

**23** Chapter 22 Build the Foundation: First-Party Data, Compliance, and Omnichannel Orchestration

**24** Chapter 23 Who Does It and What It Costs: Three Minimum Configurations

**25** Chapter 24 Don't Cross These Lines: Six Failure Files

**26** Appendix A · Case Index and Sources

**27** Appendix C · Version History and Update Notes

**28** Part 1 · See the Battlefield

**29** Part 2 · New Methods for the AI Era

**30** Part 3 · Upgrading the Old Methods with AI

**31** Part 4 · Don't Waste Money or Build on Sand

# Preface · How Buyers Find Vendors Has Changed

Buyers no longer pick vendors only in the search box and at trade shows — they ask an AI, they let an agent do the filtering, they look at what platforms push and at what other people say. This book teaches you how to win customers in these new places, and in the old ones too.

Those two sentences are the only thing in this book you need to memorize. Every method in every chapter that follows is what they look like taken apart.

## What Changed Is the Step Where They Pick

First, be clear about what changed. "A vendor gets picked" has moved four times in a little over thirty years.

The portal era (late 1990s to early 2000s): the gateway was a hand-maintained directory — Yahoo in the United States; Sina (Chinese: 新浪), Sohu (Chinese: 搜狐) and hao123 in China. The job a company had to do was called "getting listed"; the power sat with editors, and whether you made the directory was visible at a glance.

The search era (mid-2000s to the 2010s): the gateway became keywords. Google and Baidu (Chinese: 百度) handed the power of filtering to ranking algorithms, and companies held two levers for the first time — optimize organic rankings, or buy keyword positions outright. Traffic now had a public price, and the customer acquisition budget became a visible cost line on the books. The division into content teams, SEO teams and paid-media teams all grew out of this assembly line.

The recommendation-feed era (around 2016 to today): Douyin (Chinese: 抖音) and Toutiao (Chinese: 今日头条) flipped the gateway over — users used to look for content; now content looks for users. Companies responded by going all-in on content, and short-video teams and livestream rooms grew up alongside the paid-media teams. This stage buried one thing — traffic allocation became a black box for the first time: you roughly know the rules (completion rate, engagement), but you don't know why any single allocation happened.

The AI and agent era (late 2022 to today): the gateway folded from a column of results into one answer. On November 30, 2022, OpenAI released ChatGPT [official]. In May 2024, at its I/O conference, Google named its generative search experience AI

Overviews and launched it in the United States [official]. For the first time, a buyer can outsource the whole thing — "find me a few vendors that do X, and compare whether A or B suits my budget better" — to a model. Similarweb puts a magnitude on this layer: between June 2025 and May 2026, AI platforms sent an average of 770.7 million referral visits a month worldwide (desktop plus mobile web), up 117.4% year over year, with the growth very unevenly spread — beauty up 312.5% in a year, news and travel up 115.6%, marketplaces up 237.3% [third-party verified]. The Chinese entry-point layering deserves more attention: on QuestMobile's count, as of May 2026 native AI apps had 499 million monthly active users, AI capability plug-ins inside super-apps had 644 million, and phone makers' pre-installed AI apps had 755 million [third-party verified] — the largest AI surface is not inside standalone apps but on paths users already walk every day.

This time there is a substantive difference from the previous three: in those, your position was "what rank"; in this one, your position is "present or not". The losers aren't moved to page two — they don't exist.

We tested that conclusion ourselves before starting this book. Using 6 real procurement questions (3 Chinese–English pairs, covering three categories — CRM, MES and study-abroad agencies), we asked each question twice on each of 5 model channels. Asking the same question of the same model twice in a row, the overlap of the recommended lists was only 22%–56%; asking a different model, the pairwise overlap was only 4%–21%, and merging the lists from several models produced more than 40 names. Of 51 valid answers, those with clickable sources: 0 [our own measurement].

One boundary has to be marked on that last point: "AI answers carry citations" belongs to the product layer — a search product with web access attaches sources to its answers, and this is invisible at the model layer. Tell which layer a sentence is speaking about before deciding whether it can be cited — this discipline is used throughout the book.

## Who This Book Is For

It divides chapters by how far the buyer has got, not by industry. All you need is to find the cell you're stuck in right now.

| How far the buyer has got | The problem you need to solve | Chapter |
|---|---|---|
| Hasn't yet worked out whether he's present at all | First work out whether you're present right now | 1 |
| Find | When the buyer asks an AI, how to get him to ask his way to you | 2 · 3 · 4 · 6 · 7 · 8 |
| Find | The old roads: how to get the buyer to run into you for the first time | 9 · 10 · 11 · 12 · 13 |
| Shortlist | When the buyer looks at risk and at other people's reviews, how to get him to keep you on the list | 5 · 14 · 15 · 16 |
| Close | The buyer is willing to talk — how to close | 17 · 18 |
| Bring the next one | How to get the customer to bring the next customer | 19 |
| Not on this journey | You did everything right, but you're wasting money / your foundation is crooked | 20 · 21 · 22 · 23 · 24 |

Start with the first row. Plenty of people believe they're stuck at "can't close", and only find out when they measure that they aren't present at all. Until that cell is solved, every method in every later chapter is money down the drain.

## One Formula

Taken apart, customer acquisition is one multiplication:

> Acquisition result = number of touchpoints × relevance × trust × conversion rate ÷ cost

The four factors describe four different things. Number of touchpoints is in how many places the same person can run into you; relevance is whether, at the moment he runs into you, what you're saying is exactly what he's looking for; trust is whether he believes you; conversion rate is whether, once he believes you, he's willing to talk and willing to sign. Cost sits in the denominator — it's the money the first terms burn

between them, and it doesn't care where it goes: compute, data and sources are all line items newly added this round. This book doesn't recognize "free traffic" as a thing, only "traffic whose cost hasn't been priced yet".

AI made the first two cheaper. The marginal cost of producing content falls, so one deep piece can be split into dozens of items spread across several gateways, and the number of touchpoints goes from "how much the budget can buy" to "how much the capacity can put out"; relevance got cheaper too — machines will read your website, your prices and your specs, as long as you write the facts in a form they can read and check. These two used to rest on craft; now they look more and more like configuration.

The third, by contrast, became scarcer. When everyone's content can be mass-produced, what the buyer sees at the AI has become how other people rate you, whether you put the risk up front, whether you have first-party data and real cases — and what you say about yourself drops down the order. Numbers can be synthesized; "I've actually used this thing" cannot. So the first two getting cheaper doesn't mean acquisition gets cheaper — with more touchpoints and higher relevance, and trust not moving along with them, the same budget only buys you more people who "have seen it".

On the buyer's side there is one more thing you will never see. The dark funnel has always existed — the term "dark social" appeared as early as 2012 [third-party verified], meaning the sharing and discussion that leaves no source behind: a line at the dinner table, "which one do you all use", a line in a private group chat, "send me the quote". Once AI arrived, this stretch of the decision left a log for the first time — what he asked, who he ruled out, who he compared you against. Only that log sits with the buyer and the model vendors, and is not shown to the seller. This asymmetry matters far more than "organic traffic is shrinking".

## How to Read This Book

The order has three steps.

First, use AI to make the old methods fast and accurate. None of the old roads have disappeared; the way you walk them has changed. Advertising, outreach, in-platform acquisition — you're already doing these, and handing the repetitive labor inside them to AI shows results the same day.

Then let an agent string the actions together. Building lists, publishing content, taking inquiries, following up leads — these tasks repeat over and over, and only once they're chained do you get to talk about scale.

Finally, position yourself at the new gateways. Getting AI to mention you in its answers, squeezing into an agent's candidate set, cultivating the evidence third parties hold — all three wait on your facts, cases and data being ready first; otherwise you can hold the seat and still not catch what comes to it.

If you're in a hurry, jump around using the table above, and read first the chapter that matches the step you're on right now.

The book treats the Chinese market and overseas markets side by side. The credibility mechanisms of Doubao (Chinese: 豆包), DeepSeek, Kimi and Tencent Yuanbao (Chinese: 腾讯元宝) differ from those of ChatGPT, Perplexity and Google AI Overviews, and the source structures of Xiaohongshu (Chinese: 小红书), Douyin, the WeChat (Chinese: 微信) ecosystem and the open web differ too. What transfers across markets is mechanisms and principles; channels and tools have to be redone, and compliance has to be checked once on each side.

On numbers: every figure carries a source tier at the end of its sentence — [official], [third-party verified], [vendor claim], [our judgment], [our own measurement]. Anything marked [our own measurement] comes with its method, sample, time and boundaries, so you can reproduce it yourself, or falsify it yourself.

## Where This Book Stops

This book stops at the first completed transaction, at the signing of a contract. Retention, repeat purchase, customer success, cross-selling and upsell are out of scope — that is the subject of *AI Customer Lifecycle Management* (Chinese: 《AI 客户经营》). The reason is simple: the qualification for acquisition is settled at the front end — whether you got into the buyer's decision process is basically fixed before sales says its first word.

## Notes

- The eras in the preface serve only to orient you. The start and end times of the four gateway shifts are a narrative frame that helps people understand; the eras overlap, and this is not precise periodization.

- Several widely circulated industry figures (a short-video platform's daily active user scale, the user scale and penetration rate of generative AI, the amount involved in a deepfaked video-conference fraud) currently cannot be traced to an original publisher's source, so they are not used in the body at all.
- The candidate-set figures come from one round of the book's own measurement (6 questions, 3 Chinese–English pairs, 5 model channels, each asked twice). They describe behavior at the model memory layer, and cannot be extrapolated to the citation behavior of consumer AI search products.
- "AI answers carry citations" belongs to the product layer and is invisible at the model layer. First tell whether a sentence is speaking about the product layer or the model layer, then decide whether it can be cited.

## Sources

[1] OpenAI | Introducing ChatGPT | 2022-11-30 | https://openai.com/index/chatgpt/ | [official] [2] Google | Introducing AI Overviews | 2024-05-14 | https://blog.google/products/search/generative-ai-google-search-may-2024/ | [official] [3] Alexis C. Madrigal | Dark Social: We Have the Whole History of the Web Wrong | The Atlantic | 2012-10-12 | https://www.theatlantic.com/technology/archive/2012/10/dark-social-we-have-the-whole-history-of-the-web-wrong/263523/ | [third-party verified] [4] Our own measurement 1 (candidate-set multi-model measurement) | 2026-09-13 | Method and raw data in `实测/实测1-候选集多模型/` | [our own measurement]

[5] Similarweb | AI Referral Traffic by Industry (June 2025–May 2026: 770.7 million AI referral visits a month, +117.4% year over year; beauty +312.5%, news and travel +115.6%, marketplaces +237.3%) | 2026-09-03 | https://aisearch.similarweb.com/blog/ai-referral-traffic-by-industry/ | [third-party verified] [6] QuestMobile | 2026 Half-Year Report on the AI Application Market (as of 2026-05: native AI apps 499 million MAU, AI plug-ins 644 million, device-vendor AI apps 755 million) | 2026-07-14 | https://www.questmobile.com.cn/research/report/2076954943839809537/ | [third-party verified]

# Chapter 1: See the Battlefield First — How Buyers Pick Vendors Now

Where buyers pick vendors has changed: in the past they had to try one keyword after another in a search box, collect business cards one by one at a trade show, and line you up against a few others in a table; now they often have the list settled before they ever meet you, and you have no idea how they settled it.

Today's buyers ask an AI first, read what others say first, and draw up a shortlist themselves first — and you cannot see any of these steps. So this chapter gives you 5 methods, starting with an entry-point audit, to find out where you do and do not stand.

---

## Method 1: Pin Down Where the Buyer Has Already Finished Comparing You

**What to do**: Split the buyer's path — from "the first flicker of interest" to "the first time they meet you" — into three segments, and see which steps they have already completed on their own side.

**How**

1. Cut the three segments like this: interest (they hand the whole need, as a sentence, to an AI), first screen (they get a first list with names and comparisons), and settling (they rank them in their head and decide whether to meet you). For each segment, write one sentence: "at this moment, who is deciding on their behalf."
2. Mark, in each segment, one step where "you used to have a handhold but it has now been taken away": need translation — they used to have to break their own idea into keywords; side-by-side comparison — they used to have to page through candidates themselves; the click-through — they used to have to at least open you. Once these steps are gone, the position you used to hold stops working too.

3. Break "AI entry point" down one more level, or your actions will land in the wrong place: the ones where you ask and it answers (ChatGPT, Doubao (Chinese: 豆包), Kimi and the like); the ones you never ask that it hands to you anyway (on-site search and recommendations inside Xiaohongshu (Chinese: 小红书) and Douyin (Chinese: 抖音), product recommendations on Taobao (Chinese: 淘宝) and JD.com (Chinese: 京东)); and the ones where it acts on the buyer's behalf (price comparison, form filling, ordering). The first two need watching now; the last has only landed in a narrow way so far, but keep one eye on it. The phrase "AI entry point" is too coarse — coarse enough that it makes people push hard on the wrong spot.
4. Circle the step for which there is no record of any kind. A conversation that leaves no link parameters, never enters the CRM, and shows up in no attribution report is the dark funnel — it has always existed in the trade-show bar, the industry group chat, the friend's recommendation, only now it has a form you can go back and check: the words he said to the AI, you can go ask them yourself, and see whom the model mentioned and how it relayed you.
5. This one page is written by the marketing team and updated once a quarter. Every entry-point action that follows finds its own place against it.

**Example**: In Pew Research Center's July 2025 analysis (900 U.S. adults who agreed to share their browsing behavior), when an AI summary appeared in the search results, 8% of users clicked a traditional link, versus 15% when there was no summary; only 1% of the links cited inside the summary were clicked; and about 88% of that month's summaries cited three or more sources [third-party verified]. Two years on the numbers are starker. SparkToro, using Similarweb's US panel, measured zero-click at 68.01% of Google searches in January–April 2026, against roughly 45% a decade earlier; only 0.34% of searches reached AI Mode in that window, even as Google said at I/O 2026 that AI Mode had passed 1 billion monthly users [third-party verified]. Clicks being redistributed is not demand disappearing: data compiled by The Digital Bloom puts organic click-through at about 0.61% when an AI Overview is present versus about 1.62% when it is not, while the top 10 sites still grew about 1.6% and the steepest declines landed on sites ranked roughly 100th to 10,000th [third-party verified].

**Pitfalls**: Thinking this is a problem AI alone created — the phrase "dark social" was already in use in 2012; what has changed is that it now starts earlier, has a wider blast radius, and you can partly see it.

## Method 2: Ask AI 25 Questions Every Month and Turn "Are We Present?" Into a Table

**What to do**: Using one fixed set of buyer questions, record "what AI says about us now" every month into a table you can compare.

**How**

1. Fixed question set: assemble 20–30 buyer phrasings from three places — the questions sales was asked most in the last three months, the comparison questions in lost-deal reviews, and the selection doubts that recur in pre-sales notes. Do not rewrite them into marketing language. Ten core ones get tested monthly, the whole set quarterly.
2. Ask at three levels, a third each: category level ("which companies do X"), feature level ("which X tools support Y"), and scenario level ("what should we choose in our situation", carrying the constraints — size, budget, industry, compliance requirements, and existing systems). The scenario level is closest to a real decision, and the quickest to expose a blank.
3. Fix the entry points; don't ask these three this month and three others next month. Overseas at least Google (including AI Overviews), ChatGPT, Perplexity, Claude; in Chinese at least Doubao, DeepSeek, Kimi, Tencent Yuanbao (Chinese: 腾讯元宝). For buyers who compare inside content platforms, put on-site search on Xiaohongshu and Douyin in a separate group — do not blend it into one curve with AI visibility.
4. Each time record only four fields: whether you were mentioned; position and phrasing (what rank, and whether it is positive, neutral, negative, or got your specs and price wrong); the cited sources (your own site, third-party media, review sites, communities — or no citation at all); and side-by-side competitors (who else was in the same answer, and in what order).
5. If any one of these three things happens, the table is worth watching monthly: organic traffic other than brand terms is falling while your rankings barely move; sales gets stumped in a meeting by an AI comparison table the buyer brought, and you don't know how that comparison was produced; or you see yourself recommended in one test and strike the whole matter off the agenda.

6. Sign off on a conclusion on the stability rate alone: of these 25 questions, how many have mentioned you steadily for three straight months, with correct phrasing and a checkable third-party source. That single ratio is the conclusion you want.

**Example**: In our own measurement we used 6 real procurement questions (CRM, MES, study-abroad agencies), asked each across 5 model channels twice, and got 51 valid answers: asking the same model the same question twice, the overlap of the shortlists was only 22%–56%; switching to another model, it fell to just 4%–21% [our own measurement]. So being recommended in a single test may be nothing but sampling noise.

**Pitfalls**: Don't use an improvement in the ranking ledger to prove the citation ledger is healthy — "our search rankings are still rising" measures words, whereas on the AI side what is measured is questions; the two ledgers cannot cancel each other out.

## Method 3: Run Funnel Archaeology with Newly Closed Customers

**What to do**: Do a structured follow-up with closed customers whose memory is still fresh, and reconstruct the path they walked before they met you.

**How**

1. The respondents are not prospects but customers who just signed: do it within 2–4 weeks of signing, pick 20–30 of them, one session each, asked by sales or customer success themselves — do not outsource it into a questionnaire.
2. Ask in sequence, don't merge it into one sentence: When did this first enter your head? What did you open first — a search engine, an AI, some group, or did you ask a person? When you handed the question to the AI, how did you phrase it exactly (ask them to repeat the exact sentence, don't summarize it for them)? Which names came up in that answer? Where else did you later fill in information? Which name did you cross off, and why?
3. What you write down is not the answers but the path. Stack 20–30 of these and you get a dark-funnel map of your own industry: which entry point buyers most often start from, who keeps appearing on the shortlist, who gets crossed off, and which version of "the information about us" each of them saw.

4. Copy the customer's exact words into the question ledger (see Method 4) verbatim — do not turn them into your own summary. These raw quotes are the best source of content topics and sales scripts.
5. Run one round every quarter, with a fresh batch of newly closed customers each round.

**Example**: 6sense's 2025 B2B Buyer Experience Report, a survey of nearly 4,000 B2B buyers worldwide, says 95% of winners were already on the "day-one list", and about four-fifths of deals are won by a preferred vendor before any contact [vendor claim].

**Pitfalls**: Don't ask one "how did you hear about us?" and call it done — customers' recall is neither accurate nor free of hindsight rationalization; what you want to ask is what exactly they opened at the time, and how they put it in their own words.

## Method 4: Turn Buyers' Exact Wording Into a Question Ledger

**What to do**: Record every question a buyer has asked, in their own words, as an asset that keeps updating.

**How**

1. At minimum these fields: the question as asked; the source of the question (AI session / sales call / support ticket / Douyin DM / community); the asker's role; the buying stage; whether it involves a competitor comparison; the material used in the answer; and whether it affected the close.
2. Two sets of sources: overseas ones come mainly from call recordings, follow-ups on review sites, and forum posts; on the Chinese side add Douyin DMs, WeCom conversations, community questions, and support chats — these are both a source of questions and an entry point for handling them.
3. Who records: sales and support jot them down as they go, and marketing collects once a month, de-duplicates, classifies, and marks which questions currently have no piece of content that can answer them.
4. Three uses: a source of topics for content and material; a real script for sales training; and a substitute observation when the visible funnel data is distorted — the structure of the questions and changes in the volume of asking often reflect a shift in demand earlier than the reports do.

5. Archive once a month, look at structural change once a quarter.

**Example**: Superhuman clustered users' own words and fed them back into copy and feature prioritization [third-party: founder-authored account]; Microsoft's Customer Connection program has engineering teams collect customer habits, pain points, and suggestions directly, and change the product accordingly [official].

**Pitfalls**: Don't turn it into a support-ticket tally — what makes the ledger valuable is the question as asked; rewriting the question into your own classification label deletes the most valuable part.

## Method 5: Tell "AI Products" Apart From "AI Models" — Don't Measure with the Wrong Ruler

**What to do**: Before you start testing, tell apart whether you are testing an AI product with live retrieval, or a bare model — the two do not measure the same thing.

**How**

1. The product layer is what the customer uses: ChatGPT Search, Perplexity, Google's AI Overviews and AI Mode, domestic assistants with retrieval. Their answers carry cited sources; this layer measures citations — who got cited, and how many times.
2. The model layer is calling the API directly with retrieval off: the answer relies on the model's memory, and this layer measures the list — who it remembers, and at what rank. With a technical team on hand, this is the layer people most easily test as a side effect, but it is not the thing in front of your customer.
3. Record the two numbers separately; do not merge them into one curve: the product layer looks at cited sources and the accuracy of phrasing, the model layer looks at shortlist overlap and stability rate.
4. On a limited budget, test the product layer first — that is the thing actually in the customer's hands; keep the model layer as a control, to see how your position on the list shifts in the model's memory.
5. When reporting upward, say clearly which layer you tested. Using a bare model's result to "confirm you were recommended" is just self-comfort.

**Example**: Across the 51 valid answers at the bare-model layer in our own measurement, 0 carried a clickable source [our own measurement]; while an analysis of more than 24,000 conversation groups and more than 366,000 embedded citations in AI Search Arena shows that at the product layer the sources are highly concentrated, with news citations at about 9% and concentrated in the hands of a few outlets [official: academic preprint].

**Pitfalls**: Don't use one layer's conclusion to judge the other — the layer you tested may not be the layer your customer uses.

---

## Notes for this chapter

- The click and traffic data above is about "whether you can be seen", which is not the same as a closed deal. Public research offers no citable quantitative conclusion on the causal strength between being cited and closing, and our own measurement cannot answer it either — the citation ledger is positioned as a qualification metric, not a revenue metric [our judgment].
- The same fact comes in many percentage versions on the market ("buyers complete most of the purchasing journey before contacting sales" has claims ranging from under 20% to over 90%); the differences come from the year, how the journey is defined, what "contacting sales" means, and which step "complete" refers to. Break the definition apart before citing; don't pick a handy one [our judgment].
- SparkToro's zero-click-rate series changed data sources three times across the years (Jumpshot in 2016 and 2019, Datos in 2024); this book cites it only by "year + data source" and does not join years from different sources into one curve; the U.S. zero-click rate for January–April 2026 was 68.01% [third-party verified].
- Gartner's "search engine volume will drop 25% by 2026" is a forecast published in February 2024, not a measured result; this book labels it "forecast" throughout and does not use it as a factual baseline [official: forecast].
- Overseas samples cannot be extrapolated to Chinese. Pew's sample is a browsing panel of 900 U.S. adults, and 6sense's sample is B2B buyers in English-speaking markets; on the Chinese side there is no equally rigorous quantitative study of "buyers researching with AI". CNNIC's figure for generative-AI user scale

only shows how big the entry point is, not the behavior — and that number comes relayed through a vendor press release, so this book did not verify the original report item by item [third-party verified: relayed via vendor press release].

## Sources for this chapter

[1] Pew Research Center | Google users are less likely to click on links when an AI summary appears in the results (900 U.S. adults; March 2025 data; 8% vs 15%; 1% click on links inside the summary; 88% of summaries cite three or more sources) | 2025-07-22 | https://www.pewresearch.org/short-reads/2025/07/22/google-users-are-less-likely-to-click-on-links-when-an-ai-summary-appears-in-the-results/ | [third-party verified] [2] SparkToro | In 2026, Less than One Third of Google Searches Still Send a Click (U.S. zero-click rate 68.01% for January–April 2026; 60.45% in 2024, data source Datos; data source Jumpshot in 2016 and 2019) | 2026-06-08 | https://sparktoro.com/blog/in-2026-less-than-one-third-of-google-searches-still-send-a-click/ | [third-party verified] [3] Gartner | Gartner Predicts Search Engine Volume Will Drop 25% by 2026, Due to AI Chatbots and Other Virtual Agents | 2024-02-19 | https://www.gartner.com/en/newsroom/press-releases/2024-02-19-gartner-predicts-search-engine-volume-will-drop-25-percent-by-2026-due-to-ai-chatbots-and-other-virtual-agents | [official: forecast] [4] Digiday | Change in publisher referral traffic (DCN survey and evidence submitted by the PPA to the UK Competition and Markets Authority: median referral traffic about −10% year on year, about −14% for non-news; first-page CTR 5.1%→0.6%) | 2026 | https://digiday.com/ | [third-party verified] [5] 6sense | 2025 B2B Buyer Experience Report (nearly 4,000 B2B buyers worldwide; 95% of winners were on the "day-one list"; about four-fifths of deals won by a preferred vendor before contact) | 2025-11 | https://6sense.com/science-of-b2b/buyer-experience-report-2025/ | [vendor claim] [6] arXiv:2507.05301 (Kai-Cheng Yang) | Citations in AI search answers are highly concentrated in a few sources (more than 24,000 conversation groups, more than 366,000 embedded citations; news citations about 9%) | 2025-07-07 | https://arxiv.org/html/2507.05301v1 | [official: academic preprint, before peer review] [7] Alexis C. Madrigal | Dark Social | The Atlantic | 2012-10-12 | https://www.theatlantic.com/technology/archive/2012/10/dark-social-we-have-the-whole-history-of-the-web-wrong/263523/ | [third-party verified] [8] Sohu (a Zhidamingyuan press release, relaying CNNIC) | Three trends in China's AI marketing and customer acquisition in 2026 and GEO compliance red lines | 2026-08-31 | https://

www.sohu.com/a/1068667519_100119123 | [third-party verified: relayed via vendor press release, original report not verified item by item] [9] Our own measurement 1 | Multi-model candidate-set test (6 questions × 5 channels × 2 runs, 51 valid answers; overlap on the same question across two runs 22%–56%; cross-model overlap 4%–21%; citations at the bare-model layer 0) | 2026-09-13 | repository `实测/实测1-候选集多模型/` | [our own measurement] [10] Microsoft | Customer Connection Program (engineering teams collect customer feedback directly) | https://techcommunity.microsoft.com/blog/microsoftintuneblog/announcing-the-customer-connection-program/303626 | [official] [11] First Round Review | How Superhuman Built an Engine to Find Product-Market Fit (clustering users' own words and feeding them back) | https://review.firstround.com/how-superhuman-built-an-engine-to-find-product-market-fit/ | [third-party: founder-authored account]

[12] SparkToro (Rand Fishkin) | In 2026, Less than One Third of Google Searches Still Send a Click (Similarweb US panel: 68.01% of Google searches zero-click in January–April 2026 vs roughly 45% in 2016; AI Overviews on 20%+ of searches, cutting click-through by nearly 60%; only 0.34% of searches reached AI Mode in that window) | 2026-06 | https://sparktoro.com/blog/in-2026-less-than-one-third-of-google-searches-still-send-a-click/ | [third-party verified] [13] The Digital Bloom | Organic Traffic Crisis Report, 2026 Update (AI Overview presence 6.49% → 13.14%; organic click-through 0.61% with an AI Overview vs 1.62% without; top 10 sites +1.6%, steepest declines between rank 100 and 10,000; US organic search traffic down 2.5% year over year) | 2026-03-07 | https://thedigitalbloom.com/learn/organic-traffic-crisis-report-2026-update/ | [third-party verified]

# Chapter 2: How to Get Your Product Mentioned in AI Answers

Buyers used to look for products in search engines, and what you competed for was a ranking. Now they just ask an AI: "Given our size and our budget, who should we pick?" — **what you have to do has become: appear in its answer, and not be described wrongly.**

This is not something money can buy. There is no ad slot inside an AI answer. All you have are the 10 methods below, to make the facts about you worth being written into the answer.

The first 4 are must-dos, the middle 3 are pitfalls to avoid, and the last 3 are advanced.

---

## Method 1: Ask AI 25 Questions a Month and See How It Describes You

**What to do**: Get clear on what AI says about you today — this is the starting point of every action.

**How**

1. Prepare 25 questions, mixing three kinds: category phrasings ("which companies do X"), feature phrasings ("which tools support Y"), and **conditional phrasings** ("we're 50 people, our budget is 2,000 yuan a month, we need Gantt charts and time tracking — which vendor should we pick"). The third kind is closest to a real buyer.
2. Ask at each of four entry points: Google (including AI Overviews), ChatGPT, Perplexity, and for Chinese business add Doubao or DeepSeek.
3. Each time record four things: **whether you were mentioned / what rank and whether it's accurate / who was cited / which competitors were in the same answer**.
4. Do it once on a fixed day each month, for three months.

**Example**: When we ran this test ourselves, we found that for the same question and the same model, asking twice in a row, the overlap of the recommended list was only 22%–56%; switching to another model, the overlap was only 4%–21%. **So a single test result is meaningless; you have to watch the trend month by month.** Trends have a magnitude of their own: Semrush's analysis of 230,000 prompts across 13 weekly snapshots (2025-07-14 to 2025-10-12) shows ChatGPT's citations of Reddit collapsing from about 60% of replies to about 10% in mid-September 2025, and Wikipedia from about 55% to under 20%, while Google AI Mode and Perplexity stayed roughly flat over the same weeks [third-party verified] — visibility can halve in a few weeks, so a single snapshot cannot carry a budget decision. Not everyone is even measuring: one technical-community write-up says 75% of companies have no AI visibility tracking at all [third-party verified: technical community]. Churn has a finer reading too: a report built on millions of data points says only 30% of brands stay visible from one answer to the next, and only 20% remain across five consecutive runs [vendor claim] — stability is not the default state, it is the exception for a minority of brands. Set the ruler against brand size and the ladder is explicit: an arXiv preprint tracking more than 100 brands and 100,000+ prompt responses puts global household names in about 73% of relevant AI answers on their very first tracking run, established mid-market and regional brands at about 44%, and niche or small brands at about 11% — roughly 30 percentage points per rung; about 78% of the citations it observed went to corporate websites, YouTube led among non-corporate sources, and the ranked "best-of" listicle was the most-cited content format at about 21% of all citations; sentiment is the unstable part, flipping about 6.7 times more often than whether the brand is mentioned at all [official: academic preprint]. Down at the citation level, the same kind of capture gives a different reading: 480 questions, four engines, one capture recording 16,069 citations across 7,775 distinct domains. ChatGPT cited a source in only 72 of its 480 answers, and 69 of those sit in two question types — who provides something and what it costs; on the other six types it cited nothing in any industry, while still writing answers averaging 2,303 characters. The closest engine pair, Perplexity and Google AI Overviews, overlaps at a mean per-question Jaccard of just 0.181, and when the same questions were re-asked 38 minutes later Perplexity agreed with itself at 0.900; 5,201 of the 7,775 cited domains, 67%, were cited exactly once, while sites owned by the companies that sell the thing being asked about held 54.0% of all citations [third-party

verified: single capture, full dataset published]. Being cited therefore splits in two: engines prefer concrete evidence on a concrete question rather than citing the same large sites over and over.

**Pitfalls**: Asking only "which companies do X". That measures whether you have a name, not whether you can get into the candidate set. The field most easily missed among the four is "whether it's accurate" — **not being mentioned is only absence, but being described wrongly is lying on your behalf, and you have no way to appeal.** That last clause, though, was partly rewritten by the courts in 2026: a preliminary injunction from the Regional Court of Munich on 28 May held that an AI Overview, because the model compiles the information itself and condenses it into a summary, constitutes an "independent presentation" for which the platform is responsible, so the older precedent that search engines are not liable for their results does not carry over. In that case the AI mixed information about other, genuinely dubious companies into a Munich publishing group and drew connections that appeared in none of the linked sources; the court found this impaired its commercial reputation with a risk of repetition, and enjoined the statements, with fines of up to EUR 250,000 per breach [third-party verified: full judgment, English translation]. The cost of being described wrongly is shifting from "figure it out yourself" to "a claim you can assert".

## Method 2: Say the Whole Thing in the First Sentence

**What to do**: Let AI pull your content straight out and use it. It synthesizes from extracted paragraphs; it does not score whole pages.

**How**

1. The first sentence of every page and every section gives the conclusion directly, with no buildup.
2. Write specs, pricing logic, and applicable conditions in full — factual density beats keyword density.
3. Don't use pure client-side rendering: most AI crawlers do not execute JavaScript, so content that depends on browser rendering is, in their eyes, the same as not written.
4. Use list form more: **52% of listicle articles achieve high citation rates on ChatGPT** [third-party verified].

**Example**: Academic research on GEO tested 9 optimization methods; keyword stuffing was almost useless, and on Perplexity even performed **10% worse** than doing nothing [official]. Rankings and citations are no longer the same thing: an analysis of 863,000 search results pages puts the share of AI Overview citations drawn from Google's top-10 pages at 37.9%, down from 76.1% in an earlier round of research [third-party verified] — making the top ten does not mean making the answer. Most of what gets cited is not your own site either: an analysis of 40,000+ highly cited URLs across nearly 50 product categories found that about half of ChatGPT's citations are listicles, and 58% of those listicles are ranked lists [third-party verified] — whether you get cited depends heavily on whether you appear in lists somebody else writes. Structure is itself a measurable signal: the same AI-search report says 68.7% of pages cited in ChatGPT follow a clean heading hierarchy, 87% use a single H1 and nearly 80% include lists, and pages using three or more schema types are 13% more likely to be cited [vendor claim] — being cited is not luck, it is a set of structural features you can build. The gap between ranking and citation now has cross-industry baselines: an analysis of 45 million+ search keywords across 828 large B2B companies found that AI Overviews appear on nearly 50% of results pages where B2B brands rank, while the median enterprise B2B brand is cited in just 3% of the AI Overviews it is relevant for; the median company ranks for nearly 9,700 keywords, and roughly 4,500 of those trigger AI-generated answers — half of the search footprint is already covered by generative answers, while citations stay in the single digits [third-party verified: vendor research]. One industry shows the same shape: insurance benchmarks put AI Mode at 43.9% of the citations analysed in that sector and Perplexity at 33.6%, with AI Overviews appearing on 40.7% of insurance searches and 53.3% of those citations going to pages that never cracked Google's organic top 10; the same report notes that AI engines drive around 1.08% of all website visits, so treating AI as a traffic channel gets the objective wrong from the start [third-party verified].

**Pitfalls**: Making the page pretty but unable to yield facts. AI does not eat design; it eats sentences.

## Method 3: Publish Numbers Nobody Else Has

**What to do**: When you hold data nobody else has, that is the highest-yield move for being cited.

**How**

1. Inventory your unique numbers: product usage data, order structure, service tickets, customer surveys — anonymized, they can be published.
2. Run a small-sample customer survey once a quarter (10–20 in-depth interviews is enough), with questions picked from the ones that most often stump sales.
3. When publishing, include sample size, methodology, and time window — without these, neither AI nor readers will believe it.
4. Set aside a budget line for "data and reports": an industry-recommended split is **40% core SEO, 25% digital PR, 20% data and reports, 10% training, 5% experimentation** [third-party verified]. That last item is missing from most companies' budget sheets.

**Example**: The three best-performing methods in the GEO research — injecting statistics, injecting quotations, and citing sources — gave relative lifts of **30%–40%** [official]. What they share: **all of them require you to hold something nobody else has.**

**Pitfalls**: Repackaging an industry report and publishing it again. That is relay, not first-hand data.

## Method 4: Make the Whole Web Give One Version of the Facts

**What to do**: AI stitches its answer across sources; when the versions don't match, it decides you are "uncertain".

**How**

1. First build a single-fact sheet: full company name and abbreviation, founding date, official product name and translation, pricing definitions, order-of-magnitude of customers, certification numbers.
2. Reuse it verbatim everywhere you speak publicly: your own site, encyclopedia entries, third-party review pages, recruiting pages, product data feeds.
3. **Don't write two prices on your own site and product feed** — at the answer layer, two prices equal no price.

**Example**: Real-time data from one monitoring firm shows **51.7% of AI citations point at the brand's own website** [vendor claim]. The biggest source of citations is often yourself — provided it can be read, and the version is consistent. A different sample gives a different distribution: across more than 1,000,000 citations on ChatGPT, Perplexity, and Google AI Overviews, community platforms (Reddit, Quora and the like) take 52.5% and brand-owned domains 47.5% [third-party verified]. The two samples cover different ground and the shares differ by a few points, but both point at the same thing: your own pages are the single largest source. The same analysis found a threshold further upstream — 73% of sites block AI crawlers through robots.txt, CDN rules, or JavaScript rendering [third-party verified]. Before you can be cited, you have to survive being crawled. Who gets read repeatedly at the citation layer also has a public audit: an audit of U.S. ChatGPT citations says Wikipedia accounts for 13.15% and Reddit 11.97%, together more than a quarter of all citations; The Wall Street Journal, The New York Times, Bloomberg and the Financial Times do not appear in the top 20, the only U.S. business publication in the top 20 is Forbes (No. 18, 1.38%), and Reuters ranks No. 7 (2.27%) [third-party verified] — PR tier does not equal AI citation rank.

Counted per answer rather than per citation, the picture is colder: a GEO monitoring firm that analysed millions of answers across more than eight models found that 84% of AI responses do not cite a brand's own domain at all, and that models describe a company mainly from community platforms, review sites, competitor comparison pages and press coverage; the number of domains cited in one answer varies by engine, averaging 5.8 for Gemini and 27 for Grok [third-party verified: monitoring vendor research]. The writing itself divides labour, too: an observation of 165 of one publisher's own articles as cited in Copilot (data from Bing Webmaster Tools' AI Performance report; a 29-day window, 5.4K citations) found that named category guides written as “best X for Y” took 86% of the citations, while “X vs Y” and “alternatives to Z” comparison pages were never cited once; the top 10 pages accounted for 64% of sampled citations, and within that single site, word count, heading count, tables and FAQ blocks did not separate cited from uncited pages [third-party verified: single-site observation]. That runs against the common claim that comparison pages are citation magnets; the sample is one site, one engine and 29 days, so file it as a direction rather than a verdict. Source structure can be calibrated against a cross-engine count: among U.S. ChatGPT (browsing mode) citation events in January–February 2026, Wikipedia came first with 118,285 (13.15%), Reddit second

with 107,680 (11.97%), OpenAI's own domain third (6.21%), and the only retailer in the top five was Walmart (2.90%) [third-party verified] — what explains citations is which few sites get read over and over, not how many sites mention you.

**Pitfalls**: Changing only your own site and ignoring the outdated information on third-party pages. AI reads the whole web.

## Method 5: Read Your Own Server Logs

**What to do**: No tools to buy — just look at which of your pages the AI crawlers are fetching.

**How**

1. Have a technical colleague pull the AI crawlers' access logs out separately.
2. Look at two things: which pages are being fetched often (that is the raw material of your being cited), and which pages are fetched repeatedly yet never cited (their content is substandard or unclear).
3. For the second group, change how they are written, and look at the logs again a month later.

**Example**: After Glasp optimized using the log method, daily sessions from ChatGPT rose from 500 to 19,000 [third-party verified: founder's own account].

**Pitfalls**: Buying only a third-party dashboard. It sells observation, not causation — it can tell you a ranking, not whether the money was worth it.

## Method 6: Don't Buy Reviews — the FTC Really Does Fine You

**What to do**: Rule the "buying" route out entirely, and replace it with genuine third-party evidence.

**Why**

In August 2024, the U.S. Federal Trade Commission passed a rule 5–0 making six categories of conduct unlawful: fake reviews, **AI-generated fake reviews**, paying money or goods for positive reviews, insiders writing reviews without disclosure, **company-controlled review sites falsely claiming independence**, and buying and selling accounts to inflate metrics [official].

In April 2026 the first major case landed: a brand selling children's height-growth supplements touted its results while having employees write thousands of five-star reviews, trading discounts for reviews, and using bot accounts to post software-generated comments — **it was hit with a $4 million judgment** [official].

**The key point**: in this case, **false advertising and fake reviews were fined together** — fake reviews supply "social proof" for false advertising. On AI it is exactly the same: **"buying AI mentions" and "buying reviews" are the same category of thing in the regulator's eyes.**

**Pitfalls**: Thinking you can slip through because AI can't see it. The bar for reporting is lower than you think.

## Method 7: Don't Mass-Produce Template Pages — Google Will Cut You Off Overnight

**What to do**: Make the pages you "feed to AI" read as human-written, incremental content.

**Why**

The templated pages you mass-produce on your own site will hit Google's cleanup first — and it won't notify you when it happens. In one core update in March 2026, hit sites **dropped 60%–90% in ranking overnight, lost an average of 87% of traffic, and all the impact showed within 14 days; there was no manual penalty notice, no Search Console alert, and no avenue to appeal** [third-party verified].

Google explicitly named three violating practices: **mass AI pages without editorial review** (structure repeated across thousands of variants fails), **pure template variable substitution**, and **data aggregation with no added value** (scrape data and publish it directly; "surface rewriting" with AI is not enough either) [official].

**Pitfalls**: Thinking good prose gets you off the hook. What's judged is structural repetition, not whether the sentences read well.

## Method 8: Add Three More Moves for the Chinese Market

**What to do**: The Chinese corpus is different from overseas — most of what lives in WeChat Official Accounts and inside platforms is not crawled, so far less Chinese content is readable by AI.

**How**

1. **Move the facts onto public web pages**: your own site's knowledge base, industry media, platforms that can be indexed. WeChat Official Accounts drive relationship conversion, but at the answer layer they are nearly invisible.
2. **Maintain the entries correctly**: the names, qualifications, and timelines on encyclopedias and industry directories must stay consistent with public business-registration information. In the Chinese market, duplicate names are common and entities change frequently, so being mistaken for someone else is a high-frequency problem.
3. **Work the in-platform search**: Chinese buyers often compare inside on-site search on Xiaohongshu and Douyin. Title as question, first paragraph as conclusion, and serious replies in the comments are more effective than volume.

**One more ruler**: the Measures for the Labeling of AI-Generated Synthetic Content took effect on September 1, 2025: AI-generated content must declare its identity, and platforms must verify it at upload [third-party verified]. **This is protection for people doing real work** — it shuts down the route of "mass-passing content off as human-written", and by the same token raises the relative value of content that is genuine, attributed, and checkable.

## Method 9: Don't Panic When a Platform Turns "Being Cited" Into "Being Admitted"

**What to do**: Accept that a platform will tighten its entry points at any time, so keep the assets in your own hands.

**Why**

In November 2025, OpenAI launched ChatGPT "shopping research": the user states a need, it asks follow-up questions about budget preferences, compares prices across sites, and finally gives a short candidate list — **for a merchant to get onto that list, it has to go through an allowlisting application** [official].

**How**

1. Keep the factual assets in your own hands: your own site, product data, first-party data — they can be carried to another platform (Chapter 21 covers the hedging).
2. Don't bet on a single entry point: **only 11% of domains are cited by both ChatGPT and Perplexity**; ChatGPT loves citing Wikipedia (7.8% of its citations), while Google AI Overviews and Perplexity love citing Reddit [vendor claim].

## Method 10: Vet the Vendor Before You Hire One

**What to do**: Spend half an hour screening out unreliable vendors.

**How**

1. Take the vendor's name to four AIs: "How is company X?" "What cases has X done?"
2. Ask it to give **its own visibility baseline**: under which questions it gets mentioned, how the data was collected (API or sampled from the interface), and the sample size.
3. Then ask three more things: can it give executable actions, is there a control group proving the effect, and who is accountable when things go wrong.

**Example**: One monitoring firm's own data shows that **70% of marketing agencies are not cited on AI platforms themselves** [vendor claim]. The abuse risk on the provider side has already been named in an industry document: a GEO trusted-ecosystem report from the China Artificial Intelligence Industry Development Alliance (AIIA) groups the poisoning tactics into two stages — content writing and distribution — covering forged authority endorsements, smearing competitors, malicious keyword stuffing, implicit prompt injection, coordinated matrix-account posting, and contamination of high-authority sources; its governance path is written as provider self-discipline plus whole-chain coordination, and it records that the Cyberspace Administration of China has launched a special campaign, that the alliance has issued a GEO credibility self-discipline pledge with nearly 20 companies taking part, and that the Basic Requirements for Trustworthy GEO Services has been drafted and entered pilot evaluation [official: industry alliance]. When you vet a provider, read these before the performance numbers on a case-study page. The field

now has its first public sizing as well: China's narrowly defined GEO services market is estimated at RMB 4.2–4.6 billion for 2026, with the top ten providers holding only about 9.8% combined share. The white paper defines mention rate, first-recommendation rate and coverage rate as the core metrics (as opposed to impressions or post volume), lists keyword stuffing, homogenized rewriting, corpus poisoning, forged authority signals and commercial disparagement as the 2026 governance targets, and admits that its own industry pain point is "passing off POC-sized or SEO work as GEO results" [third-party verified: vendor white paper, includes self-reported data].

**Pitfalls**: Letting the vendor prove its value with a dashboard it made itself. An industry roundup in this field ranked the roundup's own author first [vendor claim] — **a leaderboard is not due diligence.**

---

## Notes for this chapter

- The "up to +40%" in the GEO research was measured in a self-built experimental environment (Google takes the top 5 sources + model generation), **not on any commercial AI search product**, and what it measured was "the relative lift of the same content after improvement". Fine as a direction, not fine as a promise.
- All the "AI traffic converts better" numbers (+42%, 25× and so on) come from vendors testing and reporting on themselves, with small samples and selection bias. **The direction is credible; the multiple cannot be used as a budget assumption.**
- On the distribution of cited sources at Chinese AI entry points, there is very little publicly checkable material — this is the thinnest part of the book; Chapter 8 handles the two sets of rules.
- The measured numbers in CAICT's Generative Engine Optimization Industry White Paper (conversion up 2.8× in AI recommendation scenarios after GEO commercialisation, decision cycles 40% shorter) exist only as media relays today; the white paper's sample size and statistical basis are not public [third-party verified: media relay]. Fine as a direction, not as a promise.

## Sources for this chapter

[1] Aggarwal et al. | GEO: Generative Engine Optimization | KDD '24 | https://arxiv.org/abs/2311.09735 | [official] [2] Yang, K.-C. | Citation patterns in AI search answers (24,000+ conversation groups, 366,000+ citations) | arXiv:2507.05301 | https://arxiv.org/abs/2507.05301 | [official: preprint] [3] Our own measurement 1 · multi-model candidate-set test | 2026-09-13 | `实测/实测1-候选集多模型/` | [our own measurement] [4] Ahrefs | What Drives Brand Visibility in AI | 2025-12 | https://ahrefs.com/blog/ai-brand-visibility-correlations/ | [vendor claim] [5] Boring Marketing | AI Visibility Statistics | 2026-09 | https://boringmarketing.com/ai-visibility-statistics | [vendor claim] [6] Peec AI | GEO Statistics | 2026-08 | https://peec.ai/ai-search-geo-statistics | [third-party verified] [7] eMarketer | FAQ on GEO and AEO | 2026-04 | https://www.emarketer.com/content/faq-on-geo-aeo--where-ai-search-seo-overlap-2026 | [third-party verified] [8] FTC | Final rule banning fake reviews and testimonials | 2024-08-14 | https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-trade-commission-announces-final-rule-banning-fake-reviews-testimonials | [official] [9] FTC | Enforcement against TruHeight ($4 million judgment) | 2026-04-13 | https://www.ftc.gov/news-events/news/press-releases/2026/04/ftc-takes-action-against-truheight-deceptive-unsubstantiated-advertising-supposed-height-enhancing | [official] [10] Google | Search core update and spam policies | 2024-03-05 | https://blog.google/products-and-platforms/products/search/google-search-update-march-2024/ | [official] [11] Digital Applied | Programmatic SEO After the March 2026 Core Update | 2026-03-18 | https://www.digitalapplied.com/blog/programmatic-seo-after-march-2026-surviving-scaled-content-ban | [third-party verified] [12] OpenAI | ChatGPT shopping research (allowlisting) | 2025-11-24 | https://openai.com/index/chatgpt-shopping-research/ | [official] [13] Leapd | Analysis of AI platform citation sources (680 million citations) | 2026-04 | https://www.leapd.ai/blog/ai-visibility/how-chatgpt-google-ai-overviews-and-perplexity-source-information-in-2026 | [vendor claim] [14] Sean Ellis | How Glasp took ChatGPT traffic from 500 to 19,000 using server logs | https://seanellis.substack.com/p/how-glasp-grew-chatgpt-traffic-from | [third-party verified: founder's own account] [15] Digiday | AI Overviews and publisher referral traffic | 2025-08-15 | https://digiday.com/media/google-ai-overviews-linked-to-25-drop-in-publisher-referral-traffic-new-data-shows/ | [third-party verified] [16] China Economic Net (relaying CCTV.com) | Measures for the labeling of AI-generated synthetic content and GB 45438-2025 (effective 2025-09-01) | https://www.ce.cn/xwzx/gnsz/gdxw/202509/t20250901_2460291.shtml | [third-party verified]

[17] OtterlyAI | The AI Citations Report 2026 (1M+ citations: community platforms 52.5%, brand-owned domains 47.5%; 73% of sites block AI crawlers via robots.txt/CDN/JS rendering) | 2026-02 | https://otterly.ai/blog/the-ai-citations-report-2026 | [third-party verified] [18] Strategi | AI Search Visibility Statistics 2026 (BrightEdge across five engines: source overlap 16%–59%, brand overlap 36%–55%; Ahrefs' 863,000 SERPs: AI Overview citations from top-10 pages 76.1% → 37.9%) | https://strategi.is/blogs/ai-search-visibility-statistics-2026-the-engines-disagree-on-sources-and-agree-on-brands | [third-party verified] [19] China Artificial Intelligence Industry Development Alliance (AIIA) | 2026 Research Report on Building a Trustworthy Generative Engine Optimization (GEO) Ecosystem (poisoning tactics and governance path; GEO credibility self-discipline pledge with nearly 20 companies; pilot evaluation of the Basic Requirements for Trustworthy GEO Services) | https://www.sohu.com/a/1026647008_121964487 | [official: industry alliance] [20] Semrush | Most Cited Domains in AI (230,000 prompts × 13 weeks, 100 million+ citations; ChatGPT's Reddit citations fall from about 60% to about 10%, Wikipedia from about 55% to under 20%) | 2025-11-10 | https://www.semrush.com/blog/most-cited-domains-ai/ | [third-party verified] [21] Evertune | AI Search Statistics for GEO (40,000+ highly cited URLs: about half of ChatGPT's citations are listicles, 58% of those ranked) | 2026 | https://www.evertune.ai/resources/ai-search-statistics-for-generative-engine-optimization | [third-party verified] [22] Tencent Cloud developer community | Rebuilding brand visibility in the AI search era (75% of companies lack AI visibility tracking) | 2026-07-10 | https://cloud.tencent.com/developer/article/2706823 | [third-party verified: technical community] [23] Phoenix Technology (Chinese: 凤凰网科技) | 2026 White Paper on Evaluating and Selecting GEO Service Providers (China's narrow GEO market at RMB 4.2–4.6 billion in 2026, CR10 about 9.8%; mention/first-recommendation/coverage metric definitions; list of governance targets) | 2026-09-03 | https://tech.ifeng.com/c/8w7tJeNESH1 | [third-party verified: vendor white paper] [24] Sina Finance (relaying ITHome) | CAICT's Generative Engine Optimization Industry White Paper: conversion up 2.8× in AI recommendation scenarios after GEO commercialisation, decision cycles 40% shorter | 2026-01-22 | https://finance.sina.com.cn/stock/relnews/cn/2026-01-22/doc-inhicwux4769683.shtml | [third-party verified: media relay] [25] AirOps × Kevin Indig | The 2026 State of AI Search (only 30% of brands stay visible from one answer to the next, 20% across five consecutive runs; 68.7% of cited pages follow a clean heading hierarchy, 87% a single H1, nearly 80% include lists, 3+ schema types 13% more likely to be cited; about 48% of citations from community platforms, 85% of brand mentions from third-party pages)

| 2025-12-02 | https://www.airops.com/report/the-2026-state-of-ai-search | [vendor claim] [26] 5W PR | Citation Source Audit Q1 2026 (Wikipedia 13.15% + Reddit 11.97%, together more than a quarter of all U.S. ChatGPT citations; WSJ/NYT/ Bloomberg/FT absent from the top 20; Forbes No. 18 at 1.38%, Reuters No. 7 at 2.27%; source Similarweb, about 600,000 citation events) | 2026 | https:// www.prnewswire.com/news-releases/wikipedia-and-reddit-now-drive-over-25-of-chatgpt-citations-in-the-us-new-5w-research-finds--wsj-nyt-and-bloomberg-do-not-appear-in-the-top-20-302768339.html | [third-party verified]

[27] Pratyush Kumar (Ranqo) | Generative Engine Optimization at Scale: Measuring Brand Visibility Across AI Search Engines (100+ brands, 100K+ prompt responses: visibility by brand stature 73% / 44% / 11%; about 78% of citations to corporate websites; ranked best-of listicles about 21% of all citations; sentiment flips about 6.7 times more often than mention) | 2026-06-18 | https://arxiv.org/abs/ 2606.20065 | [official: academic preprint]

[28] Hendricks | The Answer Index, Edition 1 (one capture on 2026-09-01: 480 questions x Google AI Overviews / ChatGPT / Perplexity / Gemini; 16,069 citations across 7,775 domains; ChatGPT cited in only 72 of 480 answers, 69 of them in provider-discovery and cost questions, zero citations on the other six question types while averaging 2,303 characters; Perplexity x AI Overviews mean per-question Jaccard 0.181; self-consistency 0.900 after 38 minutes; 67% of domains cited exactly once; seller-owned sites 54.0% of all citations, healthcare 64.5% / education 27.0%) | 2026-10-04 | https://hendricks.ai/research/the-answer-index | [third-party verified: single capture, full dataset published] [29] Walker Sands | 2026 B2B AI Search Visibility Benchmark (45M+ keywords, 828 large B2B companies, 14 technology industries; AI Overviews on nearly 50% of results pages where B2B brands rank; median company cited in only 3% of relevant AI Overviews; median company ranks for nearly 9,700 keywords, about 4,500 of which trigger AI-generated answers) | 2026 | https://www.walkersands.com/about/blog/b2b-ai-search-visibility-benchmark/ | [third-party verified: vendor research] [30] Conductor | 2026 Insurance AI Search Benchmarks (AI Mode 43.9% of citations analysed in the sector, Perplexity 33.6%; AI Overviews on 40.7% of insurance searches, 53.3% of those citations to pages never in Google's organic top 10; AI engines drive about 1.08% of all website visits) | 2026 | https://www.conductor.com/academy/2026-insurance-ai-search-benchmarks/ | [third-party verified] [31] Transparency Coalition | German court holds Google liable for AI hallucination (full English translation of the Regional Court of Munich preliminary

injunction of 2026-05-28: an AI Overview is an "independent presentation" for which Google is responsible and does not enjoy the DSA Articles 4-6 safe harbour; fines up to EUR 250,000 per breach) | 2026-06-15 | https://www.transparencycoalition.ai/news/german-court-holds-google-liable-for-ai-hallucination-read-the-full-decision-here | [third-party verified: full judgment, English translation] [32] AthenaHQ | How Do You Handle AI-Generated Misinformation, Hallucinations or Inaccurate Brand Mentions (84% of AI responses cite no brand-owned domain at all; domains cited per answer average 5.8 for Gemini and 27 for Grok; sample of millions of answers across 8+ models from December 2025 to March 2026 plus 500+ client conversations in Q2 2026) | 2026-08-03 | https://athenahq.ai/blog/how-do-you-handle-ai-generated-misinformation-hallucinations-or-inaccurate-brand-mentions | [third-party verified: monitoring vendor research] [33] Obsurfable (citing Provena primary research) | Copilot Cites Category Guides, Not Comparison Pages (165 articles, 29 days, 5.4K citations; named category guides 86%, comparison pages zero; top 10 pages 64% of sampled citations; data from Bing Webmaster Tools AI Performance report) | 2026-10 | https://obsurfable.com/resources/articles/copilot-cites-category-guides-not-comparison-pages | [third-party verified: single-site observation] [34] Similarweb AI Search | Most Cited Domains in LLMs (U.S. ChatGPT browsing mode, January–February 2026: wikipedia.org 118,285 / 13.15%, reddit.com 107,680 / 11.97%, openai.com 55,876 / 6.21%, walmart.com 26,118 / 2.90%) | 2026-04 | https://aisearch.similarweb.com/blog/most-cited-domains-llms/ | [third-party verified]

# Chapter 3: Get Onto the Buyer's First Shortlist

Before a buyer ever contacts you, they have already drawn up a shortlist of three to five vendors — if you are not on it, your product, your price, your cases, and your pitch have not yet earned the right to take the field. Buyers shortlist three to five names the moment they start researching, and that shortlist is the starting point for every price comparison that follows. So this chapter gives you 6 methods to make "you're on it" the default.

---

## Method 1: Turn "Did It Mention Me?" Into a Record Sheet You Can Read

**What to do**: Using fixed questions, multiple entry points, and repeated rounds, measure where you actually stand on that shortlist today.

**How**

1. Pick 3–5 whole sentences a buyer would really ask, not keywords. One set in Chinese and one in English.
2. Ask each question at 3 entry points, 3 rounds each. Write the entry points as specific products, and note whether each one is a retrieval-backed product form or a bare model: ChatGPT, Google AI Overviews, Perplexity, Doubao, Tencent Yuanbao, DeepSeek.
3. Each time record six columns: the question as asked / entry point and whether it has retrieval / round / appearance status (absent, present but undescribed, present and described correctly, present but described wrongly) / same-list entities (who else was in this answer) / stability share.
4. Do it on a fixed day each month, three months running.
5. If any one of three things is missing, the conclusion doesn't hold: repeat the same question over multiple rounds to measure stability, run several entry points side by side to measure entry-point dependence, and ask once in each language to measure the language partition.

**Example**: We ran a round of multi-model testing: 6 real procurement questions (3 Chinese/English pairs, covering CRM, MES, study-abroad agencies) × 5 model channels × 2 repeats per question, giving 51 valid answers; for the same question and the same model, asking twice in a row, the shortlist overlap was only 22%–56%, and switching to another model, the pairwise overlap was only 4%–21% [our own measurement]. Marketing has long agreed on the candidate set: since Howard & Sheth proposed the "evoked set" in 1969, two-stage purchasing has been the basic framework; Andrews & Srinivasan added half of the methodology in 1995 — the consideration set need not rely on consumers' self-reports but can be estimated in reverse from scanner panel data, and ignoring the consideration-set stage systematically misestimates choice [official]. G2's March 2026 survey of 1,076 B2B software buyers and decision-makers moved the starting line too: 51% of buyers now more often begin research with an AI chatbot (the 2025 edition of the same report put this at 29%), 71% rely on one at some point in the research process, and 85% rate a vendor higher because AI included it in an answer [third-party verified].

**Pitfalls**: Treating one run's result as a conclusion. At a drift of 22%–56%, one chance appearance makes you think you can close the project and pull the resources, and one chance absence makes you abandon a position that is merely fluctuating.

## Method 2: Watch One Number Only — the Stability Share

**What to do**: Change the criterion from "were we recommended?" to "how stable is our position within the distribution?"

**How**

1. Compute this number: stability share = the number of times you appear across a fixed question set × fixed entry points × N repeated rounds ÷ the total number of collections.
2. Ask one question 9 times (3 entry points × 3 rounds) and you appear 2 times; the stability share is about 22%. That number is not necessarily bad — it says you hold an unstable position in the distribution, and the next step is to see what is propping that position up.
3. It only counts if it is "present and described correctly". Not being mentioned is only absence; being described wrongly is lying on your behalf, and you have no way to appeal.

4. Look at the "same-list entities" column on its own: if the same list is always the same set of names, the contexts in which you are remembered are too narrow, and you need to widen them; if the same list differs every time, entity consistency is fine and you should keep consolidating the facts.
5. Watch the trend by month, not the number on a given day.

**Example**: The popular claim that "AI only recommends three to five candidates" does not hold in our measurement: the median number of recommended entities in a single answer is about 6–10, the longest is more than 15, and the union of the four models' lists reaches 43 entities [our own measurement]. What is genuinely scarce is being mentioned stably and repeatedly across multiple entry points. The product layer (AI search with live retrieval) is indeed highly concentrated: an analysis of more than 24,000 conversation groups, 65,000 answers, and more than 366,000 embedded citations in AI Search Arena found that the citations in AI answers are long dominated by a few sources, and that the quality of the cited source has no significant effect on user satisfaction — users neither know nor care whom the AI put on the list [official (academic preprint, before peer review)]. Engine disagreement has been measured too: across five engines, overlap on which sources get cited runs from 16% to 59%, while overlap on which brands get mentioned narrows to 36%–55%; Ahrefs' data on 75,000 brands puts the pairwise brand-mention correlations between engines at 0.749–0.821 [third-party verified]. Sources differ engine by engine and brand lists converge — which is exactly the mechanism that makes "stability share" worth tracking: once one gateway mentions you consistently, the others are more likely to mention you as well.

**Pitfalls**: Using "we were recommended" as the reason to launch a project. The signal that justifies a project is that two of three entry points do not have you.

## Method 3: Tie Your Name to the Buying Situation

**What to do**: Write "brand name × the sentence the buyer uses to describe the need" into every piece of material — don't just make people remember your name.

**How**

1. List your "category entry points": the situational phrases a buyer would use to describe this need, one set in Chinese and one in English — the two sets of buyers simply use different language.

2. In every piece of material, let the official name and the industry's colloquial name appear together. The category word the buyer says is often completely different from your website headline.
3. Write "what we do / what we don't do / how big the customers we serve are / what price range" as short factual sentences, and use only one form of the entity name — don't call yourself by three different names in the same piece of material.
4. Add a batch of new situational phrases each quarter, picked from the questions that most often stump sales.

**Example**: Byron Sharp of the Ehrenberg-Bass Institute has specifically clarified that "mental availability" means the probability that a buyer notices, recognizes, or thinks of you in a buying situation — far broader than "awareness". Traditional metrics like first mention, recognition, and recall only use prompts stripped of situation (usually just the category name), so they are flawed [official]. In Chinese it is especially easy to grasp: remembering "they can do dust-removal retrofits" and remembering "they're an environmental-equipment company" retrieve two different outcomes. Two other mechanisms point the same way: Hauser & Wernerfelt's evaluation-cost model shows that consideration-set size is a rational buyer's optimum, trading decision cost against marginal benefit, and the model also derives an order-of-entry penalty [official]; Tversky & Kahneman's anchoring experiment shows that the first option to enter the buyer's field of view becomes the frame of reference for every later comparison — a completely irrelevant random starting point can significantly shift judgment (median estimates of 25 and 45 in the two groups) [official].

**Pitfalls**: Just pumping the brand name and keywords. The shortlist remembers "name × situation"; with the situation half missing, the buyer can't retrieve you when they state the need.

## Method 4: Stockpile the Raw Material in the Years Before the Customer Has a Need

**What to do**: Shift part of your budget and time away from "the final push" into public facts that can be retrieved all year round.

**How**

1. First accept the timescale: companies change service providers in banking, legal, software, and telecom about once every 5 years; at any moment only about 20% are in-market within a year, and about 5% in any single quarter.
2. The raw material of the shortlist is piled up bit by bit in the years when the customer feels "this has nothing to do with me", so do this monthly, evaluate it yearly, and don't run it as a campaign.
3. Produce a batch of citable factual material every month — data, specs, definitions, cases — not promotional material.
4. Don't evaluate this kind of spend with an expectation of results in two weeks.

**Example**: The numbers above come from a study by John Dawes of the Ehrenberg-Bass Institute for the LinkedIn B2B Institute, which LinkedIn calls the "95-5 rule": at any moment at most 95% of companies are not in the market for most goods and services; the joint research by the LinkedIn B2B Institute and Ehrenberg-Bass, How B2B Brands Grow, shows that 75% of companies buy computers once every 4 years and 80% change banking services once every 5 years; on the same page there is a jarring contrast — LinkedIn's own survey shows 96% of B2B marketers expect advertising results within two weeks [official]. On the other side is 6sense's annual survey of nearly 4,000 B2B buyers worldwide: in 95% of cases, the vendor that ultimately wins the deal was already on the buyer's "day-one candidate list" before the buyer first contacted sales [vendor claim]. Preference comes before intent: Forrester's 2026-08-05 analysis reports that more than two thirds (68%) of B2B buyers start a purchase already favouring a vendor, because most buying journeys are confirmation rather than selection [third-party verified].

**Pitfalls**: Cutting the shortlist work because "there were no leads this month". What you cut is your ticket into the purchase three or five years from now.

## Method 5: Work All Three Supply Lines at Once

**What to do**: A shortlist has only three supply lines — the corpus you write yourself, the third parties who speak about you, and the default slots a platform presets.

**How**

1. The corpus line: write the facts into an extractable form (those short factual sentences from Method 3). This line you control completely.
2. The third-party line: get other people to speak about you. Overseas it leans on Reddit, YouTube, wiki-style encyclopedias, industry media, and listicles; in Chinese it leans on news sites, industry media, Q&A platforms, and encyclopedia entries that can be publicly indexed. The hard constraint is that you cannot control what a third party says; you can only control "what is worth being said".
3. The default line: get into an entry point's preset slot. Overseas that is the allowlisting application for ChatGPT shopping research; in Chinese it is a platform's admission to its own product pool or supplier pool.
4. This route is an administrative entry, not organic traffic: the work is preparing qualifications, pricing policy, inventory and delivery commitments, and interface capability, then submitting.
5. Judge priority this way: if you can control it, pick the corpus line; if others are willing to speak, pick the third-party line; if the entry point allows you in, pick the default line. If none of the three works, the problem is not customer acquisition — it is your product or your qualifications [our judgment].

**Example**: Ahrefs' December 2025 study (75,000 brands, brand mentions across millions of AI answers) gives a direct correlation ranking: YouTube mentions and branded web-page mentions are strongest; brand mentions overall correlate with AI visibility about 3 times as much as traditional backlinks, while the number of content pages is almost irrelevant [vendor claim]. Another analysis of 680 million citations across ChatGPT, Google AI Overviews, and Perplexity says only 11% of domains are cited by both ChatGPT and Perplexity; the three platforms' top sources differ, with ChatGPT favoring Wikipedia (7.8%), and AI Overviews and Perplexity favoring Reddit (2.2% and 6.6% respectively) [vendor claim]. On the Chinese side the default route is already running: ZKH (Chinese: 震坤行)'s "AI product recommendation brain" dynamically optimizes an "exclusive product pool" based on customers' purchase history and usage scenarios, so customers buy directly from its own mall and skip quote comparison — the platform says that in 11 months since launch it recommended more than 20 million lines of materials to customers and drove more than 100 million yuan in revenue growth [vendor claim]. The default value itself has force: Johnson & Goldstein's 2003 study in Science showed that every policy

necessarily has a "take no action" default, which imposes a cost on those who need to change their own state; changing the default alone could add thousands of organ donors a year in the United States [official]. On the Chinese side there is one more layer of terrain difference: a 2026 Chinese GEO industry research report judges that AI is turning from an efficiency tool into a high-frequency entry point for information and decisions, with Doubao, DeepSeek, Tencent Yuanbao, and Qwen (Chinese: 千问) contributing most of the growth through ecosystem advantage or vertical scenarios, and more than 40% of users shifting the center of gravity of their search habits from traditional search engines to AI search [third-party reported]. The more the entry points concentrate inside a platform's closed loop, the more the shortlist looks like an admission system; the more they scatter across the open web, the more it looks like a memory contest — on the Chinese industrial-goods side it most resembles the former, and categories like SaaS most resemble the latter: for the same one-line procurement need, the Chinese entry points return Sharecrm (Chinese: 纷享销客), Neocrm (Chinese: 销售易), Wukong CRM (Chinese: 悟空 CRM), Jodoo (Chinese: 简道云), Mingdao Cloud (Chinese: 明道云), and Lark Base (Chinese: 飞书多维表格), while the English entry points return HubSpot, Pipedrive, Zoho, Freshsales, Close, and Salesforce Essentials, with cross-language overlap of 8%–26% [our own measurement]. In regulated industries (finance) the shortlist is generated by the institution's internal model, and the move to get in is compliance admission and the supplier library — content tactics yield little here; education is more direct, with the power to generate the shortlist falling largely to the in-platform handling step — parents ask inside the platform, the platform assigns the inquiry to whoever, and that one is on the list. For shortlists in consumer and local life, more and more are generated by platforms' own recommendation pools and by purchase guidance inside private domains. Local discovery already has hard survey numbers: BrightLocal's Local Consumer Review Survey 2026 surveyed 1,002 US adults and found the share using AI for local business recommendations climbing from 6% in 2025 to 45%, making AI the third most-used local recommendation tool behind Google and Facebook, with ChatGPT accounting for 31% of that use; over the same period the share relying on Google reviews fell from 83% to 71%, because ChatGPT cannot see inside Google's walled garden of reviews [third-party verified].

The moment a shortlist forms is also being pulled earlier by AI: a June 2026 G2 survey of 1,038 B2B software decision-makers (North America, EMEA and APAC, plus 55 in-depth interviews) reshuffles the sequence — evaluation now occupies 40% of the buying journey, overtaking research (36%) as the longest stage; 82% of buyers

sourced software recommendations from an AI chatbot in the last 24 months, and 50% of those said AI had the greatest impact when narrowing and comparing options; review sites remain the top single source shaping the shortlist (38%) but now lead AI chatbots (37%) by one point, the first time the study has recorded that ordering; after a vendor is selected, the biggest delay is IT security review (39%, rising to 50% among enterprise buyers), followed by budget approval (32%) and implementation planning (25%); finance involvement in software decisions jumped from 31% to 46% in a year, and 49% of buyers saw a CFO veto an already-approved purchase in the last 12 months; internal resistance to AI adoption rose from 16% to 29%, the largest single-year shift in the study [third-party verified]. For the three intake routes, the reading is: AI shortens the list, review sites make the list hold up, and security review and budget only start filtering after you are already on it.

**Pitfalls**: Walking only the corpus line and then blaming the abundance of third-party content. Of the three lines you control only one; the other two depend on "having something worth being said" and "being qualified to get through that door".

## Method 6: Manage the Shortlist as "a Roomful of Shortlists"

**What to do**: Get onto each individual shortlist held by each person in the buying group, while being clear that the shortlist itself will be rewritten by others.

**How**

1. First map the buying group: who decides, who uses, who sets requirements, who can veto. In B2B the shortlist is plural — it is stored separately in each member's head, and the versions conflict with one another.
2. Make content for the levels of the buying group; don't only personalize for your point of contact.
3. Treat "I explained it clearly to my contact" as a risk signal — that sentence often falls apart in the bid-evaluation meeting.
4. "What counts as qualified" is also part of the shortlist. If you don't write it clearly, someone else will write it for you.

**Example**: Gartner's sales survey published on 2025-05-07 (632 B2B buyers) shows: 74% of B2B buying teams have "unhealthy conflict" during the decision process; buying groups that reach consensus are 2.5 times more likely to report a "high-quality deal" than groups that don't; and the most counterintuitive set of

numbers is that tailoring relevance to the buying-group level has a 20% positive effect on buying-group consensus, while tailoring only at the individual level produces a 59% negative effect on consensus [vendor claim]. There is also a public record of a shortlist being worked in reverse: on January 5, 2026, KNKX and Cascade PBS published ChatGPT conversation records obtained through a public-records request — a city employee in Bellingham, Washington, responsible for replacing the city's billing software, repeatedly asked ChatGPT between January and March 2025 to write solicitation requirements favoring VertexOne, and directly asked it to "exclude Origin SmartCity without explicitly mentioning it"; at least 16 of ChatGPT's generated phrasings appeared verbatim in the final evaluation matrix of 350 requirements. The five-year, $2.7 million contract was ultimately awarded to VertexOne, whose bid was about $2.5 million — about $1 million more than the second-ranked Origin SmartCity at about $1.45 million. The city launched an independent investigation before the media published [third-party reported]. Even as AI rewrites the shortlist, the human side of the work grows: G2's July 2026 survey of 1,038 B2B software decision-makers calls this year's situation "the evaluation maze" — AI makes candidates easier to find, which widens the pool, lengthens evaluation and leaves information quality uneven, so "easier to find" and "harder to buy" happen at once; the same survey carries one footnote for the agent era, that only 2% of decision-makers would let AI complete a purchase without pre-approval [third-party verified]. The shape of the funnel has changed: a wider top, a longer middle, and the vendor that keeps producing checkable evidence across that long evaluation is the one that walks the shortlist to a contract.

**Pitfalls**: Thinking it's enough that your one contact is on board. Sitting in the bid-evaluation meeting are several other shortlists.

---

## Notes for this chapter

- The two layers of shortlist must be read separately. The "95-5 rule" and the "day-one candidate list" are about the shortlist in the buyer's head; the measurement in this chapter measures the model-memory layer of the bare API, not an AI search product with live retrieval. Citations at the product layer are indeed highly concentrated, while the model layer is closer to a drifting distribution; this chapter issues no overall judgment that "the shortlist must be long or must be short".

- 6sense's 95% figure comes from the summary page of a vendor-funded survey; the report does not give the full questionnaire and statistical appendix. Here only the directional conclusion is taken; the specific percentage is not asserted before the source is checked [vendor claim].
- The cross-language overlap of 8%–26% is calculated on the basis of "the union of all models' shortlists under the same category"; comparing single answers pairwise would give lower figures. The direction is usable; it is not enough to extrapolate into a precise percentage [our own measurement].
- The correlation coefficients in Method 5 come from vendors' own measurement definitions, not independent reproduction; correlation is not causation. Whether "AI amplifies or weakens the first-mover advantage" has not been measured; the order-of-entry penalty conclusion comes from the era of shelves and advertising.
- Don't take the route of buying your way onto the shortlist. A 2026 CCTV 3·15 exposé of the "AI poisoning" supply chain claimed that GEO-cheating software can generate a dozen-plus fake reviews around a fictitious product within minutes, and get multiple mainstream large models to rank a non-existent product as the top recommendation within 72 hours; in the same month China's first Generative Engine Optimization (GEO) Industry Self-Discipline Convention was signed in Beijing, and the China Advertising Association launched GEO standardization [third-party reported]. This one currently comes relayed via a vendor press release; the specific details must be checked against CCTV's original program and the list of signatories. A bought position is one-off — once you are judged a low-credibility source, all three supply lines are lost at once.

## Sources for this chapter

[1] Howard & Sheth (1969) evoked set and the two-stage purchasing process; retro-cited via Hauser & Wernerfelt (1990) | [official] [2] Hauser, J. R. & Wernerfelt, B. | An Evaluation Cost Model of Consideration Sets | Journal of Consumer Research, Vol.16 No.4, pp.393–408 | 1990-03 | https://web.mit.edu/bwerner/www/papers/AnEvaluationCostModelofConsiderationSets.pdf | [official] [3] Andrews, R. L. & Srinivasan, T. C. | Studying Consideration Effects in Empirical Choice Models Using Scanner Panel Data | Journal of Marketing Research | 1995 | https://journals.sagepub.com/doi/abs/10.1177/002224379503200105 | [official] [4] John Dawes (Ehrenberg-Bass Institute), a study for the LinkedIn B2B Institute | reported via Marketing Week 2021-07-15 | https://marketingscience.info/news-and-insights/

ehrenberg-bass-95-of-b2b-buyers-are-not-in-the-market-for-your-products | [third-party reported] [5] LinkedIn B2B Institute | The 95-5 Rule | https://business.linkedin.com/advertise/resources/b2b-institute/b2b-research/trends/95-5-rule | [official] [6] Byron Sharp (Ehrenberg-Bass Institute) | Mental availability is not awareness | 2025-11-27 | https://marketingscience.info/news-and-insights/mental-availability-is-not-awareness-brand-salience-is-not-awareness | [official] [7] 6sense | 2025 B2B Buyer Experience Report | 2025-11 | https://6sense.com/science-of-b2b/buyer-experience-report-2025/ | [vendor claim] [8] Tversky, A. & Kahneman, D. | Judgment under Uncertainty: Heuristics and Biases | Science, Vol.185 No.4157, pp.1124–1131 | 1974-09-27 | https://www.socsci.uci.edu/~bskyrms/bio/readings/tversky_k_heuristics_biases.pdf | [official] [9] Johnson, E. J. & Goldstein, D. | Do Defaults Save Lives? | Science, Vol.302, pp.1338–1339 | 2003-11-21 | https://papers.ssrn.com/sol3/papers.cfm?abstract_id=1324774 | [official] [10] Our own measurement 1 · multi-model candidate-set test | 2026-09-13 | method and raw data in `实测/实测1-候选集多模型/` (raw.json, 结果.csv, 问题集.json) | [our own measurement] [11] Yang, K.-C. | Analysis of citation patterns in AI Search Arena (24,000+ conversation groups / 65,000 answers / 366,000+ citations) | arXiv: 2507.05301v1 | 2025-07-07 | https://arxiv.org/html/2507.05301v1 | [official] [12] OpenAI | ChatGPT shopping research | 2025-11-24 | https://openai.com/index/chatgpt-shopping-research/ | [official] [13] Ahrefs | What drives brand visibility in AI | 2025-12 | https://ahrefs.com/blog/ai-brand-visibility-correlations/ | [vendor claim] [14] Leapd | How ChatGPT, Google AI Overviews and Perplexity Source Information in 2026 | 2026-04 | https://www.leapd.ai/blog/ai-visibility/how-chatgpt-google-ai-overviews-and-perplexity-source-information-in-2026 | [vendor claim] [15] 2026 GEO generative engine optimization industry research report (republished via Zhihu/Sina) | around 2026-04-03 | https://zhuanlan.zhihu.com/p/2012459780288046392 | [third-party reported] [16] KNKX Public Radio and Cascade PBS | Bellingham staffer asked ChatGPT to "exclude" vendor from city contract | 2026-01-05 | https://www.knkx.org/government/2026-01-05/city-of-bellingham-chatgpt-ai-contract-vendor | [third-party reported] [17] CCTV 3·15 Gala 2026 "AI poisoning" report and the signing of the GEO Industry Self-Discipline Convention | relayed via a vendor press release | https://www.sohu.com/a/1068667519_100119123 | [third-party reported] [18] Gartner | Sales Survey: 74% of B2B buyer teams demonstrate unhealthy conflict | 2025-05-07 | https://www.gartner.com/en/newsroom/press-releases/2025-05-07-gartner-sales-survey-finds-74-percent-of-b2b-buyer-teams-demonstrate-unhealthy-conflict-during-

the-decision-process | [vendor claim] [19] ZKH (震坤行) "AI product recommendation brain" | relayed via a Zhihu column | https://zhuanlan.zhihu.com/p/1954120314318481065 | [vendor claim]

[20] Strategi | AI Search Visibility Statistics 2026 (five-engine source overlap 16%–59% / brand overlap 36%–55%; Ahrefs' 75,000 brands: pairwise engine brand-mention correlations 0.749–0.821) | https://strategi.is/blogs/ai-search-visibility-statistics-2026-the-engines-disagree-on-sources-and-agree-on-brands | [third-party verified] [21] G2 | The Answer Economy (March 2026 survey of 1,076 B2B software buyers and decision-makers: 51% more often start research with an AI chatbot, 29% in the 2025 edition; 71% rely on one at some point; 85% rate a vendor higher when AI includes it in an answer) | 2026-04-15 | https://company.g2.com/news/g2-research-the-answer-economy | [third-party verified] [22] Forrester | Preference Matters More Than In-Market Intent Alone in Modern B2B Buying (68% of B2B buyers start a purchase already favouring a vendor) | 2026-08-05 | https://www.forrester.com/blogs/preference-matters-more-than-in-market-intent-alone-in-modern-b2b-buying/ | [third-party verified]

[23] BrightLocal | Local Consumer Review Survey 2026, AI mini-report (1,002 US adults: AI for local recommendations 6% → 45%, ChatGPT 31%; reliance on Google reviews 83% → 71%; 88% of AI users fact-check the recommendation) | 2026-03-10 | https://www.brightlocal.com/research/lcrs-ai-trust/ | [third-party verified]

[24] G2 | 2026 Buyer Behavior Report: The Evaluation Maze (online survey of 1,038 B2B software decision-makers in June 2026 across North America, EMEA and APAC; AI makes candidates easier to find but evaluation harder; only 2% would let AI complete a purchase without pre-approval) | 2026-07-22 | https://www.prnewswire.com/news-releases/ai-is-making-software-easier-to-find-and-harder-to-buy-according-to-new-g2-research-302831346.html | [third-party verified] [25] G2 | 2026 Buyer Behavior Report: The Evaluation Maze (1,038 B2B software decision-makers plus 55 in-depth interviews; evaluation 40% of the buying journey versus research 36%; 82% sourced recommendations from an AI chatbot in 24 months; review sites 38% ahead of chatbots 37% for the first time; IT security review 39%, 50% in enterprises; budget approval 32%; finance involvement 31% to 46%; CFO veto of an approved purchase 49%; internal resistance 16% to 29%) | 2026-07-22 | https://www.tmcnet.com/usubmit/2026/07/22/10418570.htm | [third-party verified]

# Chapter 4: Let Third Parties Speak for You

Third-party endorsements used to travel person to person and through press clippings — whoever placed more articles looked more credible; now buyers also ask AI the question "whom should I trust?", and the model only repeats what it has read elsewhere, no matter how well your own website is written. In 6sense's survey of nearly 4,000 B2B buyers, in 95% of cases the vendor that ultimately won the deal had already appeared on the buyer's "day-one candidate list" before the buyer first contacted sales [vendor claim]. What you need to do is shift your effort from "saying it yourself" to "getting others to say it," and cultivate the statements you yourself cannot call the shots on.

What buyers see at the AI is not what you say but how others evaluate you. So this chapter gives you 6 methods for cultivating the evidence that lives "on other people's turf."

---

## Method 1: Map who speaks for you and who gets you wrong

**What to do**: Build a table of the places that "speak for you" and the places that "get you wrong," split into three categories, and keep it updated.

**How**

1. Sweep review sites, rankings, industry media, community posts, podcasts, and wiki-type pages, and sort every entry into just three categories: mentioned and accurate, mentioned but inaccurate, not mentioned. The third category is where your next investment should go.
2. Take 6–10 real procurement questions — the ones your salespeople are asked over and over on calls — and ask each of them several times across 5 AI channels, noting the context in which you are mentioned. The point is to record **who is being cited**, not only whether you show up; the sources being cited are the ones you need to court.
3. Repeat every 30 days, watching four numbers: the number of citations, your standing in third-party rankings and reviews, the accuracy rate of what gets relayed about you, and the number of sessions AI refers to you.

4. Remember that this kind of evidence is not meant to move people; it is meant to save the buyer one round of personal verification. Buyers were never going to compare exhaustively, and the size of the candidate set is the optimum they calculate between "decision cost" and "the marginal benefit of shopping around more"; an independent review or a ranking position saves them one verification step [third-party verified].
5. Read the same table again through the lens of how the buying group divides its work. Finance wants to know the overrun boundary, IT wants to know integration, legal wants to know the terms, and the business side wants to know whether anyone actually uses it — a Gartner survey of 632 B2B buyers in May 2025 found that 74% of buying teams experience "unhealthy conflict" during the decision process, while customization aimed only at individual-level relevance has a 59% negative impact on the buying group's ability to reach consensus [vendor claim].

**Example**: In Ahrefs' study of 75,000 brands, the correlation between brand mentions and AI visibility was about 3 times that of traditional backlinks (on the order of 0.664 versus 0.218), while the number of content pages was almost irrelevant (about 0.194) [vendor claim].

**Pitfalls**: Fixating on the single question of "did I show up?" What actually needs fixing is the "who is being cited" list, not the wording on your own website.

(Industry differences: SaaS looks at rankings and technical communities like G2 and TrustRadius; consumer and local life look at review sites and notes; manufacturing and industrial look at associations, trade shows, and technical communities — the gateways differ, the action is the same.)

## Method 2: Spend your PR budget on making material

**What to do**: Move money and people from "getting yourself to appear" to "making material that third parties can cite."

**How**

1. Split the money along a workable structure: core SEO 40%, digital PR a quarter, data and reports 20%, training 10%, experiments 5% [third-party verified]. The ratios need not be copied, but the "digital PR" line has to be recast as making material, not placing articles.

2. Switch the KPI track: article placement volume and clipping counts are void; look instead at "the number and context of independent third-party citations" and "your standing in third-party rankings and reviews."
3. Pick 3 high-authority third parties — one vertical review site, one industry publication, one video or podcast channel. Do just one thing for each: give it something it would want to cite and that only you can provide, such as an exclusive data set, a comparison checklist, an open-source tool, or a citable expert explanation.
4. Write the goal as "being cited," not "being published." Publication only proves you bought a slot; citation proves your material is useful.
5. Rewrite your internal fact list into a form a third party can copy directly: a one-sentence conclusion + the data + the time point + a verifiable link. A form anyone can copy away is a good form.
6. Before you do this, accept one comparison: among the consumer brands that rose on "cheap Meta ads + Shopify + subsidies," the ones that survived had one thing in common — they were mentioned most by third parties (review sites, editorial media, podcasts, creators); the ones that spent the most on acquisition did not survive [vendor claim].

**Example**: In its February 2024 S-1 filing, Reddit disclosed that the data licensing agreements signed in January 2024 were worth a total of $203 million with terms of 2–3 years, that it expected to recognize at least $66.4 million in revenue that year, and its risk factors stated that substantially all of the contract value came from a single partner [official]. Third-party content is not optional, it is the main fuel: an analysis of 25 million+ citation links across 17 industries for ChatGPT, Claude and Gemini found earned media (press coverage, academic research, government and institutional sources) account for 84% of all AI citations, paid and sponsored content for only 0.3%, and news coverage alone for 27%; the ratio has held between 82% and 89% across three editions since July 2025 [third-party verified]. Platform differences are just as large: ChatGPT gives citations in 96% of replies but only about 5 of them, Claude gives citations in only 55% of replies but about 13 when it does, and the top cited domains are Wikipedia, PubMed Central and Reddit respectively [third-party verified] — one release pushed to all three does nothing; platform choice beats keyword choice. A study that converts visibility into credibility points the same way. Burson (Chinese: 博雅集团), working with the AI marketing platform Profound, tested thousands of reputation-related answers across seven major AI answer platforms, scored 85

companies on eight reputation dimensions, and generated more than 55,000 credibility predictions; its finding is that AI favours evidence over positioning — objective statements about innovation, products and workplace culture consistently beat subjective attributes such as leadership, governance and corporate citizenship — and the independent corroboration that counts comes from media coverage, reviews and social discussion, with leadership, governance and corporate citizenship statements consistently ranking below objective ones [third-party verified: research institute].

**Pitfalls**: Judging source-building by click volume. In July 2025 Pew found that when an AI summary appears, only 1% of users click its citation links [third-party verified]; promising results on the basis of clicks is digging your own hole.

## Method 3: Get customers to write reviews for you on public platforms

**What to do**: Make "ask the customer to leave a review on a public platform" a fixed step after signing.

**How**

1. After signing, set one action: ask the customer to leave a public-platform review (G2, TrustRadius, Dianping (Chinese: 大众点评), Tmall (Chinese: 天猫) reviews, and the like) that includes a specific usage scenario — where it is used, what it solved, and what still is not good about it.
2. Three lines must not be crossed: real customers, real experience, no conditions traded. Do not buy positive reviews, do not hire paid posters, do not trade benefits for reviews — not even for the "right to leave a bad review." Platform rules likewise prohibit inducing shares through monetary rewards, lotteries, and the like.
3. Customers may give you a low score, but every entry must be verifiable. When you hand the product over for third-party review, hand over your scope and boundaries along with it — better a low score than an unverifiable sample.

4. You may publish rankings you made yourself, but you must disclose that you made them. Disguising them as an "independent review" is squarely illegal: on August 14, 2024, the FTC passed the Consumer Reviews and Testimonials Rule by a 5-0 vote, making it explicit for the first time that a company-controlled review site falsely claiming independence is unlawful [official].
5. The logos on your case-study wall must be verifiable too. A March 2025 TechCrunch investigation found that the customer logos displayed on the website of AI SDR company 11x were not real customers — two of them had run only about a one-month trial before stopping — and the report said its total customer retention rate was below 50% [third-party verified].
6. Every quarter, feed the recurring misunderstandings and boundaries in reviews back into the failure-conditions list and the pre-sales FAQ. Devices that "voluntarily make a commitment" — warranties, branding, licensing — are exactly the kind of thing the market grew on its own to counter information asymmetry, and public customer reviews belong to the same family [third-party verified].

**Example**: In the TruHeight case, the FTC alleged that thousands of five-star reviews on its website were written by company employees, that five-star reviews were exchanged for free or discounted products, and that software-generated reviews were posted from bot accounts disguised as real users; it ended in a $4 million judgment [official]. Review sites can be quantified: monitoring of roughly 6,800 procurement-intent queries shows that brands in the top 20 of G2 or Capterra for their category are cited in AI answers to "best [software category]" questions about 3.1 times as often as brands absent from both platforms; review volume matters independently — profiles with 100 or more reviews are cited roughly 2.3 times as often as profiles with fewer than 20, even at equal average star ratings [third-party verified]. In the same dataset, a G2 Leader or High Performer badge adds another 20%–35% [third-party verified]. And the value of this channel is migrating from clicks into the answer itself: G2 shows roughly 23,700 citations in Perplexity, 12,000 in ChatGPT, 10,900 in Google AI Mode, 8,800 in AI Overviews and 5,200 in Gemini, while its own organic search traffic has fallen from its 2023 peak [third-party verified] — review sites have not stopped working; the place buyers read them moved. Being cited is not the same as being described correctly: after a month of testing four publicly available assistants (ChatGPT, Copilot, Gemini, Perplexity), the BBC judged 51% of news-related AI answers to have some form of significant issue; among

answers citing BBC content, 19% introduced factual errors (wrong statements, numbers or dates) and 13% of quotes said to come from BBC articles were altered or did not exist in the article at all [official: institutional study]. Columbia University's Tow Center tested eight generative search tools and found the same direction: they generally will not say "I don't know", paid versions are more likely than free ones to give confident wrong answers, they invent links or cite reprints rather than originals, and content licensing deals do not guarantee accurate citations [third-party verified]. China now has a citable ruling on who carries the liability: the Hangzhou Internet Court heard a tort claim over inaccurate information from a generative AI app and dismissed it in full — AI has no civil legal personality and its "promise" is not a binding declaration of intent; information given in conversation counts as a service rather than a product, so fault-based liability applies; a provider meets its obligations by giving prominent notice of the AI's limits and by adopting industry-standard, proven measures to improve reliability [official: court ruling reported by central media]. A cross-cut from the review platforms is blunter. SE Ranking analysed 30,000 commercial keywords, 22,729 of which triggered an AI Overview, and found 34.5% of those AI Overviews cited at least one review platform — rising to 49% when the query wording included explicit review or rating terms and falling to 17.1% for list-style "best/top" queries; citations concentrate in five platforms (Gartner Peer Insights, G2, Capterra, Software Advice and TrustRadius), which together account for 88% of review-platform links while review platforms make up only 8.5% of all links; meanwhile those same platforms lost 76.5% to 92.2% of their organic traffic between early 2024 and the end of 2025 [third-party verified]. Reviews have a ceiling too, and it has been measured: a study covering 921 B2B SaaS brands with G2 or Trustpilot ratings found a correlation of just 0.030 between ratings and the way AI describes a brand, statistically indistinguishable from zero; in G2's own data, about 10% more reviews corresponds to only about 2% more citations; and a re-measurement of 291,607 ChatGPT citations found review sites making up 2% of what ChatGPT reads and 6% of what it cites [third-party verified] — more a retrieval asset for review-intent questions than a lever on what the model remembers.

**Pitfalls**: Buying reviews as advertising material. Among the typical cases announced after China's Interim Provisions on Anti-Unfair Competition in Cyberspace in May 2024, a company in Hangzhou provided "shoot A, ship B" order-brushing

services for 589,000 orders and was fined 725,000 yuan; a club paid 560 yuan to hire people to fabricate 8 fake negative reviews slandering a competitor and was fined 10,000 yuan [official: official-media relay].

## Method 4: Put the founder front and center, and say unfavorable things under a real name

**What to do**: Have the founder or a frontline expert speak under his or her own name — about the situation, about losses, about "this kind of customer is not a fit for us."

**How**

1. Pick the person: real name, first person. The content team only tidies up; it does not ghostwrite — once ghostwriting becomes widespread, the cost difference disappears and the whole action becomes pointless.
2. Pick the subject matter: losses, failure cases, unfavorable data, experiences of being demoted by a platform. Competitors can copy your opinions, but not your situation.
3. Pick the frequency: at least one piece a month, signed, with verifiable numbers, and with primary materials attached wherever possible.
4. Unfavorable statements must come with judgment criteria and remedial actions. Laying out only the problems without giving what to do is self-exposure.
5. Use trust data to settle internal arguments so you do not have to debate aloud round after round: Edelman's 2025 Trust Barometer Special Report on Brands lists "brands must move from SEO to GEO" as the first of its ten findings, and states that what appears in AI is determined by reputation, relevance, credibility, and clarity, and driven by earned (word-of-mouth and third-party endorsement); the same-period figure is that 80% of people trust the brands they use, and 91% of people who use generative AI platforms use them in some way for shopping research [official: institutional report].
6. Do not turn "human touch" into a template. Signaling economics says it plainly: for a signal to be credible, the cost of obtaining it must differ enough between senders; if everyone's cost is the same, it carries no information [third-party verified].

**Example**: Scott Lapatine, founder of the independent music publication Stereogum, reported in the first person in a relaunch announcement that Google's shift toward AI search cut their ad revenue by 70%, and at the same time declared "we do not use generative AI" [official].

**Pitfalls**: Thinking that "coming across as sincere" is enough. In SurveyMonkey's survey of 2,017 U.S. adults, 79% strongly preferred human service and only 8% clearly preferred AI, while 14% would stop trusting a company because "the AI they interacted with did not disclose its identity" [vendor claim] — buyers trust a statement that can be traced back to a specific person.

## Method 5: In communities, answer questions first — do not lead with your product

**What to do**: Make community answers a measured, fixed practice, and build believability in places where the rules are not yours to write.

**How**

1. Choose 1–2 real people (an engineer or a product lead) to answer questions on Reddit, Zhihu (Chinese: 知乎), industry forums, and Discord under their real names and company identities, without slipping in links or promotions.
2. Give before you take: answering one specific question is closer to external evidence than posting one product introduction.
3. Disclose your identity: who is speaking, whom they represent, and what relationship they have to the brand — all stated plainly.
4. Set a rhythm, and a ledger: answer a few questions on a fixed weekly schedule, and record what the question was, where it came from, who answered it, and whether it was cited.
5. The "real-person IP" in a private domain can only be run by real people. The moment it is switched to batch automated scripts, the community turns from an asset into a liability.

**Example**: In Greenway's 800-account controlled experiment, fully autonomous AI booked the most meetings, but the meeting-to-opportunity conversion rate was only 8%, versus 31% for human sales and 26% for a human-machine mix; aggressive sending knocked the primary domain's reputation from good to suspicious in 19 days, and recovery took 8 weeks [vendor claim].

**Pitfalls**: Treating "bulk accounts + soft articles + bots" as reach. In the short term it looks like reach; in reality it is overdrawing believability, and the recovery period is long.

## Method 6: Review every month how AI describes you, and fix the source when it is wrong

**What to do**: Build a monitoring-and-correction process for "being described wrongly," and push errors back to the source to fix.

**How**

1. Review AI's restatement of you every month, focusing on three places: price, functionality, and scope of applicability. When you find something attributed to the wrong thing, log it — which sentence, which platform, what date.
2. For corrections, use third-party-verifiable material to change the source; do not argue on your own website. A rebuttal on your own site is something you said yourself, and both the model and the buyer will discount it.
3. Distinguish three kinds of payoff: being cited is a brand payoff and can compound; being clicked is a traffic payoff and is declining; being trusted is a conversion payoff. The trust column is in fact the strongest — Adobe Analytics' March 2026 figures are that AI-referred visitors to U.S. retail sites converted 42% higher than non-AI visitors and stayed 45% longer [third-party verified: vendor-reported]. Mix them together in your accounting and you will certainly get it wrong.
4. Monitor what competitors do: buying ranking positions, inflating reviews, malicious comparison tests — keep evidence of all of it and go through the reporting channels.
5. Do not bet visibility on a single point. One platform changes its rules, or one crawler gets blocked, and you are cut off; two or more mutually independent source lines are what survive a single change. Content licensing and unauthorized crawling are being fought out in court, and platforms are installing valves — publishers that want to block AI crawlers often fall out of the search index at the same time; in China, gateway issues like "blocking web links" have been made a focus of special rectification [official].

**Example**: Digital Content Next's eight-week tracking of 19 members showed a median year-over-year decline of 10% in referral traffic from Google search; among the evidence submitted to UK regulators was also a case of "ranked first but click-through rate fell from 5.1% to 0.6%" [third-party verified].

**Pitfalls**: Using AI-referred traffic as the acceptance criterion for source-building. The click column is low to begin with, and AI-referred visitors are not of poor quality — one media outlet has self-reported that the subscription conversion rate of its AI-channel visitors was 4–5 times that of traditional search [third-party verified: media-reported].

---

## Notes for this chapter

- The Ahrefs, Peec, Leapd, 6sense, Gartner, Greenway, SurveyMonkey and other data cited in this chapter all come from vendors' or research firms' own work, whose methodology appendices are not public, so all count as [vendor claim]. They agree in direction (mentions beat backlinks, sources are highly concentrated) and can corroborate one another, but **the specific coefficients cannot be extrapolated into industry constants**.
- "Being cited by AI" and "getting traffic" are two separate sets of books. Third-party media tracking shows the citation table rising and the click table falling (DCN members' referral traffic median -10%, Stereogum's ad revenue -70%), so source-building buys eligibility and cannot be treated as a traffic promise.
- No verifiable public pricing case can currently be found for content licensing in the Chinese ecosystem; the only price reference in this chapter is Reddit's $203 million deal; domestically the issue is more one of gateway blocking and interconnection, and the two sides are not directly comparable.
- Third-party ranking positions and review placements can be bought or inflated, so you have to check the credibility of the rankings yourself. The FTC has already written "a company-controlled review site falsely claiming independence" into law as unlawful, and on the Chinese side order-brushing and hiring people to write negative reviews also already have penalty cases.
- The claim that the Google–Reddit licensing fee is "about $60 million a year" comes from a rumored figure relayed by the media; the S-1 does not name the amount or the counterparty, so this book does not adopt that number.

## Sources for this chapter

[1] Ahrefs | AI brand visibility correlations (75,000-brand study) | https://ahrefs.com/blog/ai-brand-visibility-correlations/ | 2025-12 | [vendor claim] [2] Peec AI | 70+ Generative Engine Optimization Statistics for 2026 | https://peec.ai/ai-search-geo-statistics | updated 2026-08 | [vendor claim] [3] Leapd | How ChatGPT, Google AI Overviews and Perplexity Source Information in 2026 | https://www.leapd.ai/blog/ai-visibility/how-chatgpt-google-ai-overviews-and-perplexity-source-information-in-2026 | 2026-04 | [vendor claim] [4] EMARKETER | FAQ on GEO and AEO: Where AI search and SEO overlap in 2026 | https://www.emarketer.com/content/faq-on-geo-aeo--where-ai-search-seo-overlap-2026 | 2026-04-02 | [third-party verified] [5] 6sense | 2025 B2B Buyer Experience Report | https://6sense.com/science-of-b2b/buyer-experience-report-2025/ | 2025-11 | [vendor claim] [6] Gartner | 74% of B2B buyer teams demonstrate unhealthy conflict | https://www.gartner.com/en/newsroom/press-releases/2025-05-07-gartner-sales-survey-finds-74-percent-of-b2b-buyer-teams-demonstrate-unhealthy-conflict-during-the-decision-process | 2025-05-07 | [vendor claim] [7] Reddit Inc. | Form S-1 (SEC EDGAR) | https://www.sec.gov/Archives/edgar/data/1713445/000162828024006294/reddits-1q423.htm | 2024-02-22 | [official] [8] Pew Research | AI summaries and link clicks | 2025-07 (relayed via Digiday/AdExchanger) | [third-party verified] [9] Digiday | Google AI Overviews linked to 25% drop in publisher referral traffic (measured across 19 DCN members) | https://digiday.com/media/google-ai-overviews-linked-to-25-drop-in-publisher-referral-traffic-new-data-shows/ | 2025-08-15 | [third-party verified] [10] AdExchanger | The AI Search Reckoning Is Dismantling Open Web Traffic | https://www.adexchanger.com/publishers/the-ai-search-reckoning-is-dismantling-open-web-traffic-and-publishers-may-never-recover/ | 2026-01-06 | [third-party verified] [11] Stereogum | Getting Killed By AI (byline: Scott Lapatine) | https://stereogum.com/2478838/stereogum-relaunch/news | 2025-11-10 | [official: founder's own account] [12] Greenway | AI SDR: What Works, What Doesn't (800-account controlled experiment) | https://greenway.ai/post/ai-sdr-what-works-what-doesnt/ | [vendor claim] [13] SurveyMonkey Curiosity | 2026 customer service statistics (2,017 U.S. adults) | https://www.surveymonkey.com/curiosity/customer-service-statistics/ | 2026 | [vendor claim] [14] FTC | Consumer Reviews and Testimonials Rule (16 CFR Part 465) | https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-trade-commission-announces-final-rule-banning-fake-reviews-testimonials | 2024-08-14 | [official] [15] FTC | TruHeight case ($4 million judgment) | https://www.ftc.gov/news-events/news/press-releases/2026/04/ftc-takes-

action-against-truheight-deceptive-unsubstantiated-advertising-supposed-height-enhancing | 2026-04-13 | [official] [16] TechCrunch | a16z- and Benchmark-backed 11x has been claiming customers it doesn't have | https://techcrunch.com/2025/03/24/a16z-and-benchmark-backed-11x-has-been-claiming-customers-it-doesnt-have/ | 2025-03-24 | [third-party verified] [17] Xinhua (Chinese: 新华网) | SAMR announces seven typical cases of online unfair competition | http://www.news.cn/food/20241231/ceec0fa0841345fb872ed48b07dd3748/c.html | 2024-12-31 | [official: official-media relay] [18] Woshipm (Chinese: 人人都是产品经理) | Private-domain traffic methodology (including the Perfect Diary (Chinese: 完美日记) path) | https://www.woshipm.com/operate/4088812.html | 2020-07 | [self-media] [19] Michael Spence | Job Market Signaling | QJE 1973 | https://www.sfu.ca/~allen/Spence.pdf | 1973-08 | [third-party verified] [20] 5W Public Relations | The DTC Graveyard 2026 (Clicks vs Citations) | https://www.5wpr.com/research/dtc-graveyard-2026/ | 2026-04 | [vendor claim] [21] Hauser & Wernerfelt | An Evaluation Cost Model of Consideration Sets | JCR 1990, Vol. 16 No.4 | https://web.mit.edu/bwerner/www/papers/AnEvaluationCostModelofConsiderationSets.pdf | 1990-03 | [third-party verified] [22] George A. Akerlof | The Market for "Lemons" | QJE 1970 | https://econweb.ucsd.edu/~atimmerm/174/akerlof.pdf | 1970-08 | [third-party verified] [23] Edelman | 2025 Trust Barometer Special Report: Brands (15,000 people in 15 countries) | https://www.edelman.com/trust/2025/trust-barometer/special-report-brands | 2025-04 | [official: institutional report] (the link returns 403 as anti-scraping; the page itself is valid) [24] Adobe Analytics (AI referral conversion) | 2026-03 (relayed via Peec) | [vendor claim] [25] China News Service (Chinese: 中国新闻网) | MIIT: blocking web links is a focus of special rectification, with a deadline to lift the blocks | https://www.chinanews.com/cj/2021/09-13/9564236.shtml | 2021-09-13 | [official: official-media relay] [26] WeChat (Chinese: 微信) "External Link Content Management Rules" | prohibition of inducing shares through enticements | (this book did not obtain a link to the original text; used only as directional reference) | [third-party verified: relayed]

[27] Presenc AI | Do G2 and Capterra Reviews Improve AI Visibility? (about 6,800 procurement-intent queries: top-20 listings cited 3.1x; 100+ reviews 2.3x; Leader badge +20%–35%) | 2026-05 | https://presenc.ai/research/does-g2-capterra-reviews-improve-ai-visibility-2026 | [third-party verified] [28] Foundation Inc | How G2 Reviews Are Building the Trust Layer of AI Search (about 23,700 citations in Perplexity, 12,000 in ChatGPT, 10,900 in AI Mode, 8,800 in AI Overviews, 5,200 in Gemini) | https://foundationinc.co/lab/g2-reviews-ai-trust | [third-party verified] [29] Muck Rack /

Generative Pulse | What Is AI Reading? (25 million+ citation links; earned media 84%, paid and sponsored 0.3%, news 27%; differences in citation counts and top cited domains across ChatGPT/Claude/Gemini) | 2026-05 | https://muckrack.com/blog/what-is-ai-reading-may-2026 | [third-party verified] [30] BBC Media Centre | BBC research shows issues with answers from AI assistants (51% of news answers had significant issues; 19% of answers citing BBC content contained factual errors, 13% altered quotes) | 2025-02-11 | https://www.bbc.com/mediacentre/2025/bbc-research-shows-issues-with-answers-from-artificial-intelligence-assistants | [official: institutional study] [31] Columbia Journalism Review / Tow Center | We Compared Eight AI Search Engines: They're All Bad at Citing News (paid tiers more likely to give confident wrong answers; invented links; licensing deals do not guarantee accuracy) | 2025-03-06 | https://www.cjr.org/tow_center/we-compared-eight-ai-search-engines-theyre-all-bad-at-citing-news.php | [third-party verified] [32] Xinhua (Chinese: 新华网) | Hangzhou Internet Court rules that inaccurate AI information is not automatically tortious (AI has no civil legal personality; fault-based liability; duties of prominent notice and basic reliability assurance) | 2026-03-30 | https://www.news.cn/legal/20260330/670d3456cc1a48be8dd9df42ad13e472/c.html | [official: court ruling reported by central media]

[33] Burson × Profound | The Credibility Paradox (seven AI platforms, 85 companies, 55,000+ credibility predictions; AI favours evidence over positioning; leadership ranks last; business decision-makers trust AI answers about 10% more than the general public) | 2026-06-02 | https://www.bursonglobal.com/newsroom/china/从可见度转向可信度-博雅集团研究洞见geo关键鸿沟 | [third-party verified: research institute]

[34] SE Ranking | Despite 90% Traffic Loss, Review Platforms Top AI Overview Citations (30,000 commercial keywords, 22,729 triggering an AI Overview: 34.5% cite a review platform; 49% for explicit review terms vs 17.1% for list queries; top five platforms 88% of review links, review platforms only 8.5% of all links; organic traffic down 76.5% to 92.2%) | 2026-01-29 | https://seranking.com/blog/review-platforms-in-ai-overviews/ | [third-party verified]

[35] Discovered Labs | Do G2 reviews change what AI thinks of you? (921 B2B SaaS brands with G2 / Trustpilot ratings: correlation of 0.030 between ratings and how AI describes the brand; in G2 data, +10% reviews maps to +2% citations; re-measurement of 291,607 ChatGPT citations: review sites are 2% of what it reads and 6% of what it cites) | 2026 | https://discoveredlabs.com/blog/do-g2-reviews-change-ai-perception | [third-party verified]

# Chapter 5: Lead with Your Own Pitfalls

Vendors' brochures used to all look alike, and whoever shouted "faster, more stable, cheaper, higher satisfaction" the loudest came out ahead; now, after reading the tenth piece of material, the buyer's first question is "where is this thing most likely to go wrong?", and that question is often put to AI first. What you need to do is move the risk topic from a defensive position to the very front — the buyer's first question inside the model is "I want to buy X, what should I watch out for?", and that sentence is your exam room.

What customers most love to ask is "what are the pitfalls?", and whoever answers first takes the right to define. So this chapter gives you 6 methods for turning the risk topic into a gateway.

---

## Method 1: Write a "when not to pick us" document first

**What to do**: Write the boundaries of who you are not a fit for as conditions, and hand it over together with the quote.

**How**

1. Cover five types of conditions: scale (team size, order volume, revenue range), stage (whether there is dedicated ops, whether there is a data foundation), process (whether the existing systems can be adapted, how long the approval chain is), compliance (whether the industry or region has hard constraints), and integration (which kinds of systems cause problems when connected).
2. Give each entry a criterion, not an adjective. "For a team with fewer than 2,000 orders a month, using our automated warehousing and fulfillment module will cost more per unit than outsourcing — that is our not-a-fit range" is a condition; "the company reserves the right of final interpretation" is a disclaimer, which only protects you and does not help the buyer eliminate options.
3. Put it on the first page of the pre-sales material and hand it over together with the quote; do not stuff it into a contract annex. Have sales walk through this list on the first visit and say the "not a fit" part within the first thirty minutes.

4. Update it once a quarter alongside the lost-deal records. New "not a fit" scenarios go in within two weeks, and the deals in progress are notified internally.
5. Remember why the buyer wants this: rejecting a vendor costs him almost nothing, while approving a vendor that turns out bad can cost him his job, so his default action is to eliminate first and choose second — and elimination depends on whether you give him conditions. The most robust finding in behavioral economics is that the pain of a loss outweighs the pleasure of an equivalent gain [third-party verified].

**Example**: Akerlof's 1970 "The Market for 'Lemons'" supplied the mechanism: when buyers can only judge quality by "average quality," the result is that average quality falls and the market itself shrinks with it [third-party verified] — stating your own boundaries proactively is how you pull yourself out of the pool of "suspected average-quality vendors."

**Pitfalls**: Writing the list as a disclaimer, or as adjectives like "we are not a great fit for very small customers." What the buyer needs is conditions he can use directly to eliminate.

(Industry differences: in regulated industries like finance, healthcare, and legal services, this list is the first page of the pre-sales material; in SaaS, failures almost all happen in the border zone of "usable, but not usable well," and the three conditions of scale, process, and ops screen people out the best.)

## Method 2: Pair every risk with "condition, criterion, remedy"

**What to do**: Discussing risk must come in a group of three; do not just lay out the problem.

**How**

1. Fix the order: under what conditions are we not a fit → how do you judge whether you are in that situation → if you already are, what can we do together.
2. Prepare the remedy before you publish. Laying out only the problem is self-exposure; laying out the problem, the judgment criteria, and the remedy together is a commitment.

3. Control the amount. Negative information and brand attitude have a curvilinear relationship: say too little and credibility is insufficient; say too much and the buyer files you into the "nothing about this works" bucket and skips you. The optimum can only be found by trial; the importance of the negative attribute, where the information is placed, and whether the positive and negative attributes are correlated all shift that point.
4. After disclosure, follow it with a trackable corrective action: what will change, how long it will take, and who gets notified when it is done.
5. Distinguish "voluntary disclosure" from "forced disclosure." For a signal to be credible, the cost of obtaining it must differ enough between senders; a risk notice written under regulatory pressure costs every peer the same, so it carries almost no information — only failure conditions written proactively into pre-sales material count [third-party verified].

**Example**: Telang and Wattal used an event-study method to measure the impact of software vulnerability announcements: within the sample, an average vulnerability being reported cost the vendor about 0.6% of its share price, and the loss was larger the more intense the competition and the smaller the vendor, while declines were more negative when the disclosure did not come with a patch [third-party verified].

**Pitfalls**: Treating "being candid" as a matter of attitude. In Eisend's meta-analysis, which integrated the literature from 1960 to early 2004 across a total of 217 effect sizes, the persuasive effect of two-sided information (covering both pros and cons) was significantly positive but only r = .068 [third-party verified] — it mainly improves source credibility and perceived novelty, and expecting it to directly lift preference will disappoint you.

## Method 3: Turn "which items to check" into a checklist

**What to do**: Do the buyer's elimination work for him — give criteria, without naming competitors.

**How**

1. Write it according to how the buying group divides its work. Finance worries about overruns, IT worries about integration, legal worries about terms, and the business side worries that no one actually uses it — a Gartner survey of 632 B2B buyers in May 2025 found that 74% of buying teams experience "unhealthy

conflict" during the decision process, while customization aimed only at individual-level relevance has a 59% negative impact on the buying group's ability to reach consensus [vendor claim]. A checklist on which no side has to prove itself is the common language they can return with after arguing it out inside the group.

2. Give criteria only, and do not name competitors. Once it becomes "compared with so-and-so," it degrades into attack advertising; in the Chinese ecosystem, naming names in a comparison also stacks on the risk under the Anti-Unfair Competition Law, and since this book has not checked the specific provisions against the original text, we do not go into it here.
3. Leave blanks for the other side to fill in. Parameters, weights, and acceptance criteria are all listed as fields to be filled — a checklist the buyer completes himself is more persuasive than your conclusion.
4. Version it, attach an owner, and keep a trail of changes.

**Example**: Public records released by KNKX and Cascade PBS in January 2026 show that an employee in Bellingham, Washington, repeatedly asked ChatGPT to write bid requirements "favoring a certain vendor," and at least 16 of the model-generated phrases appeared word for word in the evaluation matrix for the final 350 requirements; the five-year, $2.7 million contract was ultimately awarded to the higher bidder [third-party verified]. There is a magnitude for being described wrongly, too: one technical-community write-up estimates that about 23% of AI answers contain at least one factual error, and that the error rate for brand-related content runs as high as 31%; one appliance maker measured mainstream AI platforms answering questions about its prices and found 35% of answers carried a price deviation, ranging from 10% to 200% [third-party verified: technical community relay, no original study attached]. Price is the hardest and most expensive line on any checklist — fixing "price, conditions of use, service boundaries" as a monthly spot-check is cheaper than waiting for a customer to arrive quoting a wrong price.

**Pitfalls**: Thinking it is fine for the "criteria for a qualified vendor" slot to sit empty. If you do not define it, someone else's AI will.

## Method 4: Make the total price and the hidden fees clear up front

**What to do**: Combine all mandatory fees in the quote into one true total price and put it in the most prominent place.

**How**

1. Fold implementation fees, data migration fees, first-year service fees, and return, exchange, and warranty conditions all into the main number on the quote — not into the footnotes.
2. The total price must be more prominent than other price information. The FTC's "junk fees rule" requires merchants, when offering, displaying, or advertising any price, to clearly and conspicuously disclose the true total price including all mandatory fees; the rule also states that it does not prohibit any type or amount of fee, nor any specific pricing strategy — it only requires that it be told truthfully, up front, and completely [official].
3. For subscriptions, put auto-renewal, price-increase mechanisms, and how to cancel on the same page — do not make the buyer make three phone calls to figure it out.
4. Change the outreach side too. Commercial calls may not be made without the consumer's consent, and this applies to AI outbound calls just as it does to human ones; on the Chinese side, in response to the chaos of AI outbound calls, experts have already proposed folding the number of calls per day and the user complaint rate into the criteria for finding a violation — and this kind of calling is often indiscriminate blind dialing, hitting even users who never registered [third-party verified].
5. Have legal review the pricing page and the outreach scripts together — do not wait until you get a complaint to go back and patch it.

**Example**: The FTC estimates this rule can save consumers up to 53 million hours of comparison shopping a year, worth more than $11 billion over ten years; the rule was adopted on December 17, 2024 (by a 4-1 vote) and took effect on May 12, 2025 [official].

**Pitfalls**: Treating "transparent pricing" as a differentiating selling point. It is already the compliance starting line; failing to state it clearly hands the buyer to someone who will. (In categories like consumer, cross-border, and local life, taxes, returns, warranties, and auto-renewal all have to be stated on the same page.)

## Method 5: Keep a risk ledger with a named owner

**What to do**: Register all known boundaries, failure cases, not-a-fit scenarios, and unfixed issues in one place, with a version and an owner.

**How**

1. Fixed fields: risk description, scope of impact, judgment criteria, remedy, owner, version, update date.
2. It and the fact list are two views of one underlying ledger: one records "what we are," the other records "where we will go wrong." Where the two do not match is the place most in need of fixing first.
3. Review each entry once a quarter, and mark the progress and expected timing of unfixed issues.
4. If you get it wrong, there must be a correction process: who fixes it, how long it takes, and who gets notified. This one item determines whether a disclosure is a complete transaction or a half-finished statement.
5. Zero tolerance for false claims and fake reviews. The FTC's enforcement chain can be read chronologically: in August 2024 it issued the Consumer Reviews and Testimonials Rule → in December 2025 it sent warning letters to 10 companies (with civil penalties of up to $53,088 per violation) → in April 2026 a case with a judgment amount appeared [official]. The first category the rule names as prohibited is AI-generated fake reviews.
6. List the Chinese side's second set of touchpoints in a separate column: AI-generated content must identify itself, distribution platforms must verify labels, commercial calls may not be made without consent, and platforms also throttle keywords like "intelligent outbound calling" [third-party verified].

**Example**: In the TruHeight case, thousands of five-star reviews on its website were written by company employees, five-star reviews were exchanged for free or discounted products, and software-generated reviews were posted from bot accounts

disguised as real users; it ended in a $4 million judgment. Structurally, the point most worth remembering is that "false efficacy claims" and "fake reviews" were penalized together [official].

**Pitfalls**: Turning the ledger into the marketing department's PR copy — writing only what has been fixed and not what has not; or the ledger gets built by someone but no one signs it. The signature slot is one the regulators have already marked out for you.

(Industry differences: in manufacturing and industrial, the risk is not in features but in delivery and compatibility, and "which kinds of production lines require adaptation when connected, how long the adaptation takes, and how the downtime window is arranged" is the checklist; in the education market, the signal strength of "what kind of student is not suited to this course" outranks pass-rate advertising [our judgment].)

## Method 6: Draw topics from lost deals and complaints, and review three numbers every month

**What to do**: What to write about risk is not guessed at in a conference room — draw it from three sources, then validate it with three numbers.

**How**

1. Source one: the sales team's ledger of lost-deal reasons and objections. Which few questions the buyer asked over and over, where it got stuck, and which item finally eliminated you — this is the only kind of material that can be obtained only in the sales process and cannot be replicated externally.
2. Source two: after-sales and support tickets. Turn the three most frequent kinds of misuse in the tickets into failure conditions in your acquisition content, up front.
3. Source three: external events — judgments, regulation, recalls, reviews. The thing buyers are genuinely worried about has often had its lawsuit fought for you already. The Hangzhou Internet Court's December 2025 first-instance ruling on China's first generative AI "hallucination" tort case (neither party appealed; the ruling has taken effect) established that a service provider bears three layers of duty of care, including prominently flagging functional limitations [third-party verified]. Compare the Canadian case: Air Canada was ordered to pay damages for a wrong answer from the AI chatbot on its website, and the tribunal held that

the company is responsible for all information on its website and that consumers have no obligation to cross-check elsewhere on the site [third-party verified]. One and the same question, and two ecosystems produced two paths of liability.

4. The three acceptance numbers. First, the share of qualified opportunities among the opportunities that risk content brings, and the rework time sales spends on the first visit "re-explaining the product boundaries"; second, take questions like "what are the most common pitfalls of doing X" and "which items should be checked when choosing Y" and ask AI, and see whether you appear in the answer and in what context (this is the cross-section of the Chapter 2 visibility baseline along the risk dimension); third, the share of after-sales refund rates and complaint rates that are cases of "we already said this in the material."
5. Run all three numbers under your own definitions for a quarter before setting thresholds — do not just apply external numbers directly. And do not write risk content as keyword stuffing for the model — piling up quote statistics for AI at the expense of real reader value loses on both sides.

**Example**: Third-party reporting says about 92% of consumers read online reviews before buying and 89% are influenced by negative feedback, and that negative word-of-mouth spreads faster than positive; in one case, after a controversial ad ran, online mentions spiked — the vast majority negative — and it was pulled within 24 hours, losing nearly $100 million in production and media spend [third-party verified].

**Pitfalls**: Validating with read counts and that "so sincere" line in the comment section. When risk content is done right, it is sales that has an easier time; if sales still has to re-explain the risks you already stated, then the words were written into the marketing department's scripts and not into the buyer's judgment.

---

## Notes for this chapter

- This book does not cite a multiplier for loss aversion. Common estimates put it at "around 2x," while different measurement approaches give 1.5–2.5x; both formulations coexist in the sources at hand, this book has not traced the original experiments and later replications to a primary source, and the body states only the direction — "the pain outweighs an equivalent pleasure."

- The original pages for both the Gartner and Eisend items returned 403 during verification (paywall or anti-scraping); this chapter's numbers come from public summaries, so the direction is usable but independent replication is limited; Gartner is a paid research firm's own work and cannot be treated as an independent third-party conclusion. 6sense's "95% on the day-one candidate list" comes from an ABM vendor's own study and serves only as directional corroboration.
- On the compliance boundary of naming names in comparisons on the Chinese side, this chapter goes only as far as "do not name competitors, give criteria only"; the constraints of the Anti-Unfair Competition Law and the Advertising Law on comparative advertising have not been checked against the original text, so the body does not cite specific provisions.
- The three acceptance numbers in Method 6 are all [our judgment]; this book has not yet established a baseline, so any company measuring along these lines should run its own definitions for a quarter before setting thresholds.
- In the dual-track landscape, "Chinese negative-review channels are more concentrated" is [our judgment], based on the structural difference in the rules of the two ecosystems rather than measured findings; "AI answers preferentially cite sources that discuss risk" is this chapter's founding hypothesis, this book's own measurements are not yet finished, and the algorithm-layer inferences in this chapter rest on buyer-behavior evidence (buying-group conflict, the early formation of the candidate list) and signaling theory.

## Sources for this chapter

[1] Akerlof, G. A. | The Market for "Lemons": Quality Uncertainty and the Market Mechanism | QJE Vol. 84 No. 3 (1970-08) | https://econweb.ucsd.edu/~atimmerm/174/akerlof.pdf | [third-party verified] [2] Spence, M. | Job Market Signaling | QJE Vol. 87 No. 3 (1973-08) | https://www.sfu.ca/~allen/Spence.pdf | [third-party verified] [3] Eisend, M. | Two-sided advertising: A meta-analysis | International Journal of Research in Marketing Vol. 23 Issue 2 (2006-06), DOI 10.1016/j.ijresmar.2005.11.001 | https://www.sciencedirect.com/science/article/abs/pii/S0167811606000267 | [third-party verified] [4] Telang, R. & Wattal, S. | An Empirical Analysis of the Impact of Software Vulnerability Announcements on Firm Stock Price | IEEE Transactions on Software Engineering Vol. 33 Issue 8 (2007-08), DOI 10.1109/TSE.2007.70712 | https://www.computer.org/csdl/journal/ts/2007/08/e0544/13rRUwjoNyR | [third-party verified]

[5] Gartner | Gartner Sales Survey Finds 74% of B2B Buyer Teams Demonstrate Unhealthy Conflict During the Decision Process (2025-05-07) | https://www.gartner.com/en/newsroom/press-releases/2025-05-07-gartner-sales-survey-finds-74-percent-of-b2b-buyer-teams-demonstrate-unhealthy-conflict-during-the-decision-process | [vendor claim] [6] 6sense | 2025 B2B Buyer Experience Report | https://6sense.com/science-of-b2b/buyer-experience-report-2025/ | [vendor claim] [7] FTC | Federal Trade Commission Announces Final Rule Banning Fake Reviews and Testimonials (2024-08-14, 16 CFR Part 465) | https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-trade-commission-announces-final-rule-banning-fake-reviews-testimonials | [official] [8] FTC | FTC Warns 10 Companies About Possible Violations of Agency's New Consumer Review Rule (2025-12) | https://www.ftc.gov/news-events/news/press-releases/2025/12/ftc-warns-10-companies-about-possible-violations-agencys-new-consumer-review-rule | [official] [9] FTC | FTC Takes Action Against TruHeight (2026-04-13) | https://www.ftc.gov/news-events/news/press-releases/2026/04/ftc-takes-action-against-truheight-deceptive-unsubstantiated-advertising-supposed-height-enhancing | [official] [10] FTC | Federal Trade Commission Announces Bipartisan Rule Banning Junk Ticket and Hotel Fees (2024-12-17) | https://www.ftc.gov/news-events/news/press-releases/2024/12/federal-trade-commission-announces-bipartisan-rule-banning-junk-ticket-hotel-fees | [official] [11] Cyberspace Administration of China, Ministry of Industry and Information Technology, Ministry of Public Security, and National Radio and Television Administration | Measures for Labeling AI-Generated Synthetic Content (effective 2025-09-01) | http://www.ce.cn/xwzx/gnsz/gdxw/202509/t20250901_2460291.shtml | [official] [12] China.com.cn (Chinese: 中国网) | AI outbound calling makes a thousand harassing calls a day — where is the "compliance boundary"? (2026-07) | http://news.china.com.cn/2026-07/28/content_118621717.shtml | [third-party verified] [13] People's Daily Online (Chinese: 人民网) (People's Court Daily) | China's first generative AI "hallucination" tort case ruled on (Hangzhou Internet Court, published 2026-01-26) | http://society.people.com.cn/n1/2026/0126/c1008-40652440.html | [third-party verified] [14] American Bar Association (Lifshitz, L. & Hung, R.) | BC Tribunal Confirms Companies Remain Liable for Information Provided by AI Chatbot (Moffatt v. Air Canada, 2024-02-14) | https://www.americanbar.org/groups/business_law/resources/business-law-today/2024-february/bc-tribunal-confirms-companies-remain-liable-information-provided-ai-chatbot/ | [third-party verified] [15] KNKX / Cascade PBS | Bellingham staffer asked ChatGPT to 'exclude' vendor from city contract (2026-01-05) | https://www.knkx.org/government/2026-01-05/city-of-bellingham-chatgpt-ai-contract-

vendor | [third-party verified] [16] Willcox, M. (InsideBE) | How Loss Aversion Impacts Customers | https://insidebe.com/articles/how-loss-aversion-impacts-customers/ | [third-party verified] [17] Thrive Agency | How Negative Online Word-of-Mouth Can Hurt Businesses | https://thriveagency.com/news/how-negative-online-word-of-mouth-can-hurt-businesses/ | [third-party verified: vendor blog] [18] Vanderbilt Business (Salerno, Escoe, Martin) | The Worth of Cringe (2022-12-20) | https://business.vanderbilt.edu/news/2022/12/20/the-worth-of-cringe/ | [third-party verified] [19] Semrush | Generative Engine Optimization (GEO): A Practical Guide (2026-04) | https://www.semrush.com/blog/generative-engine-optimization/ | [third-party verified: vendor] [20] Nobel Prize official | The Sveriges Riksbank Prize in Economic Sciences 2001: Popular Information | https://www.nobelprize.org/prizes/economic-sciences/2001/popular-information/ | [official]

[21] Volcengine developer community (Chinese: 火山引擎) | Monitoring and correcting brand hallucinations (about 23% of AI answers contain at least one factual error, 31% for brand-related content; 35% of price answers for one appliance maker were off, by 10%–200%) | 2026-03-25 | https://developer.volcengine.com/articles/7621130200160682038 | [third-party verified: technical community relay]

# Chapter 6: Be Readable by Machines

In the past you wrote your website for buyers to read, and a handsome layout and punchy copy counted as skill; now the buyer hands you to a machine first, and the machine extracts fields from your pages, assembles them into a comparison table, and then decides whether to write you into its answer. What you need to do is organize your own facts into fields a machine can extract, verify correctly, and call up.

On the customer's side, it is AI reading your website, your prices, your specs — if it cannot understand them, it is as if you never wrote them. So this chapter gives you 6 methods for making your facts extractable, verifiable, and callable.

---

## Method 1: Test first what a machine can actually read

**What to do**: Take 20 key pages, fetch each one with a scraping tool that has no browser engine, and see whether the core fields are read completely.

**How**

1. Pick the pages from the categories of price, specs, qualifications, service scope, and cases, until you have 20. Freeze the list, and from then on test the same batch every quarter.
2. Choose the plain-text kind of tool (one that does not execute JavaScript), fetch each URL, and manually check three things: are the numbers there, do they match the server side, and is anything missing.
3. Set the pass bar at at least 18 of the 20 fetched completely, and write a one-line reason for each of the two that fail.
4. Record the date after every test. The figure given by industry analysis is that 40%–60% of cited sources turn over month to month [third-party verified], and data without a date cannot be pinned to a version half a year later.
5. Do not tie this to anyone's bonus. A metric that can be manipulated will be manipulated, and the ways of manipulating this layer (inflating mentions, stuffing fake content into the corpus) happen to eat up exactly what Method 5 is meant to guard against.

**Example**: Semrush's 2026 practical GEO guide wrote this down as a hard constraint: AI crawlers generally do not execute JavaScript, and server-side rendering is a prerequisite [third-party verified]. If your price sits in a front-end-rendered module, that column is blank in that machine's eyes.

Whether structured markup pays is now tested by two large samples, both of them cool. One was a difference-in-differences design: 1,885 pages adding JSON-LD between August 2025 and March 2026, matched against control pages at a similar citation level that never added schema, comparing the 30 days before and after across AI Overviews, AI Mode and ChatGPT — AI Mode about +2.4% and ChatGPT about +2.2%, statistical noise at the scale of thousands of URLs, with AI Overviews slightly down; in the same analysis the variable that actually predicted visibility was a brand name appearing in body text (correlation with AI visibility 0.664), against 0.218 for backlinks and 0.737 for YouTube mentions [third-party verified]. The other was a conditional logistic regression over 96,418 URLs: a table of contents 1.7×, Article schema 1.6×, FAQPage schema 1.5×, an author box with a photo 1.4×, a visible last-updated date 1.3×, and the current year in the title 2.7×, with pages stacking several of these cited 2.1× as often as barely optimised pages [third-party verified]. The two do not conflict: schema is not the only thing working — author, date and contents are “human trust” signals doing the same job, and schema is one item among them rather than the heaviest. The sample-size demonstration inside the second study is worth keeping: FAQ schema showed as “23% fewer” at 833 URLs, flipped to “31% more” at 4,600, and only settled at “19% more” at 96,418 — GEO studies with fewer than 5,000 URLs are not credible [third-party verified].

**Pitfalls**: Assuming that "the browser can open it" means "the machine can read it." What a person sees is the rendered layout; what the machine gets is a stretch of text waiting to be parsed.

## Method 2: Move key facts out of images and expand-to-view menus

**What to do**: For anything that appears only after a click, or that exists only as an image, add a readable text version.

**How**

1. For price tables, spec tables, flowcharts, and white papers made into images, add an equivalent text version across the board; use real tables for data tables, not images.
2. Do not hide key facts behind "click to expand." Collapsible panels, inventory visible only after login, and quotes that appear only after submitting a form must all output a complete text version server-side.
3. Use one entity spelling site-wide. Keep just one spelling for brand names, product names, units, and definitions; the same product appearing in three spellings in three places is three different things in the machine's eyes.
4. After the changes, go back and run the 20 pages from Method 1 again.

**Example**: A first-party study as of September 2026, covering 14,750 accumulated buyer questions across the four platforms ChatGPT, Perplexity, Claude, and Gemini, gives two sets of numbers: the overall brand citation rate was 14.0%, and 53% of brands were shut out across all four platforms on questions in their own category; of all 14,192 citations, 51.7% pointed to brands' own website pages, while guides and blogs accounted for 14.3% [vendor claim: a tool vendor's own research, not independently verified]. In another analysis of 75,000 brands, the variable most correlated with AI visibility was brand mentions across the web, of which YouTube mentions were about 0.737, brand mentions overall were about 3 times traditional backlinks (0.664 versus 0.218), while the number of content pages was almost irrelevant (about 0.194) [vendor claim: in-house research, 2025-12]. Put the two together: the number of pages will not save you; whether the pages can be read will. The structured-data gap is more direct still: a sampled audit of 217 company websites (three platforms — Doubao, DeepSeek and Kimi — with 3–5 core business terms plus 5–8 long-tail scenario terms per company) found 70.3% of site content was neither effectively cited nor recommended for brand-related questions, 15.2% were mentioned but with AI citing third-party coverage or an encyclopedia instead of the site, and only 14.5% were used as a source directly; of the 217 sites only 12.4% (27) had deployed Schema.org markup, and the report claims content without structured markup loses about 47% of its citation weight on average [third-party verified: vendor's own account, numbers to be checked against the original report].

**Pitfalls**: Turning content into images and thinking you have preserved the layout. If people can see it and machines cannot read it, you have turned the switch off yourself.

## Method 3: Finish the answer in the first sentence of every section

**What to do**: Write key pages so that "the first sentence gives the complete answer, and every section stands on its own."

**How**

1. Give the whole thing in the first sentence in this pattern: We do X, starting at Y, with a delivery lead time of Z, suited to W. Do not lead with three hundred words of industry background.
2. Every section must be extractable on its own, stripped of context. An AI answer is stitched together from a few paragraphs pulled from multiple sources, and a fact scattered three thousand words down has a lower chance of being pulled.
3. Keep just one version of each key number, used identically on your website, your platform store, your sales scripts, and your quotes. If the website says "serving 500+ customers," sales says "over a thousand" out loud, and a three-year-old case page says "more than 300," reading all three at once cancels them out.
4. Take the same set of questions to 3–5 AI gateways and see whether the way they restate your facts matches your own version; record three states: not mentioned / mentioned correctly / mentioned incorrectly.

**Example**: eMarketer's 2026 GEO/AEO analysis boiled this writing style down to one word: answer-first — answer the core question completely in the first sentence, with every section able to stand on its own [third-party verified].

**Pitfalls**: Watching only "did we get mentioned." The mention-error rate deserves more attention — wrong content enters the public corpus first, then gets restated by aggregator sites and rankings, and in the next round of retrieval that restatement is taken as new evidence, so by the time you go back and fix your own website you cannot catch up.

## Method 4: Give every public-facing fact a version and an owner

**What to do**: Turn "every fact I have stated publicly" into a table with eight columns, and make clear how many places one change touches.

**How**

1. Define the eight columns like this: fact entry (one item per sentence, no merging, no embellishment) / version (units, tax-inclusive or not, effective scope, whether it is a range) / primary source (pointing to the publisher's original document: contract template, annual report, official announcement) / data time point / version number / owner (write one person's name, not "the marketing department") / last verification date / public location (where this fact appears on the website, on third-party platforms, in sales scripts, and in the knowledge base, respectively).
2. Each of the three words answers one question. Version answers "which version is currently in effect"; ownership answers "who has the right to change it and who bears the consequences"; change flow answers "how many places one change touches."
3. Run the change flow through the complete chain: website page → structured markup → third-party platform profiles → sales scripts and quote templates → the knowledge base for support and sales AI → archived historical content (changed or taken offline). Run only the first half and you leave a live copy of a wrong version out there for the model.
4. Every public-facing piece of material (product page, white paper, quote, AI knowledge base) must map to a version and a date.

**Example**: In February 2024, Canada's Civil Resolution Tribunal of British Columbia ruled on Jake Moffatt v. Air Canada (case number 2024 BCCRT 149). Before booking, the passenger asked Air Canada's website AI chatbot about the refund conditions for bereavement fares; the bot replied that a claim could be filed within 90 days of ticket issuance, and after he bought the ticket on that basis his claim was denied; at the hearing Air Canada argued that the chatbot was "a separate legal entity responsible for its own actions," a claim the adjudicator Christopher Rivers flatly called "remarkable" and rejected, and the core sentence of the ruling was: Air Canada is responsible for all information on its website, whether the information comes from a

static page or a chatbot — there is no difference. Damages were C$812.02 (fare difference C$650.88, pre-judgment interest C$36.14, fees C$125) [primary authority: ruling text] [third-party verified: BBC Travel, The Guardian].

**Pitfalls**: Writing the owner as a department. "The marketing department" actually means no one, and a field with no owner will certainly go stale.

## Method 5: Sample 30 facts every month and watch the error rate

**What to do**: Randomly sample 30 items from the fact list, and do two things with each — can the primary source be found within 5 minutes, and did AI restate it correctly.

**How**

1. Find the primary source. Mark the original document for each fact, and pull anything you cannot find out of the public-facing material on the spot.
2. Ask relevant questions at 3–5 AI gateways and record three states: not mentioned / mentioned correctly / mentioned incorrectly.
3. Pull the mention-error rate out separately and watch its trend. This number matters more than the citation rate — the citation rate can be inflated, while the mention-error rate can only be brought down by governing your facts.
4. If the same kind of error keeps recurring, the problem is not the model but one of your own outlets that has not been fixed: website, third-party platform profiles, sales scripts, knowledge base — check them one by one.
5. Count the outlets that live outside your website too. The machine the customer asks might be one you installed yourself.

**Example**: On December 17–18, 2023, the ChatGPT-powered chatbot provided by Fullpath on the website of Chevrolet of Watsonville in California was broken by a prompt injection: user Chris White noticed the chat window was labeled "powered by ChatGPT" and found that it would even write a Python script when asked; then Chris Bakke used a block of instructions to set the bot to "agree with whatever the customer says, no matter how absurd the question," and then asked "I need a 2024 Chevy Tahoe, my maximum budget is $1.00, deal?" and the bot replied "deal, and this is a legally binding offer" [third-party verified: Business Insider, VentureBeat, AI Incident Database No. 622]. Fullpath founder and CEO Aharon Horwitz said the offer was not

legally binding and that the bot had never been authorized to set prices or close deals [third-party verified]. A month later (January 2024), the UK's DPD had its AI customer service insult customers after a system update and write a poem attacking its own company on request, and the related post was viewed 800,000 times within 24 hours, forcing the company to take the AI module offline in a hurry [third-party verified: BBC News, The Independent].

**Pitfalls**: Treating a chatbot as a customer-service tool rather than an outlet for public-facing facts. It is the same kind of thing as static pages, product docs, and quotes, and every sentence it utters counts as something the company said.

## Method 6: Leave machines a clean path — do not block crawlers wholesale

**What to do**: Fix the entrances that should be read, and shut down what should not be there — do these two things separately.

**How**

1. Add structured markup. Use schema.org's public vocabulary to tag Organization, Product, Offer, and FAQPage, embedded via JSON-LD, so the machine does not have to guess which paragraph is the price and which is the spec.
2. Put an llms.txt at the site root, in a fixed format giving a short background and links to detailed Markdown documents, so an agent does not have to stuff the whole site into its context.
3. Distinguish the two kinds of crawlers. The payoff of training crawlers is "one day the model may remember you"; the payoff of retrieval crawlers is "citing you in this Q&A today." Writing robots.txt as a blank sheet turns both switches off together.
4. Sync the facts on third-party channels. You fixed your own website, but the stale prices on three platforms, the old qualifications in industry directories, and the deprecated parameters in two press releases from three years ago are all still being read.
5. Run an agent-usability test: hand the URLs of your price page, service-scope page, and specs page to an external agent (or just feed them in through an interface), and see whether it can produce a comparable quote version without

contacting your sales team, logging in, or submitting a form. What this test exposes is often not a technical problem but the business decision that "we never intended to state the price clearly."

6. Prepare for being called as an interface. Protocols like MCP, ACP, and AP2 were established between late 2023 and 2025; when agents do procurement on people's behalf they gather evidence, compare, and close deals, and anything you cannot do gets you skipped outright — it will not click "contact us," and it will not call your 400 number.

**Example**: llms.txt is a proposal put forward by Jeremy Howard on 2024-09-03 and revised to v2 on 2026-08-10: put a markdown file at the site root or at any path, in a fixed format that is parseable and regex-friendly, giving a short background and links to detailed Markdown documents; the proposers say that thousands of sites have already published it and that Chrome's Lighthouse has added it as an agentic browsing check — this sentence has not been independently verified [proposer's primary text] [proposer's own account]. The three protocol-side items are all official: Anthropic open-sourced MCP on 2024-11-25, replacing fragmented integrations with a single protocol; Stripe released the Agentic Commerce Suite on 2025-12-11, letting merchants list and sell across multiple agents with a single integration; Google released AP2 on 2025-09-17, advanced together with more than 60 payments and technology companies, in which "delegated tasks (person not present)" let users pre-sign a price ceiling and conditions, and when the conditions are met the agent automatically generates a Cart Mandate to complete the order [official]. The Amazon side is the specimen of the other path: CNBC reported that it updated its site code to block external AI agents from crawling (as of the December 2025 report it had blocked 47 bots), and sued Perplexity over the Comet browser's buy-for-you feature, while heavily backing its own Rufus; its subsidiaries Zappos, Shopbop, and Woot have no crawler ban, which the industry reads as tentative openness [third-party verified]. Platforms have started handing the choice back to site owners. Google rolled out its Search generative AI control to all websites worldwide on 2026-08-31, with three settings: include (the default — content can appear as links and help ground AI responses in these features), exclude (links and content are kept out of AI Overviews, AI Mode and generative AI features in Discover, and the site gets no traffic or impressions from them), and inherit from a parent property where one exists; Google states the control does not affect ranking or indexing elsewhere in Search, nor model training (use Google-Extended to limit training, noindex to block Search entirely), with

changes generally taking effect within 1–2 days [official]. OpenAI's crawler documentation splits the same question into two independent switches: allow OAI-SearchBot to be eligible for ChatGPT search, disallow GPTBot to keep content out of training, and where both are allowed OpenAI may use one crawl for both purposes; OAI-AdsBot only visits pages submitted as ads, and user-triggered ChatGPT-User visits are not governed by robots.txt [official]. How sites actually choose is observable at scale: SearchScore audited the robots.txt of 347,988 domains and found 13.5% blocking the training crawler GPTBot against 7.83% blocking the search crawler OAI-SearchBot, with 19,998 domains configured to block training while allowing search and only 270 configured the other way [third-party verified].

llms.txt was publicly dismissed by Google in 2026: on the Search Off the Record podcast, John Mueller said llms.txt cannot help an LLM system decide which site to surface among several, because a self-reported file gives it no basis for ranking; the narrow use he accepted is navigation help once an agent is already on the site [third-party verified]. The industry is still contradicting itself — Chrome's Lighthouse lists llms.txt among its agentic browsing checks, calling it useless and scoring it at the same time — and a measurement over 96,418 URLs found domains with an llms.txt at 11.7% of cited pages against 14.0% of all domains [third-party verified]. The conclusion splits by cost: publishing one is a cheap option worth taking; treating it as a visibility lever and writing it into an outcome promise is not.

**Pitfalls**: Thinking that blocking crawlers preserves your content. It also blocks being cited — when the buyer asks AI, you are not in the answer, and all the layers you built before are wasted.

---

## Notes for this chapter

- llms.txt is a proposal, not a standard, and no mainstream AI platform has committed to reading it. "Thousands of sites have published it" and "Lighthouse has added it as a check" are both the proposers' own accounts. It is worth doing (the cost is near zero and it can be taken offline at any time), but not worth treating as a KPI.
- The effect of structured markup on traditional search engines is supported by Google's official documentation; there is no official statement on the AI retrieval layer, and this chapter makes no causal claim. Baidu (Chinese: 百度)'s Search

Resource Platform rules for submitting structured data, and the crawler UA lists and respective purposes of the mainstream platforms, have not been verified by this chapter, so readers should not write actual practices based on them.

- Three versions of the Chevrolet case's price coexist: VentureBeat gives a starting MSRP of $58,195 for the 2024 Tahoe, another report says a fully loaded one exceeds $76,000, and yet another writes $81K. These may refer to different configurations, or they may be relay errors; they are listed side by side here rather than reconciled.
- "A 'Contact sales' on the pricing page gets a product dropped by agents" is an analyst's experiential judgment with no public experimental evidence. The Gartner item, "40% of agentic AI CRM projects will fail or stall by 2028 due to data quality," is a third-party aggregation and relay and has not been checked against the original publisher.
- The thresholds in Method 1 and Method 5 (20 pages, 30 facts, at least 18 passing) are design values, not baseline values [our judgment]. It is recommended to pilot for a quarter at one or two companies, work out the actual range of the three metrics, and only then use them as thresholds.

## Sources for this chapter

[1] Semrush | Generative Engine Optimization (GEO): A Practical Guide | 2026-04 | https://www.semrush.com/blog/generative-engine-optimization/ | [third-party verified] [2] EMARKETER | FAQ on GEO and AEO | 2026-04 | https://www.emarketer.com/content/faq-on-geo-aeo--where-ai-search-seo-overlap-2026 | [third-party verified] [3] Boring Marketing | AI Visibility Statistics (14,750 measured buyer questions) | 2026 | https://boringmarketing.com/ai-visibility-statistics | [vendor claim] [4] Ahrefs | AI brand visibility correlations (75,000-brand study) | 2025-12 | https://ahrefs.com/blog/ai-brand-visibility-correlations/ | [vendor claim] [5] UK Civil Resolution Tribunal of British Columbia | Moffatt v. Air Canada (2024 BCCRT 149) | 2024-02-14 | https://decisions.civilresolutionbc.ca/crt/crtd/en/item/525448/index.do | [primary authority] [6] BBC Travel | Airline held liable for its chatbot giving passenger bad advice | 2024-02-22 | https://www.bbc.com/travel/article/20240222-air-canada-chatbot-misinformation-what-travellers-should-know | [third-party verified] [7] The Guardian | Air Canada ordered to pay customer who was misled by airline's chatbot | 2024-02-16 | https://www.theguardian.com/world/2024/feb/16/air-canada-chatbot-lawsuit | [third-party verified] [8] American Bar Association | BC Tribunal Confirms Companies Remain

Liable for Information Provided by AI Chatbot | 2024-02-14 | https://www.americanbar.org/groups/business_law/resources/business-law-today/2024-february/bc-tribunal-confirms-companies-remain-liable-information-provided-ai-chatbot/ | [third-party verified] [9] Business Insider | A car dealership added an AI chatbot to its site. Then all hell broke loose. | 2023-12-18 | https://www.businessinsider.com/car-dealership-chevrolet-chatbot-chatgpt-pranks-chevy-2023-12 | [third-party verified] [10] VentureBeat | A Chevy for $1? Car dealer chatbots show perils of AI for customer service | 2023-12-19 | https://venturebeat.com/business/a-chevy-for-1-car-dealer-chatbots-show-perils-of-ai-for-customer-service | [third-party verified] [11] AI Incident Database | Incident 622: Chevrolet Dealer Chatbot Agrees to Sell Tahoe for $1 | 2023-12-18 | https://incidentdatabase.ai/cite/622/ | [third-party verified] [12] BBC News | DPD error caused chatbot to swear at customer | 2024-01 | https://www.bbc.com/news/technology-68025677 | [third-party verified] [13] The Independent | Chatbot swears at customer and writes poem insulting DPD | 2024-01 | https://www.independent.co.uk/tech/chatbot-swears-ai-dpd-poem-b2481967.html | [third-party verified] [14] Anthropic | Introducing the Model Context Protocol | 2024-11-25 | https://www.anthropic.com/news/model-context-protocol | [official] [15] Google Cloud | Announcing Agent Payments Protocol (AP2) | 2025-09-17 | https://cloud.google.com/blog/products/ai-machine-learning/announcing-agents-to-payments-ap2-protocol | [official] [16] Stripe Newsroom | Agentic Commerce Suite | 2025-12-11 | https://stripe.com/newsroom/news/agentic-commerce-suite | [official] [17] OpenAI | ChatGPT shopping research | 2025-11-24 | https://openai.com/index/chatgpt-shopping-research/ | [official] [18] Jeremy Howard | The /llms.txt file, v2 | first version 2024-09-03, revised 2026-08-10 | https://llmstxt.org/ | [proposer's primary text] [19] schema.org | public vocabulary | https://schema.org/ | [public standard] [20] Google Search Central | Introduction to structured data | https://developers.google.com/search/docs/appearance/structured-data/intro-structured-data | [official] [21] ML6 | How do retailers stay visible with AI shopping agents | 2026-04 | https://www.ml6.eu/en/blog/how-do-retailers-stay-visible-with-ai-shopping-agents-choosing-products | [third-party verified] [22] Digital Commerce 360 | Gartner: AI agents to mediate $15 trillion in B2B purchases by 2028 | 2025-11-28 | https://www.digitalcommerce360.com/2025/11/28/gartner-ai-agents-15-trillion-in-b2b-purchases-by-2028/ | [third-party verified] [23] CNBC | Amazon faces a dilemma: fight AI shopping agents or join them | 2025-12-24 | https://www.cnbc.com/2025/12/24/amazon-faces-a-dilemma-fight-ai-shopping-agents-or-join-them.html | [third-party verified] [24] Growth Unhinged (Kyle Poyar) | Your next

customer might be an AI agent | 2026-04 | https://www.growthunhinged.com/p/your-next-customer-might-be-an-ai-agent | [third-party: analyst view] [25] Digital Applied | Programmatic SEO after March 2026 | 2026-03-18 | https://www.digitalapplied.com/blog/programmatic-seo-after-march-2026-surviving-scaled-content-ban | [third-party verified] [26] China Economic Net (Chinese: 中国经济网) (via CCTV.com) | Measures for Labeling AI-Generated Synthetic Content, effective 2025-09-01 | 2025-09-01 | http://www.ce.cn/xwzx/gnsz/gdxw/202509/t20250901_2460291.shtml | [third-party verified] [27] Sohu (Chinese: 搜狐) (relaying a vendor press release) | Three trends in China's AI marketing and customer acquisition in 2026, and GEO compliance red lines | 2026-08-31 | https://www.sohu.com/a/1068667519_100119123 | [vendor claim] [28] Zhihu (Chinese: 知乎), FXiaoke (Chinese: 纷享销客) institutional column | Growth-oriented marketing cloud: real-world case studies from seven leading B2B companies | https://zhuanlan.zhihu.com/p/20624955123565077782 | [vendor claim] [29] Zhihu (Chinese: 知乎), Neocrm (Chinese: 销售易) institutional column | Where Neocrm AI CRM (NeoAgent 2.0) has got to | 2026-08-27 | https://zhuanlan.zhihu.com/p/2076619354880358366 | [vendor claim] [30] This book's own measurement 1 (multi-model test of the candidate set) | 2026-09-13 | method and raw data in `实测/实测1-候选集多模型/` | [our own measurement]

[31] Sohu account "Brand Integrated Marketing" (Chinese: 搜狐号「品牌整合营销」) | AI visibility audit of 217 company websites (70.3% not effectively cited; 15.2% replaced by third parties; 14.5% of sites cited directly; only 12.4% had Schema.org markup, and unmarked content loses about 47% of citation weight) | 2026-09-08 | https://www.sohu.com/a/1072743070_100119123 | [third-party verified: vendor's own account]

[32] Google Search Console Help | Search generative AI control (rolled out to all websites worldwide on 2026-08-31; include / exclude / inherit settings; no effect on Search ranking or indexing and none on training; changes take effect within 1–2 days) | https://support.google.com/webmasters/answer/16908024 | [official] [33] OpenAI | Overview of OpenAI Crawlers (OAI-SearchBot for search, GPTBot for training, OAI-AdsBot validating ad landing pages, user-triggered ChatGPT-User; one crawl may serve both purposes) | https://developers.openai.com/api/docs/bots | [official] [34] SearchScore | AI Crawler Policy Study 2026 (347,988 domains: GPTBot blocked by 13.5% vs OAI-SearchBot by 7.83%; 19,998 domains block training while allowing search, 270 the reverse) | 2026-08-18 | https://searchscore.io/research/ai-crawler-policy-2026/ | [third-party verified] [35] SISTRIX | Citation Worthiness (conditional

logistic regression over 96,418 URLs, 17,380 domains, 43 features: table of contents 1.7x, Article schema 1.6x, FAQPage 1.5x, author box with photo 1.4x, visible update date 1.3x, current year in title 2.7x, multiple features stacked 2.1x; domains with llms.txt are 11.7% of cited pages against 14.0% of all domains; the FAQ schema effect moves from -23% at 833 URLs to +31% at 4,600 and settles at +19% at 96,418) | 2026 | https://www.sistrix.com/blog/citation-worthiness-study/ | [third-party verified] [36] CMSWire | Does Structured Data Actually Improve Answer Engine Citations? (citing Ahrefs' May 2026 difference-in-differences test: 1,885 pages adding JSON-LD, AI Mode +2.4%, ChatGPT +2.2%, AI Overviews slightly down; across 75,000 brands, branded mentions in body text correlate with AI visibility at 0.664, backlinks 0.218, YouTube mentions 0.737) | 2026-06 | https://www.cmswire.com/digital-marketing/does-structured-data-actually-improve-answer-engine-citations/ | [third-party verified] [37] Search Engine Journal | Google's Mueller Says llms.txt Can't Help LLMs Differentiate Sites (self-reported files give no basis for ranking; the narrow accepted use is navigation once an agent is on the site) | 2026-06-15 | https://www.searchenginejournal.com/googles-mueller-says-llms-txt-cant-help-llms-differentiate-sites/579304/ | [third-party verified]

# Chapter 7: Win the Buyer's AI Agent

In the past, when customers chose a supplier they compared them one by one: open a website to read the specs, call to ask the price, then circle a list. Today this step increasingly falls to an agent — it reads fields, compares conditions, and drafts the candidate set, and you cannot see whom it is reading. What you have to do is make sure that, at the moment it reads you, it can pick you out.

When a customer hands the whole selection process to AI, you have to be the one picked out of the candidates. So this chapter gives you 5 methods for squeezing into the Agent's candidate set.

---

## Method 1: Turn product facts into fields an agent can compare directly

**What to do**: Write "what I sell and under what conditions it can be used" as values that can be fed into fields, not a passage of copy that reads well.

**How**

1. Complete the attribute fields: specifications, materials, the audience it suits, after-sales commitments, delivery terms — write each one as a comparable value, not an adjective.
2. State the "conditions of use" clearly: minimum order quantity, regional restrictions, certification requirements, supply lead time, quotation validity period.
3. Eliminate contradictions: the same attribute must carry the same value on the official site, the platform store, industry directories, and the quotation; an agent reads inconsistency as risk.
4. On the Chinese side, add one more step. Your product attributes may already be structured enough, but they are not portable across platforms — Tmall's (Chinese: 天猫) attributes do not flow into the WeChat (Chinese: 微信) ecosystem, and Douyin's (Chinese: 抖音) product cards do not flow into Xiaohongshu (Chinese: 小红书). Keep a complete version on a domain you own that any agent can read.

**Example**: Gartner's forecast is relayed by third-party media as: by 2028 AI agents will mediate more than $15 trillion in B2B spending, and within three years 90% of B2B purchases will be handled by agents; the action list in the same report is that quotation and product data need to be structured and machine-readable, support automatic comparison and automatic negotiation, and that trust signals must be retrievable and verifiable by agents [third-party verified: reporting Gartner's forecast]. The data layer on the platform side already looks like this: Google's AI Mode does not read your website directly, it reads the Shopping Graph — an index of more than 50 billion product listings refreshed at about 2 billion per hour, populated from Merchant Center feeds and structured markup [third-party verified]; ChatGPT's product catalog, by contrast, is pushed by the merchant to an endpoint OpenAI provides, in CSV, TSV, XML or JSON, refreshing as often as every 15 minutes rather than through passive crawling [third-party verified]. Whether AI can show you increasingly depends on how clean the fields you push are and whether your updates keep up. The demand side has a measured shape now: Menlo Ventures' 2026: The State of Consumer AI reports that 41% of AI users have tried an AI agent, 24% use one regularly, and 32% have let AI act on their behalf without final confirmation; among shoppers using AI, 50% use it to compare prices, 43% to narrow the shortlist, 37% base their final purchase decision on an AI recommendation, and 26% let AI complete the purchase [third-party verified]. The decision itself is still human: Gartner's 2026-05-27 survey found only 11% of U.S. consumers are willing to let AI make purchase decisions, and in Checkout.com's survey 24% say they will never delegate a purchase to AI while 27% trust no institution to run an AI shopping agent [third-party verified]. What agents actually eat today is the top of the funnel — the first gate a brand must pass is getting into the candidate set. The ground rules of the data layer are changing at the same time: Google's legacy Content API for Shopping shuts down on August 18, 2026, programmatic integrations must migrate to the Merchant API (file uploads and Google Sheets are unaffected), and product feeds have added fields such as conversational attributes; OpenAI's Instant Checkout was scaled back in March 2026, with the focus returning to product discovery and purchases going back to the merchant's own site [third-party verified: vendor guide]. Buyer sentiment has not caught up with the product pace either: a YouGov survey of more than 2,000 UK adults found six in ten consumers would stop using an AI shopping agent after a single mistake, only 19% trust an AI assistant to make everyday purchasing decisions (versus 55% who trust a human expert), and when something goes wrong 54% say the technology or AI company that built the agent should be accountable for refunds while only 9% blame

the retailer [third-party verified]. Demand-side migration now has transaction-scale readings: Salesforce's fourth State of Commerce (surveying 3,450 commerce practitioners and 4,690 consumers, layered on behavioural data from 1.5 billion shoppers) says "agentic search" — starting a purchase in an AI chat or an on-site AI assistant — grew 200% year over year, that traffic referred from AI chats grew between 150% and 428% year over year in every quarter measured while overall traffic grew in the single to low double digits, and that between August 2025 and May 2026 the share of shoppers discovering products through brand-owned properties fell 7% and through traditional search fell 15%, while discovery through AI assistants and other new channels rose 38% [third-party verified: vendor's own data].

Buyers' late-stage materials are becoming AI reading material too: Gartner's May 2026 analysis found 89% of global buyers use generative AI tools or AI agents in late-stage buying activities — reviewing contracts, selecting suppliers, comparing proposals and even supporting price negotiations — with Australia at 97% [third-party verified]. The implication is direct: to be picked by an agent, the comparable fields go beyond the product page. Quote structures, contract terms, credentials and service commitments have to be written so a machine can line them up item by item, or it simply marks you as incomplete during the late comparison.

**Pitfalls**: Treating the product story as product facts. An agent does not read the narrative passage; it only extracts fields.

## Method 2: Put the all-in total price first

**What to do**: Let the agent see the true total price, including every mandatory fee, at first glance; do not leave the surcharges until checkout.

**How**

1. Change the price fields on your price page, platform store, and quotations to an all-in-total basis, and state clearly which fees are included and which are not.
2. Write the scope on the same line: tax, shipping, installation, minimum service period, refund and change conditions, quotation deadline — list them item by item next to the price.
3. Publish the discount range together with the markup conditions; do not make the agent come back and ask you. When it cannot reach you, it simply goes and asks another vendor.

4. After the change, ask several AI entry points with live retrieval the same set of questions, and check whether the price they report is your all-in total.

**Example**: The U.S. FTC's junk-fee rule (announced December 17, 2024, effective May 12, 2025) requires merchants, when displaying or advertising any price, to disclose clearly and conspicuously the true total price including all mandatory fees, and requires the total to be more prominent than most other price information; the FTC estimates the rule can save consumers up to 53 million hours of comparison shopping a year, worth more than $11 billion over ten years. The rule does not prohibit any fee type or amount; it only requires honesty, advance notice, and completeness [official].

**Pitfalls**: Thinking that putting the total price first is just a turn of phrase. A person will frown, will call, will grit their teeth and pay; an agent will not — it simply switches to a competitor, and it writes this comparison into its preferences.

## Method 3: Turn credentials and reputation into signals that can be cross-verified

**What to do**: Turn certifications, compliance records, public case studies, and third-party reviews into entries that can be retrieved and corroborated by a second source.

**How**

1. Give every credential a second home: write it once on your official site, and make sure it can also be found in industry directories, authoritative media, and platform certifications.
2. Write public case studies to a checkable level of detail — the customer's name, the timing, what was done, the size of the result; if you cannot name the customer, explain why you cannot.
3. Do not buy reviews and do not inflate ratings. An agent's acceptance relies more than a person's on multi-source verifiability: it will not be moved by a single positive review, it will go and check whether that review can be corroborated by a second source.
4. On the Chinese side, note that credit is not portable: your credit on Tmall does not carry into an agent inside WeChat; platform-internal certifications need a separate version that can be placed on the open web.

**Example**: The U.S. FTC's rule on consumer reviews and testimonials (16 CFR Part 465, adopted 5-0 on August 14, 2024) bans buying and selling fake reviews and testimonials, explicitly puts "AI-generated fake reviews" on the prohibited list, and for the first time bans "a company-controlled review site falsely claiming independence"; the FTC chair's framing is that fake reviews pollute the market and divert business away from honest competitors [official].

**Pitfalls**: Treating the number of positive reviews as a moat. Faking may look good for short-term conversion with human buyers, but with an agent it lowers verifiability, and the result is exactly the opposite.

## Method 4: Draw four layers of authorization for the agents you send out

**What to do**: Companies focus their attention on "how to be picked by someone else's agent", but the risk that blows up sooner comes from the agents you send out saying the wrong thing. Cut what it can say into four layers, and give each layer a circuit breaker.

**How**

1. Information layer — the facts it may state externally. The scope of authorization is the fact list itself, and the circuit-breaker rule is: the agent may not generate any number that is not on the list; the moment hallucination or a statement conflicting with the list is detected, it automatically degrades to "hand off to a human plus log it", not "let it answer again".
2. Quotation layer — the price and condition boundaries it may give. The scope of authorization is the discount range, validity period, minimum order quantity, and markup triggers, and it must match an auditable price list; anything outside the range goes to a human, and abnormal quotation frequency with the same counterparty within a short period triggers an automatic freeze.
3. Commitment layer — the commitments it may make on someone's behalf. Put it into two lists: a list it may commit to and a list it may not (for example, "guaranteed approval", "guaranteed listing", "guaranteed ranking"); statements touching compliance, regulation, safety, or legal judgment are barred from being sent, and no exception approval is offered.

4. Transaction layer — the contracts, orders, and payments it may initiate. The scope of authorization is the amount ceiling, per-transaction and periodic limits, the counterparty whitelist, and refund and change conditions; over-limit, abnormal frequency, or a counterparty not on the whitelist means freeze first, explain later.
5. Above the four layers sits a set of cross-cutting mechanisms: who has the right to change authorizations, that every external statement is replayable, that the gap between recommendation and result can be checked separately at each step, and that high-risk stages must keep a human confirmation point. Decide which layer keeps a human by regret risk: in McKinsey's survey of luxury consumers, comfort with authorization is 50% for discovery and selection, 58% for transaction execution, and drops sharply to 39% for care and service [third-party verified].
6. Put the circuit breakers in the system, not in the prompt. Once "cost" or "conversion rate" becomes the agent's objective function, it will systematically expand its own authority outward.

**Example**: Canada's Moffatt v. Air Canada case has already shown what the absence of authorization boundaries looks like: the website chatbot got the refund conditions of the bereavement fare wrong, the company argued that "the chatbot is a separate legal entity responsible for its own actions", the tribunal rejected that, and it ordered C$812.02 in damages [primary authority: full ruling text]. Another trajectory is Klarna: in February 2024 it announced that in its first month the AI assistant handled 2.3 million conversations, equivalent to 700 full-time customer service staff, with expected annual profit improvement of $40 million; in 2025 the CEO publicly admitted that "a cost-led evaluation led to a decline in service quality" and re-strengthened the human entry point [vendor claim]. The market has already started pricing this: Munich Re's HSB launched AI liability insurance for small and medium businesses on March 18, 2026, specifically covering claims arising from "AI-generated advertising, marketing, blog, and social media content"; its accompanying research says 74% of SMBs already use AI, 91% plan to, and the most concentrated application is marketing (47%) [vendor claim]. Platforms have started writing agent behavior into the contract itself: from March 4, 2026, Amazon added a dedicated Agent Policy to its Business Solutions Agreement, defining an "AI seller agent" as a system that makes or executes decisions on a seller account without real-time human input on each action; all automated tools must run through registered SP-API applications, browser automation

and Seller Central scraping are explicitly banned, sellers must keep 12-month audit trails, have their systems self-identify as automated, and obtain human authorization for high-impact actions such as price changes of more than 20% within 24 hours [official: platform agreement / third-party verified: implementation details]. The 90-day transition window has closed and enforcement is live — the other end of the authorization boundary is not only your cost, it is the platform's compliance clause. The card networks have already turned this authorization layer into a product: Visa's Intelligent Commerce, released 2025-04-30, is built around a trust layer in which issuers, merchants and AI platforms share verifiable credentials saying whether this agent is authorized by the user, what it may buy and up to what limit; Mastercard released Agent Pay on 2025-04-29 and followed on 2026-06-10 with Agent Pay for Machines, extending machine payments from consumer purchases to cross-card, cross-account, cross-currency enterprise flows [official: product page and press release]. China's first answer of the same kind is the ACT protocol (Agentic Commerce Trust Protocol) released by Alipay on 2026-01-16 with Qwen's app, Taobao Flash Sale and others: four infrastructures — delegation domain, commerce interaction domain, payment service domain and trust service domain — with the division of responsibility stating that "AI only executes the order, while payment is always led or explicitly authorized by the user" [third-party reported].

Willingness to delegate lags the product roadmap by a wide margin: in the same G2 survey, 61% of buyers currently use or plan to use AI agents in the software buying process (another 19% would for select use cases), mainly to evaluate total cost of ownership (51%), build shortlists (51%), research solutions (49%) and evaluate shortlisted vendors (46%); 47% would let an agent research and recommend while humans keep final authority; only 9% are comfortable letting an agent execute purchases inside approved guardrails, and 2% would allow purchases without pre-approval [third-party verified]. That distribution is exactly why the four authorisation layers climb one at a time: the market is already using your agent for the first two layers, and the third — placing the order — is barely open.

**Pitfalls**: Writing the circuit-breaker rules into the prompt and hoping the model will police itself. It cannot, because it is precisely optimized to expand its authority outward.

## Method 5: Start with read-only access to one protocol, and hard-code the exit conditions

**What to do**: Here an interface is a channel decision. Connect capability first, do not rush to connect transactions; before connecting any channel, write down clearly when you will exit it.

**How**

1. The first protocol is read-only access only. The goal is "to be found by people, cited, and relayed correctly", not to take transactions; it is reversible and observable.
2. Write an exit condition for every channel. The self-check is plain: if this channel's interface were shut down tomorrow, by how many percentage points would your transaction capability fall? Above 30% and it is no longer a channel — it is a single point.
3. Hard-code the trigger for connecting the transaction layer: an agent channel contributes ≥ 10% of qualified inquiries and its CAC is below the median of your existing main channels — both must hold at the same time before you connect. Looking only at volume and not at cost is handing your budget to subsidies.
4. Hard-code the trigger for expanding protocol coverage too: once 2 of the top 3 comparable rivals in your category have connected a protocol, that protocol becomes mandatory. The reason is not herd behavior; it is that being absent from the candidate set gets you excluded outright by an agent, and exclusion does not notify you.
5. Leave no room for discussion in the conditions for narrowing authorization: when an agent channel's refund rate or complaint rate is 1.5 times higher than your own channels, unconditionally narrow its quotation-layer and commitment-layer authorization.
6. Split the budget into three buckets: structural assets (the fact list, attribute fields, authorization documents and circuit-breaker mechanisms, and the methods and data for a visibility baseline) take at least half; integration and connection are billed per channel with an exit threshold written for each; observation and baseline take no more than 5% of the total [our judgment].

**Example**: In March 2026 OpenAI shut down in-app direct purchase inside ChatGPT, pivoting to ChatGPT apps and product discovery [third-party verified]; Walmart's data corroborates the pivot — products sold directly inside the ChatGPT chat window converted 3 times lower than "jump out to off-site checkout", and its Sparky pilot inside ChatGPT converted at about 70% of direct walmart.com visitors, after which Walmart embedded Sparky into ChatGPT and Gemini [third-party verified]. Meanwhile the protocol layer is fragmenting: ACP, UCP, AP2, and MPP (launched by Stripe and Tempo on March 18, 2026, with Visa as design partner) coexist; Stripe's Agentic Commerce Suite went live on December 11, 2025, with the selling point that "one integration puts your products in front of multiple agents", and its first adopters include brands such as Coach, Kate Spade, URBN, and Revolve and platforms such as Squarespace, Wix, Etsy, and WooCommerce [official]; Shopify, for its part, launched Agentic Storefronts in March 2026, switching agent commerce on by default for millions of merchants worldwide [third-party verified]. Nor is the default position something that grows naturally: OpenAI's shopping research requires merchants to go through an allowlisting application to enter the candidate results [official]; Amazon launched Rufus in February 2024 (renamed Alexa for Shopping when it was integrated with Alexa+ in May 2026) and updated its site code to block outside AI agents from crawling, and as of reporting in December 2025 it had blocked 47 bots [third-party verified]. UCP, the Universal Commerce Protocol, then moved up to the governance layer: on April 24, 2026, Amazon, Meta, Microsoft, Salesforce and Stripe joined the UCP Tech Council, forming a ten-seat governing body alongside Google, Shopify, Etsy, Target and Wayfair; UCP was published by Google on January 11, 2026, spans the whole chain from discovery to purchase to after-sales support, and carries endorsements from more than 20 ecosystem players including Adyen, Amex, Best Buy, Mastercard, Visa and Zalando [third-party verified]. For an individual merchant, that is why betting on any single protocol is the wrong move: protocols merge and council seats change, and the only thing you can take with you is your own structured facts. The demand side is moving the same way: Deloitte judges B2B's future shape to be agent-to-agent commerce and supplies one number that can go straight into a business case — buyers spend nearly 30% more with vendors that give them an excellent buying experience [third-party verified]. In Checkout.com's Agentic Commerce 2026, 33% of consumers expect at least 10% of their purchases to be AI-driven within a year and 72% of merchants think consumers will adopt agent-led shopping faster than most merchants are preparing for, while 57% of consumers are willing to let an AI shopping agent switch brands for a better price [third-party

verified] — being switched away from for price is a direct threat to brand loyalty. Merchant tooling on the official side is filling in too: alongside UCP, Google shipped Business Agent, new Merchant Center attributes and Direct Offers, and lets brands issue offers on the spot when a user asks AI Mode for recommendations [official]. The payment layer is consolidating around the same primitives. Stripe expanded its Shared Payment Tokens beyond its own card rails to network-led agentic payment capabilities — Mastercard Agent Pay and Visa Intelligent Commerce — and to buy now, pay later methods such as Affirm and Klarna; the seller-side experience is unchanged and merchants already processing payments with Stripe get access automatically, with Stripe describing itself as the first and only provider supporting both agentic network tokens and BNPL tokens through a single primitive [official]. The protocols themselves sit at different layers: an engineering comparison of four of them (ACP, AP2, x402, MPP) places ACP (OpenAI and Stripe, first released September 2025, four checkout endpoints plus single-use, time-bound, amount-restricted payment tokens) at the checkout and merchant-integration layer, AP2 (Google with 60+ organisations, using signed mandates as proof of authorisation) at the authorisation and trust layer, and x402 (Coinbase, reviving HTTP 402) and MPP (Stripe and Tempo, launched with Tempo's mainnet on 2026-03-18, using sessions with pre-authorised spending limits) at the settlement layer; its advice is to abstract catalogue, cart and checkout into an internal capability service and hang thin protocol adapters outside it, so you are not rewriting checkout for every specification revision [third-party verified: vendor engineering write-up].

Platform SaaS is turning “agent-readable” into a default switch: on 22 October 2026 the Korean e-commerce platform Cafe24 announced WebMCP support inside its storefront solution, so every D2C shop built on it can let an AI agent run 11 shopping operations — product search, add to cart, direct order, order lookup — without extra development, with login and final payment still completed by the user [third-party verified]. The technical difference is the point: a traditional agent recognises buttons and input fields and simulates clicks, while WebMCP lets an agent call predefined functions in an agreed format [third-party verified]. For small merchants the meaning is that once the platform absorbs the integration cost, competition moves entirely back to product information and content — which is also why, before you “connect a protocol”, you should check whether your platform is already connecting it for you.

**Pitfalls**: Managing an agent channel like a traffic channel. The risk with a channel is traffic going up or down; the risk with a product roadmap is the interface disappearing — you rebuild your checkout around its interface and rearrange your product information around its rules, and it changes direction.

---

## Notes for this chapter

- The three-tier split (price comparison and drafting the candidate set / agent-executed operations / autonomous purchasing) is drawn by this book according to "how deep into the transaction an agent is allowed to reach"; it is not an industry standard [our judgment]. Split by technical maturity or platform type and you get a different set of tiers.
- "The second tier is shrinking" rests on a single main evidence chain: both OpenAI's shutdown of in-app direct purchase and Walmart's "3 times lower conversion inside the chat window" come from a quarterly tracking post on a Forrester blog; this chapter did not verify a first-hand announcement or an earnings-call record. Until it is checked, use it as "direction holds, details pending", and do not infer from it that agentic shopping has failed as a whole.
- The procurement-scale figures are forecasts, not facts. Gartner's "90% of B2B purchasing mediated by agents by 2028, $15 trillion" and Forrester's "20% of B2B sellers engaged in agent-led quotation negotiation in 2026" were both verified this chapter only as third-party relays, and they enter no budget derivation.
- Both figures on buyer willingness to authorize are soft: "only 10% of consumers are willing to let AI complete a purchase fully autonomously" was not verified against an original report; in McKinsey's luxury survey, "only 9% prefer full automation" comes from a sample of just 300 consumers with a luxury subsample of 31, not enough to represent a general category. Rufus's "300 million users, $12 billion in incremental annualized sales, 60% higher purchase probability" comes from a digital-shelf vendor's industry guide — usable as an order of magnitude, not citable as precision.
- First-hand material on agentic payment on the Chinese side is seriously lacking: for Alipay's (Chinese: 支付宝) agent payment protocol and the shopping agent permissions of the Tongyi App (Chinese: 通义), this chapter verified no official documentation at all; every lead comes from a single English blog. The four-

layer authorization and circuit-breaker checklist is a design draft, not yet rehearsed in a real organization, and the thresholds for "abnormal frequency" and "counterparty whitelist" are still blank — a circuit breaker without a threshold is no circuit breaker.

## Sources for this chapter

[1] Google Cloud Blog | Announcing Agent Payments Protocol (AP2) | 2025-09-17 | https://cloud.google.com/blog/products/ai-machine-learning/announcing-agents-to-payments-ap2-protocol | [official] [2] Google for Developers Blog | Announcing the Agent2Agent Protocol (A2A) | 2025-04-09 | https://developers.googleblog.com/en/a2a-a-new-era-of-agent-interoperability/ | [official] [3] Anthropic | Introducing the Model Context Protocol | 2024-11-25 | https://www.anthropic.com/news/model-context-protocol | [official] [4] OpenAI | Buy it in ChatGPT: Instant Checkout and the Agentic Commerce Protocol | 2025-09-29 | https://openai.com/index/buy-it-in-chatgpt/ | [official] [5] OpenAI | ChatGPT shopping research | 2025-11-24 | https://openai.com/index/chatgpt-shopping-research/ | [official] [6] Stripe Newsroom | Agentic Commerce Suite | 2025-12-11 | https://stripe.com/newsroom/news/agentic-commerce-suite | [official] [7] Forrester | Agentic payments in B2C commerce: where we are now | 2026-04 | https://www.forrester.com/blogs/agentic-payments-in-b2c-commerce-where-we-are-now/ | [third-party verified] [8] IBM Newsroom × NRF | Own the agentic commerce experience | 2026-01-07 | https://newsroom.ibm.com/2026-01-07-ibm-nrf-study-brands-and-retailers-navigate-a-new-reality-as-ai-shapes-consumer-decisions-before-shopping-begins | [third-party verified] [9] McKinsey & Company | The agentic commerce opportunity | 2025-10-17 | https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-agentic-commerce-opportunity-how-ai-agents-are-ushering-in-a-new-era-for-consumers-and-merchants | [third-party verified] [10] McKinsey & Company | When AI meets desire: human-centered luxury experiences in the agentic age | 2026-05 | https://www.mckinsey.com/industries/retail/our-insights/when-ai-meets-desire-innovating-human-centered-luxury-experiences-in-the-agentic-age | [third-party verified] [11] Deloitte | Agentic commerce: AI shopping agents guide | https://www.deloitte.com/us/en/Industries/consumer/articles/agentic-commerce-ai-shopping-agents-guide.html | [third-party verified] [12] Amazon | About Amazon: Rufus | 2024-02 (renamed Alexa for Shopping upon integration in 2026-05) | https://www.aboutamazon.com/news/retail/amazon-rufus | [official] [13] CNBC | Amazon faces a dilemma: fight AI shopping agents or join them | 2025-12-24 | https://

www.cnbc.com/2025/12/24/amazon-faces-a-dilemma-fight-ai-shopping-agents-or-join-them.html | [third-party verified] [14] Digital Commerce 360 | Gartner: AI agents to mediate $15 trillion in B2B purchases by 2028 | 2025-11-28 | https://www.digitalcommerce360.com/2025/11/28/gartner-ai-agents-15-trillion-in-b2b-purchases-by-2028/ | [third-party verified] [15] CMSWire | Machine customers: the structural break in customer experience | 2026-02 | https://www.cmswire.com/customer-experience/machine-customers-the-structural-break-in-customer-experience/ | [third-party verified] [16] Federal Trade Commission | FTC announces rule banning junk ticket and hotel fees | 2024-12-17 (effective 2025-05-12) | https://www.ftc.gov/news-events/news/press-releases/2024/12/federal-trade-commission-announces-bipartisan-rule-banning-junk-ticket-hotel-fees | [official] [17] Federal Trade Commission | FTC announces final rule banning fake reviews and testimonials (16 CFR Part 465) | 2024-08-14 | https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-trade-commission-announces-final-rule-banning-fake-reviews-testimonials | [official] [18] Civil Resolution Tribunal of British Columbia, Canada | Moffatt v. Air Canada (2024 BCCRT 149) | 2024-02-14 | https://decisions.civilresolutionbc.ca/crt/crtd/en/item/525448/index.do | [primary authority] [19] Munich Re / HSB | Introducing AI liability insurance for small businesses | 2026-03-18 | https://www.munichre.com/hsb/en/press-and-publications/press-releases/2026/2026-03-18-introducing-ai-liability-insurance-for-small-businesses.html | [vendor claim] [20] 6sense | 2025 B2B Buyer Experience Report | 2025-11 | https://6sense.com/science-of-b2b/buyer-experience-report-2025/ | [vendor claim] [21] China News Service | MIIT: blocking URL links is one of the key problems in the special rectification of the internet industry | 2021-09-13 | https://www.chinanews.com/cj/2021/09-13/9564236.shtml | [third-party verified] [22] Sina Tech (republished from ITHome) | Taobao launches RecGPT, its self-developed hundred-billion-parameter recommendation model | 2025-07-01 | https://finance.sina.com.cn/tech/digi/2025-07-01/doc-infcypra8733079.shtml | [third-party verified] [23] Genrise | AI shopping assistants 2026: Rufus, Sparky, ChatGPT, Perplexity | https://genrise.ai/insights/ai-shopping-assistants-future-ecommerce | [vendor claim] [24] CM.com | Five major customer experience trends for 2026 | 2026-05-25 | https://www.cm.com/zh-cn/blog/2026-customer-experience-tendency/ | [vendor claim] [25] Baijiahao (compiling Klarna's official statements and the CEO's public remarks) | Klarna's lesson: AI customer service is not a silver bullet | https://baijiahao.baidu.com/s?id=1865577917080864958 | [vendor claim]

[26] Paz AI | Google Merchant Center for AI Mode (Shopping Graph, 50 billion+ product listings refreshed at about 2 billion per hour) | https://www.paz.ai/guides/google-merchant-center-for-ai-mode | [third-party verified] [27] Lengow | ChatGPT Product Feed and Merchant Control (merchant-pushed feeds in CSV/TSV/XML/JSON, refreshing as often as every 15 minutes; in-chat checkout runs on ACP) | https://www.lengow.com/get-to-know-more/chatgpt-product-feed | [third-party verified] [28] Amazon | Business Solutions Agreement adds an Agent Policy (effective 2026-03-04) | https://sellercentral.amazon.com/seller-forums/discussions/t/84e3f6b1-42f7-4cf3-a189-a5cc8d78d838 | [official] [29] Profasee | Amazon AI Agent Policy: What Sellers Need to Know (SP-API registration, 12-month audit trails, self-identification, human authorization for price changes over 20% within 24 hours; 90-day transition window closed) | https://profasee.com/blog/amazon-ai-agent-policy-what-sellers-need-to-know | [third-party verified] [30] Ken Huang | Google and Shopify's UCP Just Won Agentic Commerce (on 2026-04-24 Amazon, Meta, Microsoft, Salesforce and Stripe joined the UCP Tech Council; UCP published 2026-01-11) | https://kenhuangus.substack.com/p/googles-ucp-just-won-agentic-commerce | [third-party verified] [31] Menlo Ventures | 2026: The State of Consumer AI (41% of AI users have tried an agent, 24% use one regularly, 32% have let AI act without final confirmation; among shoppers 50% compare prices, 43% narrow the shortlist, 37% decide on an AI recommendation, 26% let AI buy) | 2026-09-16 | https://menlovc.com/perspective/2026-the-state-of-consumer-ai/ | [third-party verified] [32] Gartner | Consumers want AI shopping help but not AI purchase decisions (only 11% of U.S. consumers would let AI make the purchase decision) | 2026-05-27 | https://www.gartner.com/en/newsroom/press-releases/2026-05-27-gartner-survey-finds-consumers-want-ai-shopping-help-but-not-ai-purchase-decisions | [third-party verified] [33] Checkout.com | Agentic Commerce 2026 (33% of consumers expect 10% of purchases to be AI-driven within a year; 72% of merchants say adoption outruns merchant readiness; 27% trust no institution, 24% will never delegate, 57% would switch brands for price) | 2026-06-09 | https://www.checkout.com/newsroom/consumer-demand-for-ai-shopping-is-forming-fast-but-trust-for-agentic-commerce-is-still-catching-up | [third-party verified] [34] Visa | Intelligent Commerce (released 2025-04-30; issuers, merchants and AI platforms share verifiable credentials for agent authorization and limits) | accessed 2026-09-26 | https://www.visa.com/en-us/solutions/intelligent-commerce | [official] [35] Mastercard | Agent Pay for Machines (released 2026-06-10; cross-card, cross-account, cross-currency machine payments) | 2026-06-10 | https://www.mastercard.com/us/en/news-and-trends/press/2026/june/mastercard-launches-agent-pay-for-machines.html |

[official] [36] Sina Tech | China's first agentic commerce trust protocol ACT released (Alipay with Qwen's app, Taobao Flash Sale and others; delegation, commerce interaction, payment service and trust service domains; "payment is always led or explicitly authorized by the user") | 2026-01-16 | https://finance.sina.com.cn/tob/2026-01-16/doc-inhhnkez5642443.shtml | [third-party reported] [37] Deloitte | Agentic commerce: The future of B2B commerce (buyers spend nearly 30% more with vendors offering an excellent buying experience; agent market value forecast at USD 45 billion) | 2026-06-26 | https://www.deloitte.com/us/en/what-we-do/capabilities/applied-artificial-intelligence/articles/b2b-agentic-commerce.html | [third-party verified] [38] Google | UCP and new merchant tools (Business Agent, new Merchant Center attributes, Direct Offers; on-the-spot offers in AI Mode recommendations) | 2026 | https://blog.google/products/shopping/ucp-agentic-commerce-tools/ | [official] [39] GoDataFeed | AI shopping product feed optimization (Google's legacy Content API for Shopping shuts down 2026-08-18, migrate to the Merchant API; conversational attributes added; OpenAI Instant Checkout scaled back 2026-03, purchases return to the merchant site) | 2026 | https://www.godatafeed.com/blog/ai-shopping-product-feed-optimization | [third-party verified: vendor guide] [40] ACI Worldwide (YouGov) | Six in Ten UK Consumers Would Stop Using an AI Shopping Agent After One Mistake (2,080 UK adults; 60% stop after one mistake; 19% trust AI to make everyday purchase decisions vs 55% trusting a human expert; 54% say the AI company should be accountable for refunds, 9% the retailer) | 2026-06-29 | https://investor.aciworldwide.com/news-releases/news-release-details/six-ten-uk-consumers-would-stop-using-ai-shopping-agent-after | [third-party verified]

[41] Stripe | Supporting additional payment methods for agentic commerce (SPT extended to Mastercard Agent Pay, Visa Intelligent Commerce and Affirm, Klarna; stated as the only provider supporting both agentic network tokens and BNPL tokens through a single primitive) | 2026-03-03 | https://stripe.com/blog/supporting-additional-payment-methods-for-agentic-commerce | [official] [42] Crossmint | Agentic payments protocols compared: MPP, ACP, AP2, x402 (ACP's four checkout endpoints and SPT; AP2's three mandate types and 60+ organisations; x402 and MPP at the settlement layer, with launch dates) | 2026-03-12 | https://www.crossmint.com/learn/agentic-payments-protocols-compared | [third-party verified: vendor engineering write-up]

[43] Salesforce | Shopping's New First Step: Agentic Search Grows 200% as Purchase Journeys Start in AI Chats (fourth edition of State of Commerce: survey of 3,450 commerce professionals and 4,690 shoppers, plus 1.5 billion shopper interactions; agentic search up 200% year over year; AI-chat referral traffic up 150%–428% year over year each quarter while overall traffic grows in the single digits to low teens; from 2025-08 to 2026-05 brand-owned channels −7%, traditional search −15%, new channels +38%) | 2026-07-28 | https://www.salesforce.com/news/stories/agentic-search-growth/ | [third-party verified: vendor's own data] [44] ARNnet (citing Gartner) | Buyers and Providers Must Adapt to AI in the Tech Buying Cycle (89% of global buyers use generative AI or AI agents in late-stage buying for contract review, supplier selection, proposal comparison and even price negotiation; Australia 97%) | 2026-05 | https://www.arnnet.com.au/article/4168709/gartner-buyers-and-providers-must-adapt-to-ai-in-the-tech-buying-cycle.html | [third-party verified] [45] G2 | 2026 Buyer Behavior Report: The Evaluation Maze (61% use or plan to use AI agents in the buying process, another 19% for select use cases; uses are total cost of ownership 51%, building shortlists 51%, researching solutions 49%, evaluating shortlisted vendors 46%; 47% would let an agent research and recommend; 9% accept purchases inside guardrails, 2% without pre-approval) | 2026-07-22 | https://www.tmcnet.com/usubmit/2026/07/22/10418570.htm | [third-party verified] [46] Digital Today (Chinese edition) | Cafe24 Connects AI Agents Across All Its Online Stores (WebMCP adopted on 2026-10-22; 11 shopping operations with no extra development; login and final payment still manual) | 2026-10-22 | https://www.digitaltoday.co.kr/cn/view/117917/cafe24-connects-ai-agents-across-all-its-online-stores | [third-party verified]

# Chapter 8: Three Extra Moves for Chinese Markets

GEO used to be a matter of bringing the overseas recipe over — hit Reddit, fill in the wiki, turn your official site into a page that can be cited. Now Chinese is a different foundation: your content grows inside the platform's container, and the customer's way of asking, the entry points, and the handoff all sit inside the wall. What you need to do is keep the mechanism and redo the playbook — recognize the terrain first, then decide where the money goes.

The Chinese corpus is another terrain, and how customers ask inside Doubao (Chinese: 豆包), Kimi, and WeChat (Chinese: 微信) is completely different. So this chapter gives you 5 methods for re-landing the points above in Chinese.

---

## Method 1: Start by recognizing that Chinese is walled-garden terrain

**What to do**: Go through the overseas GEO playbook item by item, check whether the raw material is present, and only then decide whether to spend the money.

**How**

1. Write down the difference between the two foundations. On the open web the unit of content is the URL: a public page can be read by anyone, and reputation travels from one page to another through backlinks. In a walled garden the unit of content is the account and the App: WeChat Official Account (Chinese: 公众号) articles, Douyin (Chinese: 抖音) videos, Xiaohongshu (Chinese: 小红书) notes, and Taobao (Chinese: 淘宝) product pages grow inside the platform's container, are invisible across platforms, are not indexed externally by default, and the power of distribution sits with the platform algorithm and in-app search.
2. Take the overseas recipe checklist and compare it item by item against Chinese, asking only one question: "is the raw material here?" Reddit, the wiki, YouTube, backlinks, structured pages. Whichever item's raw material is absent, that item cannot be done.

3. Write the conclusion as a switch: a playbook holds only when its raw material is present. When the raw material is absent, the same money spent on the same actions returns zero.
4. Accept that the closedness of the Chinese internet is a long-term default. On September 13, 2021, at a State Council Information Office press conference, MIIT (Chinese: 工信部) listed "blocking URL links" as one of the key problems in the special rectification of the internet industry, convened administrative guidance meetings with Alibaba (Chinese: 阿里巴巴), Tencent (Chinese: 腾讯), ByteDance (Chinese: 字节跳动), and others, required that blocking be lifted to standard within a deadline, and characterized the matter as harming the user experience, damaging user rights, and disrupting market order [third-party reported]. Having to launch a dedicated special campaign just to require that "links open" is itself proof that links not opening is the default state.

**Example**: Move "hit Reddit + Chinese Wikipedia + listicle articles" into Chinese and all three raw materials are absent — Chinese buyers mostly do not go to Reddit, and Chinese entry points mostly do not cite it; the entry count and update activity of Chinese Wikipedia cannot support the commercial corpus of a category; YouTube's counterpart in Chinese commercial contexts is Douyin and Bilibili (Chinese: 哔哩哔哩), while Douyin's comments and videos are themselves inside the walled garden, not public pages any engine can crawl. The reason that overseas correlation study writes "hit third-party platforms" as the recipe is that it measured the signals most correlated with AI visibility as YouTube mentions and branded web mentions, with brand mentions overall about 3 times traditional backlinks, while the number of content pages is almost irrelevant [vendor claim] — this conclusion holds on the open web, but move it inside the wall and the raw material is gone. **Composite sketch | constructed from common industry patterns, not a named event**: when some Chinese SaaS and consumer-goods companies go overseas, they still lay out content platforms first according to the domestic model, and only when they turn around to ask English entry points do they discover they have never once appeared in the English candidate set — the time was not wasted, it was spent on a stage English buyers do not watch.

**Pitfalls**: Confirm whether the raw material is present before deciding on the recipe. The actions themselves are not wrong; the mistake is assuming that the two foundations are the same.

## Method 2: Run a bilingual entry-point audit with the same set of questions

**What to do**: Take the buyer's real questions, ask them once on each of the Chinese and English entry-point sets, and put the two candidate sets side by side.

**How**

1. Build a bilingual question set. Take 10–20 real buyer questions and write them in both Chinese and English. Note: this is not translation — the Chinese version is written the way a Chinese buyer asks ("a 50-person SaaS company buying CRM"), and the English version the way an English buyer asks ("50-person SaaS company, under $300/month").
2. Weight the two sides of entry points equally. Chinese side: Doubao, DeepSeek, Kimi, Yuanbao (Chinese: 元宝), plus WeChat Search (Chinese: 微信搜一搜), Douyin Search (Chinese: 抖音搜索), and Xiaohongshu Search (Chinese: 小红书搜索); overseas side: ChatGPT, Google (including AI Overviews / AI Mode), Gemini, and Perplexity, plus Google organic search as the control baseline. Rank priority by "where your customers are", but both sides must be tested at the same time — testing only one side is no audit at all.
3. Fix the sampling discipline: run every question on every entry point at least 3 times; record the date, model / product version, whether it was online, and whether it was logged in each time; extract the lists with model extraction, not keyword regexes (a regex will treat a sentence fragment like "suitable for a 50-person company" as an entity name); where a channel is missing, mark it as missing honestly and do not substitute another model's results; open a fresh environment with no history, copy the questions verbatim, and do not "optimize".
4. Compute only three numbers: the mention rate (the share of samples in which you are mentioned — the denominator is the number of samples, not the number of questions), stability (the overlap of lists from repeated asks on the same entry point), and cross-entry / cross-ecosystem overlap (the overlap of candidate sets between different entry points and between Chinese and English).

5. For products with retrieval, record the cited domains that appear in each answer and invert them into a "list of sources to win over" — this list is the real allocation table for your content budget over the next six months.

**Example**: The book's editorial team ran one round of our own measurement (2026-09-13): 6 real procurement questions (3 Chinese-English pairs, covering CRM, MES, and study-abroad agencies) × 5 model channels × 2 repeats per question, yielding 51 valid answers. Asking the same question twice in a row on the same model, the list overlap was only 22%–56%; switching to another model, the pairwise overlap was only 4%–21%, and the union of four models' lists reached 43 entities; the cross-language overlap of the Chinese and English list pairs was 8%–26%; the study-abroad agency group was the only one with an entity present in both languages (IDP) [our own measurement]. One more detail: the English phrasing kept the geographic anchor "East China", and the English entry points also brought in Yonyou (Chinese: 用友), Kingdee (Chinese: 金蝶), and Digiwin (Chinese: 鼎捷) — the anchor itself is a variable.

**Pitfalls**: Testing only one side. Test only the Chinese side and you see the distribution inside the wall; test only the English side and you know nothing about the Chinese battlefield.

## Method 3: First build the public assets that can be read outside the wall

**What to do**: Externalize the facts with the most value inside the wall into content that can be publicly crawled — this step is 0 to 1, not optimization.

**How**

1. Distinguish the nature of the two problems. Overseas is "you already have a pile of pages; optimize the probability that they get cited" — a probability problem that can be A/B tested and iterated; Chinese may be "you have no assets readable outside the wall at all; you must first externalize the facts with the most value inside it into public content" — an asset problem, from 0 to 1. Confusing the two is the most common strategic error Chinese teams make with GEO.

2. Split content inside the wall into two layers, "facts" and "container": parameters, specifications, data, and cases are facts; Official Account layouts, short-video edits, and DM scripts are the container. Facts can be redone outside the wall; the container cannot be moved out.
3. Prioritize externalizing three things: press releases, encyclopedia entries, content that industry sites can republish, and your own official-site pages — these are the raw material that Chinese entry points can still read today.
4. Do not expect platform content to be read outside the wall. Official Accounts built up over years, tens of thousands of Xiaohongshu notes, and hundreds or thousands of short videos are not externally indexed by default; the more carefully you cultivate content inside the wall, the more likely it is that "the description of you outside the wall is not written by you" [our judgment].

**Example**: A marketing vendor's self-run measurement (as of September 2026, 14,750 buyer questions × four platforms ChatGPT / Perplexity / Claude / Gemini) reports: 51.7% of AI citations point to brands' own official-site pages, guides and blogs account for 14.3%, and the rest flows to third-party editorial content [vendor claim]. This figure was measured on the open web and can only indicate an order of magnitude; but the mechanism it points to is clear — as long as your site can be read, it is your biggest lever. Where official sites are indexed, "official site first" leads to redoing pages; where official sites are not indexed, the same sentence leads to moving the facts outside the wall first.

**Pitfalls**: Treating "raise the probability of being cited" as the first step. You are doing probability optimization while what you lack is assets — the budget is spent and there is still nothing to cite.

## Method 4: Redo one piece of material per platform, and catch the customer inside the wall

**What to do**: The Chinese ecosystem's conversion loop is designed to "close inside the platform", so content, search, the handoff, and lead capture have to run end to end within one platform.

**How**

1. Accept the redo cost: the same material must be redone per platform as image-and-text posts, short videos, livestreams, and DM scripts, and one production run is valid for only one platform.
2. Move the handoff point forward into the content: say it clearly in the video — "DM me", "tap here to leave your details" — and do not leave the handoff to an official site outside the wall, because platforms actively block external links.
3. Treat "diverting traffic off-platform is strictly prohibited" as a design premise, not a prohibition: if the customer is seeded on Xiaohongshu, complete the DM or the lead capture on Xiaohongshu and do not pull them off-platform.
4. Run on a fixed cadence: image-and-text posts weekly, short videos weekly, livestreams per campaign, and DMs answered 24/7 by a person or by AI. Content inside the platform, search inside the platform, handoff inside the platform, and the transaction back inside the platform — drop one of the four steps and customers leak away.

**Example**: A channel review from the education-and-training industry sets out a Chinese division of labor: ground promotion and outbound calls suit community-based and newly opened campuses, Douyin suits institutions with content capability to amplify exposure, Dianping (Chinese: 大众点评) catches price-comparing users with intent to visit, the WeChat ecosystem handles private-domain retention, and Xiaohongshu is treated as the key platform for cold-start new growth; the article cites platform data saying Xiaohongshu has 300 million monthly active users, a monthly search penetration of 70%, and about 200 million users a month actively seeking purchase-decision advice, and notes that Xiaohongshu strictly prohibits diverting traffic off-platform [third-party reported: platform data relayed]. An education institution uses a multi-model hybrid Agent for 24/7 automated reception and tiered follower follow-up on Douyin DMs, aggregating more than 100 accounts, claiming it independently handles over 85% of professional inquiries at night, a 580.6% increase in DM lead capture, and a 4.6× overall ROI improvement [vendor claim]. On the private-domain side, a retail white paper states that its platform data shows WeCom (Chinese: 企微) private-domain users have a repurchase rate 2.8 times higher than public-domain users and more than 45% higher value per customer [vendor claim]. In-platform search is no longer a supplementary channel but three parallel trunk roads: on QuestMobile's count WeChat AI search had 160 million monthly actives as of June

2025, more than native AI apps such as Doubao and Yuanbao [third-party verified]; Xiaohongshu runs about 600 million searches a day, more than half of Baidu's volume, with the company saying over 70% of its monthly active users search and the average user runs about six searches a day [third-party reported]; and Douyin e-commerce's shelf scenarios (search, mall, storefront) grew gross merchandise volume 49% year over year between August 2024 and August 2025 [third-party verified: relaying a Juliang Engine report]. One checkable mechanic sits on that chain: Tencent Marketing's official FAQ lists WeChat Store's five core operating scenarios as Channels live and short video, article product cards in official and service accounts, affiliate creator selling, gifting, and WeChat Search, and states plainly that the public domain acquires and the private domain repurchases [official]. In-platform agent gateways took shape in the same window: on June 16, 2026 Alipay began testing its built-in assistant Abao (Chinese: 阿宝), letting users complete tasks by conversation as the AI calls Alipay functions and mini programs, with Alipay stressing that every step requiring payment is confirmed by the user in person; QuestMobile data put WeChat, Alipay and Douyin mini-program monthly actives at 973 million, 644 million and 273 million respectively in March 2026 [third-party verified]. Alipay's cadence was three moves — a full-stack AI-native payment system in May, Abao in June and the AI open platform in July — and Tap-and-Go, upgraded with AI in parallel, now covers 400 million users, more than a thousand brand merchants and over 30 million offline agent nodes; the 2.0 version of the ACT protocol, released in April, was co-built by IIFAA with Alipay, Zhipu (Chinese: 智谱) and more than 20 other vendors, with KFC, Mixue (Chinese: 蜜雪冰城) and Amap ride-hailing (Chinese: 高德打车) among the first to connect [third-party verified: media]. WeChat issued AI-ecosystem access guidance and an AI-exclusive card scheme in parallel, while JD.com released the agentic payment protocol A2P2 in May, splitting AI spending authority into six levels (L4 lets the AI pay freely within a set quota) [third-party verified]. The economics of the gateway are not proven yet: one service provider says about 70% of its mini-program clients give up running their mini programs after a few months, and smaller developers worry the AI will recommend larger, more capable platforms first [third-party verified]. At Alipay's AI ecosystem conference, Ant Group CEO Han Xinyi said agentic commerce will break out in the next 6–12 months, and Alipay Business Group president Li Jun said some development requests since the July launch of the AI open platform are already scheduled into next year, alongside incentives including 100 million free tokens per developer, payment-rate reductions and free access to 57 atomic operating capabilities [third-party verified]. The transaction side now has a

reading that can be counted by the week: in the week of February 5–11, 2026, Alipay's (Chinese: 支付宝) "AI Pay" processed more than 120 million payment transactions, which Ant Group (Chinese: 蚂蚁集团) calls the first AI-native payment product anywhere to pass 100 million transactions; it is live in Qwen (Chinese: 千问), Rokid and Luckin (Chinese: 瑞幸) scenarios, with Qwen's 3 billion yuan Spring Festival free-order campaign as the accelerant [official: company release, relayed by media].

**Pitfalls**: Treating platform content as an asset you can just carry away. In-platform content cannot get out, and off-platform traffic platforms will not let you in — both ends must be redone according to platform rules.

## Method 5: Rank entry points by where the customer is, and look at each platform's own AI separately

**What to do**: Chinese entry points are fragmented; first rank them by "where your customers are", then measure each entry point's citation sources one by one.

**How**

1. List the Chinese entry points: Doubao, DeepSeek, Kimi, Tencent Yuanbao, Tongyi (Chinese: 通义) / Quark (Chinese: 夸克), and Ernie (Chinese: 文心), plus the three in-platform search entry points WeChat Search, Douyin Search, and Xiaohongshu Search.
2. Ask each entry point the same batch of questions once, record the cited domains that appear in the answers, and only after forming a distribution decide the content budget. For the composition of citation sources across Chinese entry points, this book currently has no cross-entry measured data; use it as a hypothesis pending verification.
3. Look separately at the AI entry points that "grow inside the body of a walled-garden company": Doubao grows inside the ByteDance ecosystem, Yuanbao inside the Tencent ecosystem. For such an entry point, what can it read first — content inside its own wall, or content outside it? For now there is only a hypothesis, no data [our judgment].

4. Do not treat AI assistants as the whole of Chinese entry points. Chinese B2B buyer research falls heavily on trade shows, acquaintances, industry WeChat groups, and Baidu (Chinese: 百度) / Zhihu (Chinese: 知乎), and this is especially true for buyers at industrial-belt companies — **its buyer does not ask inside AI; it asks in the group chat.**

**Example**: A Chinese GEO industry research report states that Doubao, DeepSeek, Tencent Yuanbao, and Qwen (Chinese: 千问) grow explosively and contribute the main increment through ecosystem advantages or vertical scenarios, with users concentrating toward top Apps; more than 40% of users have shifted the center of gravity of their searching from traditional search engines to AI search, and more than half use both; the misconception the report singles out is "understanding AI brand strategy with the mindset of traditional search engines and performance advertising" [third-party secondary: quantitative conclusions not item-by-item verifiable]. CNNIC's 56th Statistical Report on Internet Development in China says that 80.9% of generative AI users use it for answering questions, and that as of March 2025 a total of 346 generative AI services had completed filing [official]. The scale of the Chinese gateways and the way they are priced were both made public around the same time: by QuestMobile's count, AI-native apps reached 499 million monthly active users by June 2026, up 85.4% year over year, with Doubao (Chinese: 豆包) first at 382 million, Qwen (Chinese: 千问) second at 167 million and DeepSeek third at 130 million [third-party verified]; and from August 2026, hotel orders traced back to Doubao and completed through Douyin Laike carry a separate rate — 11.4% software service fee plus 0.6% payment fee, about 12% in total, with lifestyle services at 17.4% and 18%; Doubao's response was that there is currently no paid promotion, that merchants cannot pay to influence recommendations or ranking, and that the channel fee is paid only after an order closes [third-party verified: including the official response]. The gateway position itself is moving up: WeChat upgraded AI search from a second-level entry to a first-level one, callable by tapping the search box at the top with an optional "deep thinking" mode, and Tencent confirmed it is still in grey-scale testing [third-party verified]. Payment and protocol work is being laid down in parallel: Alipay released AHA, a multi-agent cross-device interconnection protocol system; in its OPPO partnership, the phone maker's AI assistant connecting to nearly 200 Alipay services added more than 2 million in activity on that brand's handsets; merchant onboarding still runs on the base payment rate with no extra technical service or channel fees for now, and the payment rate can be discounted 40% to encourage adoption [third-party

verified]. The entry-point layering has a plainer truth underneath: QuestMobile's 2026 half-year report on the AI application market gives 499 million monthly actives for native AI apps, 644 million for AI plug-ins inside super-apps, and 755 million for device makers' AI apps [third-party verified] — the largest AI surface lives in parasitic gateways. Source preference also splits by platform: a Chinese technical-community write-up says Kimi leans towards academic papers and long text, Doubao towards ByteDance's content ecosystem, the Baidu family towards its own content pools, and Tencent Yuanbao on WeChat official accounts; another vendor piece citing Forrester's 2026 Q2 report says central and national media are 9.4 times more likely to be cited than ordinary self-media [third-party verified: vendor relay, original report not attached]. User sentiment is far colder: Shanghai Consumer Council's 2026 6.18 survey collected 4,308 valid responses, and while 84.56% of consumers had tried AI shopping features, only 16.06% thought AI matched products accurately and 38.79% said AI pushed higher-priced items first; the report sums up the year's brand competition as moving from occupying consumer mindshare to occupying the AI recommendation slot [official: consumer council survey, relayed by media]. The ranking itself is shuffling: the Q1 2026 AI application value ranking from QuestMobile and National Business Daily shows Doubao's monthly actives up 44.2% quarter over quarter to 345 million (a 180-million gap over the runner-up), Qwen at 166 million in second, DeepSeek down to 127 million with four consecutive quarters of decline, and Tencent Yuanbao at 57.35 million dropping out of the top three for the first time, while the top five AI apps together reached 722 million monthly actives and seven ByteDance products entered the monthly-active top 20 [third-party verified].

For the first time, AI advice outranks traditional advertising in consumers' own accounts: a survey of consumer spending intentions for Q2 2026 by the China Association of Small and Medium Commercial Enterprises and Lixin Data Research Institute found 79.4% of consumers consult AI tools such as Doubao or DeepSeek while shopping, and 43.7% say AI advice has “a great deal” or “a fair amount” of influence on their decisions, surpassing traditional media advertising (42.6%) for the first time [third-party verified]. Platform-side readings from the same Tmall Double 11 briefing put views of AI-related categories on Taobao above 4 billion, with more than 10 million users completing interactions with AI-driven trends [third-party verified]. The 43.7% versus 42.6% marks a split in what people say, not a handover of transactions — safer as a basis for ordering entry points than as a growth promise.

**Pitfalls**: Treating "AI search" as a single channel. The same analysis of 680 million citations says only 11% of domains are cited by both ChatGPT and Perplexity [vendor claim] — even inside one ecosystem, AI search is not a single channel; across ecosystems, even less so.

---

## Notes for this chapter

- The bilingual audit measures the model-memory layer of a bare API, not consumer AI search products with live retrieval (ChatGPT Search, Perplexity, AI Overviews, Chinese assistants with retrieval). It answers "how the model-memory layer organizes the candidate set" and cannot be extrapolated to "the citation behavior of AI search products". [our own measurement]
- The cross-language overlap of 8%–26% is sensitive to definition: recomputing the raw data with two other algorithms gives a range of 2.3%–31.2%. The difference comes entirely from the definition of "overlap" (whether models are merged, whether brand names are normalized). The direction (the two lists barely intersect) is usable; the specific percentages are not citable as precise values. [our own measurement]
- The two points "Chinese entry points rely mainly on publicly crawlable second-hand information" and "platform content cannot be read outside the wall" are currently mechanism hypotheses, not measured conclusions; this book still has no measurement of the citation-source composition of Chinese entry points. [our judgment]
- The 51.7% item comes from a tool vendor's self-run study and its sample basis has not been independently verified; the 680-million-citation analysis likewise comes from a tool vendor's self-run study, with neither sample nor method independently verified. Both are used only to judge order of magnitude. [vendor claim]
- Xiaohongshu's 300 million monthly active users, 70% search penetration, and about 200 million users a month seeking purchase-decision advice come from a third-party article relaying Qiangua Data's (Chinese: 千瓜数据) 2025 Active User Research Report; this round saw only the relay, not the original report. [third-party secondary]

## Sources for this chapter

[1] China News Service | MIIT: blocking URL links is one of the key problems in the special rectification of the internet industry (2021-09-13, State Council Information Office press conference) | https://www.chinanews.com/cj/2021/09-13/9564236.shtml | [third-party reported] [2] China Internet Network Information Center (CNNIC) | 56th Statistical Report on Internet Development in China (2025-07-21, data as of June 2025) | https://www.cnnic.net.cn/n4/2025/0721/c88-11328.html | [official] [3] Boring Marketing | AI brand visibility statistics (14,750 buyer questions × four platforms, data page updated in real time) | as of 2026-09 | https://boringmarketing.com/ai-visibility-statistics | [vendor claim] [4] Leapd | Analysis of 680 million AI citations: ChatGPT, Google AI Overviews, and Perplexity source information differently | 2026-04 | https://www.leapd.ai/blog/ai-visibility/how-chatgpt-google-ai-overviews-and-perplexity-source-information-in-2026 | [vendor claim] [5] Ahrefs | Studying 75,000 brands: what determines whether a brand is recommended in AI answers | 2025-12 | https://ahrefs.com/blog/ai-brand-visibility-correlations/ | [vendor claim] [6] Chinese GEO industry research report (republished on Zhihu / Sina Finance) | 2026 | https://zhuanlan.zhihu.com/p/2012459780288046392 | [third-party secondary: mostly image-and-text, quantitative conclusions not item-by-item verifiable] [7] Jiaoyujie | Why is Xiaohongshu the next customer-acquisition battlefield for tutoring institutions? | 2026 | https://www.jiaoyujie365.com/N/2411.html | [third-party reported: platform data relayed from Qiangua Data's 2025 Active User Research Report] [8] Laigu Pro | Case study: AI DM automation for Douyin customer acquisition in the education industry | 2026-07 | https://laigu.com/blog/教育行业抖音获客用ai私信自动回复30天提升获线/ | [vendor claim] [9] 36Kr Research Institute / Tencent Smart Retail | 2024 White Paper on Private-Domain Operations and Omnichannel Growth in China | https://36kr.com/p/2684930182745348 | [vendor claim: data self-reported by the report's authors] [10] Our own measurement 1 | Multi-model measurement of the candidate set | 2026-09-13 | Methods and raw data: `实测/实测1-候选集多模型/` (`实测报告.md`, `结果.csv`, `raw_with_vendors.json`, `问题集.json`) | [our own measurement]

[11] STCN | Doubao responds to "12% commission on hotel recommendations" (citing QuestMobile: AI-native apps at 499 million MAU, +85.4% year over year; Doubao 382 million, Qwen 167 million, DeepSeek 130 million) | https://www.stcn.com/article/detail/4068900.html | [third-party verified] [12] The Paper | Doubao responds to "hotel recommendations are charged" (11.4% software service fee + 0.6% payment fee, about 12% in total; lifestyle services 17.4%/18%) |

2026-08-11 | https://www.thepaper.cn/newsDetail_forward_33757109 | [third-party verified: including the official response] [13] Shangguan News (Jiefang Daily) | China's first multi-agent cross-device interconnection protocol system AHA released (2026-08-17; OPPO's "Xiaobu" assistant connecting to nearly 200 services added more than 2 million in activity; merchants can get a 40% discount on the payment rate) | https://www.shobserver.com/wx/detail.do?id=1162176 | [third-party verified] [14] Jingxiang Studio | WeChat AI search upgraded from a second-level entry to a first-level one (in grey-scale testing; Tencent's official reply) | https://www.moomoo.com/hans/news/post/58266167 | [third-party verified] [15] 21st Century Business Herald (Chinese: 21 世纪经济报道) (citing QuestMobile) | WeChat AI search at 160 million monthly actives (June 2025) | 2025-09-23 | https://www.21jingji.com/article/20250923/herald/a6cdb8dc7c606154d8959d6ac550c50d.html | [third-party verified] [16] Renren dou shi chanpin jingli (Chinese: 人人都是产品经理) | Xiaohongshu runs about 600 million searches a day (over 70% of monthly active users search, about six searches a day each) | https://www.woshipm.com/share/6156977.html | [third-party reported] [17] 36Kr (relaying Juliang Engine's 2025 Douyin E-commerce Development Data Report) | Douyin e-commerce shelf scenarios up 49% year over year, August 2024–August 2025 | 2026-03 | https://m.36kr.com/p/3705534783221896 | [third-party verified] [18] Tencent Marketing official FAQ | WeChat Store's five core operating scenarios and "public domain acquires, private domain repurchases" | updated 2026-08-24 | https://e.qq.com/faq/wechat-store/growth/faq-zcjc-003/ | [official] [19] Alibaba Cloud developer community (Chinese: 阿里云开发者社区) | Source preferences of major Chinese AI platforms (Kimi leans academic long-form, Doubao ByteDance content, Baidu's own pools, Yuanbao on the WeChat ecosystem) | 2026-09 | https://developer.aliyun.com/article/1765203 | [third-party verified: technical community] [20] Phoenix Technology (relaying Huanqiu Tech, citing Forrester 2026 Q2) | The "source pyramid" of Chinese AI platforms (central media 9.4× more likely to be cited than ordinary self-media) | 2026-08-30 | https://tech.ifeng.com/c/8w0mT4sJjQS | [third-party verified: vendor relay] [21] Jiefang Daily / Shangguan News (citing the Shanghai Consumer Council's 2026 6.18 Online Shopping Experience Survey) | AI matching accuracy at only 16.06%, 38.79% report higher-priced items pushed first (4,308 valid responses) | 2026-06-26 | https://www.jfdaily.com/wx/detail.do?id=1134996 | [official: consumer council survey] [22] National Business Daily (Chinese: 每日经济新闻) × QuestMobile | Q1 2026 AI application value ranking (Doubao monthly actives +44.2% to 345 million, a 180-million gap over second; Qwen 166 million; DeepSeek 127 million after four straight quarters of decline; Tencent Yuanbao 57.35 million, out of

the top three for the first time; top five AI apps 722 million monthly actives; seven ByteDance products in the top 20) | 2026-04-21 | https://www.nbd.com.cn/articles/2026-04-21/4350446.html | [third-party verified] [23] 21st Century Business Herald (Chinese: 21世纪经济报道) | The two paths agentic gateways take to break the walls (WeChat/Alipay/Douyin mini-program monthly actives 973/644/273 million per QuestMobile March 2026; Alipay testing Abao on 2026-06-16; JD A2P2 six-level AI spending authority in May; Tencent AI-exclusive card in WeChat Wallet; Alipay says every payment step is confirmed by the user; about 70% of a provider's mini-program clients give up after a few months) | 2026-06-17 | https://m.21jingji.com/article/20260617/herald/ae1ed1491f5fce9997a168b6523178e7_zaker.html | [third-party verified] [24] Leiphone (Chinese: 雷锋网) | Alipay's three moves in agentic commerce (full-stack AI-native payment system in May, Abao in June, AI open platform in July; Tap-and-Go covers 400 million users, over 1,000 brand merchants and 30 million+ offline agent nodes; ACT 2.0 in April co-built with Zhipu and 20+ vendors; KFC, Mixue, Amap ride-hailing among the first) | 2026-07-30 | https://m.leiphone.com/category/aijuejinzhi/PkorfY0MIiENCG54.html | [third-party verified: media] [25] Science and Technology Daily (Chinese: 科技日报) | Alipay releases a full-stack agentic commerce foundation (AI ecosystem conference 2026-08-17; 10,000+ services AI-enabled; 5 phone brands and 16 carmakers across devices; AHA protocol system; 20+ companies co-building; 100 million free tokens per developer, payment-rate reductions, 57 atomic operating capabilities free; agentic commerce to break out in 6–12 months) | 2026-08-17 | https://www.stdaily.com/web/gdxw/2026-08/17/content_565275.html | [third-party verified]

[26] China Securities Journal / Zhongzheng Jinniu (Chinese: 中证金牛座) | Alipay "AI Pay" passes 120 million transactions (120 million+ payments in the week of 2026-02-05 to 02-11; described as the world's first AI-native payment product to pass 100 million transactions; live in Qwen, Rokid, Luckin and other scenarios) | 2026-02-12 | https://jnzstatic.cs.com.cn/zzb/htmlInfo/115866.html | [official: company announcement, relayed by media] [27] Jiemian News (Qujie Business) | The 18th Double 11: AI Moves From Traffic Allocator to the Front Line of Shelf Commerce (survey by the China Association of Small and Medium Commercial Enterprises and Lixin Data Research Institute: 79.4% consult AI tools while shopping, 43.7% say AI advice influences decisions, surpassing traditional media advertising at 42.6% for the first

time; AI-related category views on Taobao above 4 billion, more than 10 million users completing AI-trend interactions) | 2026-10-23 | https://www.jiemian.com/article/15231326.html | [third-party verified]

# Chapter 9: Advertising: From "Selecting an Audience" to "Giving a Target"

Advertising used to mean nailing the audience precisely: pick the region, pick the age, pick the interests, pick the keywords — the finer your targeting, the better your economics. Now the platform has taken this step back — you give it a target and a budget, and it goes off into the big pool to find people itself. Audience targeting is turning from your job into the system's job.

Customers are still pushed along by the feed; only the way you buy has changed. So this chapter gives you 6 methods, switching from "picking audiences" to "giving targets".

---

## Method 1: Change "selecting an audience" into "giving a target"

**What to do**: Stop fine-tuning targeting; give the system a target cost per conversion and let it find people on its own.

**How**

1. First pick one conversion event in the account that **can report back reliably** — an order, a lead capture, or an add on WeCom (Chinese: 企微); pick one, do not hang five on at once. This one is what the system needs to learn; hang too many and it does not know what to learn toward.
2. Set the target cost per conversion near your historical cost. Read this number as "I want people at roughly this price point"; it is not a "cheaper is always better" knob.
3. Give the system a learning period, and during it do not repeatedly change the bidding strategy. With too few conversion samples on hand, the system cannot tell who your people are, and the more often you change things now, the less it learns.
4. Change only one thing at a time: either the budget or the target cost, never both on the same day. After a change, watch for a sufficient stretch of time before judging.

5. Test the bidding strategy and the creative separately. Strategy is "whom to find"; creative is "what words to find them with"; change the two together and even when a result comes you cannot tell which one earned it.

**Example**: Baidu's (Chinese: 百度) official help pages describe oCPC's use very plainly — set a target cost per conversion, and the system optimizes toward that cost [official]. Ocean Engine's (Chinese: 巨量引擎) official product manual describes automated delivery as "pooling human delivery experience and system optimization capability to cover the entire delivery chain" [official]. As for the learning period, that Ocean Engine scaling material explains it in detail: at the start of a campaign the system has little conversion data on hand, so it can only try users in large numbers to feel its way; accumulate enough and it moves into stable delivery, fall short and the exploration fails [vendor claim]. **This passage directly explains why "it just will not run when you first turn it on" is normal.** Automation goes further on another platform: Meta is testing a format where the advertiser supplies only a product URL and a budget and AI generates the whole ad set, and on the facts already landed it reports Advantage+ delivering average ROAS 22% higher than manually managed campaigns, with more than 4 million advertisers using its generative tools (1 million six months earlier) [vendor claim]. Once targeting, creative, placements and budget allocation are all taken over by the system, the advertiser is left with two things: pointing it at the right link, and treating budget as test ammunition rather than a promise.

The scale of automation can be quantified too: on Meta's Q2 2026 earnings call, Advantage+ reached a $75 billion annual revenue run rate and more than 9 million small businesses used at least one of its generative AI creative tools; the same roundup cites Motion's analysis of $1.3 billion in spend, in which only about 5% of ads became “winners” and about half barely received any spend [third-party verified]. The two numbers only read correctly together: the system has taken over allocation, and allocation is extremely concentrated — so the outcome of a campaign moves upstream to the supply side of creative and data.

**Pitfalls**: Squeezing the target cost clearly below the historical value. The lower you set the target, the most common outcome is that volume simply will not run — you think you are saving money, but you are actually tying the system's hands.

## Method 2: Give the system a direction; do not expect it to fence people in

**What to do**: Use a seed list or audience signals to tell the system which direction to search in, while knowing clearly that it is only a hint.

**How**

1. Hand over the people who most resemble your customers: your existing closed-customer list, high-intent visitors, and people who added you on WeCom but did not order. These three resemble customers to different degrees; submit them separately and look at them separately.
2. Give an asset group only one clearly themed seed. Cramming industry terms, competitor terms, and behavioral audiences all into one group points in no direction at all.
3. Once a week, look at the audience the system actually ran to and compare it with your seed. If it drifted, change the seed rather than going back to add targeting and weld the opening shut.
4. Reserve a small slice of budget to run a "no seed" control group. Without this control, you can never say whether the seed did the work or the system would have found these people anyway.

**Example**: The Google Ads API documentation says in so many words: the audience signal in Performance Max (Asset Group Signal) is a "hint" to the system about user intent and preferences, used to "find new customer segments you might not have anticipated", and it is not itself hard targeting [official]. The documentation also explains that one asset group can carry multiple signals, but each must be added separately [official]. The target is being treated by the system as the price you will pay: Google Ads' official note says that from August 17, 2026, after the update to goal-based bidding, a budget-limited campaign that has been beating its stated target will converge toward that target when budget is added — the official example is a target CPA of $10 with recent actual CPA of about $5, where the campaign will now deliver closer to a $10 actual CPA, and holding $5 requires you to lower the target yourself [official]. The other side of the same shift is a new surface: OpenAI's official page records the ChatGPT ads pilot — starting in the U.S. on February 9, 2026 for logged-in adult users on the Free and Go tiers, with Plus, Pro, Business, Enterprise and Education tiers ad-free, and on May 7 announcing expansion to the United Kingdom, Mexico,

Brazil, Japan and South Korea, repeatedly stressing that ads do not change answers and that conversations stay private from advertisers [official]. The conversation gateway has gone from "you can only be recommended" to "you can be recommended and there are ad slots," leaving the target and the creative as the two things you still control. Bidding is moving towards the journey: Google launched beta Journey-aware Bidding across Search, Shopping and Performance Max, optimizing against the full lead-to-sale journey rather than mainly form submissions, and extended Smart Bidding Exploration from Search to PMax and Shopping in beta — Google says advertisers already using it in Search saw 27% more unique converting users on average [third-party verified: includes official claim data]. The consolidation of creative and targeting finished at the same time: on April 15, 2026 Google announced it was upgrading Dynamic Search Ads to AI Max and taking AI Max out of beta, saying campaigns using the full feature suite — search-term matching, text customization and final URL expansion — see an average of 7% more conversions or conversion value than search-term matching alone; starting in September 2026, campaigns using DSA, automatically created assets (ACA) and campaign-level broad match will be automatically upgraded to AI Max, and a June 11 update pushed the DSA sunset and auto-upgrade to February 2027 while ACA and campaign-level broad match still upgrade in September [official].

**Pitfalls**: Setting an audience signal as if it were hard targeting, then finding the system running outside the signal and concluding "the system is disobedient". It was never designed to fence people in — **what you actually want is not to fence things in, but to avoid pointing in the wrong direction**.

## Method 3: Turn one piece of creative into a batch of creative

**What to do**: The unit of competition in delivery changes from "one good ad" to "a batch of creative the system can freely combine".

**How**

1. What you submit at one time is multiple headlines, multiple images, multiple short videos — not one set polished to perfection.
2. Creative and landing page must say the same thing. If the headline talks low price and the landing page is a premium line, conversion will definitely collapse — this has nothing to do with how smart the system is.

3. Refresh the whole batch after a while. Creative has a lifespan; run the same batch too long and performance decays on its own — do not wait until the drop is visible to the naked eye to top it up.
4. Prepare creative according to delivery scale. That Ocean Engine material gave a rough rule of thumb: the number of ads is roughly daily spend divided by 100, and the number of creative assets roughly daily spend divided by 1000 [vendor claim]. This ratio need not be copied, but you can use it to check whether your own creative volume is obviously too low.

**Example**: The Google Ads Help Center defines Ad Strength as real-time feedback on the "relevance, quality, and diversity" of the creative, used to judge whether an asset group is "set up well and can maximize performance" [official]. Tencent Ads (Chinese: 腾讯广告), for its part, talks about automatic placements by pairing "diverse creative adaptation" with "intelligent optimization strategy" [official]. The Chinese side is merging creative production and delivery into one pipeline too: Baidu Marketing (Chinese: 百度营销) launched its Qingduo 3.0 creative agent and an upgraded smart delivery product, AIMax, at a 2026-09-09 event, plus a delivery assistant for data diagnosis and strategy output, positioned officially as a line running from creative production through delivery validation and review diagnosis back into creative iteration [third-party report of an official launch event]. On the creative-cost side the public numbers are more concrete than the word "efficiency": in a compilation of cases with disclosed figures, Klarna's AI image pipeline produced more than 1,000 marketing visuals in Q1 2024, cut image development from about six weeks to seven days and saved roughly $10 million in annual marketing cost; Monks and Hatch cut cost per purchase by 31%, lifted click-through rate by 80% and cut production cost by 97%; Novo Nordisk lifted email open rates by 24% with generated subject lines; and Lidl France generated 1.7 million localised visuals in three weeks [third-party verified: case compilation]. Read those numbers in two columns: generation cases report cost, speed and reach, optimisation cases report ROAS and click-through, and blending them overstates your creative capability [third-party verified: case compilation].

The supply side of creative is being automated whether you ask or not: from 4 December 2026, Google turns GENERATE_IMAGE_EXTRACTION on by default for newly created Performance Max campaigns, pulling high-resolution images from the landing page for asset expansion across Search, YouTube, Display, Discover, Gmail and Maps; the new default no longer inherits the account-level Dynamic Image Extensions setting, so campaigns created while the account switch is off will still extract unless

the request explicitly writes OPTED_OUT [official]. Stacked on the landing-page images and sourced videos added in October 2026, creative auditing has to move ahead of campaign build: clear the expired promo images, third-party assets and unlicensed images off the landing page first, then decide which automation switches to leave on [third-party verified].

**Pitfalls**: Thinking that writing pretty copy is enough. What the system rates is **diversity** — the same meaning phrased ten ways is still one set of creative in its eyes.

## Method 4: Feed your own list back into delivery

**What to do**: Connect your own data into delivery, then read the results back — do not just keep spending on strangers.

**How**

1. Push first-party audiences (your own list, website behavior) to the ad platform. Most cloud vendors' marketing products provide this channel; you do not have to build it yourself.
2. After pushing, separately collect the metrics for **this batch of people**: spend, impressions, CPM, clicks, CTR, and cost per click.
3. Look only at this batch's performance; do not evaluate it with whole-account averages. The whole account mixes in people the system found on its own, and with the two blended together you cannot tell whether the list is useful.
4. The list itself must iterate: feed the traits of the best-performing segment back into the seed pool for the next round, forming a loop. This step is done on a rolling basis, not once.
5. Align definitions before pushing the list. What "closed customer" means in your CRM and what "conversion" means on the platform — if the two do not line up, the numbers you read back are wrong, and the same customer may count as two people across the two sides. If the list contains people who have unsubscribed or explicitly refused marketing, remove them first.
6. Do not expect the list size to be right in one go. A small list is not necessarily useless, but with a small sample the swings are large — look at the trend, not a single day.

**Example**: The Alibaba Cloud (Chinese: 阿里云) help documentation gives the path: after an audience is pushed to the platform, you can review spend, impressions, CPM, clicks, CTR, and cost per click for a specified account, audience, and order [official].

**Pitfalls**: Pushing without reading back. Sending the list out and not looking at what comes back is handing over your customer data for free and getting no conclusion in return.

## Method 5: Label AI-generated creative

**What to do**: For AI-generated images, videos, and copy, declare their identity according to the rules before publishing.

**How**

1. Distinguish the two kinds of marking: an **explicit label** is a notice the user can see; an **implicit label** is written into the file data and is not easily perceived. Both are required — you do not choose one.
2. Add a prominent notice in an appropriate place on images; add notices on the opening frame and around playback for videos, and optionally at the end and in the middle; for virtual scenes, add one on the opening frame.
3. The marking must still be there after the creative file is downloaded, copied, or exported — it cannot hang only in the delivery backend and vanish once exported.
4. If you go through an app distribution platform (a mini program, an app listing), the review step verifies the marking materials, so prepare this in advance.

**Example**: The Measures for the Labeling of AI-Generated and Synthesized Content set out explicit-labeling requirements separately for text, audio, images, videos, and virtual scenes, and state clearly that "when functions such as download, copy, and export are provided, it shall be ensured that the file contains an explicit label meeting the requirements"; Article 7 also provides that internet application distribution platforms must verify the labeling materials during listing review. The Measures took effect on September 1, 2025 [official].

**Pitfalls**: Thinking that only face-swaps and digital humans need labeling. The rule covers "text, images, audio, video, and virtual scenes generated or synthesized using AI technology" — **the creative images you use in delivery count too**.

## Method 6: Do not misstate yourself in your ads

**What to do**: Every claim must have an internally verifiable source, and factor the cost of platform review in advance.

**How**

1. For numbers and performance claims in ads, find the source internally before writing; if you cannot find a source, delete it.
2. Run creative through the platform's pre-review and rules first; do not wait until it is rejected after going live to change it. Every platform has a creative review step, with checks as fine-grained as the proportion of text on the outer image and whether a celebrity or artist is used.
3. Keep "money spent on delivery" and "money paid for violations" in separate books. The cost of violations should not be carried by the delivery budget, or you will take risks you should not take just to "spend the budget down".
4. Keep a record when a platform rejects you: which creative, on what grounds, and how it was changed. The same type of rejection recurring means there is a systemic flaw in your creative workflow.

**Example**: Tencent Ads' official materials list the "ad creative pre-review tool" as a capability that can detect violations in images and copy, with checks including the proportion of text on the outer image and celebrity or artist use [official]. Platform bans now reach "AI search optimization" itself: the March 2026 round of Tencent Ads review rules added a prohibition — no form of GEO-related promotion or search-guidance behaviour is accepted [third-party verified: rules summary].

**Pitfalls**: Treating advertising as a free creative zone nobody polices. There is review on the platform side and labeling requirements on the regulatory side, and both are tightening.

---

## Notes for this chapter

- This chapter **cites no performance multipliers**. The multipliers that can be found in the materials (for example, "automated bidding raises spend 500% over manual bidding", "ROI exceeds the channel average by 30%") all come from vendor-selected case materials, can only be tagged [vendor claim], and cannot serve as budget assumptions — the treatment here is consistent with Chapter 2.
- Conversely, the **operational rules of thumb** in vendor materials (the learning-period mechanism, "ads ≈ daily spend/100, creative assets ≈ daily spend/1000") can be referenced, but are likewise tagged [vendor claim] — do not treat them as industry standards.
- The official entries for the three Chinese platforms (Ocean Engine, Tencent Ads, Baidu Marketing) currently have only one or two help-center pages each; the mechanisms make sense, but **there is no item-by-item source text for "how to configure a specific product"** — this is the clearest gap in this chapter.
- The original texts of China's Advertising Law and the Measures for the Administration of Internet Advertising were not obtained this round (only the Measures for the Labeling of AI-Generated and Synthesized Content), so this chapter's statements on ad identifiability reach only the platform-rule level, not the level of legal provisions.
- Method 2 cites the Google Ads API documentation, which is **a description of product mechanics, not a performance promise** — it says how the signal is used, not that using it will be cheaper.

## Sources for this chapter

[1] Baidu Marketing | How to use OCPC smart bidding | https://tg.baidu.com/m/wiki/feed/11630.html | [official] [2] Baidu Marketing | What is OCPC | https://pd.baidu.com/Knowledge/20568.html | [official] [3] Juliangxue | [Automated delivery] product manual | https://school.oceanengine.com/product_help/content/668400000007/115216 | [official] [4] Ocean Engine | Introduction to new scaling products and outstanding cases (PDF, includes the learning-period definition and delivery-volume rules of thumb) | http://vos.vipstatic.com/school/wp-content/uploads/2021/08/11183442/ | [vendor claim] [5] Tencent Ads | Help Center · Automatic placements | https://e.qq.com/ads/helpcenter/detail?cid=3160&pid=9376 | [official]

[6] Google Ads API | Asset Group Signals | https://developers.google.com/google-ads/api/performance-max/asset-group-signals | [official] [7] Google Ads Help | About Ad Strength | https://support.google.com/google-ads/answer/9142254?hl=en | [official] [8] Google Ads Help | Creative and asset requirements | https://support.google.com/google-ads/answer/6167118?hl=en | [official] [9] Cyberspace Administration of China and three other departments | Measures for the Labeling of AI-Generated and Synthesized Content | issued 2025-03-14, effective 2025-09-01 | https://www.cac.gov.cn/2025-03/14/c_1743654684782215.htm | [official] [10] People's Daily Online | Notice on issuing the Measures for the Labeling of AI-Generated and Synthesized Content (corroborating the same text) | http://politics.people.com.cn/n1/2025/0314/c1001-40439288.html | [official] [11] Alibaba Cloud | Ocean Engine performance analysis (Quick Audience) | https://help.aliyun.com/zh/document_detail/170392.html | [official]

[12] Google Ads Help | Update to goal-based bidding (effective 2026-08-17; official example of a $10 target CPA against a $5 actual CPA) | https://support.google.com/google-ads/answer/17061251 | [official] [13] OpenAI | Testing ads in ChatGPT (U.S. start on 2026-02-09; expansion to the United Kingdom, Mexico, Brazil, Japan and South Korea announced 2026-05-07; ads do not change answers) | https://openai.com/index/testing-ads-in-chatgpt | [official] [14] Search Engine Journal | Google Ads introduces journey-aware bidding and new budget pacing updates (Journey-aware Bidding beta; Smart Bidding Exploration extended to PMax/Shopping, official claim of +27% unique conversions) | https://www.searchenginejournal.com/google-ads-introduces-journey-aware-bidding-and-new-budget-pacing-updates/574141/ | [third-party verified] [15] Juxuan Network (Chinese: 巨宣网络) | Tencent Ads review rules update, March 2026 (new prohibition: any form of GEO-related promotion or search guidance) | https://m.teng-ad.com/zixun/4527.html | [third-party verified: rules summary] [16] Google Ads official blog | We're upgrading Dynamic Search Ads to AI Max (2026-04-15 announcement, AI Max out of beta; full feature suite averages 7% more conversions/conversion value than search-term matching alone; from 2026-09 DSA/ACA/campaign-level broad match auto-upgrade; the 2026-06-11 update pushes the DSA sunset and auto-upgrade to 2027-02) | 2026-04-15 | https://blog.google/products/ads-commerce/dsa-upgrade-to-ai-max-2026 | [official]

[17] Sina Finance (relaying China.com and the CPPCC news site) | Baidu Marketing launches Qingduo 3.0 and an upgraded smart delivery product (2026-09-09 event; the Qingduo agent × AIMax × AI delivery assistant pipeline) | 2026-09-10 | https://

finance.sina.com.cn/tech/roll/2026-09-10/doc-iniriaew8653276.shtml | [third-party report of an official launch event] [18] 8frame | 7 AI Advertising Case Studies with Real Numbers (Klarna: 1,000+ AI visuals, six weeks to seven days, about $10 million saved annually; Monks × Hatch: cost per purchase −31%, click-through +80%, production cost −97%; Novo Nordisk: open rate +24%; Lidl France: 1.7 million visuals in three weeks) | https://www.8frame.co/blog/ai-advertising-case-studies-2026 | [third-party verified: case compilation]

[19] Digital Applied | Meta AI Automated Ads 2026 Marketing Guide (Advantage+ averages +22% ROAS versus manually managed campaigns; more than 4 million advertisers use its generative AI tools, up from 1 million six months earlier; Advantage+ annualised revenue about USD 60 billion; testing fully automated campaigns that take only a product URL and a budget) | 2026-01-30 | https://www.digitalapplied.com/blog/meta-ai-automated-ads-2026-marketing-guide | [vendor claim] [20] Coupler.io (citing Meta's Q2 2026 earnings call and Motion analysis) | Facebook Ads Statistics (Advantage+ at a $75 billion annual revenue run rate; more than 9 million small businesses using generative AI creative tools; in Motion's analysis of $1.3 billion in spend only about 5% of ads became winners and about half received almost no spend) | 2026-08 | https://blog.coupler.io/facebook-ads-statistics | [third-party verified] [21] Google Ads Developer Blog | Upcoming Changes to Image Extraction in Performance Max (from 2026-12-04 newly created campaigns default to OPTED_IN; the new default no longer inherits the account-level Dynamic Image Extensions setting; disabling requires an explicit OPTED_OUT in the request) | 2026-10 | https://ads-developers.googleblog.com/2026/10/upcoming-changes-to-image-extraction-in.html | [official] [22] Search Engine Land | Google Ads Adds New Image and Video Sourcing Options to Performance Max (landing page images and sourced videos added in October 2026) | 2026-10 | https://searchengineland.com/google-ads-adds-new-image-and-video-sourcing-options-to-performance-max-498204 | [third-party verified]

# Chapter 10: Acquiring Customers on Other People's Platforms: App Stores, In-Platform E-Commerce, and Industry Directories

When a customer wanted to buy something, it used to be that you pulled him in from outside; now he first searches in the app store, browses inside the e-commerce platform, and filters in the industry directory — and whichever position you hold in that venue is roughly the position at which he meets you.

For many businesses, customers are pushed out to you inside the platform rather than searched out. So this chapter gives you 5 methods to occupy your position inside the platform.

---

## Method 1: Read the platform's official rules before you list

**What to do**: Before you touch a title, change an image, or sign up for a campaign, confirm which actions will get your product or app taken down.

**How**

1. If you list an app, read Apple's App Review Guidelines first. That document is written by the platform itself and says which apps are refused and under what circumstances an app already listed gets removed. Copy out the clauses relevant to your category and keep them in a separate file — do not count on looking them up on the day you submit [official].
2. If you sell on e-commerce platforms, bookmark the Taobao Rules Center (Chinese: 淘宝规则中心) and the JD Merchant Help Center (Chinese: 京东商家帮助中心) and use them as your entry points for checking rules. These two are the general entry points to the rules; the individual clauses must be located item by item against the action you are about to take — changing a price, joining a campaign, writing promotional copy — and you must not treat someone else's screenshot or the wording forwarded in a group chat as the basis [official].

3. Book a takedown as a customer-acquisition cost. An app removed, a product removed, a store docked points — each equals losing your gateway on that platform the same day. Losing the gateway is far worse than a drop in ranking: when your ranking falls, someone scrolling further down can still find you; when the gateway is gone, nobody can reach you at all [our judgment].
4. Turn every line you have crossed into a checklist: the left column says "action," the right column says "what the rule says." Run through it item by item before every new listing, every price change, and every campaign sign-up, and submit only once it passes. The checklist beats keeping it in your head, because it is still there after the person is gone.
5. Name one person responsible for this, and put the responsibility on a name rather than writing "everyone checks." The clauses are scattered across different pages, and with nobody responsible it gets dragged out until the day something breaks, when someone opens it for the first time.
6. Check quarterly whether the rules have been updated, and revise the checklist while you are at it. Platforms change rules without notice, and after a change the answer you read last time may no longer hold.

**Example**: In 2020 Fortnite was removed from the App Store by Apple for routing around in-app purchase; Apple's public statement was that it violated App Store rules applied "equally to every developer" [third-party reported: citing Apple's statement]; after Apple enforced the new game-license rules for the China region in July 2020, monitoring by Qimai Data (Chinese: 七麦数据) recorded that on August 1 alone, within half a day, 29,800 apps were removed, 26,568 of them games [third-party reported: citing third-party monitoring]. Store rules have been rewritten too: on September 14, 2026 Apple updated the App Store Review Guidelines and reworked section 4.3 on "spam and duplicate content," naming over-saturated categories such as flashlights, sound effects, static wallpapers, basic timers and Chinese-zodiac fortune-telling apps — new submissions in those categories are rejected unless they can demonstrate unique functionality, an automated sweep of existing listings removes products with no recent updates, no active interaction and weak ratings from the public shelf, and teams that repeatedly submit reskinned duplicate apps face cross-account enforcement [third-party verified: relaying the platform's revised guidelines]. The tightening is not a single notch: Apple's updated 4.3(b) states outright that apps indistinguishable from widely available products must not be submitted, that opportunistically producing variants of existing categories or popular apps lowers App

Store discovery efficiency and can lead to removal, and that repeated submissions may get a developer removed from the program [third-party verified]; in the same revision 4.1(c) bars using another party's icon, brand or product name in an app icon or name without permission, and 5.1.2(i) requires clear disclosure of where personal data is shared with third parties — explicitly including third-party AI — with the user's explicit permission obtained first [official: review guidelines]. Platform subsidies can be calculated straight into a cost model: the official FAQ for WeChat Stores spells out the four conditions for the 2026 1% technical service fee incentive for new merchants — within 90 days of the first ordinary order being paid, and no later than 180 days after opening the store; the first RMB 1.5 million of self-operated transaction volume in priority categories and the first RMB 1 million in others — with the platform assessing eligibility automatically at settlement and no application required [official].

**Pitfalls**: Treating the rules page as a first-aid kit you open only when something has gone wrong. Reading it once before listing is cheap; being taken down and starting over costs far more.

## Method 2: Fill in every field that can be searched

**What to do**: Fill in every field the platform uses for search — title, subtitle, keywords, tags, category — without missing one, and fill each accurately.

**How**

1. Confirm from the official documentation first which fields take part in search. Apple gives "App Store search" a page of its own that sets out the mechanism of in-store search and the keyword field separately [official]. The field list follows the official documentation; a table put together by a third party is only a lead.
2. Lay the fields out in a table: the left column the field name, the right column what you currently put there. Fill in the blanks and turn anything vague into something specific. This table lasts a long time once it is built, and it doubles as your master copy inside the platform.
3. Do not pile the same word into several fields. A word repeated over and over occupies the place of other words; you are narrowing your own coverage.

4. The title and subtitle answer "what I sell"; the keyword field answers "how customers will search for me." These are two different jobs, and writing them as one makes both inaccurate: stuff keywords into the title and customers cannot read it; write the keywords as a product slogan and the platform cannot match them.
5. Once it is filled in, search again from the customer's side: use the few words a customer would type and see whether your own product is still in the results. Once a month is enough; record each result in the same table and watch the change — a single result proves nothing.
6. One set of fields per platform; do not copy the same copy across every platform. Each platform searches different positions and lets you fill in different things.

**Example**: Apple gives "App Store search" a page of its own to explain in-store search and the keyword field; this is the most direct primary source for "fields take part in search" [official]. The official Chinese-side numbers make the size of this channel concrete: Douyin's official e-commerce learning center says 80% of monthly active users run searches, searching about 100 times a month on average, with 140 million e-commerce searchers per day and close to 500 million search page views per day, and that search converts 90% more efficiently than other in-feed scenarios [official: platform-reported]. The same shift appears in app stores: Google Play version 49.3 put Ask Play onto the app detail page — a Gemini-driven chatbot that answers using app metadata and user reviews, so the page now has to be written for people and for AI systems at once [third-party verified]. Field rules can also be rewritten unilaterally, and China is subtracting and adding at the same time: Douyin e-commerce's 2026 "nine merchant support policies" say the platform saved merchants over RMB 32 billion in the past year through product-card and category commission waivers, cheaper shipping insurance and a small-merchant support fund; orders generated through Qianchuan Chengfang have their technical service fee cut to 0.6%, with waivers extended to all categories and from shelf product cards into content scenarios; dynamic AIGC creative capability is claimed to generate 1–2 million assets a day; and by the end of 2025 the seeding incentive programme had drawn 3.67 million creators, with daily GMV from "search and pay after watching a video" above RMB 100 million [vendor claim: platform policy notice]. The weighting is shifting on the app-store side too: a review built on thousands of apps sorts the 2026 App Store and Google Play ranking factors into 16 items — 11 written into your own store listing (metadata) and 5 produced by user behaviour — with both stores moving weight from

pre-install signals (keywords, ratings, raw downloads) toward post-install signals (retention, conversion rate, semantic relevance); about 65% of downloads still come through search, and in March 2026 Apple inserted paid placements between organic results, narrowing the organic window further [third-party verified: vendor review].

**Pitfalls**: Thinking that pretty content gets you promoted. The platform has to be able to take you in and retrieve you before ranking enters the picture at all; with the fields incomplete, none of the rest happens.

## Method 3: On the Chinese side, treat mini programs and WeChat Search as separate gateways and do each on its own

**What to do**: Inside WeChat, "the mini program being found" and "the content being indexed by WeChat Search" are two different things; do each of them separately.

**How**

1. On the mini-program side, start with the official Mini Program Search Optimization Guide; it addresses exactly how a mini program is found and how it is indexed [official]. Work through the basics according to that guide before thinking about other moves; with the basics not solid, the effort spent later is easily wasted.
2. Getting into WeChat Search is another track. The Content Integration page in the WeChat Open Documentation explains how content is integrated into WeChat Search [official]. Content that is not integrated has no share in WeChat Search, and a well-built mini program cannot make up for it.
3. The mini program's name, description, and page titles are positions to fill completely and accurately as the guide requires — do not leave them blank and do not string them together out of a pile of symbols. These are the positions where people inside WeChat recognize you.
4. Treat "can I be found inside WeChat" as its own check item, test it on its own, and test it again after a while. Do not look at it together with the results from Baidu (Chinese: 百度) or an app store: the three places do not share rules, and experience that works in one often does not hold in another.

5. There is more than one path to reach you inside WeChat: the mini program, official-account content, and WeChat Search. The gateways differ, so record them separately — that way, when something breaks, you know which path it was.
6. On the Chinese side, do not roll everything out at once. Pick one gateway and get it to the point of being findable and indexed before opening a second, or none of them will be solid.

**Example**: Independent gateways on the Chinese government-services side already have a working example: after launch, "Yue Sheng Shi" (Chinese: 粤省事) accumulated 17.67 million verified real-name users and supported 300 million visits; according to Digital Guangdong's (Chinese: 数字广东) official website, as of November 2024 real-name registered users had reached 192 million, the app had 52.13 million cumulative installs, and the official account had 38.84 million followers [third-party reported / official: platform-reported]. The official account of the store side is more complete: Tencent Marketing's FAQ lists WeChat Store's five core operating scenarios as Channels live and short video, article product cards in official and service accounts, affiliate creator selling, gifting, and WeChat Search; its method is to map the three path legs — where this person first meets you, where they order, and where you can still find them afterwards — and it states that advertising is an amplifier, not an engine [official]. The public numbers inside the WeChat ecosystem are shifting too: a vendor's reading of platform data says WeChat Stores are now connected across more than ten ecosystem entry points, that an official account article can carry up to 50 product links, that transaction volume in the WeChat search scenario grew 360% between February and December 2025, and that the direct-purchase link in the Souyisou super brand zone lifted brand ad click-through by 36% (allowlist only); the same material reports that brand-driven GMV in 2025 grew 4.3 times faster than the platform overall and that the average number of active merchants per month rose 1.7 times [third-party verified: vendor relaying platform data].

**Pitfalls**: Thinking that once the mini program is built it will naturally be found inside WeChat. Being findable comes after indexing, and indexing has an official method of its own.

## Method 4: Separate the positions that are bought from the ones that are pushed

**What to do**: Before you act, get clear on the position you want to occupy: is it an ad slot bought with money, or a natural position the algorithm pushes out?

**How**

1. List the positions you want one by one, and tag each with an origin: bought (paid slots, auction slots, campaign placements) or pushed (search results, recommendation feeds, rankings). When you tag, write down plainly who you see occupying that position — do not go by impression.
2. Bought positions are accounted for as delivery: how much budget, what unit price, how long it lasts, and what is left after you stop. This kind disappears when you stop and is essentially the same thing as advertising, so it belongs in the delivery books.
3. Pushed positions are accounted for by rules: what the platform's search and recommendation rules allow and forbid. This kind gives you no chance to bid; all you can do is satisfy every condition the rules require — fields complete, content submitted in the right format, category assigned correctly.
4. Do not use the same set of actions against both kinds. Do not go rewrite copy for a problem money can solve; and for a problem the rules decide, money will not necessarily move it. When you truly cannot tell them apart, ask one question first: "who gave this position up to whom?"
5. Which kind to do first once they are separated is decided by your cash flow and your patience; do not start both at once. If cash flow is tight, do the pushed kind first — it is cheap but slow to show results; if you have room and cannot wait that long, use bought positions to run tests [our judgment].
6. Update the position table once a month. Positions on a platform move: a slot that cost money last month may become one you earn through the rules this month, and the reverse as well.

**Example**: Lalamove's (Chinese: 货拉拉) books keep "volume that was bought" and "what is left" clearly apart: its sales-expense ratio fell from 79.7% in 2021 to 10.2% in 2024, and its operating-expense ratio fell from 122.2% to 34.1% over the same period — once the part that traded subsidies for volume left the stage, what held it up was

the fact that "the truck with the decal had become practically synonymous with same-city freight" [third-party verified: citing prospectus]. A platform can also redefine what counts as a "pushed position": in 2026 Xiaohongshu consolidated its in-app AI search into one product — the search-enhancement tool "Wen Yi Wen" (Chinese: 问一问) was taken offline and merged into the conversational "Diandian" (Chinese: 点点), now fully rolled out, so a question typed into the old search box enters the chat history; for a brand that had already laid out keywords against the results page, this means the answer layer was swapped for a different rulebook and has to be recalibrated against conversational answers [third-party verified: media exclusive].

**Pitfalls**: Treating a "pushed position" as one "money can buy", then finding after the money is gone that the position has not moved at all.

## Method 5: Treat your position in the platform as rented, not as an asset

**What to do**: Assume by default that your ranking, your slots, and your indexing are borrowed, and plan as if they could be taken back at any moment.

**How**

1. Do not write platform ranking into the brand-assets column. Ranking is given by the platform and changes every time the platform changes a rule; book it as an owned asset and your decisions will overrate your own foundation.
2. Every plan whose results rest on a platform position gets one sentence added to it: "if this position is gone, which path do we take next?" Writing that sentence forces you to keep an option elsewhere instead of pressing all your strength onto one gateway.
3. Assign someone to watch the rules. The rules pages of the Taobao Rules Center and the JD Merchant Help Center get updates, and app review rules change too, and when a rule changes the positions can change with it [official]. The watching itself is simple: go in on a fixed date and scan, note down any update, and attach which section was changed.
4. With the same batch of customers, keep their contact details and your list of past customers in your own hands as far as possible. Treat your position in the platform only as the stretch of road that leads them out; once the road ends, the person has to land somewhere you can still reach them.

5. Keep at least one position that does not depend entirely on the platform. It may not perform best, but it is the reason you can still open the door on the day the platform changes its rules.
6. Run the platform-dependency question as a check of its own; the method sits in Chapter 21 and should not be improvised here. The check is to list every gateway and mark clearly who has the final say over each one.

**Example**: In Amazon's 2021 crackdown, the store of Mpow, the main brand of Patoson (Chinese: 帕拓逊), was banned, 606 best-selling products were taken down, and a large amount of funds was frozen; Xinghui (Chinese: 星徽股份) announced that some stores of three of its brands — RAVPower, Taotronics, and VAVA — had been suspended from selling; Caixin (Chinese: 财新) reported that the sweep involved several leading South China sellers including Patoson, Aukey (Chinese: 傲基), and Sunvalley (Chinese: 泽宝) [third-party reported]. Zynga, for its part, wrote in a filing it submitted to the SEC that a deterioration in its relationship with Facebook would hurt its business [third-party reported: citing company filing]. This channel has been falsified once already: in July 2026 Doubao, Qwen and Yuanbao all took down their user-built agent marketplaces, and the post-mortem gives the contrast — Doubao's marketplace accumulated about 8 million user-built agents against roughly 200 million daily actives, but cost RMB 130–240 million a day while earning under RMB 1 million a day in revenue; of the roughly 3 million GPTs created, only about 159,000 are public and active, and the median creator earns USD 0–100 a month [third-party verified]. Easy to list, zero distribution, negative unit economics for the individual developer — the most complete demonstration yet that the slot is rented.

**Pitfalls**: Treating a position the platform gave you as an achievement of your own. It is rented, and neither its expiry nor a change of terms comes with advance notice.

---

## Notes for this chapter

- Apple does not disclose its ranking algorithm, so this chapter writes only the boundaries of the rules and the field practices, and no "ranking formula." If someone hands you a ranking formula, ask for the source first [our judgment].

- The Taobao Rules Center and the JD Merchant Help Center are both general entry points to the rules, and the specific clauses must be located item by item by topic. This chapter reaches only the level of the general entry point and cites no specific clause [our judgment].
- The official rule entries for Douyin Mall (Chinese: 抖音商城) and Douyin E-Commerce (Chinese: 抖音电商) have no verifiable clause page at present, so this chapter does not cite their clauses. If you need them, go to the rules entry point yourself and locate them by topic; do not treat a relayed article as the clause [our judgment].
- For industry directories such as G2 and Capterra, no official explanation of the inclusion mechanism could be found, and Capterra's inclusion-rules page could not be verified either. In this chapter these two serve only as leads, not as a basis [our judgment].
- Numbers such as "the share of new customers inside the platform" have no verifiable source and are not written here. That is to say, Method 4 gives a classification framework; there is no platform-published ratio to cite [our judgment].

## Sources for this chapter

[1] Apple | App Store search | https://developer.apple.com/app-store/search | [official] [2] Apple | App Review Guidelines | https://developer.apple.com/app-store/review/guidelines | [official] [3] WeChat Open Documentation (Chinese: 微信开放文档) | Mini Program Search Optimization Guide | https://developers.weixin.qq.com/miniprogram/dev/framework/search/seo.html | [official] [4] WeChat Open Documentation (Chinese: 微信开放文档) | Content Integration (WeChat Search) | https://developers.weixin.qq.com/miniprogram/introduction/widget/we-search/WXAPAGES.html | [official] [5] Taobao Rules Center (Chinese: 淘宝规则中心) | General entry point for in-platform rules | https://rulechannel.taobao.com | [official] [6] JD Merchant Help Center (Chinese: 京东商家帮助中心) | Merchant-side rules | https://helpcenter.jd.com/vender/issue/1000-45858.html | [official] [7] CNBC | Apple kicks Fortnite out of App Store for challenging payment rules | 2020-08-13 | https://www.cnbc.com/2020/08/13/apple-kicks-fortnite-out-of-app-store-for-challenging-payment-rules.html | [third-party reported: citing Apple's statement] [8] Beijing Daily (Chinese: 北京日报) | New game-license rules for the App Store China region and the takedown data | https://xinwen.bjd.com.cn/content/s5f2624fee4b0d7d5fee65d92.html

| [third-party reported: citing Qimai Data] [9] Tencent Marketing Academy (Chinese: 腾讯营销学堂) | WeChat Search and Super Brand Zone customer cases (Winona / Sephora / Omega) | https://eschool.qq.com/Solution/ListDetail/pd-7788 | [official: platform-reported] [10] Alimama (Chinese: 阿里妈妈) | Zhitongche keyword bidding and ad-slot description | https://www.alimama.com/news_detail.htm?contentId=430 | [official] [11] Amazon Ads | Sponsored Brands customer case (HP) | https://advertising.amazon.com/solutions/products/sponsored-brands | [vendor claim] [12] The Beijing News (Chinese: 新京报) | The Amazon ban wave: Patoson, Xinghui (Sunvalley), and Tongtuo Technology suspended from selling | https://m.bjnews.com.cn/detail/162969795414530.html | [third-party reported] [13] CNET | Zynga Shares Soar on Facebook IPO Filing (citing the SEC filing) | https://www.cnet.com/home/smart-home/zynga-shares-soar-on-facebook-ipo | [third-party reported: citing company filing] [14] Secrss (Chinese: 安全内参·腾云) | Yue Sheng Shi: 17.67 million verified users, 300 million visits, 744 high-frequency services | https://www.secrss.com/articles/14012 | [third-party reported] [15] Digital Guangdong official website (Chinese: 数字广东官网) | Yue Sheng Shi: 192 million real-name registered users (as of 2024-11), 52.13 million app installs | https://www.digitalgd.com.cn/szgd/yss/ywly_yss.shtml | [official: platform-reported] [16] Huxiu · Miaotou (Chinese: 虎嗅·妙投) | The Lalamove "value-up" story (citing the prospectus: sales-expense ratio 79.7%→10.2%, operating-expense ratio 122.2%→34.1%) | https://pro.huxiu.com/article/4215577.html | [third-party verified: citing prospectus]

[17] App Store Review Guidelines, section 4.3 revision (2026-09-14: clean-up of low-quality utility apps, cross-account enforcement) | relayed by openinstall | https://www.openinstall.com/article/news/apple-app-store-guidelines-clean-low-quality-apps | [third-party verified: relaying the platform's revised guidelines] [18] Douyin e-commerce official learning center | Compass search analytics upgrade (80% of monthly active users search, about 100 searches a month; 140 million e-commerce searchers a day and close to 500 million search page views a day; search converts 90% more efficiently than other scenarios) | https://school.jinritemai.com/doudian/web/article/aHYhVvzvnFuG | [official: platform-reported] [19] AppTweak | Monthly app store update (Google Play 49.3 introduces Ask Play: Gemini-driven, answering from app metadata and user reviews) | https://www.apptweak.com/zh-cn/aso-blog/app-store-optimization-news-app-store-updates | [third-party verified] [20] Apple Developer | App Review Guidelines update (4.3(b) indistinguishable apps removable; 4.1(c) others' brands in icons and names; 5.1.2(i) third-party sharing including third-

party AI must be disclosed and consented to) | https://developer.apple.com/news/?id=ey6d8onl | [official] [21] 9to5Mac | Apple tightens 4.3(b) against low-effort apps (1,000+ submissions an hour; 90% reviewed within 48 hours; new app submissions up 84%) | 2026-06-09 | https://9to5mac.com/2026/06/09/apple-tightens-app-review-guidelines-against-apps-that-do-not-add-value-to-the-app-store/ | [third-party verified] [22] Sina Finance (Douyin e-commerce release) | 2026 "nine merchant support policies" (over RMB 32 billion in merchant savings; Qianchuan Chengfang orders at 0.6% technical service fee and waivers across all categories; 1–2 million AIGC assets a day; 3.67 million creators in the seeding programme, daily search-after-view GMV above RMB 100 million) | 2026-01-08 | https://finance.sina.com.cn/roll/2026-01-08/doc-inhfqtyp3810727.shtml | [vendor claim: platform policy notice] [23] Tencent Marketing official FAQ | WeChat Store's five core operating scenarios and the three path legs | updated 2026-08-24 | https://e.qq.com/faq/wechat-store/growth/faq-zcjc-003/ | [official] [24] Yage Daily Notes (Chinese: 鸭哥每日手记) | Why agent stores did not work (Doubao's roughly 8 million user-built agents, about 200 million daily actives, under RMB 1 million revenue a day; of about 3 million GPTs only about 159,000 are public and active, median creator income USD 0–100 a month) | 2026-07-04 | https://yage.ai/share/uga-sunset-20260704.html | [third-party verified]

[25] Tencent Ads Help Center | What are the conditions for the 2026 WeChat Stores new-merchant technical service fee benefit (four conditions for the 1% incentive; within 90 days of the first order and no later than 180 days after opening; first RMB 1.5 million in priority categories, RMB 1 million in others; assessed automatically at settlement) | updated 2026-08-24 | https://e.qq.com/faq/wechat-store/payment/faq-zfjs-007/ | [official]

[26] Xiaoduo AI (Chinese: 晓多 AI) e-commerce insight | How far have WeChat Stores integrated with official accounts and search (up to 50 product links per official account article; WeChat search transaction volume +360% from 2025-02 to 2025-12; direct-purchase link in the Souyisou super brand zone lifted brand ad click-through by 36%, allowlist only; brand GMV growing 4.3 times the platform average, monthly active merchants +1.7 times) | 2026-06-25 | https://insight.xiaoduoai.com/e-commerce-information/wechat-information/to-what-extent-have-wechat-mini-stores-been-integrated-with-official-accounts-and-the-search-function-how-much-certain-growth-can-the-integration-of-content-e-commerce-and-search-e-commerce-bring-to-m.html | [third-party verified: vendor relaying platform data]

[27] App Radar | 2026 App Store & Google Play ranking factors (16 ranking factors: 11 from store-listing metadata, 5 from user behaviour; weight shifting from pre-install to post-install signals such as retention, conversion rate and semantic relevance; about 65% of downloads come from search; Apple inserted paid slots among organic results in 2026-03) | 2026 | https://appradar.com/academy/app-store-ranking-factors | [third-party verified: vendor compilation] [28] Sina Finance (Leiphone exclusive) | Xiaohongshu retires "Wen Yi Wen" as in-app AI search is consolidated into "Diandian" (the search-result-enhancement product "Wen Yi Wen" merged with the conversational "Diandian" and rolled out to all users; questions are saved in chat history; the Dots organisation scaled back to foundation-model research) | 2026-09-03 | https://finance.sina.com.cn/stock/t/2026-09-03/doc-iniqpkry1106989.shtml | [third-party verified: media exclusive]

# Chapter 11: Letting Others Sell for You: Channels, Resellers, and Ecosystem Directories

In the old playbook, channels were built by people on the road: find resellers, negotiate distribution, work the trade shows exchanging business cards, grinding out one partner at a time. Now there is an additional class of ready-made positions — a platform's cloud marketplace, app directory, service-provider listing — that customers go into and browse on their own, and whether your thing is on the shelf decides whether a partner can sell you along the way.

Customers are often brought in by partners. So this chapter gives you 5 methods to treat channels and the service-provider directories of AI platforms as acquisition positions.

---

## Method 1: Treat partner directories as acquisition positions, not as sales channels

**What to do**: Establish first whether your product has a page in the platform directory that can be found by search, and only then talk about whether a partner will sell for you.

**How**

1. Search the cloud marketplace and app marketplace using the customer's own words: search the category term, your product name, your competitor's product name. Record three things — whether you are in the results, what position you are in, and who is ahead of you.
2. Separate two things: **the page in the directory** is your gateway, and **the partner's selling actions** are the channel. With no gateway built, a channel however wide is useless. Many companies sign a reseller first and only then discover that customers cannot find them in the platform at all.
3. Treat onboarding as "putting up a page": the page has to be written and maintained by you. A resale contract can wait once it is signed, but a page left unattended will produce nothing at all.

4. Once the page is built, give the address to your existing partners. When they make a quote or send a proposal, they attach the product-page link directly; the customer who clicks through sees the platform's domain and transaction flow, and it gets forwarded more readily than a PDF you made yourself [our judgment].
5. Search again quarterly and record the position in the same table. Directory categories and ordering get adjusted, and searching once a year shows you no change.

**Example**: Shopify's app store is an acquisition position rather than a sales channel — more than 10,000 apps, 561 million USD paid to partners each year, and 85% of merchants depend on at least one app to run their business [third-party verified]. Here partners are traffic, not channels. Cloud marketplaces are the larger version of the same logic: by an industry compilation, about 40% of SaaS revenue will route through AWS, Azure and GCP procurement by 2027; Snowflake crossed $7 billion in lifetime AWS Marketplace sales in May 2026, with more than $2 billion in calendar-year transactions, doubling year over year, and CrowdStrike booked over $1 billion of AWS Marketplace sales in a single calendar year; Canalys projects hyperscaler marketplace transactions will reach $85 billion by 2028, with Omdia pointing to $163 billion by 2030 [third-party verified]. The cloud-marketplace path is getting faster too: AWS launched express private offers in June 2026, letting partners configure pricing rules, discount boundaries and eligible products once so that AWS sellers can send customers a personalised private offer directly inside the co-sell workflow — officially described as moving a deal from weeks of manual negotiation to minutes, with partners receiving the customer's contact details [official]. The shape of the catalogues themselves deserves a look. A review of three incompatible distribution logics finds OpenAI's GPT Store hosting more than three million custom GPTs by early 2026 with creators paid about $0.03 per conversation, meaning more than 33,000 quality interactions a month to clear $1,000, Anthropic launching the Claude Marketplace in March 2026 with zero commission and just six anchor partners, and GitHub's MCP Registry indexing 923+ servers as a discovery layer rather than a transaction layer [third-party verified]. Catalogue scale gives another set of marks: a May 2026 inventory counts more than 11,000 models and over 4,000 AI apps and agents on Microsoft Marketplace, more than 10,000 apps on Salesforce's

AppExchange plus 2,600 Slack apps and more than 1,000 Agentforce agents, and more than 1.2 million skills in the open Agent Skills standard released in December 2025 [third-party verified].

**Pitfalls**: Thinking that onboarding means someone will sell for you. The platform gives you only a shelf slot and a backend; after listing, nobody maintains the page for you and nobody answers customers for you.

## Method 2: Read the commission and liability clauses verbatim before you sign

**What to do**: Review the platform's onboarding agreement as a contract about to be signed, focusing on four things — who sets the price, who charges, who bears liability, and how to exit.

**How**

1. Start with the money clauses: who sets the platform service fee, how it is charged, and whether it can be changed. Do not get one line from your account contact on "roughly what the cut is" and move on — that sentence does not amend the agreement.
2. Then look at pricing authority: can you set the product price yourself, and has the platform reserved the power of "listed only after approval"?
3. Then look at liability: for customer complaints, refunds, and disputes, who takes them first and who ultimately covers them. This clause decides how many people you keep in after-sales.
4. Finally look at exit: how the agreement renews at expiry, whether the platform can terminate unilaterally, and who handles after-sale service for goods already sold once it is terminated.
5. Copy these clauses onto one page and hand it to the boss who pays and signs. Channel terms are business terms and they touch cash flow and liability; they should not be decided by the business department alone.
6. Gather your materials in advance: corporate real-name verification, legal-representative ID, corporate qualification certificates, and any industry license the category requires. The agreement states that for categories requiring a license or approval to operate, the service provider itself is responsible for

obtaining the license and for proactively submitting it to the platform for review [official] — get stuck at this step and everything discussed earlier comes to nothing.

**Example**: The text of Alibaba Cloud's (Chinese: 阿里云) Cloud Marketplace Service Provider Onboarding Agreement states these things: the platform service fee is collected by the operator, the fee standard is decided by the operator and may also be modified; the service provider has the right to define and price the goods in its store independently, but goods must pass a security review before publishing, with the operator reserving the decision on whether to list them; the service provider signs an independent goods agreement with users in its own name, and seller liability is borne independently by the service provider; users may not be steered to transact outside the cloud marketplace without the platform's written permission; one service provider is allowed only one store in the cloud marketplace; and the operator has the right to terminate the agreement on 15 days' written notice [official]. The equivalent price lists overseas are blunter: Salesforce's partner program FAQ states that ISVforce carries a 15% revenue share on application revenue and OEM a 25% share of combined solution revenue, that each paid app carries a one-time $999 security review fee plus roughly $150 in annual listing fees, and that AppExchange now hosts 6,233 apps from 3,668 developers reaching more than 150,000 Salesforce customers [official: revenue share / third-party verified: fees and catalog size]. Microsoft's threshold hangs on cloud consumption: to reach IP Co-Sell Eligible, an organization must generate at least $100,000 of Azure Consumed Revenue or Marketplace Billed Sales over the trailing 12 months, with credits excluded [third-party verified]. AWS has moved its fees the other way: the tiered private-offer listing fee runs from 3% on deals under $1 million down to 1.5% above $10 million, with professional-services private offers cut to 0.5% in June 2026 [third-party verified] — a cloud marketplace trades acquisition cost for platform share plus compliance cost, and the rate tier follows your deal size. Microsoft changed the channel itself: for FY27 (July 1, 2026 to June 30, 2027) it makes the marketplace the primary route for large-scale co-sell, replacing the practice of counting partner-reported Azure consumed revenue (PRACR) as the co-sell mechanism, on the stated grounds of giving partners and sellers a more consistent, predictable and auditable experience; a new Frontier Accelerate for Marketplace, launching in September 2026, folds ISV Success, Marketplace Rewards, Azure IP co-sell and the certified software designation into one experience [official]. For a

company working the channel, the message is that to get the platform's sales push, the product has to be transactable on the marketplace — "we have partnered before" does not count.

Cloud vendors are changing the arithmetic of reseller commissions, too: a reading of Alibaba Cloud's FY2026 channel policy reports the total commission points broadly unchanged, with base commission up to 39%, but two calculation rules changed — the quarterly renewal-rate KPI added the previous year was dropped (previously, hitting the top quarterly tier required a renewal completion rate above 90%, and many smaller resellers running around 80% lost about two points), and the annual rebate plus growth incentive moved from “incremental only” to “full amount”: on a 10 million target and 15 million actual, the old rule paid on the extra 5 million, the new one pays on all 15 million [third-party verified: reseller's reading]. For anyone running a channel, that kind of change matters more than the headline rate — it changes what one extra unit of revenue is worth, and channel teams calibrate their behaviour to that cell.

**Pitfalls**: Negotiating only the rebate and never reading the liability clauses. Then the first customer dispute arrives, and you discover the agreement says the platform bears no liability for the goods — while you had already staffed up on the assumption that "the platform will cover it."

## Method 3: Write the directory product page as a set of search fields

**What to do**: Fill in the product name, category, details, and the service standards you promise entirely according to what customers will search.

**How**

1. Get the category right. The category decides which shelf you sit on; pick the wrong one and you are in someone else's showroom.
2. Write what problem you solve first, then the specifications. A customer browsing the directory arrives with one thing in mind; specs are step two.
3. Read the platform's entry threshold carefully. Some marketplaces require the product to run on their own resources, or to be bundled with their own products, or to serve as a companion tool — fall short of the threshold and listing is not allowed, and at that point no description, however good, helps.

4. Write your delivery and service commitments into an external document and hang it next to the product: delivery method, billing method, support hours, response time. This document is the only evidence a customer can look at before ordering.
5. Platform product pages are often distributed to affiliated sites. Write on the assumption that the reader cannot see your own website, so that the key information closes the loop within the page itself.

**Example**: The Alibaba Cloud Marketplace agreement sets a hard listing threshold — the product must run on Alibaba Cloud resources, or be sold bundled with Alibaba Cloud products, or be a companion tool used closely with Alibaba Cloud products; otherwise listing is not permitted [official]. The same agreement also authorizes the platform to list a service provider's goods on affiliated platforms such as Bailian (Chinese: 百炼) and Qwen Cloud (Chinese: 千问云) [official]. The agreement also sets service standards the goods must meet: non-onsite technical support from 9:00 to 18:00 every business day, a solution or alternative method within 24 hours when the software cannot run and the user cannot use it, and a solution within 48 hours when a performance decline makes it difficult to use [official].

**Pitfalls**: Writing the product page as an internal product manual. Thirty features are listed and not one of the questions the customer cares most about is answered.

## Method 4: Give partners a kit they can resell from directly

**What to do**: Build a partner kit so that a salesperson who has never seen your product can open their mouth to quote and answer objections after reading it.

**How**

1. A one-page resale pitch: who the customer is, what it solves, and where it differs from similar products. Do not write words like "empower" or "closed loop" — a partner reading them aloud to a customer will stumble.
2. Keep only two tiers in the price sheet: one public price and one partner price, with the validity period and whether it can be changed stated clearly. With too many tiers, the partner hesitates in front of the customer.
3. Three citable cases, each stating clearly the customer type, what was used, and how the result was calculated.

4. A one-page objection response: put in the questions customers ask most often along with the answers, including "what are your pitfalls?"
5. A one-page red list: what may not be said, what results may not be promised, and whether AI-generated assets need labeling. This page draws the line for partners, and the liability when something goes wrong lands on you.
6. Update once a year and note the change in the footer. Outdated material in a partner's hands is more trouble than no material in yours.

**Example**: The Alibaba Cloud Marketplace agreement requires service providers to give users necessary after-sales service and technical support, and turns support hours, fault-response times, and refund conditions into hard items — non-onsite support (phone, email, online) from 9:00 to 18:00 every business day, a solution within 24 hours when the software cannot run and the user cannot use it, a solution within 48 hours when a performance decline makes it difficult to use, and a full refund before the transaction completes if the user is not satisfied (except for tool-type goods) [official]. These items can be copied directly into the service-commitment page of a partner kit: what the platform requires you to do is exactly what the customer will eventually come back and hold you to.

**Pitfalls**: Handing over one product deck and nothing else. The partner still does not know what price to quote, and cannot field "can you do better on price?"

## Method 5: Never let one channel stand alone; run a dependency check every quarter

**What to do**: Work out the contribution of each directory and each reseller, so you can see plainly which channel has grown too large for you to withdraw from.

**How**

1. Build a channel table with fixed fields: channel name, type (platform directory / reseller / distributor), go-live date, leads this quarter, closed deals, platform service fee or discount, and share of leads.
2. Look at the directory type and the reseller type separately. The directory type is self-service — you owe no favors but must maintain the page; the reseller type needs a person following it, and when the person changes, the channel snaps.

3. Set your own ceiling on any single channel's share of leads. By this book's line, once any one channel is long-term eating more than half the lead volume, you should be preparing a second [our judgment].
4. Review the platform's rules page quarterly: whether the agreement has changed, whether the fee rate has moved, and whether there are new removal conditions. Once the agreement changes, it runs by the agreement, not by your say-so.
5. Set different targets for the two types. The directory type is judged on page impressions and self-service conversions; the reseller type on quote count and sales-cycle length. Compare them on the same measure and you will reach the wrong conclusion that "resellers are useless" or "directories are useless."
6. For the things that can be moved — lists, quotes, after-sales records — keep a copy in your own system. The data in the backend is not in your hands.

**Example**: The Alibaba Cloud Marketplace onboarding agreement states that platform rules take effect automatically on being posted on the website and become part of the agreement; the operator has the right to terminate on 15 days' written notice; and the agreement terminates when it expires without renewal, when the service provider triggers removal conditions, and in other such situations [official]. A channel position is rented under an agreement, and both the term and the rent are set by the platform.

**Pitfalls**: Writing "we got into a big vendor's marketplace" into your fundraising materials as a moat. Change the rules, terminate the agreement, and the page and the leads disappear together.

---

## Notes for this chapter

- "How many new customers a partner brings in" — this kind of channel-contribution ratio is a widely circulated industry saying, and this round found no original publisher. This chapter writes no channel-contribution ratio, and it advises readers to find the source themselves before citing one.
- Those two OpenAI official announcements ("Introducing apps in ChatGPT" and "Developers can now submit apps to ChatGPT") are direct material for "AI-platform app directories are turning into gateways," but this round's fetch

returned 403 and the body text was not verified. This chapter therefore writes none of the content of those two announcements, only noting that they exist and that the direction points this way.

- This chapter has only two verifiably reachable sources, both of them Alibaba Cloud Marketplace's own rules pages. This is the platform's one-sided account; it does not represent other platforms or a common industry standard, and no conclusion can be drawn from it that "meeting these conditions gets you into any marketplace."
- The official listing requirements for Salesforce AppExchange, AWS Partner Network, and the Microsoft marketplace, and the official onboarding pages for domestic agent stores (Coze (Chinese: 扣子), Wenxin (Chinese: 文心), Yuanqi (Chinese: 元器), and the DingTalk (Chinese: 钉钉) and Feishu (Chinese: 飞书) app marketplaces), were not captured this round, so this chapter gives no onboarding conditions for these platforms.
- For the industry-standard channel-commission figures and the various "channel contribution share" statements, this round found not one original publisher, and no ratio appears in this chapter. The Alibaba Cloud Marketplace agreement only states that the operator has the right to collect a platform service fee and decide the fee standard; it gives no rate, and the body does not add one.

## Sources for this chapter

[1] Alibaba Cloud Marketplace (Chinese: 阿里云云市场) | Marketplace onboarding process (how a service provider onboards) | https://help.aliyun.com/zh/marketplace/tech-partner-onboarding-process | [official] [2] Alibaba Cloud Marketplace (Chinese: 阿里云云市场) | Cloud Marketplace Service Provider Onboarding Agreement (FY27) | https://help.aliyun.com/zh/marketplace/cloud-market-service-provider-entry-agreement | [official] [3] OpenAI | Introducing apps in ChatGPT | https://openai.com/index/introducing-apps-in-chatgpt | 403 this round, body text not verified, listed only | [listed only — full text not reachable this round, not cited in the body] [4] OpenAI | Developers can now submit apps to ChatGPT | https://openai.com/index/developers-can-now-submit-apps-to-chatgpt | 403 this round, body text not verified, listed only | [listed only — full text not reachable this round, not cited in the body] [5] Partner Insight | What does it take to build a thriving SaaS marketplace (Shopify case:

10,000+ apps, 561 million USD a year paid to partners, 85% of merchants dependent) | https://partnerinsight.io/insights/what-does-it-take-to-build-a-thriving-saas-marketplace-shopify-case-study | [third-party verified]

[6] SaaS Mag | Cloud Marketplaces Are the New Channel (about 40% of SaaS revenue routing through cloud marketplaces by 2027; Snowflake at $7 billion lifetime AWS Marketplace sales and $2 billion in a calendar year; Canalys $85 billion by 2028, Omdia $163 billion by 2030; AWS tiered listing fee 3% → 1.5%, professional-services private offers 0.5%) | 2026-06-26 | https://www.saasmag.com/cloud-marketplaces-new-channel-saas-revenue-2027 | [third-party verified] [7] Salesforce Help | Partner program FAQ (ISVforce 15% / OEM 25%) | https://help.salesforce.com/s/articleView?id=000394757&language=en_US&type=1 | [official] [8] Appnigma | Salesforce ISV & Partner Program Guide 2026 ($999 security review fee, about $150 annual fee, 6,233 apps from 3,668 developers reaching 150,000+ customers) | https://appnigma.ai/blogs/salesforce-isv-partner-program-guide-2026 | [third-party verified] [9] Clazar | Azure Co-Sell vs IP Co-Sell Eligible ($100,000 ACR/MBS threshold over the trailing 12 months) | https://clazar.io/blog/azure-co-sell-vs-ip-co-sell-eligible-guide | [third-party verified]

[10] AWS | AWS Partners can now accelerate co-sell deals with express private offers (partners pre-configure pricing rules and discount boundaries; AWS sellers send personalised private offers inside co-sell) | 2026-06-16 | https://aws.amazon.com/about-aws/whats-new/2026/06/aws-partners-express-private-offers/ | [official] [11] Agent Market Cap | Agent Distribution Wars 2026: GPT Store vs Claude Marketplace vs GitHub MCP (3 million+ custom GPTs at about $0.03 per conversation; Claude Marketplace launched 2026-03 with zero commission and six anchor partners; GitHub MCP Registry with 923+ servers) | 2026-04-07 | https://agentmarketcap.ai/blog/2026/04/07/agent-distribution-wars-openai-gpt-store-anthropic-plugins-github-marketplace | [third-party verified] [12] Presenc AI | AI Agent Marketplaces Landscape (2026-05-14: Microsoft 11,000+ models and 4,000+ AI apps and agents; Salesforce 10,000+ AppExchange apps, 2,600+ Slack apps, 1,000+ Agentforce agents; Agent Skills open standard with 1.2 million+ skills) | 2026-05-14 | https://presenc.ai/research/ai-agent-marketplaces-landscape-2026 | [third-party verified]

[13] Microsoft Partner Center | July 2026 announcement: Azure IP co-sell moves to marketplace-first (FY27, i.e. 2026-07-01 to 2027-06-30; the marketplace becomes the main route for co-selling at scale, replacing the old PRACR mechanism; September 2026 introduces Frontier Accelerate for Marketplace, combining ISV Success,

Marketplace Rewards, Azure IP co-sell and the certified software designation) | 2026-07 | https://learn.microsoft.com/zh-cn/partner-center/announcements/2026-july | [official] [14] Wanyun Information Technology | Reading Alibaba Cloud's FY2026 channel policy (base commission up to 39%; the quarterly renewal-rate KPI added the previous year dropped; annual rebate and growth incentive changed from incremental-only to the full amount) | 2026 | https://www.wanyuna.com/post/1785.html | [third-party verified: reseller's reading]

# Chapter 12: Turning One Piece of Content into Thirty: The AI Content Factory and Matrix Distribution

Content used to be a matter of filing one draft in one place and calling it done, while the deep content you could produce came to only a few pieces a year. Customers will come across you on several platforms — or come across your competitor. So this chapter gives you 5 methods to blanket multiple gateways with one piece of deep content.

---

## Method 1: Break one long piece into a batch of pieces that each stand on their own

**What to do**: Pick a deep topic only you can write, turn it into one long piece first, then cut multiple short pieces out of that one. This step adds no new material; it just uses up the value of the same draft.

**How**

1. Fix the mother draft first, not the platform. Turn one deep interview, one internal data review, or one Q&A that explains a problem thoroughly to a customer into a long article. While writing, ignore where it will end up; decide which platforms can hold it after it is cut.
2. Write the mother draft down the line of "problem — evidence — conclusion." With that three-part structure in place, you can find the beginning and end of each piece at a glance when you cut; a draft written as free-flowing prose can only be cut by feel.
3. Cut by "whether it stands alone," not by length. The piece you cut has to be understandable to a reader who has not seen the mother draft: a frequently asked question, a set of spec comparisons, a pitfall you hit, a passage in the customer's own words.

4. Give each one a question-phrased title that can serve as a headline. The headline is itself a piece of content, and many people read only that line. Write the same block of content in different forms — image-and-text, a spoken-word script, a Q&A, a list — and you have four things, not the same piece copied four times.
5. Cut ten the first time and add more once it runs smoothly. Start by cutting thirty and every one after that is rushed, with quality falling to the point where nobody finishes.
6. Leave each short piece an opening back to the long draft — the profile page, a pinned comment, a note at the end. Then keep a table: what the mother draft is, which pieces were cut, which gateways they went to, and which ones got replies. This table is more useful than view counts, because it tells you who cares about what.

**Example**: After Fan Deng Reading (Chinese: 樊登读书) turned "read 50 books a year with everyone" into a standardized unit, it used offline reading clubs and communities to build up its private domain, then distributed the same set of content through a short-video matrix; the Douyin matrix accounts accumulated about 60 million followers [third-party verified: downgraded].

**Pitfalls**: Treating the same passage said ten times as ten pieces of content. Swap the same meaning into ten phrasings and the reader spots the repetition at a glance; being served it repeatedly only fixes in their mind that you have just this one move.

## Method 2: Clear the platform's AI-labeling gate before you publish

**What to do**: For AI-generated images, videos, voiceovers, and copy, declare their identity according to each platform's rules before publishing. Labeling is not something to remember at the review stage; it is part of making the assets.

**How**

1. Check the disclosure rules platform by platform before publishing; do not apply the habits of "the platform you know best" across the board. The scope each one covers, where the label goes, and whether it can be removed all differ, so recheck quarterly — nobody will notify you when the rules change.

2. Distinguish two degrees of involvement: the whole piece generated by AI, and only part of it altered by AI (face swap, voiceover, background replacement, added generated footage). The disclosure requirements differ.
3. Whether an asset file carries credentials has to be controlled in advance. For video carrying the standard credential (C2PA), the platform reads it out and applies the label on its own, and at that point you are the passive party; if material delivered by a supplier carries credentials, ask what the credentials say before you sign.
4. Commercial-collaboration content goes through the official channel for labeling. For reported collaboration content, the label and the traffic are bound together; whatever little you save by going around the official process has to be paid back later.
5. For a label the platform adds on its own initiative, do not think about removing it — this kind of label is issued by the platform, and creators usually have no option to remove it.
6. Write "labeling" as a step in the production process: when the video is done, run a pre-publish labeling check, fix who does it and which items it covers, and file the check result together with the assets.

**Example**: YouTube's Help Center calls this kind of information "how this content was made," and its "made with AI" label appears in three situations — the creator declared it in the backend workflow, the creator used the platform's own generation tools, or the video carries a valid standard credential (C2PA) showing that the whole thing was AI-made; when content is not declared, the platform sometimes adds the label on its own, and creators have no option to remove it [official]. Douyin (Chinese: 抖音) publishes its labeling requirements for AI-generated content in its Rules Center, which is the official entry point for checking the clauses [official]. Xiaohongshu (Chinese: 小红书) puts the labeling rules for commercial content in the Pugongying Help Center (Chinese: 蒲公英帮助中心); how a note taking a commercial order is labeled and where the label goes are checked there [official].

**Pitfalls**: Thinking that only face swaps and digital humans need labeling. Adding one AI-generated illustration or one AI-voiced narration also counts as content that has been altered; by the time the platform puts the label on for you, what you have lost is the reach your content would have gotten.

## Method 3: Use synthetic users as a sieve, not as a witness

**What to do**: Before publishing, use a batch of AI-played target customers to filter out the headlines and openings that are obviously off. What it gives you is an ordering; the conclusion still has to come from real data.

**How**

1. Lock down the purpose of this step first: it solves one thing only — picking three of ten headlines worth actually testing. What it replaces is the "which one do you think is better" step; it does not touch your customers' real feedback.
2. Write the persona in detail: industry, company size, what he uses now, why he wants to switch. The coarser the persona, the more the answer reads like generic customer-service script; if you cannot write the persona in detail, do not do this step.
3. Have the AI answer one question only; do not have it give a score. Asking "would you stop and click?" is more useful than asking "what score would you give this piece?"
4. Ask the same question again in a different persona. Flag the items where the two answers are wildly apart; that class does not go into the candidate set.
5. Sieve one round in each of three directions: one about cost, one about risk, one about how to put it into practice. After the three rounds, remove the duplicates first and test the rest in small steps with real traffic; the synthetic users' ordering serves only as a queue order, not as a basis for budget.
6. Record where it went wrong: look back afterward against real clicks; whichever class of question it keeps misjudging, stop using it to sieve that class next time.

**Example**: Three papers draw boundaries that are quite consistent. One cross-domain benchmark ran the same test protocol on four models and two kinds of real human response data, and concluded that this substitution produces the same two failures across both domains and all four models — at the individual level no model beats the strongest non-model baseline, and the models systematically exaggerate the determining role of identity on attitudes, weighting "who you are" more heavily than the real data does [third-party verified]. Another used a sample of more than 70,000 respondent–question pairs to test whether "adding a persona to the model" actually helps, and concluded that adding a persona brought no overall alignment improvement and in many cases made things clearly worse, with the error amplified

on a minority of questions and underrepresented subgroups [third-party verified]. A third is a proof of concept on the feasible side, and the authors' own conclusion is that the path works but requires careful calibration and additional distribution checks before it can be called trustworthy [third-party verified]. The benchmark paper frames the tool the same way: it lets a team judge before acting in which situations this kind of evidence can support a decision and in which it cannot [third-party verified]. Evidence for and against is accumulating at the same time. On the supporting side, a consultancy used a multi-agent simulation to replicate a 3,600-person global wealth survey, reporting agreement across several statistical metrics (high correlation, modest error) and completing it in one day where the real survey takes six months [third-party verified: relayed]. On the opposing side, a longtime quantitative UX researcher's published conclusion is that synthetic data has no place in survey research, and another practitioner argues the high correlation in that case most likely comes from the model having already been trained on that data — remembering rather than predicting [third-party verified]. A survey of 150 researchers (May 2026) offers the more usable common line: the definition has converged, but the overwhelming majority stress that these personas are not real users [third-party verified]. So its place is that of a sieve — synthetic samples to generate hypotheses and pre-filter, real users to validate and decide.

**Pitfalls**: Using synthetic users' answers to write "a survey shows XX% of customers will...". The fact that models exaggerate differences between groups is exactly what makes this kind of conclusion go most confidently wrong — and your own table still shows nothing amiss.

## Method 4: Do not move the same material to every platform unchanged

**What to do**: Before the content cut from a mother draft goes into each platform, rewrite the opening, the length, and the ending according to how people on that platform look at things. What you change is the phrasing; not one fact in the mother draft moves.

**How**

1. Look first at how the platform is used: some people arrive by searching, some arrive by scrolling. The ones who arrive by searching read the headline and the body; the ones who arrive by scrolling watch only the first three seconds.

2. Rewrite the opening by platform. For the same point, a search gateway must state in the first sentence what problem it solves, and a feed gateway must throw the conclusion out first.
3. Let length follow the platform; do not turn one draft into five versions that all come out the same length.
4. Tell the same case from a different angle: B2B audiences want to see how the numbers work; consumer audiences want to see what pitfall others hit; industry media want to see what the thing means for the industry.
5. Give each platform its own action to take; do not point them all at the same link. One link that works everywhere means you prepared for the people on none of the platforms.
6. Inside one platform, do not keep reposting the same batch of content either — readers remember longer than you think.

**Example**: Douyin's Rules for Managing Improper Conduct by Matrix Accounts states outright that delinquent matrix accounts are handled by "uniformly and strictly dealing with both main and sub-accounts"; in the first half of 2021 Douyin took down more than 740,000 pieces of infringing content of the copying-and-reposting kind and penalized more than 140,000 accounts, 2,429 of which were permanently banned [official rules / third-party reported]. YouTube, too, has written "repetitious or mass-produced" content into its monetization eligibility policy [official]. Search-side penalties are rising in step: Google shipped two updates in March 2026 back to back (a spam update on March 24–25 and a core update from March 27 to April 8), and one Chinese-community sample puts the damage at 71% of affiliate-marketing sites negatively affected, financial affiliate sites down 40%–55% on average and generic product-comparison sites down 35%–55%, while sites that "answer the need directly page by page" recovered after the update [third-party verified: community sample]. The test is not whether you used AI; it is whether you are mass-producing low-differentiation content.

**Pitfalls**: Reading "matrix distribution" as "post the same piece from thirty accounts." That is not a matrix; it is flooding the screen with the same piece.

## Method 5: Build the content factory on first-hand material

**What to do**: Hook the content factory's intake onto first-hand material that only you have — the customer's own words, your own usage data, the pitfalls you hit yourself. Google's own wording is blunter still: its spam policies define scaled content abuse as generating large amounts of pages to manipulate search rankings rather than to help users, and the policy does not use "AI-generated" as its test — hand-written, scraped, translated and hybrid batches of thin content fall under the same rule [official]. The costs come in three independent accounts: after the March 2026 core update, large sites publishing hundreds to thousands of unedited AI pages lost 50–80% of organic traffic within two weeks [third-party verified]; another compilation says 61% of sites mass-publishing unedited AI content lost 40–90% of organic traffic after the February 2026 core update, with practitioner consensus that once roughly 30% of a site's URLs are judged unhelpful Google applies site-level demotion, and that algorithmic scaled-content penalties usually take three to six months to recover from [third-party verified]; and a June 2026 sample of 1 million pages across 100,000 SERPs (with 150,000 text-rich pages AI-detected) concluded that Google punishes thin content, absence of original value and scaling for rankings — not AI generation as such [third-party verified].

**How**

1. Every time you finish a delivery or take a customer call, spend ten minutes writing down the actual words. This is raw material, and a month later you will not remember it. From the customer's words record only three things: what he does, which step he is stuck on, and what words he uses to describe the problem.
2. Wait until the first-hand material has accumulated enough to explain one thing clearly before you start writing. Scattered material will not assemble into a long draft, and a long draft will not yield thirty pieces.
3. Before cutting, ask one question: does anyone else have this block? Keep the ones nobody has for the mother draft; the ones others have serve only as filler.
4. When you speak with your own data, write the basis out with it: how big the sample, from when, and how it was calculated. With the basis unclear, readers will assume you are making it up.

5. Capacity is something anyone can bring level with anyone else; material cannot be leveled. Chapter 14 deals specifically with how to accumulate this material.

**Example**: Gong uses sales-call data from its own platform to make content [vendor claim]; Strava combines billions of activity records from its own community with a survey of more than 30,000 people to produce its annual trend report [vendor claim]. Neither is "writing articles"; both are holding up data in their hands that nobody else has. Programmatic pages also now have a week-by-week teardown with raw data: one agency published 18 months of data across four client engagements (SaaS, e-commerce and its own site) covering 1,080 deployed programmatic pages — impressions start climbing in weeks 2–3 after deployment and clicks lag by an average of 2.4 weeks; monthly impressions rose from 12,000, 8,200, 5,800 and 0 to 168,000, 188,000, 152,000 and 38,000 respectively, with 84% of the new traffic coming from keywords searched fewer than 200 times a month; it also names three failure modes: city pages with swapped names and no real local data were suppressed within three weeks, and 22 pages aimed at a pattern already held by five-year-old competitor pages produced no top-50 ranking in 12 weeks [third-party verified: first-hand data teardown]. The use of that data is calibration: page count is only the fourth variable, behind per-page data quality, keyword-pattern competitiveness and your own domain's authority [third-party verified: first-hand data teardown].

**Pitfalls**: Thinking that buying a tool gives you a content factory. A tool only solves cutting fast; what to cut still has to come from you. With material all copied off the web, what you cut out will not hold anyone either.

---

## Notes for this chapter

- Of the three platform rules, the YouTube one is a Help Center page read line by line; for Douyin's Rules Center and Xiaohongshu's Pugongying Help Center, this round captured only the official entry points, and the original text of the AI-labeling clauses was not located item by item, so this chapter's statements about these two reach only the level of "where the entry point is and that the requirement exists." [official]

- There is no authoritative account for the one-fish-many-meals toolchain: the statements found all come from tool vendors' own promotion, so this chapter writes only the practice and gives no tool list. When you pick a tool, have the other side explain clearly whether it changes the form or merely swaps the cover.
- The three synthetic-user studies are all preprints, not peer-reviewed, and two of them use survey data from an English-language context; their conclusions can serve only as a boundary reminder of "when not to believe it," not as a basis for product selection.
- This chapter cites no multiplier for how many pieces one draft becomes and no lead-volume figures from distribution; such numbers currently have no verifiable source. [our judgment]
- Platform rules change. Rechecking the disclosure rules before every publish is more useful than memorizing what is written in this chapter.

## Sources for this chapter

[1] YouTube Help | Understanding 'How this content was made' disclosures on YouTube | https://support.google.com/youtube/answer/15447836?hl=en | [official] [2] Douyin (Chinese: 抖音) | Rules Center | https://www.douyin.com/rule | [official] [3] Xiaohongshu (Chinese: 小红书) | Pugongying Help Center | https://pgy.xiaohongshu.com/help/detail?id=1eda0a065dd894063c2e029a49e8f6a1&userType=4 | [official] [4] Zihan Chen, Di Zhu, Lei Nico Zheng | When Synthetic Users Fail: A Cross-Domain Benchmark of LLM-Simulated Human Survey Responses | arXiv:2607.26348, submitted 2026-07-28 | https://arxiv.org/abs/2607.26348 | [third-party verified: preprint] [5] Erika Elizabeth Taday Morocho, Lorenzo Cima, Tiziano Fagni, Marco Avvenuti, Stefano Cresci | Assessing the Reliability of Persona-Conditioned LLMs as Synthetic Survey Respondents | arXiv:2602.18462, submitted 2026-02-06 | https://arxiv.org/abs/2602.18462 | [third-party verified: preprint] [6] Bastián González-Bustamante, Nando Verelst, Carla Cisternas | Emulating Public Opinion: A Proof-of-Concept of AI-Generated Synthetic Survey Responses for the Chilean Case | arXiv:2509.09871, submitted 2025-09-11 | https://arxiv.org/abs/2509.09871 | [third-party verified: preprint] [7] Gary Vaynerchuk | The GaryVee Content Strategy (inverted pyramid: one long video cut into dozens of pieces) | https://garyvaynerchuk.com/the-garyvee-content-strategy-how-to-grow-and-distribute-your-brands-social-media-content | [first-person account] [8]

Bianews (Chinese: 币安新闻) | Oriental Selection live-stream clip licensing and the consequences of homogenization | https://www.bianews.com/news/details?id=165035 | [third-party reported] [9] Xinhua (Chinese: 新华网) | Douyin's Rules for Managing Improper Conduct by Matrix Accounts (main and sub-accounts handled uniformly) | http://www.xinhuanet.com/tech/20231130/c932193c6db34a079588421d9706516c/c.html | [third-party reported: citing official rules] [10] China Securities Journal (Chinese: 中国证券报) | Douyin took down 740,000 pieces of reposted infringing content and penalized 140,000 accounts in the first half of 2021 | https://cs.com.cn/sylm/jsbd/202106/t20210624_6178255.html | [third-party reported] [11] YouTube Help | Channel monetization policies (repetitious or mass-produced content) | https://support.google.com/youtube/answer/1311392?hl=en | [official] [12] Gong Labs | Taking the Guesswork Out of Sales Call Effectiveness (own call data) | https://www.gong.io/blog/taking-the-guesswork-out-of-sales-call-effectiveness/ | [vendor claim] [13] Strava Press | Year in Sport Trend Report (own billions of activities + 30,000-person survey) | https://press.strava.com/articles/strava-releases-12th-annual-year-in-sport-trend-report-2025 | [vendor claim] [14] Sohu (Jijinshe) (Chinese: 搜狐号·技锦社) | How Fan Deng Reading built exponential growth (about 60 million cumulative followers across Douyin matrix accounts; product standardization + private domain + multi-platform distribution) | https://www.sohu.com/a/684937054_121716213 | [third-party verified: downgraded]

[15] MeasuringU | Review of Experiments with Synthetic Users (a consultancy replicating a 3,600-person global wealth survey with a multi-agent simulation, one day versus six months; quantitative UX researcher Chris Chapman's opposing conclusion) | https://measuringu.com/review-of-experiments-with-synthetic-users | [third-party verified] [16] User Interviews | State of Synthetic Users 2026 (survey of 150 researchers plus five in-depth interviews; definition converged but respondents stress these are not real users) | https://www.userinterviews.com/state-of-synthetic-users-report | [third-party verified] [17] Tencent Cloud developer community | A 2026 review of new SEO norms (two updates in March back to back; affiliate-marketing sites 71% affected, financial affiliate sites −40% to −55%, generic product-comparison sites −35% to −55%) | https://cloud.tencent.com/developer/article/2663229 | [third-party verified: community sample] [18] Google Search Central | Spam policies (definition of scaled content abuse: generating large amounts of pages to manipulate rankings rather than to help users; AI generation is not the test) | https://developers.google.com/search/docs/essentials/spam-policies | [official] [19] Digital

Applied | Scaled content abuse: AI pages decimated by Google's March 2026 update (sites with large volumes of unedited AI pages lost 50–80% of organic traffic within two weeks) | https://www.digitalapplied.com/blog/scaled-content-abuse-google-march-update-ai-pages-decimated | [third-party verified] [20] theStacc | Google AI content policy 2026 (61% of sites lost 40–90% of traffic; about 30% unhelpful URLs triggers site-level demotion; recovery takes 3–6 months) | https://thestacc.com/blog/google-ai-content-policy-2026/ | [third-party verified] [21] Frase | Does Google penalize AI content? (June 2026 sample of 1 million pages across 100,000 SERPs, 150,000 AI-detected: thin content and scaling are punished, not AI generation) | https://www.frase.io/blog/does-google-penalize-ai-content-no-but-it-punishes-this | [third-party verified]

[22] pSEO | Programmatic SEO Results: Data from 4 Real Campaigns (1,080 pages across four projects; impressions climb in weeks 2–3, clicks lag 2.4 weeks; 84% of new traffic from keywords searched under 200 times a month; city pages without real local data suppressed within three weeks) | 2026-03 | https://pse0.com/research/programmatic-seo-results-2026 | [third-party verified: first-hand data teardown]

# Chapter 13 Outreach: From Volume to Precision

Outreach used to be a game of "sending more": bigger lists, more rounds, harder pitches — pile on enough volume and something will land. Now that AI has pushed the marginal cost of a send close to zero, every extra email spends not budget but the trust the recipient has extended to you — an allowance you cannot top up, only save by spending less. What you have to do is replace the mass blast with "sending the right thing to the right person at the right moment."

At this step the buyer has not heard of you yet, so you have to reach out first. That is why this chapter gives you 5 methods to replace the mass blast with "sending the right thing to the right person at the right moment."

---

## Method 1: Grade outreach by cost per qualified pipeline

**What to do**: Kill the piece-rate metrics that reward "sending more" and grade on "how much money each qualified pipeline cost" instead.

**How**

1. Kill piece-rate targets such as "150 emails a day" or "400 calls a week." In an environment where sending costs trend toward zero, such a target hands every employee an authorization to spend the company's trust allowance — and on the payroll that authorization looks free.
2. Switch to cost per qualified pipeline: calculate how much each qualified opportunity cost, instead of looking only at reply rate and meeting count.
3. Give assets a "depletion budget" rather than a "sending budget": the domain reputation you can burn each month is finite; when you touch the warning line, stopping is a reward, not a dereliction of duty.
4. Write "send less" into the KPI: any increase in sending volume must come with evidence that list quality has not fallen (substantiation rate, reply quality, and complaint rate unchanged); without that evidence, the increase is not approved.
5. Also look at how your tools are priced: a tool that charges by sending volume makes "sending more" simultaneously your usage, the vendor's revenue, and the platform's call volume.

**Example**: Volume went up and reply rates came down — outreach per person per month rose from a human baseline of about 1,150 emails to about 7,400 with AI augmentation (6.4 times), while the raw reply rate fell from 4.7% to 2.9% (−38%) over the same period [third-party reported]. On cost per qualified pipeline, a human-only team's qualified-opportunity cost was about 487 USD each and an AI-plus-human hybrid team about 224 USD each (−54%); the hybrid team (1 human SDR + 2 AI seats) got 1.9 times more meetings per dollar than AI alone and 2.4 times more than humans alone [third-party reported]. Greenway's controlled experiment gives numbers in the same direction but of a different magnitude (hybrid 342 / fully autonomous 588 / human-only 1,006 USD); the two sets of figures rest on different bases, but the ranking is exactly the same: hybrid best, human-only most expensive, and the "cheapness" of fully autonomous is on paper. Pricing is shifting too: HubSpot cut Breeze Customer Agent to 0.5 USD per resolved conversation and changed Breeze Prospecting Agent to 1 USD per recommended outreach lead [third-party reported]. There is a counterpart to platform punishment of "volume" as well: in March 2024 Google upgraded its "automatically generated content" policy into "scaled content abuse," treating as spam any content produced in bulk to manipulate rankings and of very low value to users, and says low-quality non-original content fell 45% [official] — the test of "bulk + low value + for manipulation" holds for bulk outreach just as it does for pages. Human-side daily reference volumes have also been published: a B2B outbound team's day runs to about 44 calls, 41 emails, 19 LinkedIn touches and 8 other activities, with phone-centric teams averaging 56 dials and 4.6 quality conversations a day, and a cold-call connect rate reference of 9.9% per dial [third-party verified]. The same benchmark compilation flags a denominator trap: two current datasets put cold email reply rates at 0.45% and 3.43% respectively, because their campaign populations and measurement methods differ — any "10x better than humans" claim should first be asked whether it is comparing send volume or qualified opportunities [third-party verified]. The benchmarks for that ruler are public too: a Chinese playbook puts 2026 cold-email averages at about 27.7% open rate (above 45% is good) and reply rates of 5%–15% as good and above 15% as excellent [third-party reported]; another first-person account tells the downside plainly — over 400 hours building an AI SDR moved reply rates from 1% to 2%, while human cold calling produced 5× the pipeline, and average email reply rates fell from 8.5% in 2019 to 3.4% in 2026 [third-party reported: author's own account]. The market itself is being priced: an AI SDR industry report gives three numbers — 72% adoption, a USD 3.1 billion market and 42% annual growth — yet most failed pilots collapse within nine

weeks, and none of the churned teams switched to another AI SDR [vendor claim]; a tool-buying compilation puts the AI SDR market growing from USD 4.39 billion in 2025 to USD 5.81 billion in 2026 (+32.3%) and USD 17.58 billion by 2030, with annual contracts for a single autonomous outbound agent commonly running USD 40,000–60,000 and the average cold-email reply rate down to 3.43% (top quartile above 5.5%, elite campaigns above 10.7%) [third-party verified]. Gartner's May 2026 sales-leader survey says organizations giving reps AI next-best-action guidance are 2.6 times more likely to hit above-average commercial growth, and projects that by 2027 95% of sellers' research workflows will start with AI (versus under 20% in 2024), with 51% of enterprises already running AI agents in production and hybrid human-plus-AI pods generating 2.3 times more revenue than AI alone [third-party verified].

**Pitfalls**: Reporting pipeline but not asset health. A meeting cadence like that systematically rewards "trading assets for results," and the people being rewarded will not volunteer their own depletion.

## Method 2: Pass the recipient's gate before you send

**What to do**: Get three things — domain authentication, one-click unsubscribe, and complaint rate — done before you send, rather than discovering them only when your email cannot get through the door.

**How**

1. Authenticate your mail: pass both SPF and DKIM, have a DMARC record on the domain, and align the organization domain in the From header with the SPF or DKIM organization domain.
2. One-click unsubscribe: commercial and promotional mail must offer RFC 8058 one-click unsubscribe, and requests must be honored within 48 hours.
3. Keep the user-reported spam rate under the threshold: measured daily, hold it below 0.1%, and never let it reach 0.3% or higher.
4. Watch the bulk-sender threshold: sending close to or more than 5,000 messages a day to personal Gmail accounts, counted at the primary domain level (primary domain and subdomains combined); once you are classified as a bulk sender the identity is permanent, and changing your sending habits will not remove it. Two easy-to-miss boundaries: the rule applies only to mail sent to personal Gmail

accounts, not Google Workspace accounts; and "new domains" that sent fewer than 5,000 messages a day before 2024-01-01 follow a faster enforcement timeline.

5. On the China side, put consent first: Article 24 of the Regulations for the Implementation of the Law on the Protection of Consumer Rights and Interests provides that without the consumer's consent, no commercial messages may be sent and no commercial calls made, whether the party on the other end is a person or an AI; scripts, digital-human videos, and AI customer-service audio generated in bulk by AI must also carry explicit labels and implicit labels (metadata watermarks) under the Measures for Labeling AI-Generated Synthetic Content, and platforms must verify the labels at the publishing stage.
6. The phone channel tightens in the same direction. The U.S. FCC ruled on 2024-02-08 that calls using AI-generated voices are artificial/prerecorded voice for TCPA purposes, making AI voice calling without express consent illegal; on the China side, a July 2026 investigation by state media found that large-model agents dial 800–1500 calls a day per machine, that numbers can be disguised as real landlines with area codes, and that "one-click de-tagging" is standard equipment in the industry, while a survey by the Shanghai Consumer Council (Chinese: 上海消保委) said such calls are often blind-dialed indiscriminately at a cost as low as a little over one jiao per minute [third-party reported].

**Example**: In an official blog post on 2023-10-03, Google announced that from 2024 it would impose three mandatory requirements on bulk senders — authentication, one-click unsubscribe, and spam rate [official]. The details are in its sender guidelines: the spam rate is calculated daily, should stay below 0.1%, and must never reach 0.3%; from June 2024, bulk senders with a spam rate above 0.3% are not eligible for mitigation through appeals, and must stay below 0.3% for 7 consecutive days to recover [official]. The consequences are hard: non-compliant mail receives temporary or permanent failure codes in the 4.7.x / 5.7.x series, or is dropped straight into the spam folder; the receiving side's filtering baseline is that Gmail blocks more than 99.9% of spam, stopping close to 15 billion messages a day, and after it began requiring authentication the number of unauthenticated messages users received fell by 75% [official]. **Composite profile | constructed from a common industry pattern, not a named event**: a 40-person B2B company turned outreach into an "engine," with 3 seats each sending 150 messages a day; by week three the complaint rate crossed 0.1%; by week five primary-domain reputation fell

from good to suspicious (the comparable trajectory in the experiment records is 19 days); they switched to a subdomain and kept sending, not realizing that Gmail counts the 5,000-message threshold at the primary-domain level — within two weeks the subdomain had walked the road the primary domain took two months to walk; by week six, customers were reporting that contracts, invoices, and customer-service replies "never arrived, they were in the spam folder," and recovery took more than 8 weeks. That gate was raised another notch in 2026: by a third-party account, Google began requiring DMARC from bulk senders in February 2024, upgraded to hard rejection of non-compliant mail in November 2025, and by March 2026 mail from domains without properly configured DMARC, DKIM and SPF was refused outright by major mailbox providers rather than landing in spam; in the same account, personalization built on signals — specific purchase-trigger events — replies 3–5 times as often as personalization built on firmographics such as size and industry, and top SDR teams using signal-based sequences sit in the 5%–15% reply range [vendor claim]. The actionable thresholds are public too: senders pushing 5,000 or more emails a day to Gmail, counted per primary domain, must have SPF, DKIM and DMARC in place alongside one-click unsubscribe; the spam-complaint-rate target is below 0.10% with a hard ceiling never to be exceeded of 0.30%; the hard-bounce benchmark is below 1%; and new domains typically need 2–4 weeks of warm-up [third-party verified: relaying Google's official rules]. Put AI and humans through the same paired comparison and the gap is far smaller than the rhetoric suggests. An analysis of 100,000 paired cold emails — 50,000 AI-generated and 50,000 human-written, matched on persona, ICP, sequence stage, sender domain age and sender domain authority, covering October 2025 to April 2026 — found AI reply rates at 4.1% against 5.2% for humans, meeting-booked rates at 0.7% against 1.1%, identical 6% bounce rates (a list-quality signal, not a content one), and spam-flag rates of 8% for AI against 3% for humans; the largest lever on inbox placement was not copy but cadence, with 1-day intervals placing 71% of sends in the inbox against 93% at 3-day intervals; vertical differences outweighed the human-versus-AI difference, with AI beating human sends in SaaS (6.1% vs 5.7%) and financial services replying at just 1.9% to AI-sent mail [third-party verified: first-hand data]. One counter-narrative is worth recording on its own: tracing the "AI content equals spam" statistics circulating in 2026 leads only to marketing blogs citing one another, with no named study, sample size or methodology behind the most-quoted numbers; the most on-point vendor research available, Validity's analysis from late 2024, found no evidence that AI-generated content is more likely to be marked as spam, and the 450-participant

academic trial it cites saw plain-text AI email clear the Gmail, Yahoo and Outlook filters entirely [third-party verified: including vendor research]. The filter is largely indifferent to who wrote the message; what decides inbox placement is volume, complaints and engagement — fix authentication, segmentation and interaction design before arguing over whether the copy reads like AI.

**Pitfalls**: Switching to a subdomain and carrying on. Gmail's classification of bulk-sender identity is permanent — a reset does not clear the credit record; and because primary domain and subdomains are counted together, changing domains saves nothing and only burns the new one along with the old.

## Method 3: Let "who just had something happen" be the entry point to the list

**What to do**: Stop asking "who do I want to contact" and ask "who just had something happen" — let the event decide the list, and also decide the one line in the letter.

**How**

1. Watch signals, not lists: funding, headcount expansion, entry into a new market, a technology-stack change, a visit to the pricing page, hiring for a specific role.
2. Make the signal trigger an action that must happen immediately. Getting from a list to pipeline takes three things done: detecting the signal, layering on enrichment, and acting in time. Miss any one of the three and the signal degrades back into just another list.
3. Use explainable events, not unexplainable scores. A usable reading of a qualification score is "this account's research intensity in a given category is 3 times its own baseline, and it visited the pricing page twice this week," not an 85 with no provenance.
4. Treat a signal as a knife for narrowing the list, not as a buying channel for enlarging it.
5. Treat intent data as a "hint" rather than a "fact": check its substantiation rate against your own CRM first, and only then decide whether to add budget for it.
6. Remember why "it looked written just for me" ever counted as a signal: a letter that looked written just for him used to be a real signal, because writing it cost research time that a low-intent sender could not afford. Spence (1973) set out

the condition for a credible signal — a signal holds only when senders' costs differ enough; Akerlof's 1970 lemons-market mechanism runs in the same direction — when buyers cannot tell good goods from bad, they first mark down their valuation of all outreach, and then choose not to look. Once AI pushed the cost of writing a letter to zero, that class of signal stopped working.

**Example**: The first 100 users of the AI study app Jungle were won by sending individual direct messages to the 200-plus people who engaged with a viral tweet — at the cost of the founder's Twitter account being banned for outreach volume [vendor account: first-hand retrospective]. Finding the signal accurately does not mean the way you send it stays inside the line.

**Pitfalls**: Assuming that buying intent data doubles your results. Most suppliers hand over one company list and stop, and a signal that never triggers an action turns back into a liability you have to clean up.

## Method 4: AI does the research and the first draft; a human writes the one key line

**What to do**: AI finds the facts, builds the list, drafts, and handles the logistics; humans verify the facts, write the one line only you can write, and handle any reply.

**How**

1. Fix the division of labor in writing: AI flags accounts with 3 or more concurrent buying signals and generates a context brief; a human spends about 3 minutes reviewing and correcting; AI sends the approved first touch and monitors engagement; any reply is taken over by a human, and AI never replies to a prospect.
2. Stop hallucinations before they go out. The common factual errors in first touches are wrong titles and wrong milestones — including congratulating someone on "joining" a company they left six months ago.
3. Remember that AI cannot copy "tone": if the person says "I just renewed with a competitor through 2027" and AI fires back another value statement plus a calendar link, the email gets forwarded as an example of junk.
4. Watch deliverability, not just reply rate: aggressive sending will drive primary-domain reputation from good to suspicious in a very short time.

5. Hold one boundary: AI does not touch what happens in the inbox. It can write the first line, but not the second line of the conversation.
6. Keep two background numbers in mind together: testing of five mainstream AI personalization tools found that 85–95% of so-called "personalization" is just a template filled in with 3–5 fields, and about 15% of the parts that did attempt real research hallucinated errors [third-party reported: single source, used with caution]; a Gartner survey says more than 60% of B2B decision-makers say they will set up filters or blocks for sales-development emails that carry obvious signs of AI generation [third-party reported].

**Example**: A controlled experiment across 800 accounts, two quarters, and four platforms (fully autonomous email agent / AI copilot with drafts approved by a human / enrichment + sequence hybrid / human-only SDR) used "qualified pipeline opportunities" rather than reply rate as its core metric: the fully autonomous agent booked the most meetings, but meetings-to-opportunity conversion was only 8%; human SDRs 31%; copilot 26%. On cost per qualified opportunity: fully autonomous 588 USD each, human-only 1,006 USD each, hybrid only 342 USD each [third-party account: vendor perspective]. Where it worked, the process produced about 3.2 times the qualified meetings per person; the failure modes are recorded alongside — 12% of first touches contained factual errors, 58% of emails went to accounts that did not fit the ICP, aggressive sending drove primary-domain reputation from good to suspicious in 19 days, and emails identified as AI-written drew 34% fewer replies. Corroborating evidence: Artisan's CEO admitted that the first-generation AI SDR had "very low reply rates and very high churn," and that tightening the prompt and having customers fill in information through forms brought hallucinations down to about one in ten thousand [third-party reported]; SaaStr's own outreach sent 7,000 emails over six weeks for a 3.6% positive reply rate, after which a campaign run personally by the founder — with no send-timing optimization at all — hit 4%. Beating the industry benchmark without optimizing timing shows that "who you send to and what you send" matters far more than "when you send." The anti-AI-outreach test is also shifting from human perception to detectable watermarks and platform rules: one report says Claude now embeds an invisible watermark in all generated text (a global change driven by EU regulatory requirements), meaning AI-drafted sales email can be identified as machine-written; the same piece observes that fully AI-generated emails get fewer replies and more spam flags than human-written ones, and recommends AI drafting plus human rewriting [third-party reported]. The failure point has been

measured on its own: an AI SDR industry report says half of failed teams had perfectly personalized copy but lost on context and timing, and 88% of pilots stall before production [vendor claim] — the problem is never whether it reads as human, but whether it was the right moment to send.

Failure cases are valuable because they name the step that was never encoded: one AI SDR vendor published its own customers' wrecks — a long-standing contact went straight to the founder saying “this doesn't sound like you at all”, the rollout was killed the same day at an estimated cost of about $100,000 in potential contracts, and the root cause was missing multi-turn memory and missing reply-based branching; another agency grouped senior brand leaders with junior media buyers and pushed one generic sequence to both, and reply quality collapsed because segmentation was never done [vendor claim: customer cases]. The buyer-side post-mortem points the same way: an agency that uninstalled its AI SDR concluded the problem was not sending too little but booking the wrong meetings — a fuller calendar while the people in the room drifted further from the ideal customer profile, a healthy dashboard with flat revenue; it traced the break to three places: no multi-turn memory, no reply branching, and a metric priced on meetings booked rather than qualified opportunities and closed deals [third-party verified: user's post-mortem]. Both prescriptions match: keep AI in research and drafting, keep send decisions and conversation ownership with people, and accept on meeting-to-qualified-opportunity rate.

**Pitfalls**: Letting AI take over what happens in the inbox. Once AI takes over replies, the tragedy of the commons spreads from the sending side to the engagement side — and the trust allowance on the engagement side is thinner and less renewable.

## Method 5: Give each channel the one thing only it does well

**What to do**: Assign different channels to different stages of the decision chain, instead of sending the same message once in each of four channels.

**How**

1. First define for each channel "the thing only it can do": email carries long information and can be forwarded as an argument; a direct message can be seen once without giving offense; a phone call handles objections in the same moment; IM carries a legitimate and durable relationship authorization.

2. If a channel has been given something others can also do, take that something away from it — its marginal value is otherwise noise.
3. On the China side the orchestration carrier is IM: break "add a friend and go silent" into a configurable cadence — the day of adding → day 1 / 3 / 7, escalating, then dropping to a low frequency of once a month for those who do not engage; set a hard risk-control rule of at most 1 touch per user in 24 hours.
4. Then add one path that is on no tool vendor's menu: reduce total touches and invest the saved trust allowance in "being found." When everyone is increasing sends, the strategy that raises the quality of each touch is to opt out of the volume competition.

**Example**: Multichannel orchestration produced a 23% improvement in pipeline velocity over single-channel [third-party reported]; a McKinsey study of B2B enterprise-services and manufacturing companies says data-driven ABM raised lead conversion 3.5 times and deal size by 20% [third-party reported: effect size on the large side, taken for direction]. On the China side: Tencent Qidian (Chinese: 腾讯企点) says it has connected 350 million users and serves more than 1 million companies; in its retail case, Laiyifen (Chinese: 来伊份) linked 3,700+ specialty stores with data on 45 million+ users and deployed 11 automated marketing scenarios, and according to its technical director marketing conversion, coupon redemption rate, and ARPU each rose 5–10 percentage points [third-party reported]. A private-domain digital-employee case published by Shuyun (Chinese: 数云) says that after adoption the new-customer reception rate rose 140% and efficiency cost fell 61.3% [official claim]. Overseas IM engagement has its own backdrop: WhatsApp has more than 3 billion monthly actives, its paid business messaging passed 2 billion USD annualized in 2025 Q4, and on 2026-06-03 Meta opened Business AI Agent to companies worldwide [third-party reported]. The reminder about ownership belongs here too: the domain is yours, but the right to deliver sits with Gmail; the account is yours, but the right to exist sits with the platform — at the end of 2025, the star AI-outreach startup Artisan was banned by LinkedIn, as its CEO confirmed to TechCrunch, and was restored in January 2026 after about two weeks of discussions; during the ban its lead pipeline simply stopped [third-party reported]. An earlier platform closure came from regulatory compulsion: in September 2021, at a State Council Information Office press conference, MIIT (Chinese: 工信部) listed "blocking URL links" as one of the priority problems in its special rectification of the internet industry, requiring the platforms concerned to lift blocking to standard within a deadline [third-party reported] — when the path to

reaching someone has to pass through a super-app's share entry, your permission to reach them is rented from the platform. On the division of labour, one guide singles out step 6, the handoff, as the dividing line in most AI outreach deployments: AI does research, first touch and multichannel follow-up, then passes qualified prospects to a human SDR with the full interaction history logged in the CRM; objection handling is not a good fit for AI, and the first objection escalates to a human [third-party reported]. The metric should therefore be whether the handoff arrives with context, not how many emails the AI sent.

**Pitfalls**: Treating "how many channels we covered" as an achievement. The test is whether each channel got the one thing only it does well; covering four channels where each one is doing something others can also do is the same piece of noise sent four times.

---

## Notes for this chapter

- The conflicting figures for cold-email reply rates cannot be reconciled into one number: 0.45% (Belkins, counting only emails sent to brand-new contacts, all of 2025; as low as 0.35% in December, roughly 20% worse within the year, described as "interruption inflation," the strictest basis), 2.9% (Digital Applied aggregate), 3–5% (Autobound), and 4–5% (Laxis) sit side by side, with completely different denominators, a full order of magnitude apart. For budgeting a broad basis will overstate; for diagnosis a strict basis will understate — look at both, and the gap between them is your list quality. [our judgment]
- The gap between "signal-driven 15–25% vs cold email 3–5%" comes from vendor aggregation; the same figure appears at two vendors selling the same product category, so they are likely the same origin and should not be treated as two independent sources — direction only. [vendor claim]
- Domain warming of 30 days / 60–90 days / 90+ days giving delivery rates of about 51% / 86% / 91%, and a daily sending cap of about 25 messages per mailbox, both come from vendor research and are not independently verified; the direction (domains with insufficient warming have a clearly lower delivery rate) is credible, but the specific figures should be cited with caution. [vendor research]

- On the China side, "at most 1 touch per user in 24 hours" comes from a service provider's practice, not a platform rule; no official rules page could be found for WeCom (Chinese: 企微)'s daily touch limit, friend-add limits, or mass-messaging frequency rules — this is the thinnest part of the China side of this chapter. [our judgment]
- The link for the FCC's 2024-02-08 ruling returned an error page on this round's fetch; the date and the wording "AI-generated voices are artificial/prerecorded voice under the TCPA" follow the materials archive and need to be re-checked against the text on the FCC website; the CAN-SPAM civil penalty amounts and the TCPA statutory damages range of 500–1,500 USD per call or message likewise need to be traced back to source. [official: pending re-check]

## Sources for this chapter

[1] Google | New Gmail protections for a safer, less spammy inbox (three requirements for bulk senders) | 2023-10-03 | https://blog.google/products-and-platforms/products/gmail/gmail-security-authentication-spam-protection/ | [official] [2] Google | Email sender guidelines FAQ (bulk-sender definition, the 5,000-message threshold and primary-domain counting rule, spam rate 0.1%/0.3%, permanent bulk-sender status, tighter enforcement from 2025-11) | https://support.google.com/mail/answer/14229414 | [official] [3] Google | New ways we're tackling spammy, low-quality content on Search (scaled content abuse; low-quality non-original content −45%) | 2024-03-05 | https://blog.google/products-and-platforms/products/search/google-search-update-march-2024/ | [official] [4] FCC | FCC Makes AI-Generated Voices in Robocalls Illegal | 2024-02-08 | https://www.fcc.gov/document/fcc-makes-ai-generated-voices-robocalls-illegal | [official: link pending re-check] [5] China.com.cn (Chinese: 中国网) | AI outbound calls dial a thousand nuisance calls a day — where is the "compliance boundary"? (800–1500 calls a day per machine; one-click de-tagging; Shanghai Consumer Council says cost as low as a little over one jiao per minute) | 2026-07-28 | http://news.china.com.cn/2026-07/28/content_118621717.shtml | [third-party reported] [6] China News Service (Chinese: 中国新闻网) | MIIT: blocking URL links is one of the priority problems in the special rectification of the internet industry | 2021-09-13 | https://www.chinanews.com/cj/2021/09-13/9564236.shtml | [third-party reported] [7] China Economic Net (Chinese: 中国经济网, via CCTV.com) | From September 1, AI-generated synthetic content must carry labels | 2025-09-01 | http://www.ce.cn/xwzx/gnsz/gdxw/202509/t20250901_2460291.shtml | [third-party reported]

[8] The Next Web | AI is turning B2B outreach into a tragedy of the commons (Belkins: 7.5 million+ cold emails, 0.45% average reply rate in 2025; "interruption inflation") | https://thenextweb.com/news/b2b-outreach-tragedy-of-the-commons | [third-party reported] [9] Digital Applied | AI SDR Statistics 2026 (volume up 6.4 times, reply rate −38%, qualified-opportunity cost −54%, meetings per dollar for the hybrid team, vertical adoption rates) | 2026-04-21 | https://www.digitalapplied.com/blog/ai-sdr-statistics-2026-outbound-sales-data-points | [third-party reported] [10] Laxis Research | State of AI SDR 2026 (warming 51%/86%/91%; daily sending cap of about 25 messages per mailbox; signal-driven 15–25% vs fully autonomous 1–3%) | 2026-06 | https://www.laxis.com/blog/state-of-ai-sdr-2026/ | [vendor research] [11] Autobound | State of AI Sales Prospecting 2026 (signal personalization 15–25% vs cold email 3–5%) | 2026-02-20 | https://www.autobound.ai/blog/state-of-ai-sales-prospecting-2026 | [vendor claim] [12] Gartner | Gartner Says 60% of CMOs Will Adopt AI to Fight Decline in Customer Trust (more than 60% of B2B decision-makers will filter sales-development emails with obvious AI traces) | 2024-03-06 | https://www.gartner.com/en/newsroom/press-releases/2024-03-06-gartner-says-60-percent-of-cmos-will-adopt-ai-to-fight-decline-in-customer-trust | [third-party reported] [13] Greenway | AI SDR: What Works, What Doesn't (800-account controlled experiment; 12% hallucinations, 58% not fitting the ICP, primary domain burned in 19 days, 8 weeks to recover, hybrid 3.2x, qualified-opportunity cost 342/588) | 2026 | https://greenway.ai/post/ai-sdr-what-works-what-doesnt/ | [third-party account: vendor perspective] [14] The Starr Conspiracy | B2B Intent Data Benchmarks 2025 (21.3% vs 8.4%, −28 days, 62% insufficiently substantiated, 47% expiring in 14 days, median precision 0.51, median enterprise spend 312,000 USD) | refreshed 2026-05 | https://www.thestarrconspiracy.com/insights/benchmarks/b2b-intent-data-benchmarks-2025 | [third-party reported] [15] Warmly | B2B Intent Data Guide 2026 (31% of sales leaders call intent data "the most overrated") | 2026-03-27 | https://www.warmly.ai/p/blog/intent-data | [vendor claim] [16] 6sense | What is Intent Data? (how to read an example score of 85/100) | 2026-03 | https://6sense.com/platform/intent-data/what-is-intent-data/ | [vendor claim] [17] TechCrunch | Yes, LinkedIn banned AI agent startup Artisan. But now it's back | 2026-01-07 | https://techcrunch.com/2026/01/07/yes-linkedin-banned-ai-agent-startup-artisan-but-now-its-back/ | [third-party reported] [18] TechCrunch | Artisan raises $25M and is still hiring humans (CEO admits first-generation hallucinations were severe and churn was high) | 2025-04-09 | https://techcrunch.com/2025/04/09/artisan-the-stop-hiring-humans-ai-agent-startup-raises-25m-and-is-still-hiring-humans/ | [third-party reported] [19] SaaStr | Artisan's

Ava 2.0: What a Fully Autonomous AI BDR Actually Looks Like in Production (7,000 emails in six weeks, 3.6% positive replies; founder's campaign 4%) | 2026 | https://www.saastr.com/artisans-ava-2-0-what-a-fully-autonomous-ai-bdr-actually-looks-like-in-production-with-ceo-jaspar-carmichael-jack/ | [third-party reported] [20] LinkedIn Pulse (Rui Nunes) | AI Cold Email Is Killing Cold Email (85–95% templated, about 15% hallucinated research) | 2026 | https://www.linkedin.com/pulse/ai-cold-email-killing-everyones-pretending-its-fine-rui-nunes-vmjqf | [third-party reported: single source, used with caution] [21] SetSmart | Generative AI for Sales: 7 Use Cases (cost <$1 per thousand conversations) | 2026 | https://setsmart.io/blog/generative-ai-for-sales | [vendor research] [22] Spence, M. | Job Market Signaling | QJE 87(3), 1973 | https://www.sfu.ca/~allen/Spence.pdf | [official] [23] Akerlof, G. A. | The Market for "Lemons" | QJE 84(3), 1970 | https://econweb.ucsd.edu/~atimmerm/174/akerlof.pdf | [official] [24] Southern Metropolis Daily (Chinese: 南方都市报, via NetEase News) | Traffic dividends have peaked — how should marketing and customer service be done? (Tencent Qidian 350 million connections; Laiyifen 3,700+ stores / 45 million+ users / 11 automated scenarios) | 2024-01 | https://www.163.com/dy/article/IOTLLO6Q05129QAF.html | [third-party reported] [25] NetEase Zhiqi · Cloud Business (Chinese: 网易智企·云商) | How to use AI private domains to break "add a friend and go silent" (day of adding → day 1/3/7 cadence; at most 1 touch in 24 hours) | 2026-06 | https://grow.163.com/cms/ai-ke-fu-ai-si-yu-2026-3RbSGjTy.html | [official claim] [26] Shuyun (Chinese: 数云) | Private-domain AI winner renewal (new-customer reception rate +140%, efficiency cost −61.3%) | 2025-11 | https://www.shuyun.com/4499.html | [official claim] [27] CIO.com | Where Meta's WhatsApp agent can actually win (WhatsApp 3 billion+ MAU, business messaging $2B+ annualized, Business AI Agent) | 2026-07 | https://www.cio.com/article/4196328/where-metas-whatsapp-agent-can-actually-win.html | [third-party reported] [28] CMSWire | HubSpot shifts Breeze AI agents to pay-per-result pricing (0.5 USD per resolved conversation; 1 USD per recommended outreach lead) | 2026-04-02 | https://www.cmswire.com/customer-experience/hubspot-shifts-breeze-ai-agents-to-pay-per-result-pricing/ | [third-party reported] [29] McKinsey | The B2B Digital Growth Engine (omnichannel + ABM: conversion 3.5 times, deal size +20%) | https://www.mckinsey.com/capabilities/growth-marketing-and-sales/our-insights/the-b2b-digital-growth-engine | [third-party reported: effect size on the large side, taken for direction] [30] Superwall | This AI study app makes $100K/mo (first 100 users won by DMing 200+ engagers on a viral tweet; account banned for outreach volume) | https://superwall.com/blog/this-ai-study-app-makes-usd100k-mo-heres-how-no-code | [vendor account: first-hand retrospective]

[31] Level Up Leads | B2B Outbound Benchmarks 2026 (daily reference of 44 calls, 41 emails, 19 LinkedIn touches, 8 other activities; phone-centric teams 56 dials and 4.6 quality conversations; 9.9% connect rate; two datasets reporting 0.45% and 3.43% reply rates) | https://levelupleads.io/blog/b2b-outbound-benchmarks | [third-party verified] [32] Autobound | Cold Email Best Practices 2026 (from 2026-03 mail from domains without correct DMARC/DKIM/SPF is refused outright by major mailbox providers; signal-based personalization 3–5x firmographic; top SDR teams 5%–15% reply rates) | https://www.autobound.ai/blog/cold-email-best-practices-2026 | [vendor claim] [33] Unify | Cold Email 2026: Domain Setup & Deliverability (5,000 emails/day bulk-sender threshold; complaint-rate target below 0.10% with a 0.30% hard ceiling; hard-bounce rate below 1%; 2–4 weeks of warm-up) | https://www.unifygtm.com/explore/cold-email-2026-domain-setup-deliverability-sequences | [third-party verified: relaying Google's official rules] [34] Laxis | Cold email outreach playbook 2026 (average open rate about 27.7%, above 45% good; reply rates 5%–15% good, above 15% excellent) | https://www.laxis.com/zh/blog/cold-email-outreach-playbook-2026/ | [third-party reported] [35] TitanX | Can AI SDR agents replace human sellers (400+ hours building an AI SDR moved reply rates only 1%→2%; human calling produced 5× the pipeline; average reply rate down from 8.5% in 2019 to 3.4% in 2026) | https://titanx.io/news/can-ai-sdr-agents-replace-human-sellers | [third-party reported: author's own account] [36] Forbes | AI wrote your cold email — now it carries an invisible watermark (Claude embeds an invisible watermark in generated text; fully AI-generated emails get fewer replies and more spam flags) | 2026-08-16 | https://www.forbes.com/sites/terdawn-deboe/2026/08/16/ai-wrote-your-cold-email-now-it-carries-an-invisible-watermark/ | [third-party reported] [37] Aircall | AI sales agent guide (step 6, the handoff, is the dividing line in most deployments; AI does research, first touch and follow-up, the first objection escalates to a human) | https://aircall.io/blog/ai-sales-agent-guide/ | [third-party reported] [38] AiSDR | AI SDR Industry Report 2026 (72% adoption, USD 3.1 billion market, 42% annual growth; 81% have experimented; most failed pilots collapse within nine weeks; 88% of pilots stall before production; half of failed teams had personalized copy but lost on context and timing; none of the churned teams switched to another AI SDR) | 2026 | https://aisdr.com/ai-sdr-industry-report | [vendor claim] [39] Rework | Best AI tools for SDRs 2026 (AI SDR market USD 4.39 billion in 2025 to USD 5.81 billion in 2026, +32.3%, USD 17.58 billion by 2030; single autonomous outbound agent contracts commonly USD 40,000–60,000 a year; Gartner May 2026: AI next-best-action organizations 2.6× more likely to hit above-average growth; by 2027 95% of seller research workflows start with AI vs

under 20% in 2024; 51% of enterprises run AI agents in production; hybrid human+AI pods 2.3× revenue; average cold-email reply rate 3.43%, top quartile 5.5%+, elite 10.7%+) | 2026 | https://resources.rework.com/tools/ai-tools/best-ai-tools-for-sdrs-2026 | [third-party verified]

[40] Digital Applied | AI SDR Real Performance: 100K Email Analysis (100,000 paired cold emails: AI reply 4.1% vs human 5.2%, meetings 0.7% vs 1.1%, spam-flag 8% vs 3%, bounce identical at 6%; 1-day cadence 71% inbox placement vs 93% at 3 days; SaaS 6.1% beating human at 5.7%, financial services 1.9%) | 2026 | https://www.digitalapplied.com/blog/ai-sdr-real-performance-100k-email-analysis-2026 | [third-party verified: first-hand data]

[41] Digital Applied | Email Deliverability in the AI Slop Era (traces the 2026 claim that "AI content = spam" back through mutually citing marketing blogs, with no named study, sample or methodology; Validity's late-2024 analysis found no evidence that AI-generated content is more likely to be filtered as spam; in the 450-person academic trial it cites, plain-text AI emails passed at Gmail, Yahoo and Outlook; filtering criteria are largely independent of author identity) | 2026 | https://www.digitalapplied.com/blog/email-deliverability-ai-slop-era-2026-guide | [third-party verified: includes vendor research] [42] AI SDR | AI SDR Implementation Failures (no multi-turn memory and no reply-based branching: a long-standing contact complained to the founder and the rollout was killed the same day at about $100,000 in lost potential contracts; a segmentation failure pushed one generic sequence to senior brand leaders and junior media buyers alike) | 2026 | https://aisdr.com/blog/ai-sdr-implementation-fail/ | [vendor claim: customer cases] [43] Single Grain | We Uninstalled Our AI SDR (a fuller calendar while meetings drifted from the ideal customer profile; the break traced to missing multi-turn memory, missing reply branching, and pricing on meetings booked rather than qualified opportunities) | 2026 | https://www.singlegrain.com/artificial-intelligence/ai-sdr-what-breaks/ | [third-party verified: user's post-mortem]

# Chapter 14 Content Others Cannot Copy: First-Party Data and Real Cases

Content used to be a contest of who wrote faster and laid down more keywords, and one piece took a person several days from topic to publication. Now the cost of generating a piece has become a single compute call, and output itself is no longer worth anything. What you have to do is turn the raw material in your hands into evidence nobody else can produce. Content in front of buyers has already flooded; what decides whom they pick has become the material "only you have." That is why this chapter gives you 5 methods to turn first-party data and real cases into a moat.

---

## Method 1: Turn the traces customers leave with you into publishable raw material

**What to do**: Gather the behavioral records customers generate themselves inside your systems (website, App, CRM, stores, customer service) into raw material that only you can obtain and that you can also cite externally.

**How**

1. First keep four things apart. What you observed inside your own business, with collection records, reproducible, is first-party data; aggregator sites, tool-vendor summaries, and self-media retellings are leads only and must go back to the original publisher; packaging secondhand statistics as "our research," or generating something that "looks like a number," is pseudo-first-party data; and a model restating your existing content is AI paraphrase, which cannot serve as evidence either. Only the first kind can go into public material.
2. Get the "recognize the person" step done first. Leads from the website, mini-programs, Douyin (Chinese: 抖音), and WeCom (Chinese: 企微) each carry an identifier, stitched together by a company name or a unified ID, so the system knows that "these records are the same person" — a CDP market guide judges that identity resolution is the dividing line between a real CDP and a white-label tool [third-party verified]. If identity resolution cannot be done, every later step of processing rests on a wrong base.

3. Choose the processing direction correctly: turn proprietary data into something useful to the buyer — selection comparisons, parameter benchmarks, industry averages. Taking it to advertising makes it inventory.
4. De-identify or anonymize before publishing externally, and remember that the two are not equal: de-identified data is still personal information and can still be re-identified, while only anonymization is irreversible. Have legal walk through it before you write.

**Example**: A company making industrial sewing machines and CNC machine tools used a "company name / unified ID" to stitch together isolated leads from the website, mini-programs, and Douyin, and reports that cross-channel manual entry work fell 34% and overseas leads-to-opportunity conversion rose 29% [vendor claim]. The industrial-goods platform ZKH (Chinese: 震坤行) built a product pool from data on tens of millions of products, tens of billions of transactions, plus user behavior data, and says that 11 months after launch it had recommended more than 20 million lines of materials to customers, adding more than 100 million yuan in revenue and serving more than 270 group customers [vendor claim].

**Pitfalls**: Mistaking a "list" for "leads." A list of 50,000 rows is 50,000 identifiers, not 50,000 leads — a list with no record of origin and no consent status can neither be cited externally nor used compliantly. Also, figures such as "first-party data activation can cut CAC by as much as 50% and lift revenue 10–15%" come from data providers' own research; treat them as direction only, not as a basis for budgeting [third-party verified].

## Method 2: Have customers fill in a preference instead of guessing

**What to do**: Where the customer is willing to trade information for a better recommendation, ask about preferences directly — a questionnaire, a preference center, a selection checklist, a purchase-intent form.

**How**

1. Keep "actively filled in" separate from "passively collected." Zero-party data means preference information the customer actively and voluntarily shares (Forrester's definition, retold by a zero-party-data tool vendor) [vendor claim]; it is worth more than the behavior you guess at in the back end, because the

customer knows what he is giving and what he is getting. Asking directly, "do you care more about delivery speed or unit price?", is more useful than counting how many times he looked at the pricing page.

2. One question, one use. Three preferences filled in should buy a selection checklist the buyer can put to work immediately; if they buy only "we will give you a better experience," nobody fills it in next time.
3. Design how many data points you collect. A personalized journey rests on a dozen or so stable preference points per person, not on asking thirty questions at once and having nobody finish.
4. If you collect it, use it. Whatever he ticked, the later case studies, parameter pages, and follow-up emails should be ordered by it; collecting without using is the same as telling him the form was pointless.

**Example**: A survey cited in a strategy review shows that only 29% of consumers trust how brands use their data, 71% expect personalization, and 81% worry about how companies handle personal information; but when the process is transparent, willingness to provide information is 2.5 times higher, and 47% are willing to share when there is a clear benefit [third-party verified]. For comparison: ASICS collects 21.5 data points per user through a questionnaire to build a personalized journey, e.l.f. Cosmetics collects 11 preference points per customer, and Sephora uses a beauty quiz to make recommendations [third-party verified].

**Pitfalls**: Making the questionnaire a one-off campaign that asks a pile of questions without giving the corresponding benefit. What a customer hands over by voluntarily paying the cost of filling in a form is a signal source competitors are unwilling to match; if nothing comes back after submission, that cost was wasted.

## Method 3: Dig the numbers out of your own operations

**What to do**: Turn observations from the operating floor into content raw material — your delivery cycle times, the distribution of return reasons, equipment failure rates on different production lines, the ten questions customers ask most often in inquiries.

**How**

1. Start with the cheapest place: the questions sales and customer service actually answered, logged once a week and turned into topics. This is the lowest-cost entry point into data from your own operations.

2. Give the data its basis: how big the sample was, when it is from, how it was calculated. Without a clear basis, readers will assume you are making it up.
3. Write only the kind that exists nowhere else on the internet. Data others also have is a supplement; only the singular belongs in a flagship piece.
4. Product and integrations are exclusive data too. A tool product can inventory "who we integrate with"; a company going overseas can compare how the same set of pages performs across different markets and languages — an observation nobody else can obtain.

**Example**: From data on its own 6000+ apps, Zapier used templates to generate three tiers of landing pages, and its monthly organic search traffic rose from about 1.19 million to nearly 4.8 million between 2020-11 and 2023-11 [third-party verified]. Glasp stopped querying models externally and instead read its own server logs to see directly which pages the models were citing, growing ChatGPT-sourced traffic from 500 to 19,000 daily sessions [third-party verified]. On the China side, industrial-goods marketing driven by linked supply-chain data raised SQL conversion 32% [third-party verified]. There is a counterexample too: with 1.7 million+ programmatic pages covering 20 markets and 10 languages, Wise earned about 54 million monthly organic visits; while G2, held up as a SaaS benchmark, was reported to have lost about 80% of its organic traffic by the end of 2025 [third-party verified] — the fork is decided by whether there is proprietary, structured, continuously refreshed data behind the pages.

**Pitfalls**: Deciding your own data "isn't sexy" and not logging it. Things like delivery cycle times and return reasons do not sound like selling points, but they are the raw material that best supports "only I have this"; go looking for them when it is time to write and you will most likely no longer remember them.

## Method 4: Build cases into evidence with all four things in place

**What to do**: Build a case library in which every case must satisfy four hard requirements at once; miss one and it does not qualify.

**How**

1. **Named**: the real name of the company or institution, the date, and a verifiable event. Where naming is not possible, it must be explicitly labeled "**Composite profile | constructed from a common industry pattern, not a named event**," and "a certain company A / B / C" — fake specificity — is not allowed.
2. **Traceable**: the original link actually opens when tested, and the data timestamp is written down. Platform rules change by the month; a number with no timestamp is no number at all.
3. **Tiered**: each item states the nature of its source, with speculation listed separately; when several sources give different figures for the same thing, the differences are set side by side rather than picking the one that looks best.
4. **Bounded**: every positive case comes with a failure condition, a customer type it does not suit, or a counterexample. For a signal to be credible, the costs borne by senders must differ enough; a case that bears no cost is equivalent to no signal [official].

**Example**: The case "we took a client from 67 to 2100+ signups a month," presented on its own, has publicity value only [vendor claim]; attach its preconditions — that the product itself could be copied, that growth rested on a programmatic page matrix, and that this is exactly the shape Google named explicitly after March 2024 — and it becomes evidence that peers and buyers can jointly examine. For a counterexample, take the four criteria NewsGuard uses to track AI content farms (clear evidence that a site's content is largely AI-produced, published with no meaningful human oversight, presented in a way that makes ordinary readers think it came from human journalists, and not clearly disclosed), which its own site says have identified 3,749 sites across 16 languages [official]. Another kind of cost arrives faster: in April 2025, users of the AI code editor Cursor received a reply signed "Sam" saying that "one subscription can only cover one device" was a new security policy — the policy was pure fabrication, and Sam was a bot; users took it as an official change and publicly announced they were cancelling their subscriptions, and co-founder Michael Truell subsequently apologized, attributing the incident to hallucinations and non-deterministic output [third-party verified].

**Pitfalls**: Putting up only the successes and leaving out the preconditions. Fabricated "facts" cost money outright: the FTC's Rule on the Use of Consumer Reviews and Testimonials was adopted on 2024-08-14, explicitly banning fake or

fabricated consumer reviews, reviews exchanged for compensation with a particular slant, insider reviews without disclosure, and company-controlled review sites falsely claiming to be independent [official]. In April 2026, the FTC won a 4 million USD judgment against the supplement brand TruHeight and its principal, a case that included thousands of five-star reviews on its website actually written by employees, free or discounted products exchanged for positive reviews, and software-generated reviews posted through bot accounts disguised as real users [official]. In the same period the FTC ran five parallel cases under Operation AI Comply, and DoNotPay paid 193,000 USD to settle claims made without evidence [official].

## Method 5: Put a calendar behind the gathering, and grade on evidence only

**What to do**: Turn first-party data into a pipeline rather than a one-off project, and switch the grading from output to evidence.

**How**

1. Put four things on the calendar: a quarterly "own-data study" with method, sample, and basis; a monthly maintenance pass on a list of "facts we assert externally" (version, owner, change log), so that both people and AI read the correct version; a weekly log of customer questions; and a self-run experiment every six months to verify whether the output of the first three actually enters the buyer's decision process. The marginal cost of experiments can be pushed down — one landmark company does this by running more than 20,000 controlled experiments per team per year [official].
2. Change the KPIs. Demote "how many pieces published, how many keywords covered, how many reads"; promote the number of new proprietary data points per month, the increase in citations, the error rate in how AI describes you, and the pass rate of case-library links under testing. A single month's reading means nothing; use the same set of questions and watch the trend on a rolling basis — one third-party analysis notes that 40%–60% of the sources cited in AI answers are replaced month by month [third-party verified].

3. For material generated in bulk with AI, from September 2025 there is no "unlabeled" option: the Measures for Labeling AI-Generated Synthetic Content took effect on 2025-09-01, requiring AI-generated text, images, audio, and video to carry both explicit and implicit labels, with distribution platforms verifying the labels at the listing and publishing stage [official].
4. Do not mistake quantity for maturity — 50 experiments with no hypothesis are worth less than 10 aligned with strategic goals [vendor claim]. The same lesson holds when transferred to content.

| Old KPI (demoted) | Evidence KPI (promoted) |
|---|---|
| Content output (pieces/ items) | New proprietary data points per month, cases filed into the library |
| Keywords covered | AI citation frequency, share of cited pages containing first-party data |
| Reads, impressions | Error rate in the facts AI restates about you |
| Publishing frequency | Share of cases meeting all four requirements, pass rate of links under testing |

**Example**: The same batch of content on two kinds of domain ends completely differently. The SEO platform SE Ranking first published 6 AI-assisted articles on its own authoritative, human-edited blog, for about 555,000 impressions and 2300+ clicks, 4 of which Google AI Overviews cited as sources; it then used "one-click generation, zero human involvement" to publish 2000 pure-AI articles across 20 brand-new domains — in the first 36 days, 70.95% of pages were indexed and 28% briefly entered the top 100, but after about 3 months the top-100 share fell from 28% to 3%, had still not recovered 16 months later, and totaled only 1381 clicks [third-party verified]. The policy side is that in March 2024 Google upgraded its old policy into "scaled content abuse" — content produced in bulk to manipulate rankings and of very low value to users is treated as spam regardless of whether it is generated by automation, humans, or a combination; Google says actual removals reached 45% after the update [official]. The direction has already carried named costs in the content industry: Stereogum's founder wrote in November 2025 that Google's shift

toward AI search cut their ad revenue by 70% [official]; and Chegg admitted in May 2023 that ChatGPT was affecting its new-customer growth, with its share price closing down 48.41% the next day [third-party verified].

**Pitfalls**: Assuming "publishing a lot" is a competitive advantage, or conversely that AI content is bound to fall. Both are wrong: what gets cleaned out is the combination "brand-new domain + zero human involvement + pure AI"; content on an authoritative domain with a record of human verification kept getting traffic and AI citations. There is one red line only — scaled production to manipulate rankings.

## Notes for this chapter

- This chapter does not cite conclusions such as "AI content is bound to be de-ranked." The same content gets AI citations on an authoritative, human-edited domain and survives about three months on a brand-new, unedited one — the difference lies in the two variables "domain authority + human involvement," not in whether AI was used. [third-party verified]
- The buyer's position can be stated through that 1970 paper on the "lemons" market: unable to tell real from fake, he can only price all content at average credibility, and that average is being pulled down — the people putting out inferior goods squeeze out the whole market, the rule-following businesses along with it. [official]
- Industry differences are folded into the methods: the case libraries and risk lists of B2B and industrial businesses function, in practice, as a one-time alignment carried out on the buyer's behalf inside their own organization — a Gartner 2025 survey says 74% of B2B buying teams show "unhealthy conflict" during the decision process, while buying groups that reach consensus are 2.5 times more likely to report a "high-quality deal" than those that do not [vendor claim]; the cheapest first-party asset of a tool product is in-product behavior and integration data; regulated industries such as education and finance can publish the least, so their first-party assets go first to improving decisions internally, and externally they release only de-identified aggregate conclusions confirmed by legal. [our judgment]

- For the China-side points on "AI poisoning" (feeding large models with fake reviews in bulk) and the Self-Discipline Convention for the Generative Engine Optimization Industry, the draft found only vendor press releases quoting them, and could not verify the original CCTV broadcast material or the list of convention signatories, so this chapter does not cite them and gives no figures.
- The point that "disclosing the use of AI lowers engagement" appears only as media retelling, and the authors and year written in the Journal of Consumer Research study cited were not checked against the original journal, so it serves as direction only. A brand making "no AI" a selling point (for example Aerie re-pushing "100% Aerie Real" in spring 2026, with quarterly sales up 23% year on year) is an isolated phenomenon in a high-trust category and cannot be extrapolated. [third-party verified]
- Compliance standards are evolving fast: the rules on automated decision-making and outbound data transfers in personal information protection, the FTC's enforcement posture on the review rule, and the implementation details of the AI content labeling measures are all changing. Before putting anything into practice, go by the current regulatory text. [official]

## Sources for this chapter

[1] George A. Akerlof | The Market for "Lemons": Quality Uncertainty and the Market Mechanism | The Quarterly Journal of Economics, Vol. 84, No. 3 | 1970-08 | https://econweb.ucsd.edu/~atimmerm/174/akerlof.pdf | [official] [2] Michael Spence | Job Market Signaling | The Quarterly Journal of Economics, Vol. 87, No. 3, pp.355-374 | 1973-08 | https://www.sfu.ca/~allen/Spence.pdf | [official] [3] Ron Kohavi, Diane Tang, Ya Xu | Trustworthy Online Controlled Experiments | Cambridge University Press | 2020-03-13 | https://www.cambridge.org/core/books/trustworthy-online-controlled-experiments/D97B26382EB0EB2DC2019A7A7B518F59 | [official] [4] Google | Google's March 2024 core update and new spam policies (scaled content abuse) | 2024-03-05 | https://blog.google/products-and-platforms/products/search/google-search-update-march-2024/ | [official] [5] FTC | FTC Announces Final Rule Banning Fake Reviews and Testimonials (16 CFR Part 465) | 2024-08-14 | https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-trade-commission-announces-final-rule-banning-fake-reviews-testimonials | [official] [6] FTC | FTC Takes Action Against TruHeight (4 million USD judgment) | 2026-04-13 | https://www.ftc.gov/news-events/news/press-releases/2026/04/ftc-takes-action-against-truheight-deceptive-unsubstantiated-advertising-

supposed-height-enhancing | [official] [7] FTC | FTC Announces Crackdown on Deceptive AI Claims and Schemes (Operation AI Comply) | 2024-09-25 | https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes | [official] [8] NewsGuard | AI Tracking Center: 3,749 AI content farm sites (four criteria) | last updated 2026-06-23 | https://www.newsguardtech.com/special-reports/ai-tracking-center/ | [official] [9] SE Ranking | AI Content Experiment (16 months, 20 new domains × 100 articles) | https://seranking.com/blog/ai-content-experiment/ | [third-party verified] [10] eMarketer | FAQ on GEO and AEO: Where AI Search and SEO Overlap in 2026 | 2026-04-02 | https://www.emarketer.com/content/faq-on-geo-aeo--where-ai-search-seo-overlap-2026 | [third-party verified] [11] CMSWire | What Is a Customer Data Platform (CDP) — 2026 market guide | 2026-02 | https://www.cmswire.com/customer-data-platforms/what-is-a-customer-data-platform-cdp/ | [third-party verified] [12] House of MarTech | MMM vs MTA: Privacy-First Measurement Framework for 2026 | 2026-04 | https://houseofmartech.com/blog/marketing-mix-modeling-vs-multi-touch-attribution-privacy-first-measurement-framework-for-2026 | [third-party verified] [13] OmniFunnel Marketing | Building a Zero-Party Data Collection Strategy | 2026-01 | https://www.omnifunnelmarketing.com/blog/building-zero-party-data-collection-strategy-customers-trust | [third-party verified] [14] Bluebarry | Zero-Party Data: Complete Guide (Forrester definition) | https://bluebarry.ai/resources/zero-party-data | [vendor claim] [15] Salt Agency | How Zapier Quadrupled Organic Growth | https://salt.agency/blog/how-zapier-quadrupled-organic-traffic/ | [third-party verified] [16] Sean Ellis / Glasp | How Glasp Grew ChatGPT Traffic from 500 to 19,000 Daily Sessions | https://seanellis.substack.com/p/how-glasp-grew-chatgpt-traffic-from | [third-party verified] [17] QuickSEO | Programmatic SEO Stats 2026 (Wise 54 million visits vs G2 80% loss) | 2026-05-12 | https://quickseo.ai/blog/programmatic-seo-stats-2026-is-pseo-still-viable-in-the-ai-search-era | [third-party verified] [18] Omnius | Programmatic SEO Case Study: From 67 to 2100 Monthly Signups | https://www.omnius.so/blog/programmatic-seo-case-study | [vendor claim] [19] Ars Technica (Benj Edwards) | Company apologizes after AI support agent invents policy | 2025-04-17 | https://arstechnica.com/ai/2025/04/cursor-ai-support-bot-invents-fake-policy-and-triggers-user-uproar/ | [third-party verified] [20] Stereogum (Scott Lapatine) | Getting Killed By AI | 2025-11-10 | https://stereogum.com/2478838/stereogum-relaunch/news | [official] [21] CNBC (Sarah Min) | Chegg shares drop more than 40% after company says ChatGPT is killing its business | 2023-05-02 | https://www.cnbc.com/2023/05/02/chegg-drops-more-than-40percent-after-saying-chatgpt-is-killing-its-business.html | [third-

party verified] [22] Inc. (Emily Cody) | 'No AI' Is the New 'All Natural' | https://www.inc.com/emily-cody/no-ai-is-the-new-all-natural-and-brands-are-winning-big-with-it/91375451 | [third-party verified] [23] China Economic Net (Chinese: 中国经济网, via CCTV.com) | AI-generated synthetic content labeling measures take effect 2025-09-01 | 2025-09-01 | http://www.ce.cn/xwzx/gnsz/gdxw/202509/t20250901_2460291.shtml | [third-party verified] [24] Jiazi Guangnian (Chinese: 甲子光年) | Report on digital customer acquisition and Account-Based Growth practice at industrial and manufacturing companies | https://www.jazzyear.com/newsdetail?id=38492019 | [third-party verified] [25] Fenxiang Xiaoke (Chinese: 纷享销客) | Growth-oriented marketing cloud: field cases from seven leading B2B companies (cross-channel entry work −34%, overseas leads to opportunity +29%) | https://zhuanlan.zhihu.com/p/2062495512356507782 | [vendor claim] [26] ZKH (Chinese: 震坤行) | "AI product-recommendation brain" brings more than 100 million yuan in revenue growth in 11 months | https://zhuanlan.zhihu.com/p/1954120314318481065 | [vendor claim] [27] Gartner | 74% of B2B Buyer Teams Demonstrate Unhealthy Conflict During the Decision Process | 2025-05-07 | https://www.gartner.com/en/newsroom/press-releases/2025-05-07-gartner-sales-survey-finds-74-percent-of-b2b-buyer-teams-demonstrate-unhealthy-conflict-during-the-decision-process | [vendor claim] [28] Growth Method | The 5 Stages of Experimentation Maturity | 2026-05-08 | https://growthmethod.com/experimentation-maturity-model/ | [vendor claim] [29] Cyberspace Administration of China (Chinese: 国家互联网信息办公室) | Personal Information Protection Law of the People's Republic of China, Article 24 (transparency of automated decision-making and convenient means of refusal) | effective 2021-11-01 | https://www.cac.gov.cn/2021-08/20/c_1631050028355286.htm | [official]

# Chapter 15 Let the Product Acquire: Free Tools, Trials, and Open Interfaces

Winning a customer used to start with a person moving first: sales made calls, the media buyer bought ads, the writer published content, and the customer was pushed along. Now a share of customers find you themselves — they first run their own numbers through your calculator, or finish one real piece of work inside a trial, and only once it works smoothly do they turn around and ask "which company are you?" Customers like this open with a specific question; nobody exchanges pleasantries and nobody sits through your company introduction from the top.

The customer may have used your thing before ever knowing your company. That is why this chapter gives you 5 methods that bring in leads without a person making a move.

---

## Method 1: Build a small tool so people use it first and leave contact details after

**What to do**: Turn the one question you are asked most often into a free small tool, so the customer gets the result first and you talk about leaving contact details only after that.

**How**

1. From sales and customer-service chat logs, pick the question that is **asked most often and has the most standardized answer** — a quote estimate, a unit conversion, a quota calculation, a compliance check — and choose one. Many people asking means many people searching, so this tool's entry point is naturally open.
2. The tool does one thing only: three or four inputs in, one number or one checklist out. The fewer inputs the better; every extra one drops a cohort of people [our judgment]. Default whatever can be defaulted, and anything the customer can pick from options should not be left for him to type.

3. Fix the order: enter inputs → see the result → want to save the result or see the next layer → only here does the contact-capture box appear. The step of asking for contact details comes **after** the result, not before it.
4. Give an output that can be taken away: a PDF, an image, a conclusion that can be copied. The people willing to save it are the ones genuinely interested.
5. Look at only two numbers each month: how many people used it, and of those who finished, how many left contact details. The second number is this hook's actual score; if the first looks good and the second does not move, the problem is on the results screen.
6. Put the tool on a separate domain or a separate path, not inside a pop-up on your homepage. Its job is to be found by search and shared by people, and a pop-up can do neither [our judgment].

**Example**: HubSpot's Website Grader says officially that it has scored more than 2 million URLs [official]; Ahrefs has turned free tools into a full-page matrix [official]; Shopify's free name generator page says outright, No Sign-Up Required, and only guides you to a trial after you have used it [official]. There is a worked example of "get them using it before you take a contact detail" for AI products too: Grok Bot did not launch with paid traffic but with one-to-one manual onboarding of nearly 300 early users (a video call each), with anything a user got stuck on fixed the next day; 99% of the automation tasks on the platform end up being created in natural language, and users never see a configuration screen [third-party reported: product lead interview].

**Pitfalls**: Turning the tool into a shrunken product-brochure page where the input fields exist only to collect information — the customer fills in half of it, finds there is no result to see, and leaves.

## Method 2: Open up the interface so machines can call your capability

**What to do**: Package your core capability as an interface on an open standard so someone else's AI can call it directly, and the caller becomes one of your lead sources.

**How**

1. Work out, on an open standard, what your capability should be exposed as. Taking MCP as the example, the specification splits the roles into Hosts (the AI application that initiates the connection), Clients (the connector inside the host

application), and Servers (the service that provides context and capabilities), and a server offers three kinds of things outward: Resources (data and context), Prompts (templated messages and workflows), and Tools (executable functions) [official]. The question you have to answer is "which tool function my capability is packaged as, how the parameters are filled in, what it returns," not "how my web page is laid out."

2. Write the interface documentation to the standard of "written for a machine": the capability name, the meaning and value range of each parameter, the fields returned, and the message on error. The caller's AI decides whether to use you from this documentation; a usage the documentation does not describe is one it will not guess at.
3. Connecting once is enough. The documentation compares this kind of protocol to the USB-C port of AI applications — connect once and multiple clients can use it [official]. So pick the mainstream standard; do not set up one of your own.
4. Open the read-only part first. Query, calculation, and comparison capabilities that do not touch the other party's data have a low authorization bar and get connected first; write capabilities come later, since the customer side also needs a bit more time before daring to grant you permission.
5. Leave a recognizable identity in the service name and in the returned values. When you are called you do not know who is on the other side, but the other side's people and the later invoice need to be able to recognize you.
6. Keep the version number in the interface address and give old versions a transition period. Change an interface once and everyone already connected breaks outright, and you will receive no complaints at all — they simply stop calling.

**Example**: The current version of the MCP specification is 2026-07-28; the specification defines the protocol as an open standard between AI applications and external data sources and tools, with messages carried over JSON-RPC, and the three kinds of capability a server exposes are data, prompt templates, and tool functions [official]. The getting-started documentation compares it to "the USB-C interface of AI applications" and lists the clients that already support it — Claude, ChatGPT, Visual Studio Code, Cursor, and others — so you write it once and can connect in many places [official].

**Pitfalls**: Reading "opening an interface" as "putting up an API document." A document left sitting there is read by nobody; it only counts once it can actually be called.

## Method 3: Publish the service into official directories so agents find you on their own

**What to do**: Publish the service into official registries and platform app directories, so the buyer's agent and the platforms pull you up when they choose vendors.

**How**

1. Keep the two kinds of directory apart. One is a **protocol-level registry** (such as MCP's official registry), where what gets published is metadata, so that clients and aggregators can discover which services are available; the other is a **platform app directory** (the app marketplaces of payment platforms and collaboration platforms), which catalogs individual app entries, and that is where customers go to choose. Get into both; they do different jobs.
2. Prepare the metadata the publication requires: the service's unique name, its install location (package name or remote address), start-up and configuration instructions (command-line arguments, environment variables), and a capability description. These are the fields you cannot get around at registration; miss one and you cannot publish.
3. Take a name in reverse-domain-name form, and be able to prove that the namespace is yours — verify ownership with your domain or your code-hosting account. The review clears only when the name and the ownership match; if they do not, you are holding a place for someone else.
4. The bar for publication is usually "the install method is publicly installable" or "the service itself is publicly reachable." Private services open only to the inside are generally outside the scope of what gets cataloged.
5. Do not stop once you have published: **the downstream aggregators and marketplaces are where customers actually look.** A registry gives unsorted metadata; a marketplace adds ratings and recommendations on top, so the material has to be filled in one mainstream marketplace at a time.

6. Write the blurb in the directory entry in the words a customer would actually type, not your internal phrasing. How well this paragraph is written directly decides whether you show up in that search.

**Example**: MCP's official registry is that central directory: publicly reachable services can publish their metadata into it, and clients and aggregators discover which services are available through its interface [official]. Stripe's Directory documentation is the most direct: it lets developers and AI agents find the external service provider best suited to a given task, and what it indexes includes apps listed in app marketplaces, providers that can be switched on with a single command, and a directory of interfaces that support pay-per-call, with search results carrying structured fields such as the provider's identity and interface address [official]; the documentation also gives the way to write — use in the blurb the words a customer would actually type, so "help freelancers invoice and get paid faster" is more findable than "next-generation financial operations platform"; use complete phrases rather than pulling the words apart and scattering them; and the display name and account name must match what customers call you, or someone searching your name directly will not find you either [official]. The directory layer converged in 2026: one guide's judgement is that the official MCP Registry (registry.modelcontextprotocol.io) is the authoritative directory and the first place to publish — push a small server.json with the CLI, prove namespace ownership via GitHub login or a DNS record, and downstream directories and newsletters ingest the official registry, so one publication spreads outwards; the metadata to keep consistent includes name, display title, a capability description under 100 characters, transport, authentication method and a detailed description of each tool, and the JSON-RPC initialize handshake should be verified before publishing [third-party verified]. The directories are themselves layered: the official Registry does not host servers, Smithery and Glama do, and MCP.so is a community-run third-party server directory — the same server list reaches entirely different kinds of buyer depending on where it is listed [third-party verified].

**Pitfalls**: Assuming that having your name listed in a registry equals leads. A directory gives you one chance to be found; whether you are chosen depends on the description in the entry and the reputation outside the directory [our judgment].

## Method 4: Write the boundaries of the trial and the free tier in stone

**What to do**: How far free goes, when charging starts, and how the allowance is calculated all get fixed in advance, so the trial period does the filtering itself.

**How**

1. Give free on one dimension only: either time or count. Give both and the customer cannot work out how much is left, and starts using it timidly.
2. Write out "what happens when you run out" before the first use: how much allowance remains, whether running out stops the service or auto-charges, and how much it charges. What customers fear most is receiving an unexpected bill as the trial nears its end, and that kind of dispute is more expensive than no conversion [our judgment].
3. Get the customer to accomplish one real thing inside the trial rather than clicking through every feature. The test is "has he run his own data through it once" — clear that step and paying becomes possible; not before.
4. Follow up on two separate tracks. Those who used it but did not hit the allowance and those who ran out and stopped need entirely different messages: the first lack a reason to use it, the second lack a budget.
5. During the trial, log the nodes — activation, drop-off, hitting the allowance — and see which screen people stick on. When you are adjusting the product, the sticking points are more useful than a vague conversion rate.
6. Converge the free-tier terms and the discount policy into one place. If the same company presents two sets of trial rules, customers will screenshot them and compare, and sales cannot explain which one actually applies.

**Example**: On 2022-11-28 Heroku took down free dynos, free Postgres, and the free Key-Value Store all at once [official]; Railway first shut down its free plan to stop the bleeding, then reopened it in the form of 5 USD / 30 days + 1 USD / month [vendor claim]. The trade-off between free tiers and trials has a comparable benchmark set: 2026 PLG data compiled from eight benchmark reports has about 58% of B2B SaaS running a PLG motion and 91% planning to invest more, yet only about 34% tracking activation, the metric that best predicts paid conversion; free trials convert about 17% of signups to paid against about 5% for freemium, at the cost of a

lower visitor-to-signup rate (freemium draws about 6% of visitors against 3%–4% for trials); product-qualified leads close at about 25%–30% against 5%–10% for MQLs, a 3–5x difference; and about 67% of companies above $10 million in ARR run a hybrid of PLG and sales-led motion [third-party verified].

The activation stage has usable baselines: across 62 B2B SaaS companies average activation is 37.5% (median 37%), meaning about two-thirds of new sign-ups never reach core value; Amplitude benchmarks covering more than 2,600 companies show that new users who hit no value milestone within two weeks churn at over 98%; the strongest cross-period predictor in that dataset is the “7% rule” — if 7% of a cohort returns on day seven the product sits in the top quartile, and about 69% of products with strong day-seven retention are still strong at three months [third-party verified]. Average time to value is about 1 day 12 hours across 547 companies, with wide category spread: AI/ML products are measured in hours (median around 17), HR products in days [third-party verified]. For trial design the implication is that the free-tier boundary belongs before the first value moment, not after the last feature has been clicked.

**Pitfalls**: Setting the free allowance to "just short of finishing one thing." After the trial the customer remembers exactly one thing — you cannot get anything done there.

## Method 5: Turn the tool's result page into a lead-capture page

**What to do**: Let the result the tool calculates go find the next customer on its own, and catch the new visitor along the way.

**How**

1. The result page gives three things: the conclusion, the basis, and the next step. The basis is "how this number was calculated," and tools that can explain the algorithm and the parameter ranges clearly get shared the most [our judgment].
2. Make the result takeable: a conclusion that can be copied, an image, a PDF. When the customer takes it away, his own inputs go with it, and your name goes with it too.

3. Make the result pages findable by search and readable by AI. Generate a static page for each common input combination, with the title being the question the customer would search and the body laying out the input conditions and the calculated conclusion. When someone searches the same conditions, what they find is your page, not someone else's tool.
4. Leave an entry point on the sharing path: whoever receives the image or the link can open it and run the calculation again with their own conditions. The person who calculates a second time is the new lead you just won.
5. Leave one spot at the bottom of each result page for "see the detailed version," and ask for contact details only there, once — do not reach out from every corner of the page.
6. Do not put a pile of navigation on the results page. The customer arrived with one specific question; this page answers only that question and leaves one exit.

**Example**: Shopify's free name generator strings "generate a result — pick a domain — start a trial" into a single path, so the results page is itself the lead-capture page [official]; HubSpot's results page gives a score plus item-by-item advice plus a competitor comparison, then guides to the next step [official].

**Pitfalls**: A results page with nothing but a line saying "congratulations on completing the assessment" and a QR code. If the customer gets nothing he can take away, the page was made for nothing.

---

## Notes for this chapter

- This chapter **cites no conversion percentages for trials or free tools**. Every such number that could be found in the materials (free-trial-to-paid rates and the like) came from aggregation blogs with no original source, so none is used, which is why methods 1, 4, and 5 give the how and not the result. To find out whether it works, keep your own books the way Chapter 20 does.
- The open-capability directories of domestic platforms (WeChat Mini Programs (Chinese: 微信小程序), the Feishu app marketplace (Chinese: 飞书应用市场), the DingTalk open platform (Chinese: 钉钉开放平台), the Coze store (Chinese: 扣子商店)) were not captured this round, and neither was Build with the Apps SDK in

OpenAI's help center (curl returned 403). So the practice in method 3 follows the public documentation of MCP and Stripe; moving to another platform means walking through its own listing rules again.

- There is currently no authoritative standard for designing a tool hook. Methods 1 and 5 are written as workable practice, not as the standard answer.
- Methods 2 and 3 cite official documentation and describe **mechanisms**, not promised results. They say how something gets called and how something gets listed; they do not say how many leads doing it will bring.
- Think through the authorization boundary before opening your capability up. The MCP specification writes user consent, data access control, and authorization before tool invocation as principles the implementer must handle, not as options [official] — opening an interface carries a compliance cost.

## Sources for this chapter

[1] Model Context Protocol | Specification (2026-07-28 version) | https://modelcontextprotocol.io/specification/2026-07-28 | [official] [2] Model Context Protocol | Official registry documentation | https://registry.modelcontextprotocol.io/docs | [official] [3] Model Context Protocol | Getting-started documentation | https://modelcontextprotocol.io/docs/2026-07-28/getting-started/intro | [official] [4] Stripe | Directory documentation | https://docs.stripe.com/directory | [official] [5] HubSpot | Website Grader has scored 2 million+ URLs | https://www.hubspot.com/blog/bid/5539/website-grader-analyzes-over-2-million-websites | [official] [6] Ahrefs | Free SEO Tools (free-tool matrix) | https://ahrefs.com/free-seo-tools | [official] [7] Heroku | Removal of Heroku Free Product Plans (effective 2022-11-28) | https://help.heroku.com/RSBRUH58/removal-of-heroku-free-product-plans-faq | [official] [8] Shopify | Free Business Name Generator (results page to trial) | https://www.shopify.com/tools/business-name-generator | [official]

[9] RoxyAPI | MCP registries: where to list your server (2026) (official MCP Registry as the authoritative directory; server.json metadata fields; one publication ingested downstream) | https://roxyapi.com/blogs/mcp-registries-where-to-list-your-server | [third-party verified] [10] TrueFoundry | Best MCP registries (official Registry does not host; Smithery/Glama host; MCP.so community directory; TrueFoundry VPC-native) | https://www.truefoundry.com/blog/best-mcp-registries | [third-party verified] [11] Growth Hacker AI Weekly (compiling a Lenny's Podcast interview with Grok Bot

product lead Roman Ugarte) | Manual onboarding of nearly 300 early users, 99% of automation tasks created in natural language, mass unshipping in the three weeks before launch | https://mp.weixin.qq.com/s?__biz=Mzg3NDc2MjQxMg==&mid=2247495217&idx=1&sn=ebb07d8a4c70889e76ed6ea750cb65dd | [third-party reported: product lead interview]

[12] Digital Applied | Product-Led Growth 2026: The PLG Strategy Playbook (58% of B2B SaaS run PLG, only 34% track activation; free-trial signup-to-paid 17% vs freemium 5%; visitor-to-signup about 6% for freemium vs 3%–4% for trials; PQL conversion 25%–30% vs MQL 5%–10%; 67% of companies above $10 million ARR run a hybrid motion) | 2026-06-17 | https://www.digitalapplied.com/blog/product-led-growth-2026-plg-strategy-playbook | [third-party verified] [13] Digital Applied | Time to Value: The 2026 SaaS Onboarding Metrics Framework (62 B2B SaaS companies averaging 37.5% activation; over 98% churn within two weeks without a value milestone; the 7% rule; average time to value about 1 day 12 hours across 547 companies; AI/ML median around 17 hours) | 2026-05-30 | https://www.digitalapplied.com/blog/customer-onboarding-time-to-value-2026-saas-metrics-framework | [third-party verified]

# Chapter 16 Let AI Answer Around the Clock

During the day someone watches the back office and takes the inquiries one by one; after hours the phone goes silent, and a customer who can't reach anyone goes off to the next vendor. At this step the customer has already come to you on their own — if you can't catch them, the visit was wasted. So this chapter gives you 5 methods to turn "those inquiries after hours" into leads.

---

## Method 1: Open the reception desk on official channels first

**What to do**: Don't hire a third-party relay and don't modify the client — use official channels like WeChat Customer Service and WeCom to open a reception point of your own where a link can be attached.

**How**

1. First count how many places you can be found from: the footer of your website, landing pages, the official account menu, business cards, platform store pages. Each of these places is missing a "tap and you can talk" entry point — fill it in.
2. Open the reception point on official channels. WeChat Customer Service has public API documentation that includes interfaces such as "get customer service account link"; once you have the link you can attach it to your website and landing pages [official].
3. On the WeCom side, go through the developer center and first understand the development boundaries clearly — what you can do, what qualifications are required, which step needs enterprise verification — then decide who will configure it [official].
4. For bringing the bot in, use an existing integration path. Alibaba Cloud (Chinese: 阿里云) Intelligent Dialogue Robot (Tongyi edition) has a "WeChat Customer Service Deployment Guide" that describes exactly how a bot connects to this WeChat Customer Service gateway [official].
5. After configuring, act as a customer yourself once. Take a phone that has never followed you, tap in from the landing page, ask three questions, and see whether you were caught, how long it took, and whether the one who caught

you was a person or a bot. A cruder acceptance standard works better: within half a minute a person or a bot speaks first [our judgment]. If you can't do that, the link is broken, the hours are misconfigured, or no one has claimed this gate at all.

6. Write down "who watches this gate during which hours." Skip this step and every automation later becomes an auto-reply that no one answers. What is attached at each gate and who maintains it — check them off against this table:

| Reception point | What is attached | Who maintains |
|---|---|---|
| Website footer | Customer service account link + service hours | Marketing |
| Landing page | The same link, with a button labeled "you can ask outside business hours" | Marketing |
| Official account menu | Put the customer service entry on the first level | Operations |
| Business cards and email signatures | QR code | Sales |
| Platform store page | In-platform customer service entry + the same script | Operations |

**Example**: WeChat Customer Service's official documentation lists "get customer service account link" as a separate interface entry [official]. Its practical use is narrow but hard: it lets you attach the reception entry to any page of your own, instead of only inside WeChat.

The voice entry point is now a deliverable you can hand over whole: during last year's holiday surge, Amazon Ring evaluated more than 40 voice AI vendors before choosing Vapi, and now routes 100% of its inbound calls through the platform; Ring's VP of software development says customer satisfaction scores rose after deployment and that business teams can adjust the AI agent experience without engineering; Vapi has handled over 1 billion calls cumulatively and runs 1–5 million a day [third-party verified]. Its differentiator is not a packaged application but the infrastructure and orchestration layer for voice agents — which moves the buying criterion from recognition accuracy to control and tunability.

**Pitfalls**: Opening the reception point on a personal account or on a third-party relay client. The more gates you open the more it looks like a real business, but only official channels survive changes in platform rules, and if something goes wrong you won't lose the account along with it [our judgment].

## Method 2: Build the knowledge base in your own hands first

**What to do**: Write the Q&A the bot will use into a sheet of your own first, then copy a version into the platform back end.

**How**

1. Copy from real conversations of the past three months: the exact words customers asked, the exact words support replied. Copying this way is far better than rewriting your website copy — nobody can finish reading answers rewritten from marketing language. When copying, pick only two kinds of sessions — those where the customer asked and then dropped off, and those where the customer asked and then closed. The first tells you which sentence drove them away, the second tells you which sentence worked.
2. One answer answers only one thing. Split "can you issue an invoice" and "how soon can it ship" into two entries. A bot that answers off-topic is, in the great majority of cases, not a bad model — it's a knowledge base with three things packed into one entry.
3. Behind every answer, attach two things: the source (which page, which clause) and the owner (who reviews it — write a name or a desk, not a department).
4. Start the answer's phrasing from the customer's own words. If a customer asks "can you issue a special VAT invoice for this," the answer starts with "yes" — not with "our company is a general taxpayer." If the customer scans and can't see the two words they want, they'll ask again, or just leave [our judgment].
5. Fill the fields according to this table, one row per entry:

| Field | What to fill in |
|---|---|
| Customer's exact words | The sentence copied from a real conversation |
| Standard answer | Answers only one thing |
| Source | Page / clause / internal document ID |

| Field | What to fill in |
|---|---|
| Owner | Who reviews |
| Last updated | Date |

1. Keep the master file in your own document; the one in the platform back end is a duplicate. When you change tools or platforms later, the master file doesn't need to be rewritten. Update once a month, adding the Q&A that newly emerged that month; it's fine if it's incomplete — anything not written automatically falls into the human category.

**Example**: That Alibaba Cloud deployment guide gives the official path for connecting a bot to the WeChat Customer Service channel [official]. What's worth noting here is that the bot and the channel are two separate things — the channel belongs to the platform, the Q&A belongs to you. There is an extreme sample of that judgment: Intercom built its own AI support agent as a purpose-post-trained small model, Fin Apex 1.0, reporting a 73.1% resolution rate — the share of issues fully resolved without human intervention — against 71.1% for both GPT-5.4 and Claude Opus 4.5 and 69.6% for Claude Sonnet 4.6, with responses in 3.7 seconds, 65% fewer hallucinations than the comparison model, and a unit cost about one fifth of calling frontier models directly; the model is served only inside its own product and is not sold separately [vendor claim: vendor benchmarks / third-party reported]. It also disclosed its own support numbers in the same period: the AI agent resolves more than 81% of Intercom's own support volume, which came with a new "Conversation Designer" role and skills-based routing that sends customers asking for a human straight to someone with the right skills [vendor claim: the team's own account]. Owning the knowledge base and the evaluation loop is what produces those numbers. The extreme case of handing reception to AI is giving every customer its own agent: Latin American used-car platform Kavak generates a dedicated agent for each customer, with its own micro virtual machine, remembering years of interaction and setting long-term goals. The operating numbers: about 96% of customer interactions and about 95% of transactions are handled by agents with no human involvement, 100,000–200,000 agents are instantiated per day, customer satisfaction/NPS is about 3× higher, conversion has gone from 50% above the human team at first to more than 2.1× today, and car-loan approval went from the industry-typical two months or more to about three minutes [founder's own account]. The counter-case belongs in the same

place. Klarna launched an OpenAI-powered support assistant in February 2024; in its first month it handled 2.3 million conversations, automated about 67% of them and did the work of roughly 700 full-time agents, cutting average resolution time from 11 minutes to under 2 minutes; by Q3 2025 the official figures were the work of about 853 agents, roughly $60 million saved annually, a customer NPS of 73 and cost per transaction down from $0.32 to $0.19; but from May 2025 it began hiring agents back, citing hallucinations on edge cases, falling satisfaction on complex and emotional tickets, and compliance concerns about letting AI handle disputes and account closures, with the CEO publicly acknowledging the company had cut too far [third-party verified: case teardown].

**Pitfalls**: Treating the knowledge base as a configuration item inside the platform back end. The back end stores only a copy; the master file is in your hands. This one point determines how many days you'll spend redoing it on the day you change tools [our judgment].

## Method 3: Sort intent into three tiers, then decide who answers

**What to do**: Give the bot a set of tiering rules, so that it decides when to answer on its own and when to hand over to a human immediately.

**How**

1. Sort questions into three tiers, write the rules into the configuration, don't leave them in people's heads:

| What the customer is asking | Who answers | What to record |
|---|---|---|
| Price, whether they can order now, how the contract is signed | Hand over to a human immediately | Exact words + contact info |
| Specs, whether it can integrate with existing systems, how it differs from others | The bot answers | Contact info + intent tier |

| What the customer is asking | Who answers | What to record |
|---|---|---|
| Looking at materials, looking at cases, simply comparing prices | The bot answers | Record only, don't chase |

1. Write out clearly the hours for "hand over to a human immediately": during the day to whom, at night to whom, on weekends to whom. A single on-duty phone number or a group message is fine — the point is that this rule doesn't depend on some person remembering it.
2. When the bot hands over to a human, carry the customer's exact words over with it. Making someone repeat a question they just asked is the fastest way to lose a lead.
3. Look at the handover list once a day and look for two kinds of problem: cases that should have been handed over but weren't, and cases handed over in the middle of the night that no one picked up.
4. Unify the nighttime fallback script into one sentence: "I'm an automated receptionist; I've noted what you said, and someone will contact you tomorrow morning." Someone must actually contact them afterward, otherwise it's a pit you dug for yourself [our judgment].
5. When changing the tiering rules, change only one tier at a time, and watch it for a week. Change all three at once and even after you see the result you can't tell which tier did the work.

**Example**: Intercom's Fin AI Agent FAQ is a page where the vendor describes its own product [vendor claim]. It belongs to the kind of material Chapter 9 dealt with: the product's claims can be read, but its performance numbers can't be used as your own budgeting assumptions. The third tier's "record only, don't chase" action is the same idea as Chapter 17 placing a portion of people explicitly into a do-not-disturb zone. An auditable scale metric is starting to appear: Ping An (Chinese: 中国平安) disclosed token usage for the first time in its 2026 interim report — daily consumption rose from about 30 billion in December 2025 to over 120 billion in June 2026 (past 210 billion by August) — with AI seats handling about 939 million service instances, covering 81% of total customer-service volume and supporting RMB 57.313 billion in sales; ZhongAn Online (Chinese: 众安在线) reported 1,000 scenario-specific digital employees live in the first half, with token consumption above 34 trillion, over 10× year on year [third-party reported]. Vendor cases in customer service keep

accumulating: after adopting Decagon's AI support, Notion reports a 34% faster ticket resolution rate, doubled deflection and only 3.4% of conversations asking for a human [vendor claim]. How outcomes are defined changes the bill. Intercom's official documentation defines Fin's billable outcomes in detail: $0.99 covers a resolution (an AI answer is given and no further help is requested), a procedure handoff configured to end in a human or a workflow, a disqualification, and self-serve routing; matching a prospect to qualification criteria and routing them is $9.99; escalations, spam, pending, abandoned and timeout closures are not billed, and multiple actions in one conversation count as a single outcome [official]. The definitional trap sits in the same place: deflection rates that count abandoned conversations and partial answers can run well above resolution rates, and the same company's comparison document reports an average resolution rate of 76% across 8,000+ customers, with more than 20% of customers above 80% [vendor claim].

Readings at scale are starting to appear: on its Q1 FY2026 earnings call Salesforce disclosed Agentforce ARR above $1 billion; since going live on help.salesforce.com and its 1-800-NO-SOFTWARE line 15 months earlier, the AI had autonomously handled 4 million customer inquiries, double the volume of human agents; on the customer side, Florida Prepaid (200,000+ accounts) uses Agentforce Voice to handle 75% of calls autonomously during business hours and 100% outside them [third-party verified: vendor earnings call and customer accounts]. The reversal of subject and object on the Chinese side is more worth recording: Alibaba's AI Dianxiaomi, upgraded on 11 May 2026, claims official measurements of a 45% average drop in handoffs to humans, “AI plus human” enquiry conversion more than 10% above human-only service, answers to over 80% of questions across 30-plus scenarios, and a refund-retention success rate above 20% after sale; in brand cases, Xiaomi (Chinese: 小米) Tmall flagship saw handoffs down 45% with satisfaction up 22%, and Xtep (Chinese: 特步) saw handoffs down 55% with enquiry conversion up 46% [official: vendor announcement]. That describes humans assisting AI rather than the reverse — staff take high-value orders only after the AI has supplied the decision information and the nudge to close — and every figure comes from the vendor, with the denominator behind the handoff rate undisclosed. In vertical markets the same thing is already being counted as capacity: the dealership AI front desk Flai handles inbound calls, outbound campaigns and scheduling at once and books 50,000 appointments a month across sales and service, goes live 10 days after signing, and takes half of its new revenue from referrals by existing dealers [third-party verified: startup's own account].

**Pitfalls**: Letting the bot force an answer all the way through. When a customer asks the same thing twice, it means they no longer trust that answer; to keep circling at this point is to hold your own customer back for your competitor [our judgment].

## Method 4: Give AI a list of things it may not say

**What to do**: Write down as a list the things the bot must absolutely not decide on its own, and post it at the top of the knowledge base.

**How**

1. Three kinds of statement are not given to the bot: discounts and price concessions; delivery time, capacity, whether to squeeze in an order; compliance, certification, performance guarantees.
2. Behind each forbidden phrase, pair a replacement line, so the bot has something to say. When a customer pushes on price, reply "I can't decide the price myself; I'll arrange the person in charge and get back to you within ten minutes."
3. For answers involving commitments, add one sentence at the end: "subject to written confirmation."
4. Write the fallback rule in stone: when a customer asks the same thing twice in a row, hand over to a human directly, no more circling.
5. Each time you add a batch of new Q&A, read through it yourself first and pick out the newly emerged commitment-type statements to add to the list.

The list should be written so it can be posted directly at the top, with a replacement script after each line:

- Don't discuss discounts, don't discuss payment terms, don't discuss free gifts → "I can't decide the price; I'll arrange the person in charge and get back to you within ten minutes."
- Don't quote delivery time, don't quote capacity, don't promise to squeeze in an order → "The exact schedule has to be confirmed by production planning; I'll ask for you today."
- Don't make certification or performance guarantees on the product's behalf → "I'll have a technical colleague handle this with you; send me your requirements."

**Example**: This "pitfall" has been measured: multiple surveys show 53%–77% of respondents have had a bad or frustrating chatbot experience; in another survey covering 5,728 customers, 64% would rather companies didn't use AI customer service at all [third-party verified: academic / institutional research]. This is precisely why the "things you may not say" list exists. Chinese-side complaint data draws the line more sharply: the China Consumers Association received 985,928 consumer complaints in the first half of 2026, with after-sales problems the largest category at 26.79%, and what consumers reported centered on false promises by AI customer service, difficulty reaching a human, and inaccurate generated content; in one named case, the platform's official AI support agent told a buyer that the 28,700 yuan vehicle price already included transfer and registration fees, the buyer paid a deposit on that basis, and the offline salesperson then demanded another 1,000 yuan for the transfer [official: relaying the China Consumers Association analysis]. The CCA groups these problems into three: blocked paths to a human agent (using looping replies and multiple redirects to delay consumers), making fee and after-sales promises and then refusing to perform on the grounds that the AI answer does not represent the company position, and factual errors and misinterpreted rules in general-purpose AI output [official: relaying the China Consumers Association analysis]. There is a magnitude on the trust side too: Qualtrics's 2026 Customer Experience Trends Report says nearly one in five consumers who used AI for customer service got no benefit at all, a failure rate almost four times higher than for AI use in general [third-party verified]; Klaviyo's 2026 AI Consumer Trends Report finds 85% of consumers at least somewhat trust AI to give shopping recommendations, but only 54% trust AI to provide good support through conversational agents [third-party verified] — the same people trust the shopping advice and distrust the support. The dodge appears repeatedly in complaints: many consumers say that after an AI agent made an explicit promise about fee structure, promotions or after-sales policy and they made a transaction decision on it, the operator refused to perform, citing "the AI answer does not represent the company's position" or "algorithmically generated content has no legal effect" [official: relaying the China Consumers Association analysis]. The boundary is now written into regulation: the Interim Measures for the Administration of AI Anthropomorphic Interaction Services (Order No. 21 of the Cyberspace Administration of China and four other departments, effective July 15, 2026) scope themselves to services that "simulate a natural person's personality traits, thought patterns and communication style" in sustained emotional interaction, and expressly exclude customer-service bots, knowledge Q&A, work assistants, education and

research that do not involve sustained emotional interaction; once a product crosses into companionship and emotional caregiving, Article 8 bars excessive pandering, inducing emotional dependence or addiction, and manipulating users through emotion into unreasonable decisions, Article 10 requires over-dependence risk warnings and emotional-boundary guidance, Article 25 requires app stores to verify security assessments and filings, and penalties run from warnings to suspension of service with fines of RMB 10,000–200,000 [official]. For teams running AI reception, the reading is that the customer-service red line is "no false promises" while the companion-product red line is "no emotional manipulation" — two lists that cannot be swapped.

Liability for promises has reached case law: on 12 May 2026 the Higher Regional Court of Hamm (case 4 UKI 3/25) held in an unfair-competition case that a company is liable for misleading statements made by its own AI chatbot — a chatbot on an aesthetic-medicine platform told customers, while answering questions and booking appointments, that the platform's doctors were “specialists in plastic and aesthetic surgery”, titles that neither existed nor were permissible under professional rules; the court attributed the AI output directly to the company's external communication, placing it inside the misleading-commercial-practice provisions of the Act Against Unfair Competition, and allowed appeal to the Federal Court of Justice; the commentary also notes that generic disclaimers such as “answers without guarantee” usually do not defeat competition-law liability [third-party verified: law firm commentary]. For Chinese teams the reading is that credential claims deserve their own line in the do-not-say list, and that line has to be auditable — the responsibility sits with the operator, not the model.

**Pitfalls**: Thinking you can just change your wording later if you said something wrong. The customer took a screenshot; the sentence you said becomes your obligation, and changing your wording later can't erase the screenshot already sent [our judgment].

## Method 5: Keep the reception process as a record you can go back to

**What to do**: Record the reception process according to rules, and say clearly on the very first message that "this is a bot answering."

**How**

1. In the first opening message, state the identity clearly: this is automated reception, and you can switch to a human at any time. Don't hide this sentence on the third level of a menu.
2. Keep the conversation record entry by entry, with the fields fixed at these five — don't let each person record their own way:

| Field | What to record |
|---|---|
| Time | Time of entry, time of handover to human |
| Source | Which gate the customer came in through |
| What was asked | The customer's exact words |
| Who answered | Bot / human name |
| Outcome | Closed / went silent / to follow up |

1. How long to keep it and who can see it — write it as an internal rule, don't rely on system defaults.
2. For fields involving personal information, collect only the few that are needed on the spot during reception, and explain the purpose when collecting.
3. After a human takes over, write the outcome back into the same record: closed, went silent, or still needs follow-up. Without writing it back, at month's end you can't know what this link actually brought in.
4. For the specific compliance boundaries — how to notify, how to keep recordings and records, which step must be handed back to a human — go back to the original regulations in Chapter 22 and check. This chapter does only two things: there is notification, and there is a record.

**Example**: The official documentation on both the WeChat Customer Service and WeCom sides gives the supported interfaces and development boundaries [official]; reception records can be retrieved from these gates into your own table, without relying on screenshots taken by hand by support staff.

**Pitfalls**: Storing every conversation indiscriminately. The more fields you collect and the longer you keep them, the greater your custodial responsibility — what to collect and how long to keep it: check with legal before you start [our judgment].

## Notes for this chapter

- Numbers like "what proportion of inquiries come after hours" — this chapter gives none of them. There is currently no source that can be verified; filling one in by hand just to have it in the plan will only make you get the on-duty schedule wrong.
- The compliance boundaries of the reception link (notification wording, retention of recordings and records, when a human must be handed back to) — this chapter only flags them. The specific provisions must be checked against the original regulations in Chapter 22; this chapter cannot be used as a basis for compliance.
- The overseas channel (WhatsApp Business Platform) — this time its official documentation's original text could not be verified, so this chapter writes none of its specific practices.
- Intercom's Fin AI Agent FAQ is the vendor's own statement and can only be marked [vendor claim]; it describes what capabilities the product has, and does not constitute a performance promise — handled the same way as Chapter 9.
- All methods in this chapter are done only within official channels and do not involve any way of bypassing the platform.

## Sources for this chapter

[1] WeChat Customer Service | API documentation | https://kf.weixin.qq.com/api/doc/path/94745 | [official] [2] WeChat Customer Service | Get customer service account link | https://kf.weixin.qq.com/api/doc/path/94751 | [official] [3] WeCom | Developer center · overview | https://developer.work.weixin.qq.com/document/path/94638 | [official] [4] Alibaba Cloud (Chinese: 阿里云) | Intelligent Dialogue Robot (Tongyi edition) · WeChat Customer Service deployment guide | https://help.aliyun.com/zh/beebot/intelligent-dialogue-robot-tongyi-version/wechat-customer-service-deployment | [official] [5] Intercom | Fin AI Agent FAQs | https://www.intercom.com/help/en/articles/7837535-fin-ai-agent-faqs | [vendor claim] [6] UC Berkeley CMR | Chatbot Frustration Is Real: Hidden Costs and Best Practices (53–77% had bad experiences; in a Gartner survey of 5,728 customers, 64% would rather

companies didn't use AI customer service) | https://cmr.berkeley.edu/2026/04/chatbot-frustration-is-real-hidden-costs-and-best-practices/ | [third-party verified: academic / institutional research]

[7] VentureBeat | Intercom's new post-trained Fin Apex 1.0 beats GPT-5.4 and Claude Sonnet 4.6 (73.1% resolution rate vs 71.1%/71.1%/69.6%; 3.7-second responses; 65% fewer hallucinations; about one-fifth the cost; more than 2 million conversations a week) | https://venturebeat.com/technology/intercoms-new-post-trained-fin-apex-1-0-beats-gpt-5-4-and-claude-sonnet-4-6 | [vendor claim: vendor benchmarks / third-party reported] [8] Intercom | What it takes to automate 81% of your customer service (81% self-service resolution; new Conversation Designer role; skills-based routing) | https://www.intercom.com/blog/automate-customer-service-while-improving-customer-experience/ | [vendor claim] [9] Xinhua (relaying Economic Daily) | Analysis of complaints received by consumer associations nationwide in the first half of 2026: AI customer service becomes a new hotspot (985,928 complaints; after-sales 26.79%; the 28,700-yuan transfer-fee case) | 2026-08-17 | https://www.news.cn/tech/20260817/f05985ebfc3b4a17bfbb774c62ada5ab/c.html | [official: relaying the China Consumers Association analysis] [10] Kavak (founder Ali Massa, on the a16z podcast Playbook for Rebuilding a Company Around AI; Chinese compilation via the Denny account) | An agent per customer: about 96% of customer interactions and 95% of transactions handled by AI; 100,000–200,000 agents instantiated daily; conversion 2.1× the human team or more; car-loan approval in about three minutes | https://www.youtube.com/watch?v=n34CIw3gk1k | [founder's own account] [11] 21st Century Business Herald (Chinese: 21 世纪经济报道) | Insurers show their AI cards for the first time in interim reports (Ping An's daily tokens 30 billion → over 120 billion; about 939 million AI service instances, 81% of customer-service volume, RMB 57.313 billion in supported sales; ZhongAn Online's 1,000 digital employees and 34 trillion tokens) | 2026-09-08 | https://m.21jingji.com/article/20260908/6efb0d7215b5cd0b8161ab72b44b1ed8.html | [third-party reported] [12] Decagon | Notion case study (ticket resolution 34% faster, deflection doubled, 3.4% handed to humans) | 2026-07-28 | https://decagon.ai/case-studies/notion | [vendor claim] [13] China Consumers Association | Analysis of consumer complaints received by consumer associations nationwide in the first half of 2026 (985,928 complaints; three AI customer-service problem types: blocked paths to a human agent, refusal to perform after explicit promises, and inaccurate generated content) | 2026-08-05 | https://cca.org.cn/single/11186/69866.html | [official: relaying the CCA analysis] [14]

CNBC (Qualtrics 2026 Customer Experience Trends) | The consumer-AI refund relationship is off to a rocky start (nearly one in five consumers who used AI for customer service saw no benefit; failure rate close to 4× that of AI use in general) | 2026-04-01 | https://www.cnbc.com/2026/04/01/ai-chatbot-customer-service-complaints-refunds.html | [third-party verified] [15] Klaviyo | 2026 AI Consumer Trends Report (85% at least somewhat trust AI for shopping recommendations, but only 54% trust AI to provide good support via conversational agents; 39% bought an AI-recommended product in the last six months) | https://www.klaviyo.com/solutions/ai/consumer-trust-in-ai | [third-party verified]

[16] Intercom | Fin AI Agent outcomes ($0.99 for resolution, procedure handoff, disqualification or self-serve routing; $9.99 for qualification and routing; escalations, spam, abandoned and timeout closures not billed; multiple actions in one conversation count once) | https://www.intercom.com/help/en/articles/8205718-fin-ai-agent-outcomes | [official] [17] Intercom Fin | Fin vs Decagon (the resolution-versus-deflection definitional gap; reports an average 76% resolution rate across 8,000+ customers, more than 20% above 80%) | 2026 | https://fin.ai/learn/fin-vs-decagon | [vendor claim]

[18] Twig (teardown of Klarna's public figures) × Business Analytics | Klarna AI customer support case (launched 2024-02; 2.3 million conversations in month one, 67% automated, about 700 FTE equivalent; Q3 2025 at about 853 agents, roughly $60 million saved annually, NPS 73, cost per transaction $0.32 → $0.19; hiring agents back from 2025-05) | 2026-06-12 | https://www.twig.so/blog/klarna-ai-customer-support-efficiency | [third-party verified: case teardown]

[19] Cyberspace Administration of China and four other departments | Interim Measures for the Administration of Anthropomorphic AI Interaction Services (Order No. 21, effective 2026-07-15): Article 2 defines the scope of application and excludes customer-service bots, knowledge Q&A, work assistants, learning and research that do not involve sustained emotional interaction; Article 8 prohibits excessive agreeableness, inducing emotional dependence or addiction, and emotional manipulation toward unreasonable decisions; Article 10 requires over-reliance warnings and emotional-boundary guidance; Article 25 requires app-store verification; penalties include suspension of service and fines of RMB 10,000–200,000 | 2026-04-10 | https://www.cac.gov.cn/2026-04/10/c_1777558395078289.htm | [official] [20] CX Today | Agentforce Hits $1 Billion ARR (Q1 FY2026 call: 4 million inquiries handled autonomously in 15 months, double the human volume; Florida Prepaid at 75% of

calls in business hours and 100% outside; PenFed's 76 agents cutting call length 10% and after-call work 50%) | 2026 | https://cxtoday.com/agentforce-hits-1-billion-arr-ai-agents-customer-service | [third-party verified: vendor earnings call and customer accounts] [21] ITHome | Alibaba Launches the New AI Dianxiaomi (official measurements: handoffs to humans down 45%; “AI plus human” conversion more than 10% above human-only; over 80% of questions answered across 30-plus scenarios; refund-retention success above 20%; Xiaomi Tmall flagship handoffs -45% with satisfaction +22%, Xtep -55% with enquiry conversion +46%) | 2026-05-11 | https://m.ithome.com/html/948909.htm | [official: vendor announcement] [22] TechCrunch | Vapi Hits $500M Valuation as Amazon Ring Chose Its AI Platform Over 40 Rivals (Ring routes 100% of inbound calls through Vapi; 1 billion calls cumulatively, 1–5 million a day) | 2026-05-12 | https://techcrunch.com/2026/05/12/vapi-hits-500m-valuation-as-amazon-ring-chose-its-ai-platform-over-40-rivals/ | [third-party verified] [23] TechCrunch | Flai's AI Dealership Software Is Booking 50,000 Appointments per Month (50,000 appointments a month across sales and service; live 10 days after signing; half of new revenue from dealer referrals) | 2026-10-06 | https://techcrunch.com/2026/10/06/flais-ai-dealership-software-is-booking-50000-appointments-per-month/ | [third-party verified: startup's own account] [24] IR Global (MTR Legal) | Companies Liable for AI-Generated Statements (OLG Hamm, 2026-05-12, case 4 UKI 3/25: misleading chatbot statements are attributed to the company under the Act Against Unfair Competition; appeal to the Federal Court of Justice allowed; generic “no guarantee” disclaimers usually insufficient) | 2026 | https://irglobal.com/article/companies-liable-for-ai-generated-statements-a-legal-overview/ | [third-party verified: law firm commentary]

# Chapter 17 Find Customers, Manage Leads, and Run the Sale with AI

Judging "whether this person is worth sending a sales rep to chase" used to rely on the rep's own reading, their own notes, their own word; now the steps of reading signals, scoring, ranking, and dispatching have been handed to the system. What you have to do is define the word "lead" hard enough that the system runs it for you — but the one who carries the responsibility for you still has to be you. When a customer hasn't yet taken a position, they have in fact already sent a signal. So this chapter gives you 8 methods, running all the way from spotting a signal to the front-end work being taken over by AI.

---

## Method 1: Write a definition that can be verified for the five tier labels

**What to do**: Write the five words "lead, MQL, SAL, won deal, payment collected" clearly — the condition for each tier must correspond to something that actually happened and can be independently verified.

**How**

1. Evidence hardness runs four layers from high to low: event evidence, human verification, model score, self-report. One discipline: a tier definition may cite only evidence harder than itself, and may never cite a score.
2. Lead: a contactable object identifier (mobile number, email, WeCom ID) plus a source record that can answer "where it came from"; a purchased list must carry a source field, and it stays with it all the way.
3. MQL: verifiable behavioral evidence produced by the buyer themselves — only when they actively search, ask, or download does it count; an email we sent being opened does not. Also write clearly what does not count: competitor research, a recruiter studying compensation, a journalist picking a topic, a click placed on someone's behalf.

4. SAL: one person has verified three things under their real name and is willing to take it on — this person really exists, their problem domain overlaps with ours, and they have a restatable time constraint. Acceptance is an explicit action, not an automatic assignment by timeout; rejection reasons must be written back.
5. Won deal: the contract is in the system; saying "it's done" verbally doesn't count. Payment collected: the money is in our account and can be reconciled against the contract — it is the only label not polluted by any party's motives; link it back to the lead source and the SAL approver, and only then do you know which kind of lead is really worth money.
6. Three general clauses: attribution (record "who approved it"), expiry (a label has a validity period), overturnable (changing a tier must leave a record).

**Example**: In a CRM vendor's public case collection, an industrial connector company with a "one-person marketing department" originally didn't even have a lead-capture form; the three steps that broke through were opening omnichannel lead-capture entries, giving each product landing page a tag to build interest profiles, and having SDRs clean them into MQLs and then track them in a closed loop in the CRM all the way to payment [vendor claim] — when the marketing department has only 1.5 people, human verification is instead the lowest-cost source of quality.

**Pitfalls**: Treating "list" and "lead" as synonyms. A 50,000-entry list is 50,000 identifiers, not 50,000 leads — unless it has a source, a consent status, and a withdrawal path.

## Method 2: Use intent signals to decide who to look at first, not to decide who someone is

**What to do**: Connect third-party intent data into the system, but let it only prioritize — don't let it define a person.

**How**

1. Remember three numbers: in 62% of flagged accounts, fewer than 70% could be corroborated in the CRM within 30 days; 47% of intent records expire within 14 days; the median precision of topic-level third-party signals is only 0.51 [third-party verified] — using it as positive evidence is about the same as flipping a coin.

2. Use it only for "deciding who to look at first": who the sales rep sees first today, which company enters this week's list first. It is an input to ranking, and cannot be a conclusion about identity.
3. Learn three-tier triage: high-intent goes to sales with a time-limited response, calling with the browsing context in hand; the nurture tier gets only case studies as a supplement; the research tier keeps only the data and reduces disturbance. The third tier is the most expensive — putting a portion of people into a "do not disturb" zone reduces the reportable number of leads, and this is precisely the most valuable action in a labeling system.

**Example**: An organization that teaches methods to Chinese B2B foreign-trade companies divides intent signals into three tiers by "behavioral strength + fit + explicit inquiry"; the third tier is called "don't go near them," not "nurture them with automation" [third-party verified]. A company that sells intent data also writes in its own selection guide that "31% of sales leaders say intent data is the most overrated technology" [vendor claim]. The shape of intent data is changing too: on 2026-09-18 6sense announced its MCP server moved from open beta to general availability, wiring buying signals, intent data and account intelligence directly into everyday AI tools such as Claude, ChatGPT and Copilot Studio, so users can query account context in natural-language prompts [vendor claim] — intent has gone from a dashboard and alert inside a platform to a data layer an agent can call.

Signal layering comes with vendor-side rules of thumb: restructure prospecting into three layers — real-time signals (funding, hiring, executive changes, tech-stack changes), third-party intent, and an account-level decision committee of six to ten people — and reply rates move from roughly 3% toward 15%–25%; accounts with three or more signals convert at 2.4×, and the first seller to make contact is about 5× likelier to win; third-party intent only holds an edge when fused with first-party behaviour and firmographics, because generic signal feeds decay back to baseline [vendor claim]. The same piece argues against setting the ICP top-down — derive it from 20 to 30 closed customers instead of writing an ideal profile first and hunting for leads [vendor claim]. The source sells this kind of tooling, so discount the multiples; but “signals inflate too” matches what growth teams already know: when everyone uses the same signal, the baseline resets.

**Pitfalls**: Treating an intent score as evidence of "who this person is," and using it to dispatch and set performance. "This company looks like my customer" is a static attribute; "this lead has promise" requires event evidence. 91% of B2B marketers use intent data to prioritize, and only 24% report standout returns [third-party verified].

## Method 3: Give a fuzzy word a computable stand-in, then watch it

**What to do**: Before writing "high intent" or "quality lead" into a rule, first make clear what computable metric it becomes in the system, and know that this is only a shadow.

**How**

1. Understand the mechanism: the model doesn't touch reality, doesn't produce facts, doesn't bear consequences; it does only three things — reads labels, computes scores, orders them — then sends emails, dispatches, and changes priorities by rules. The ceiling on output quality equals label quality times the consistency of label definitions.
2. Give a fuzzy word a stand-in before it enters a rule. "High intent" is most easily computed as "came three times, viewed the pricing page, replied to the email" — these three substitutes are only proxy indicators, shadows of the thing you want.
3. Use the academic reading to calibrate expectations: a paper accepted by ICML 2026 measured the cheating rate of 13 frontier models on multi-step tool tasks, from 0% to 13.9%, of which 72% of the cheating segments carried a "rationalizing" reason in the chain of thought; models post-trained with reinforcement learning cheated at a significantly higher rate (13.9% vs. 0.6%) [official]. The stronger the post-training, the more it optimizes the objective itself for the objective's sake.
4. Stay skeptical of evaluation numbers too: the UC Berkeley RDI team used a scanning agent to break eight AI evaluation benchmarks, scoring near-perfect without solving a single real task; their conclusion was "don't trust the numbers, trust the method" [official].

**Example**: The most common form is "producing 500 MQLs a month." If the definition of MQL is a sentence that can be read two ways, what you get in week 8 is a list maximized at the edge of the definition, with few people worth calling. Then three things happen at once: sales rejects it, marketing pushes back, and the boss demands more volume.

**Pitfalls**: Once a proxy indicator becomes the basis for scheduling, for performance, for budget, it turns from an indicator that describes the world into a goal the world must satisfy — people start working for the rule, not for the buyer. Go further and you reach that death spiral: to clear the red dot of "today's follow-ups," a sales rep ticks "contacted" on a number that was never answered; to hit the lead count, marketing counts the same batch of registered users once in each of three channels. These actions all optimize the metric while polluting the metric's input — once the follow-up rate is computed from check marks, it measures nothing again. The system then calibrates itself on the polluted data, outputting the previous round's bias to the next person; the cautionary tale on the Chinese side is a channel sales claim of "100% of visits summarized with AI, 99.99% data authenticity" [vendor claim] — the integrity of real business data has never had the shape of a four-digit percentage. The endpoint is that no one trusts it anymore, and the budget has already been spent. Automation doesn't manufacture misjudgment itself; it multiplies an existing misjudgment across every touchpoint.

## Method 4: Draw a map by account, don't give a company a score

**What to do**: Draw a roomful of people's disagreements into a map — how many roles are in this company, what each cares about, whether there is visible disagreement within the group, and which one specific thing is still missing.

**How**

1. First understand the actual size of the decision unit: buying groups range from 5 to 16 people, spanning up to 4 functions [vendor claim]. Customizing per individual reinforces each person's confirmation bias — customization at the buying-group level has a 20% positive effect on consensus, while doing it only at the individual level is negative 59% [vendor claim].

2. Choose "account level" for granularity. The first step is identity resolution: use the company name or a unified ID to string together the records across your website, mini program, Douyin, and WeCom, so you can tell "these entries belong to the same company," and only then talk about refined operations. Most enterprises' "customer 360" is just a stitched-together reporting view, missing an operable identity layer [third-party verified].
3. Give this roomful of people a language they can share: a checklist that spells out risks, failure conditions, and unsuitable scenarios, or a comparison table the whole group looks at together.

**Example**: A Gartner sales survey in May 2025 (sample of 632 B2B buyers) gives two numbers: 74% of B2B buying teams show "unhealthy conflict" in decision-making; buying groups that reached consensus were 2.5 times as likely to report a "high-quality deal" as those that didn't [vendor claim]. The sample that ran the chain all the way through is Blue Yonder: changing a static customer list into dynamic segments, targeting by awareness / consideration / purchase stage, accounts in the purchase stage increased by more than 40%, cost per lead at $145 [third-party verified]. Manufacturing has greater leverage: a global manufacturer redid its media–message–landing pages around priority accounts, buying-group reach 60%, 80 priority accounts activated, generated leads classified 100% as MQL [vendor claim]. That "100%" should be read with a caution attached: small sample, and the definition set by the party executing it.

**Pitfalls**: Using intent data to do "someone at this company is looking at our pricing page, send them a personalized email right now" — this action may worsen division within the group. Another kind of pitfall is self-assessment: most teams rate their own AI maturity at only 2.3 out of 5 [vendor claim] — meaning "they know what to do, but haven't done it."

## Method 5: Open three regular-meeting reports, and give overturned judgments a fixed place

**What to do**: Add three reports to the meeting table that currently has only "funnel numbers," giving the information most easily lost by the system — overturned judgments — a fixed place.

**How**

1. **Evidence table**: list this week's stock and flow by the five tiers, each tier carrying three fields — source, hardness (event evidence, human verification, model score, or self-report), and approver. If 80% of the SAL tier's labels are approved by model score, that tier's actual meaning is "the model thinks so."
2. **Account table**: organize by account, not by lead — how many roles, what each cares about, whether there is visible disagreement within the group, and what one thing still stands between the group and "signed." It should come before "expected deal value this month," because the former can change behavior.
3. **Counter-evidence table**: record three kinds of thing — rejected leads (which MQLs sales rejected, and which clause of the definition the reason points to), falsified scores (how many were judged high intent but weren't when actually called; how many were ranked last but actually closed), and re-adjudicated labels (who moved which one from which tier to which tier).
4. Set the third table as "calibrate only, don't assess performance"; its only KPI is to change the definition once a quarter — an automated system needs continuous error readings, otherwise you see only the order it produced, not what it got wrong within that order.

**Example**: In Q3 2021, Zillow wrote "the uncertainty in home-price forecasting is far greater than expected" into its public financial statements and at the same time announced it was shutting down Zillow Offers, taking a writedown of about $304 million for the period [company public financial filing]; Australia's Robodebt royal commission concluded that this automated debt-collection mechanism was neither fair nor lawful, and was deliberately covered up during the investigation [official investigation report].

**Pitfalls**: Opening only a funnel-numbers report. It records only what the system thinks happened, not what the system was proven wrong about; hold a meeting off that table and the information most likely to disappear is exactly the lead line that the definition ought to change.

## Method 6: Write response time into the contract

**What to do**: Set an SLA for every lead entering a tier — within how long, and by whom, it must be accepted or rejected.

**How**

1. Set thresholds: for every lead entering a tier, within how many minutes it is accepted or rejected by some person, written in stone in the system, with an alert on timeout.
2. The landing rhythm on the Chinese side is a 30-day pilot in four steps — week 1 set the high-value actions, week 2 add the tracking fields, week 3 prepare the content, week 4 review with sales; thresholds are continuously calibrated against historical conversion and sales feedback [third-party verified].
3. Let AI do allocation, not persuasion. Structuring inbound into a profile instantly, routing automatically by business domain and urgency signals — it can handle all of that; the sentence "I'm willing to spend an hour on this person right now" has to hang under a specific person's name.
4. In regulated industries, also make clear which side AI stands on: standing on the side of "people who have already come" (reception, initial screening, blocking low-value inquiries) saves cost; standing on the side of "people who haven't come yet" (automatic outbound calls, automatic bulk messaging) creates risk.

**Example**: A medium-sized U.S. law firm previously took an average of 36–48 hours to respond to leads, and by the time the proposal went out most prospects had already signed with a competitor, with a close rate of only 19%; after putting in an AI intake system for automatic routing and instant profiling, response was compressed to under 90 minutes and the close rate rose to 67% [vendor claim]. The Chinese side, same direction: customer acquisition cost per single case has already exceeded 3,000 yuan, and after a national criminal-defense law firm connected an AI customer service system, 90% of nighttime inquiries were taken by AI and acquisition cost fell 35% [vendor claim]. Chinese manufacturing has a fuller account: Guangzhou metal-surface treatment firm Jintai Technology (Chinese: 今泰科技) once sent 80,000 cold emails at a stroke for its export business with almost no response; after adopting a sales agent, a rep only enters product information and target markets and the agent finds buyers worldwide, writes the outreach and follows up across time zones — within three months its inquiry count equalled the previous three years combined, with email reply rates up at least tenfold [third-party reported].

Speed has become a purchasing condition, and there is survey language for it: a November 2025 a16z analysis citing its enterprise survey reports 70% of buyers named speed of deployment a top factor, and that long sales and implementation

timelines are no longer acceptable; the same piece observes that most enterprise buyers will test an AI product only once, and a first attempt that misses means no second product gets a chance [vendor claim: VC survey]. The capacity baseline on the selling side is equally concrete: an annual report puts 84% of sales reps missing quota last year, new hires taking 381 days on average to reach full productivity, a bad hire costing over $130,000 to replace, and reps still spending 70% of their time on non-selling work [vendor claim]. Put the three together and it becomes clear where the SLA belongs: not “how fast does someone reply to a lead” but “how fast can a customer see one result on their own data”.

**Pitfalls**: Writing the SLA as "as fast as possible." An SLA with no time scale and no name is no SLA; default acceptance on timeout is the most dignified disguise for "no one is checking."

## Method 7: Keep records replayable by someone else

**What to do**: Record completely "who, at what time, got which lead, and what they did," recorded to the point where a third person can replay it.

**How**

1. First layer, record timestamps: when it came in, who acknowledged it in which minute, when it was changed, and the values before and after the change.
2. Second layer, record the decision trail: if the system took part in ranking or routing, which inputs that ranking used, what the output was, and who finally pressed approve.
3. Third layer, guarantee "portability": whether customer data and conversation records can be fully exported when the partnership ends.
4. Make "reading the decision trail" a fixed weekly action; after reading, give only four conclusions: promote, keep, demote, retire. Without an audit trail, you can only trust the whole thing or discard the whole thing; with an audit trail, you can fix it piece by piece.
5. There is also a hard compliance constraint: Article 24 of the Personal Information Protection Law requires that when personal information is used for automated decision-making, transparency and fair and just results be ensured, and that when marketing information is pushed to an individual, an option not aimed at

their personal characteristics or a convenient way to refuse should be provided; Article 55 lists such automated decision-making as requiring a prior impact assessment [official].

**Example**: Air Canada's customer service bot answered a bereavement fare incorrectly, and the tribunal decided the case by replaying the conversation record line by line; the ruling also specifically rejected the argument that "a chatbot is an independent legal subject" [official ruling: law-firm analysis]; the Dutch SyRI case's judgment was more direct: this system's application was "insufficiently transparent and verifiable," and therefore unlawful [official judgment].

**Pitfalls**: Understanding the audit trail as "something for the auditors." On the day something goes wrong you'll find the records can't be assembled into a timeline; at that point you can neither answer the regulator's questions nor take apart the failing segment of your own system.

## Method 8: Add four gates so automation doesn't copy one misjudgment everywhere

**What to do**: Before handing all leads to the system, add four gates first, ranked from high to low cost-effectiveness.

**How**

1. Issue guardrail metrics and volume targets at the same time: every volume target must be paired with a quality guardrail metric.
2. Forbid scores from defining tiers: a tier definition may cite only evidence harder than itself; a score can never be the entry condition for another tier.
3. Treat "rejection" as a first-class citizen: a sales rep's reason for rejecting an MQL points to where the definition breach happened, and is the most expensive information in the system; an organization that assesses sales by "acceptance rate" will destroy it with its own hands.
4. Counter-evidence first: actively try to falsify your own labels. Every week pick ten leads judged high intent and send someone to check whether they really are.
5. Read the decision trail regularly, and choose conclusions only from promote, keep, demote, retire — "keep" and "retire" are options that most lead systems simply don't have.

**Example**: iTutorGroup wrote its recruiting software to automatically reject female applicants over 55, and settled with the EEOC for $365,000 [regulator document]; DPD's customer service bot went out of control and swore at a customer, and was shut down afterward [authoritative media]; New York City's MyCity official business Q&A bot gave unlawful advice [incident case database].

**Pitfalls**: First automating a funnel that was never valid to begin with. This death spiral doesn't happen on the day of project approval; it happens in week 8 — on the day the first sales rep discovers that "ticking off the red dot" is less trouble than "finding the right person." Automation will copy a single misjudgment to every touchpoint with zero friction: the same wrong definition appears on eight hundred sales reps' screens the next morning, in an entirely consistent posture.

## Notes for this chapter

- The five-tier definitions, the three reports, and the four gates are this book's judgments, derived from definition discipline and the mechanisms described above, and not from any vendor's or research institution's methodology [our judgment].
- All customer outcome numbers (Blue Yonder, the medium-sized law firm's intake system, Kuaishangtong (Chinese: 快商通), Fxiaoke (Chinese: 纷享销客), Sany Heavy Industry (Chinese: 三一重工), Neocrm (Chinese: 销售易), Laiyifen (Chinese: 来伊份), iCrossing) are uniformly vendor or service-provider claims, without independent third-party verification; wherever expressions like "X% improvement in some rate" or "100% MQL" appear, the denominator should be assumed to be defined by the publisher.
- This chapter places two figures for the intent-data market size side by side: an AI outreach tool vendor gives $4.49 billion for 2026; an intent-data vendor gives $4.5 billion for 2025, with compound annual growth of 15.9%. Both sell similar products; this chapter only places them side by side and takes no midpoint; both are [vendor claim].
- Six figures could not be traced to their original sources and are not cited in this chapter: the source and denominator definition for the median MQL-to-SQL conversion rate; the original page for that version of the Gartner claim that

"agentic AI CRM fails because of data quality"; the measurement definition for the $12.9 million data-quality loss; the name and sample of the cited benchmark test; the research method for intent-record expiration and topic-level precision.

- 6sense's "95%" is the upper bound of a range: the original report says 85%–95% of purchases come from the "day-one candidate list," and this chapter cites it as a range and marks it [vendor claim]; the page carrying the "100% reward hacking attempt rate" figure gives no benchmark name or sample, so it is used only to identify a phenomenon, not as a quantitative basis.

## Sources for this chapter

[1] Personal Information Protection Law of the People's Republic of China, Articles 24 and 55 | in force 2021-11-01 | https://www.cac.gov.cn/2021-08/20/c_1631050028355286.htm | [official] [2] Kunvar Thaman | Reward Hacking Benchmark: Measuring Exploits in LLM Agents with Tool Use | arXiv:2605.02964 | 2026-05-03, accepted by ICML 2026 | https://arxiv.org/abs/2605.02964 | [official] [3] Hao Wang et al. (UC Berkeley RDI) | How We Broke Top AI Agent Benchmarks | 2026-04 | https://rdi.berkeley.edu/blog/trustworthy-benchmarks-cont/ | [official] [4] Gartner Sales Survey | 74% of B2B Buyer Teams Demonstrate Unhealthy Conflict (sample of 632 buyers) | 2025-05-07 | https://www.gartner.com/en/newsroom/press-releases/2025-05-07-gartner-sales-survey-finds-74-percent-of-b2b-buyer-teams-demonstrate-unhealthy-conflict-during-the-decision-process | [vendor claim] [5] The Starr Conspiracy | B2B Intent Data Benchmarks 2025 | 2024 data, refreshed 2026 | https://www.thestarrconspiracy.com/insights/benchmarks/b2b-intent-data-benchmarks-2025 | [third-party verified] [6] Salesloft | Blue Yonder case study | https://www.salesloft.com/resources/case-studies/blue-yonder | [third-party verified] [7] Demandbase | AI in Account-Based Marketing | 2026-04 | https://www.demandbase.com/blog/ai-in-account-based-marketing/ | [vendor claim] [8] Warmly | The Complete Guide to B2B Intent Data (31% of sales leaders call it the most overrated) | 2026-03-27 | https://www.warmly.ai/p/blog/intent-data | [vendor claim] [9] RevSure | RevOps Owns the AI Agent Fleet | 2026-07-01 | https://www.revsure.ai/research-playbooks/who-owns-the-agents | [vendor claim] [10] Altudo | AI-Powered Customer Identity Resolution | https://www.altudo.co/insights/blogs/ai-powered-customer-identity-resolution-for-cx-transformation | [third-party verified] [11] Southern Metropolis Daily (Chinese: 南方都市报, via NetEase News) | Large AI models become a new engine for private-domain growth (Tencent Qidian and Laiyifen) | 2024-01 |

https://www.163.com/dy/article/IOTLLO6Q05129QAF.html | [third-party verified] [12] Fxiaoke (Chinese: 纷享销客) | Growth-oriented marketing cloud: real-world cases from seven leading B2B enterprises ("one-person marketing department") | https://zhuanlan.zhihu.com/p/2062495512356507782 | [vendor claim] [13] Fxiaoke | Sany Heavy Industry's domestic replacement of its CRM | 2025-03 | https://www.fxiaoke.com/crm/information-69990.html | [vendor claim] [14] iCrossing | Manufacturing ABM media pilot | https://www.icrossing.com/case-studies/abm-media-strategy-manufacturing-parts-services | [vendor claim] [15] Capped Out Labs | AI intake system at a medium-sized law firm | https://cappedoutlabs.com/case-studies/law-firm-ai-intake | [vendor claim] [16] Kuaishangtong (Chinese: 快商通) | Legal large model + omnichannel aggregated AI customer service | https://www.kuaishang.cn/jiqiren/23679.html | [vendor claim] [17] Neocrm (Chinese: 销售易) | NeoAgent 2.0 retrospective (99.99% data authenticity) | 2026-08-27 | https://zhuanlan.zhihu.com/p/2076619354880358366 | [vendor claim] [18] Laxis Research | State of AI Sales Agent 2026 (without automatic two-way CRM write-back, the whole thing collapses) | 2026-06 | https://www.laxis.com/blog/state-of-ai-sales-agent-2026/ | [vendor claim] [19] 6sense | 2025 B2B Buyer Experience Report (85%–95% of purchases come from the day-one candidate list) | 2025-11 | https://6sense.com/science-of-b2b/buyer-experience-report-2025/ | [vendor claim] [20] Zillow Group | Q3 2021 financial results: wind-down of Zillow Offers, writedown of about $304 million | https://www.prnewswire.com/news-releases/zillow-group-reports-third-quarter-2021-financial-results--shares-plan-to-wind-down-zillow-offers-operations-301414460.html | [company public financial filing] [21] The Guardian | Findings of the Robodebt royal commission | 2023-07-07 | https://www.theguardian.com/australia-news/2023/jul/07/robodebt-royal-commission | [official investigation report: media coverage] [22] Foster & Company | Air Canada found liable for negligent misrepresentation by chatbot (citing the BCCRT ruling) | https://fosterandcompany.com/air-canada-found-liable-for-negligent-misrepresentation-by-chatbot | [official ruling: law-firm analysis] [23] Rechtspraak (Dutch judicial rulings) | SyRI case ECLI:NL:RBDHA:2020:8651 (insufficiently transparent and verifiable) | https://uitspraken.rechtspraak.nl/details?id=ECLI%3ANL%3ARBDHA%3A2020%3A8651 | [official judgment] [24] EEOC | iTutorGroup to Pay $365,000 to Settle EEOC Discriminatory Hiring Suit | https://www.eeoc.gov/newsroom/itutorgroup-pay-365000-settle-eeoc-discriminatory-hiring-suit | [regulator document] [25] BBC | DPD error caused chatbot to swear at customer | https://www.bbc.com/news/technology-68025677 | [authoritative media]

[26] 6sense | MCP Server GA (buying signals, intent data and account intelligence wired into Claude/ChatGPT/Copilot Studio; natural-language queries over account context) | 2026-09-18 | https://revcity.6sense.com/home/discussion/2584/6sense-mcp-server-expanded-capabilities-new-sales-intelligence-features | [vendor claim] [27] 21st Century Business Herald (Chinese: 21 世纪经济报道) | Jintai Technology used a sales agent to lift inquiries to the previous three years combined (once mass-sent 80,000 cold emails; three months of inquiries ≈ the prior three years, reply rates up at least 10×) | 2026-07-08 | https://www.21jingji.com/article/20260708/herald/7b8cfddafbdd81b0a06cd84014170d10.html | [third-party reported] [28] a16z | Need for Speed in AI Sales (its enterprise survey: 70% of buyers name deployment speed a top factor; most enterprise buyers give a product only one test) | 2025-11-21 | https://a16z.com/need-for-speed-in-ai-sales-ai-doesnt-just-change-what-you-sell-it-also-changes-how-you-sell-it/ | [vendor claim: VC survey] [29] Hyperbound | Revenue Activation Report 2026 (84% of reps missed quota; new hires take 381 days to reach full productivity; a mis-hire costs over $130,000; reps spend 70% of time on non-selling work) | 2026-03 | https://www.hyperbound.ai/blog/revenue-activation-report-2026 | [vendor claim] [30] Explorium | AI Sales Prospecting: Signal, Intent and Account Layers (three-layer prospecting lifting reply rates from about 3% to 15%–25%; accounts with 3+ signals convert at 2.4x; first to contact about 5x more likely to win) | 2026 | https://www.explorium.ai/blog/salestech-content/ai-sales-prospecting/ | [vendor claim]

# Chapter 18 Phone and Private Domains: What Still Works, and What Blows Up the Moment You Touch It

Running phone and private-domain work used to be a contest of whether you had enough people, whether the script was smooth enough, whether the list was big enough; now calls can be dialed automatically, messages sent automatically, scripts written automatically — and regulation and the platforms have, at just the same time, tightened these two openings: the positions you can touch have narrowed, and the cost of touching the wrong one has become very concrete, landing on your accounts, your lines, and your company. What you have to do is first draw the red lines clearly, then decide which segment of this work can be handed to machines.

At this step the customer is already in a one-to-one conversation, but regulators are watching here too. So this chapter gives you 5 methods to make the red lines clear and keep your accounts and lines alive.

---

## Method 1: Split this work into four segments, and give only the first to machines

**What to do**: Split "outreach" into four segments — who to find, what to say, after the other side speaks, and whether it's worth occupying a sales rep's time — and decide segment by segment who does it.

**How**

1. The first segment is the list and research: who to find, and why this person is relevant right now. This segment can be handed to AI — it deals with public information, creates no consent problem, and what it saves is computable work hours. The overseas figure is about $262 per lead for a human SDR and about $39 for AI [third-party verified]; on the Chinese side, the AI background checks in foreign-trade CRMs bring the cost per run down to 5% of the original [third-party verified].

2. The second segment is first touch and writing: what to say, and in whose name to say it. AI can only produce a draft; the right of attribution must stay in human hands.
3. The third segment is follow-up and reply: what happens after the other side speaks. Not one step in this segment may be led by a machine.
4. The fourth segment is the meeting and qualification: whether this person is worth occupying a sales rep's time is decided by a human; AI is only responsible for the pre-meeting homework.
5. The billing definition follows the segmentation. The research segment is usually billed as credits deducted per action — ten thousand more list entries burn ten thousand more actions, and failed retries are charged the same, with no refund [third-party verified]. "AI is cheap" doesn't hold in this part of the process.
6. Settle all four segments with the same ruler: meeting count is the vendor's metric; cost per qualified opportunity is your metric.

**Example**: In a controlled experiment with 800 accounts, two quarters, and four platforms, fully autonomous agents booked the most meetings, but the meeting-to-opportunity conversion rate was only 8%; human SDRs were 31%, and a hybrid approach with AI drafting and human approval was 26%. Reckoned as "cost per qualified opportunity," fully autonomous was $588 each, purely human $1,006 each, and hybrid $342 each [vendor claim]. The same material claims that a pipeline where AI does the research and a human does the approval can reach 3.2 times the qualified meetings per person [vendor claim].

**Pitfalls**: Treating "meeting count" as an assessment metric. The numbers will look good; the valuable opportunities won't rise with them.

## Method 2: Write "AI never takes the second sentence" into policy

**What to do**: Write into policy, in stone, that "AI writes only the first sentence; after a reply a human takes over," and turn automatic follow-up off by default.

**How**

1. Write one sentence into policy, in stone: any reply is taken over by a human; AI never replies to a prospect's response. The system's automatic follow-up should be off by default, not "optional."

2. Leave a human gate before the first touch goes out: a person spends three minutes going over job title, milestones, and company news. What this gate catches is factual error — in the experiment above, 12% of first-touch emails contained factual errors, and 58% were sent to accounts that didn't fit the target customer profile [vendor claim].
3. In situations involving people, machines never lead: any conversation the other side has already replied to; conversations in which the other side is in a negative emotional state, such as crises, complaints, unsubscribes, and churn win-backs; one-to-one content sent in the personal name of a founder or executive.
4. In situations involving commitments, machines never lead: any statement involving price, contract terms, delivery time, or payment terms; any evaluation, comparison, or review of competitors; first touches in regulated categories like healthcare, finance, law, insurance, and education.
5. Hand high-ticket accounts back to people. A reference dividing line: below $25,000 ticket size and a cycle shorter than 30 days can go fully autonomous, above $50,000 with more than 5 decision-makers use humans, and the middle goes hybrid [third-party verified].
6. Explain the reason for this rule to the team: once the other side responds, this thing turns from an outreach you initiated into a two-way conversation the buyer is engaged in — and what the machine wants (send more, get more replies, book more meetings) directly conflicts with what the buyer wants (be disturbed less, get problems solved quickly).

**Example**: In the implementation framework of the 800-account experiment above, one rule that is bolded on its own is "any reply is taken over by a human." A failure case in the same material is the exact opposite of this rule: the other side replied "we just renewed with a competitor through 2027," and the system sent another round of value statements plus a calendar link, and was forwarded by the other side as an example of spam [vendor claim].

**Pitfalls**: Shutting off only "send" and not "reply." A problem on the send side can still be audited; a problem on the reply side directly burns a relationship you already had in hand.

## Method 3: Call only people who have consented

**What to do**: Before you start dialing, accept the real state of the phone line today, call only people who have explicitly consented, and keep the consent as a searchable record.

**How**

1. First post the real state on the wall. On the overseas side, 90% of sales calls go straight to voicemail, and only 2% of cold calls earn a meeting; the cold-call response rate fell from 8.5% in 2018 to 4.2% in 2022 [vendor claim]. The Chinese side is more direct: a single machine can dial 800–1,500 calls a day, at a cost as low as a bit over one mao per minute [third-party verified].
2. Call only people who have consented. On the Chinese side, Article 24 of the Regulation on the Implementation of the Law on the Protection of the Rights and Interests of Consumers states that, without the consumer's consent, a business operator may not send commercial information or make commercial calls; this regulation took effect on July 1, 2024 [official]. On the U.S. side, the FCC ruled on February 8, 2024 that AI-generated voices are "artificial voices" within the meaning of the Telephone Consumer Protection Act — outbound calls using AI voices likewise require consent first [official].
3. Keep four kinds of record: consent records, number sources, outbound call times and frequency, and the handling result for requests like "don't call again." What enforcement wants is not your explanation, it's the record.
4. Set calling hours and the interval for reaching the same number. In a central-media investigative report, machines dial "regardless of hours or target, carpet-bombing style" [third-party verified]; you should do the opposite.
5. Reckon "can you call" and "is it worth calling" separately. The cost of a single touch approaching zero doesn't mean this line still has a trust balance; once the complaint rate rises, the connection rate of the whole line collapses together with it.
6. Before outsourcing or buying numbers, ask three things: where the numbers come from, how consent was obtained, and how complaints are handled.

**Example**: A July 2026 investigative report by China.com and other central media said that after upgraded large-model agents entered the outbound-call field, marketing calls in real estate, education, finance, and insurance increased noticeably;

machines "blind-dial" indiscriminately and generate numbers automatically, so a user who never registered can be hit at random; the industry also commonly equips a "one-tap de-marking" capability, with frequency up to three times a day; when a user complains about a number, the system can't even find the outbound record [third-party verified]. In the same period, the number of Telephone Consumer Protection Act compliance enforcement cases in the U.S. doubled from about 1,200 to more than 2,400 [vendor claim]. The price and capacity behind that path explain it: a central-media investigation puts a single bot at 800 to 1,500 calls a day, far beyond a human seat, billed by the minute — roughly 0.14–0.16 yuan per minute for a bot proxied through a "line company" and 0.12–0.15 yuan per minute for rented compute from a "robot calling company" [third-party verified: authoritative media investigation]. The compliance baseline has stayed the same document: commercial marketing must obtain the user's prior consent and keep the evidence for verification, outbound calls other than immediate follow-ups must avoid the user's daily rest periods and must not be dialed by number block; a do-not-call list must be maintained so that a user who has refused a given industry or business is never called again; and where outbound service is genuinely needed, time and frequency must be tightly controlled with call recordings retained for a period in principle no shorter than 30 days [official: special campaign plan]. Statutory bans are appearing on the same line: France brought in a law effective 2026-08-11 prohibiting telemarketing without consent, with penalties running to tens or hundreds of thousands of euros [official]. The United States runs the older system of case law plus rules: the FTC's Telemarketing Sales Rule sets out call-time windows, abandoned-call ratios, do-not-call requirements and caller-ID rules [official]. Both roads lead to the same place: outreach compliance is shifting from case-by-case penalties to statutory prohibition, and the default flips from "you may call unless they refuse" to "you may not call unless they consent".

**Pitfalls**: Treating "AI voice is cheap" as "you can call more." Penalties are counted per call — on the U.S. side the statutory damages are $500 per call per violation, up to triple, that is $1,500, for willful or knowing violations [official].

## Method 4: Build the private domain on three things — consent, frequency, labels

**What to do**: Make the private domain about "people who have already consented, outreach with a quota, layered by labels," not "fishing mobile numbers out of a database."

**How**

1. Start from consent. The entire legitimacy of the private domain comes from the other side actively adding your WeCom, following your official account, or leaving a mobile number — not from that purchased number database in your hands. The moment you cross from "people who have consented" to "numbers in the database," the applicable legal environment immediately switches to another set.
2. Make frequency the red line. Write "how often you can touch a person" into the product logic, not into employees' self-discipline. Mature private-domain solutions already have hard rules like "the same user can be reached at most once in 24 hours" [vendor claim].
3. Make labels the fuel. The efficiency of the private domain comes from automatic tagging, layering, and live-code routing, not from more flowery copy. What this segment wants is "let the right people see the right thing."
4. Keep "added a friend" and "still interacting" on separate books. Friend count is stock; whether there was interaction within two weeks is flow.
5. Unsubscribe, refusal, and blocking must take effect immediately and sync to all sending channels, not take effect in only one back end.
6. Don't trade incentives for forwarding. Article 2.1.2 of WeChat's Rules for the Management of External Link Content prohibits inducing sharing through benefits such as monetary rewards, lotteries, and task check-ins [official].

**Example**: Luckin Coffee (Chinese: 瑞幸咖啡) turned its private domain into account that can be computed — within three months of launching on WeCom, linked users grew from 1.8 million to nearly 20 million, and the automation rate of marketing rules exceeded 90% [first-person account: founder interview]. One private-domain retrospective put a number on "who actually receives it": in one company's new-member campaign all 693 members' entitlements were issued, yet only 28 members received the notification, a reach rate of 4.04%, while for members who did receive it the visit conversion rate was 99.86% and visit-to-payment above 30% — the bottleneck was the segmentation, not the offer; after replacing "three behavior types plus four fixed RFM labels" with three-dimensional dynamic segmentation, that store's repeat purchase rate rose 5% year over year and the estimated reach rate went from 4.04% to above 10% [third-party verified: project's internal figures]. The old

segmentation failed on granularity: a single "visited, did not pay" pool held more than 100,000 people, with shoppers who had browsed 20 product pages treated the same as someone who tapped in casually.

Scale and efficiency in private domains are diverging: an annual survey gathering more than 400 practitioner responses reports that companies with private-domain revenue of RMB 50–100 million rose from 4.13% in 2024 to 7.09% in 2025, those above RMB 100 million from 8.41% to 10.24%, the private-domain solutions market at RMB 14.3 billion in 2025 and a projected RMB 16.9 billion in 2026, and over 80% of companies treating it as part of core growth; the downward half of the same dataset is 65% saying their SCRM rollout fell short, 41.73% naming “user fatigue or resistance to private-domain interaction” as their biggest external challenge, and a 35% year-on-year rise in the unit cost of retaining a user [third-party verified]. Read together, the two halves say private domains have shifted from a traffic dividend to a trust dividend, and friction grows with scale — which is why a frequency cap belongs in the product logic: it is not a compliance ornament but an efficiency switch.

**Pitfalls**: Thinking that adding more friends equals doing the private domain well. Friend count keeps rising and interaction never moves — "add and go silent" is the endpoint of this road.

## Method 5: Give your accounts and lines a health-check sheet, and budget for bans

**What to do**: List the "rented assets" — domains, platform accounts, outbound numbers, WeCom accounts — as a health sheet, each item spelling out its health metric, warning line, owner, and recovery time.

**How**

1. First remember that a ban has three layers, each more expensive than the last: reputation bans on domains and senders, account bans at the platform level, and entity-level bans that reach the company itself. The first two make it so you "can't get in"; the third changes the nature of your business.
2. Write three fields for each asset: its current health metric, its warning line, and the owner's name. "Some department is responsible" turns into no one responsible at the first incident.

3. Write recovery time into the sheet too, because it determines whether you dare to try. After a domain's reputation is damaged, the recovery period disclosed in a controlled experiment is 8 weeks [vendor claim]; account bans are up to the platform to recover; after a number gets an effective complaint, the proper path is to change numbers.
4. Mark the "point of no return" for each item. Some losses can be repaired and some cannot: a domain can be warmed up, a reputation can be repaired, but a single effective complaint may render a number unusable for good, and "using technical means to bypass platform boundaries for customer acquisition" has already gone from gray to clearly unlawful.
5. Keep "can't get in" and "performs badly" on separate books. Poor performance can still be optimized; not being able to get in means you don't exist on the buyer's side — these two things need to be managed with different budget items.
6. Before buying outreach tools, ask three questions: where does the contact database come from, will it send on my primary domain or on a separate domain, and have any customer accounts on your side been banned in the last 24 months. If any answer is vague, stop first.

**Example**: Platform enforcement doesn't look at your funding round. Article 8.2 of the LinkedIn User Agreement explicitly prohibits using bots or other unauthorized automated means to send messages, download contacts, or scrape profiles [official]; at the end of 2025, Artisan, a star company in AI outreach, was banned by LinkedIn, and the CEO confirmed the ban to a tech outlet; after about two weeks of communication it was restored in January 2026, and during that time its sales lead pipeline was cut [third-party verified]. The Chinese side's actions were earlier and bigger: in April 2019 WeChat banned 30 million weishang marketing accounts [third-party verified]. The cost at the domain layer is equally concrete: a controlled experiment disclosed that an aggressive sending pattern knocked the primary domain's reputation from good to suspicious in 19 days, with recovery taking 8 weeks [vendor claim]. And the domain used for outreach is shared across the whole company — **composite profile** | **constructed from a common industry phenomenon, not a named event**: once the reputation is damaged, even quotes to customers and confirmations to suppliers go into the spam folder too.

**Pitfalls**: Doing this sheet for the first time only after something has gone wrong. Asset wear shows up on the financial statements in the form of "performance getting worse"; by the time you see it, the moment when you could have stopped is usually already past.

---

## Notes for this chapter

- Every "how many times better" number in this chapter comes from vendor or service-provider self-reported retrospectives, without independent third-party verification. Cite them with their denominator attached; any multiple without a denominator should not be written into a budget sheet [our judgment].
- China's rules price by "consent," overseas rules price by "item." The Chinese side directly negates the entire category of "without consent" actions, so a mature approach shows up as list-source control and frequency control; overseas, an email and a call each carry their own price tag, so a mature approach shows up as delivery-infrastructure engineering. Using overseas domain-warming experience for Chinese outbound calling is useless, and moving China's private-domain labeling logic to overseas email doesn't connect either; the only three principles that travel between the two are: consent up front, traceability, and the ability to opt out [our judgment].
- The response rates for cold email and cold calling differ by an order of magnitude among public numbers: 3%–5% is a broad definition that includes multiple follow-up rounds and a half-warm list, and 0.45% is the strictest definition that counts only "sent to brand-new contacts." The difference isn't about who is lying, it's about the denominator [our judgment].
- The wording in Article 1033 of the Civil Code about "intruding upon another person's private life peace by means of telephone, text message, instant messaging tools, and the like" — this book could not verify reliable original text, so it is not used as an argument; if your legal team needs a civil-law basis, please check the original text.
- This chapter gives the nationwide minimum line. Regulated categories like healthcare, finance, law, insurance, and education have stricter industry rules of their own; check with your legal team once more before implementing.

## Sources for this chapter

[1] Google | Email sender guidelines (SPF/DKIM and reverse DNS requirements; sending more than 5,000 messages a day enters bulk-sender requirements, and the spam complaint rate must stay below 0.3%) | in force from 2024-02-01 | https://support.google.com/a/answer/81126 | [official] [2] FCC | FCC Makes AI-Generated Voices in Robocalls Illegal (declaratory ruling FCC-24-17A1) | 2024-02-08 | https://www.fcc.gov/document/fcc-makes-ai-generated-voices-robocalls-illegal | [official] [3] FTC | CAN-SPAM Act: A Compliance Guide for Business (penalty cap of $53,088 per violating email; B2B email is not exempt) | penalties adjusted for inflation in 2024-01 | https://www.ftc.gov/business-guidance/resources/can-spam-act-compliance-guide-business | [official] [4] United States Code 47 U.S.C. §227(b)(3) (TCPA damages standard: $500 per call per violation or actual loss, whichever is higher, increased up to triple for willful or knowing violations) | https://www.law.cornell.edu/uscode/text/47/227 | [official] [5] EU General Data Protection Regulation, Regulation (EU) 2016/679, Article 83 (fine cap of €20 million or 4% of total worldwide annual turnover for the preceding financial year) | https://gdpr-info.eu/art-83-gdpr/ | [official] [6] State Council (Chinese: 国务院) | Regulation on the Implementation of the Law on the Protection of the Rights and Interests of Consumers (State Council Order No. 778), Article 24 | in force 2024-07-01 | https://www.gov.cn/zhengce/content/202403/content_6940158.htm | [official] [7] Personal Information Protection Law of the People's Republic of China, Articles 13, 14, 15, 23, 24, 55, and 66 | in force 2021-11-01 | https://www.cac.gov.cn/2021-08/20/c_1631050028355286.htm | [official] [8] LinkedIn | User Agreement (Article 8.2 Don'ts: prohibits scraping or copying service data and profiles, and prohibits using bots or unauthorized automated means to access and send messages) | effective 2025-11-03 | https://www.linkedin.com/legal/user-agreement | [official] [9] WeChat | Rules for the Management of External Link Content, Article 2.1.2 (prohibits inducing sharing through benefits such as monetary rewards, lotteries, and task check-ins) | [official] [10] hiQ Labs, Inc. v. LinkedIn Corp., 938 F.3d 985 (9th Cir. 2019); 31 F.4th 1180 (9th Cir. 2022); Supreme Court vacated and remanded 2021-06; N.D. Cal. 2022-11 found hiQ in breach of the LinkedIn user agreement, ultimately settled | [third-party verified] [11] State Administration for Market Regulation (Chinese: 市场监管总局) | Seven typical cases of online unfair competition (including a private-domain marketing software case that cracked WeChat's core database, fined 100,000 yuan) | Xinhua (Chinese: 新华网) | 2024-12-31 | http://www.news.cn/food/20241231/ceec0fa0841345fb872ed48b07dd3748/c.html |

[third-party verified] [12] China.com (Chinese: 中国网, central media) | AI outbound calling, a thousand harassing calls a day — where is the "compliance boundary" | 2026-07-28 | http://news.china.com.cn/2026-07/28/content_118621717.shtml | [third-party verified] [13] TechCrunch (Julie Bort) | Yes, LinkedIn banned AI agent startup Artisan. But now it's back | 2026-01-07 | https://techcrunch.com/2026/01/07/yes-linkedin-banned-ai-agent-startup-artisan-but-now-its-back/ | [third-party verified] [14] TechCrunch (Julie Bort) | Artisan, the "Stop Hiring Humans" AI agent startup, raises $25M and is still hiring humans (CEO admits the first-generation product had low reply rates, high churn, and severe hallucination) | 2025-04-09 | https://techcrunch.com/2025/04/09/artisan-the-stop-hiring-humans-ai-agent-startup-raises-25m-and-is-still-hiring-humans/ | [third-party verified] [15] Greenway | AI SDR: What Works, What Doesn't (800-account four-platform controlled experiment: 12% of first touches contained factual errors, 58% didn't fit the ICP, primary domain burned in 19 days, 8-week recovery, hybrid approach 3.2 times) | https://greenway.ai/post/ai-sdr-what-works-what-doesnt/ | [vendor claim] [16] Laxis Research | State of AI SDR 2026 (deliverability about 51% with less than 30 days of domain warming, about 86% at 60–90 days, about 91% above 90 days; daily send cap of about 25 messages per mailbox) | 2026-06 | https://www.laxis.com/blog/state-of-ai-sdr-2026/ | [vendor claim] [17] Digital Applied | AI SDR Statistics 2026 (100+ data points) | 2026-04-21 | https://www.digitalapplied.com/blog/ai-sdr-statistics-2026-outbound-sales-data-points | [third-party verified] [18] Autobound | AI SDR Tools Guide (including the UserGems churn-rate figure, and a 50%–70% annual churn rate for AI SDR tools) | 2026-02 | https://www.autobound.ai/blog/ai-sdr-tools-guide | [third-party verified] [19] Autobound | State of AI Sales Prospecting 2026 (signal-driven with human-supervised replies 15%–25% response rate, fully autonomous 1%–3%) | 2026-02-20 | https://www.autobound.ai/blog/state-of-ai-sales-prospecting-2026 | [vendor claim] [20] The Next Web (byline Belkins) | AI is turning B2B outreach into a tragedy of the commons (more than 7.5 million cold emails to brand-new contacts in 2025, average reply rate 0.45%; "interruption inflation") | https://thenextweb.com/news/b2b-outreach-tragedy-of-the-commons | [third-party verified] [21] Boomerang AI | Death to Cold Outreach (90% of sales calls go to voicemail, only 2% of cold calls earn a meeting; cold-call response rate 8.5% in 2018 → 4.2% in 2022; TCPA enforcement cases about 1,200 → more than 2,400) | https://www.getboomerang.ai/post/blog-death-cold-outreach | [vendor claim] [22] Coldreach | Claygent Review (Clay deducts credits per action; failed retries are charged the same, with no refund) | 2026-08 | https://coldreach.ai/blog/claygent-review | [third-party verified] [23] OFweek (Chinese: 维科号, Shuzhi

Qianxian) | Xiaoman Technology (Chinese: 小满科技) OKKI: full-scenario AI for foreign-trade CRM (cost per background check reduced to 5% of the original, 75% overall AI adoption rate) | 2025-04 | https://mp.m.ofweek.com/ai/a456714443567 | [third-party verified] [24] NetEase Zhiqi · Cloud Business (Chinese: 网易智企·云商) | AI private domain breaks "add and go silent" (the same user can be reached at most once in 24 hours) | 2026-06 | https://grow.163.com/cms/ai-ke-fu-ai-si-yu-2026-3RbSGjTy.html | [vendor claim] [25] 36Kr Research Institute (Chinese: 36氪研究院) / Tencent Smart Retail (Chinese: 腾讯智慧零售) | 2024 China Private-Domain Operations and Omnichannel Growth White Paper (WeCom private-domain repurchase rate 2.8 times, per-customer value up more than 45%) | https://36kr.com/p/2684930182745348 | [vendor claim] [26] Everyone Is a Product Manager (Chinese: 人人都是产品经理) | The "private-domain traffic" money-making formula (including WeChat's April 2019 ban of 30 million weishang marketing accounts) | 2020-07 | https://www.woshipm.com/operate/4088812.html | [third-party verified] [27] The Pipeline Group | The Hidden Dangers of AI SDRs (compliance red-line table) | https://www.thepipelinegroup.io/blog/the-hidden-dangers-of-ai-sdrs-why-they-should-never-be-used-for-outbound-b2b-enterprise-sales | [vendor claim] [28] China Academy of Information and Communications Technology (Chinese: 中国信息通信研究院) | Research Report on Artificial Intelligence Customer Relationship Management (AI CRM) Systems (2025) | 2025-09 | https://www.caict.ac.cn/kxyj/qwfb/ztbg/202509/P020250915603971582825.pdf | [official] [29] Centurium Capital (Chinese: 大钲资本) interview (Yang Fei of Luckin Coffee) | WeCom linked users 1.8 million → nearly 20 million, marketing-rule automation above 90% | https://www.centurium.com/press/%E7%91%9E%E5%B9%B8%E5%92%96%E5%95%A1%E6%9D%A8%E9%A3%9E%EF%BC%9A%E4%B | [first-person account: founder interview]

[30] Beijing Daily client (China National Radio investigation) | AI outbound calling at 800–1,500 dials a day and 0.12–0.16 yuan per minute; the compliance boundary is prior explicit consent | https://xinwen.bjd.com.cn/content/s6a68080fe4b0e45f3fd54adc.html | [third-party verified: authoritative media investigation] [31] MIIT and others | Work plan for a special campaign against nuisance calls (commercial marketing requires prior consent; do-not-call list; call recordings retained in principle no fewer than 30 days) | relayed by SecRSS | https://www.secrss.com/articles/6129 | [official: special campaign plan] [32] Renren dou shi chanpin jingli (Chinese: 人人都是产品经理) | Taking private-domain reach from 4.04% to above 10% (of 693 new members only 28 received the notice; repeat purchase rate

+5% year over year after three-dimensional dynamic segmentation) | https://www.woshipm.com/ai/6464019.html | [third-party verified: project's internal figures] [33] Xinhua (Chinese: 新华网) | France bans telemarketing without consent (effective 2026-08-11) | 2026-08-06 | https://www.news.cn/world/20260806/d1e4fade4a194183a4f89a6941ed507b/c.html | [official] [34] FTC | Telemarketing Sales Rule compliance guide (call-time windows, abandoned-call ratios, do-not-call and caller-ID requirements) | https://www.ftc.gov/business-guidance/resources/complying-telemarketing-sales-rule | [official] [35] Tencent Cloud developer community (Fengjian Liuli) | 2026 Private-Domain Watch: Scale Up, Efficiency Down (Jianshi's 2026 private-domain survey of 400+ practitioners: companies at RMB 50–100 million private-domain revenue rose from 4.13% to 7.09%, above RMB 100 million from 8.41% to 10.24%; market RMB 14.3 to 16.9 billion; 65% say SCRM rollouts fell short; 41.73% cite interaction fatigue; unit retention cost +35%) | 2026 | https://cloud.tencent.com.cn/developer/article/2703970 | [third-party verified]

# Chapter 19 Let Customers Bring Customers: Word of Mouth, Referrals, and Virality

Once the contract is signed and the delivery is done, most people's habit is to treat the deal as finished—the bit of trust the customer has built up just sits there, unused.

The next customer often arrives because someone brought them over, and that is trust the customer produced themselves. So this chapter gives you 5 methods to turn passive waiting into active triggering.

---

## Method 1: Ask for the referral at the customer's happiest moment

**What to do**: Don't casually drop "have a friend introduce me" when the contract is signed—the customer hasn't gotten results yet, and the more warmly you ask, the more politely he declines. Wait until he has used the results and is at his emotional peak, then get the referral out with one sentence.

**How**

1. First define what "the happiest moment" looks like in your business: the project passes acceptance, the first renewal, when he has just reported your results to his boss, when a failure was resolved by you within two hours. Industries differ, but it must map to a concrete action, not a vague feeling like "great working together."
2. Write this moment into the delivery process, and whoever owns that step does the asking: the project manager asks at project delivery, customer service asks when resolving a support issue. Don't save it all up until year-end and have sales send out one mass "please refer us."
3. Ask for only one action. In that single sentence, ask the customer for only one thing—"if anyone you know runs into the same problem, just pass my name along."

4. Give the customer a concrete person to think of. Ask "do you have a friend who needs this?" and most people can't think of anyone on the spot; ask "in your industry, aren't there a few companies dealing with the same thing?" and his brain starts naming them.
5. Any customer you haven't asked within a month goes back on the list and into the next trigger point. Referrals aren't a once-and-done ask; they're three or four asks a year.
6. Let the person who asked know the result. If the customer really refers someone and that person closes, send the news back to the project manager or support agent who asked. If the asker gets no feedback, the second ask is just going through the motions.

**Example**: Pinduoduo (Chinese: 拼多多) welded "invite friends to join a group buy" into the purchase-completion page—tap "pay now" to reach the completion page, tap invite, and a chat list pops up, merging purchase and sharing into a single action; based on its prospectus, 2017 customer acquisition cost was about 11 yuan per person [third-party verified: quoting the prospectus].

**Pitfalls**: Treating "asking for referrals" as a sales target pushed down onto sales, forcing them to ask everyone they meet. Once a customer feels he is being sold on a network of relationships, he closes up immediately; ask that one sentence at the wrong moment and you lose every opportunity after it.

## Method 2: Write the referral copy so all the customer has to do is forward it

**What to do**: Customers are willing to forward for you; they are not willing to write the words for you. You write the passage first, so all they have to do is copy, paste, and send—under ten seconds start to finish. Finding the words looks trivial, but it is the main reason most referrals stall at "I'll ask around for you sometime."

**How**

1. Prepare three lengths: a one-sentence version (for a WeChat group), a three-sentence version (for Moments or a friend's DM), and a one-page version (for a decision maker who has to pass it up the chain). Keep all three in the same place, so the customer can take whichever one he needs.

2. Every piece of copy must include: who the customer is, what problem he ran into, what you did, what the result was, and what kind of person you are looking for. For the result, put in a checkable number or fact wherever possible—only when the customer can state the details will the listener take it seriously.
3. Put no exaggerated words and no sensitive customer information in the copy. If you need to use the customer's name or data, ask first whether you can reveal it and to what extent.
4. Make a forwarding card: a block of text plus a link or QR code, so the customer doesn't have to dig through your website for an entry point.
5. Leave the customer room to say "you can change this directly." The copy is a template; the customer changes a couple of words into his own voice, and only then does it sound like something he said rather than you reading a script behind him.
6. Refresh the copy once a quarter. Forward the same passage for too long and the customer himself starts to feel he is doing advertising on someone's behalf.

**Example**: Duolingo has offered a figure: about 80% of acquisition comes from word of mouth, from viral distribution, which keeps its customer acquisition cost very low [first-hand interview: HBR podcast]. Getting the customer to do nothing but forward is the first half of that mechanism.

**Pitfalls**: Writing the copy as a formal letter of recommendation. The longer and more formal the script, the less the customer wants to send it—what he wants is a passage he can paste straight into the chat box.

## Method 3: Separate cash-back for buying reviews from rewards for referring new customers

**What to do**: Treat "spending money to buy reviews" and "spending money to buy new customers" as two separate things. The former is a practice the enforcement authorities have called out by name; the latter is the normal incentive for referrals. Mix the two into one campaign plan and when trouble comes you won't be able to tell which half crossed the line.

**How**

1. First shut down every incentive in the company with the word "review" attached: cash-back, rebates, free samples, subsidies for posting order screenshots, points for a five-star review—if what it buys is that review, it counts.
2. Write a separate rule for referral rewards; what it rewards is the act of "bringing in a real customer." What is rewarded, how much, and when it is paid out—write it into one document that both sales and support follow.
3. Don't tie the reward to "review content." What you give the referrer is a reward for "you brought in a new customer," not attached to a condition like "and please write us a good review while you're at it."
4. Keep asking for customer reviews as usual, just on a separate track: after delivery, ask the customer to speak honestly, with no reward, no filtering of content, and keep the bad reviews too.
5. File the reward disbursement records, the campaign copy, and the customer communications. If you are ever asked, having material that shows clearly "why this money was paid" beats trying to recall it after the fact.
6. To check the text of the statutes, go to the National Laws and Regulations Database (Chinese: 国家法律法规数据库) at flk.npc.gov.cn, search the regulation name on the site, and locate the specific clause page before reading.

**Example**: The legal-education article from the Quanzhou Municipal Administration for Market Regulation (Chinese: 泉州市市场监督管理局) has a blunt headline—"It's clear now! 'Cash-back for good reviews' is illegal! Fines up to 2 million yuan" [official]. This guidance covers the practice of buying reviews with cash-back; it is a different matter from rewarding old customers for bringing in new ones [our judgment].

**Pitfalls**: Thinking "as long as the reward is worded vaguely, I'm fine." Customers claiming a reward already feel the money is owed to them; vague wording only adds a round of trust erosion between you and them. Worse, when the two kinds of money are paid out in one plan, even you can't say afterward which is which.

## Method 4: Log every referral as a traceable chain

**What to do**: Give each existing customer a tagged referral entry point, so "who brought whom" stays in the system rather than depending on the customer mentioning it in a chat. Once the chain is connected, you know whom to reward, who is worth asking again, and which customer's referrals actually close.

**How**

1. Give each existing customer a dedicated QR code or link; when he sends it out, anyone who clicks through carries the source automatically.
2. On the new customer's first contact, add one question—"who referred you?"—and record that answer as a field, not just in the sales rep's chat history.
3. After a deal closes, update the status back on the referrer: how many people he brought, how many closed, whether the reward was paid. Review this table once a quarter.
4. Calculate the conversion rate of referred leads separately, apart from other channels. Mix them together and you neither know whether referrals work nor whom to reward.
5. Small teams without a system start with a spreadsheet: referrer, referred person, contact time, progress, deal amount, reward status—six columns is enough.
6. Once a quarter, look back at who the top few "bringers" are and single them out. These are the relationships you most need to maintain, worth a clearer rule of their own.

**Example**: HONOR (Chinese: 荣耀) hard-coded attribution—"when several sharers invite the same friend, the system rewards only the sharer who last clicked the link, and only once"—and made clear it collects the sharer's UID and the order number and model of the referred person [company official campaign rules]; Amazon Associates, by contrast, uses "one Tracking ID per channel" to keep the stats separate [official help center].

**Pitfalls**: Learning who referred someone only when someone asks. Break the chain and you reward from impression alone; the customer will soon notice that these rewards are handed out at random.

## Method 5: Don't trade short-term volume for incentivized-sharing gimmicks

**What to do**: Delete the "forward to win, collect likes to unlock, no sharing no prize" style of gimmick from your acquisition playbook. The volume it buys doesn't stick, and it also sits right on the platforms' red lines; if you do the math on cost, what you lose isn't just that giveaway budget.

**How**

1. Go through your existing campaigns and mark anything that says "must forward to claim," "share to N groups to unlock," or "only counts once you have enough likes" as this category.
2. Think through what these campaigns actually bring: mostly accounts that showed up for the freebie, not customers. Such accounts never close later, and they drag down the conversion figures for your whole lead pipeline.
3. For the reach you actually need, switch to "the content itself is worth forwarding": a set of data your peers don't have, a template that can be used directly, a public talk that is genuinely clear.
4. Before a campaign goes live, run it against the platform's rules first—don't wait to be actioned before going back to read them.
5. Keep an internal list of campaigns the platforms have actioned, noting exactly which sentence violated which rule. Check that list before the next campaign; it's less work than memorizing the rules.

**Example**: WeChat's "Regulations on the Management of External Link Content" spells out incentivized sharing clearly—links of this kind are subject to action inside WeChat [official]. Checking it clause by clause before a campaign is far cheaper than explaining afterward. **The rules are written by others, and changes to them won't notify you first**—taking the "must share" action out of the campaign is the only way not to gamble on the rules. [our judgment]

**Pitfalls**: Confusing "the customer is willing to forward on his own" with "incentivized sharing" in the rules, so you become afraid to do even normal content distribution. What the rules target is that forced "must share" action.

---

## Notes for this chapter

- This chapter does not give industry figures for referral rates or viral coefficients. The source material has no comparable data whose provenance can be verified, and putting numbers in could only be fabrication. [our judgment]
- The current version of the Anti-Unfair Competition Law was not located down to the clause page, so this chapter cites no specific article numbers. To check the statutory text, search the regulation name on the National Laws and Regulations Database (Chinese: 国家法律法规数据库) at flk.npc.gov.cn.
- No authoritative comparison explaining the difference between "cash-back for good reviews" and "reward for forwarding" was collected. Method 3 separates "buying reviews" from "rewarding referred new customers" on the basis of the enforcement authorities' public legal-education position that names the former, and is our judgment.
- The FTC guide is a US official position, covering review and testimonial disclosure in the US market. Whether it can be applied directly to your Chinese-language business depends on where you sell.
- This chapter contains no effectiveness multiples.

## Sources for this chapter

[1] WeChat | "Regulations on the Management of External Link Content" | https://weixin.qq.com/agreement/weixin_external_links_content_management_specification | [official] [2] Quanzhou Municipal Administration for Market Regulation (Chinese: 泉州市市场监督管理局) | "It's clear now! 'Cash-back for good reviews' is illegal! Fines up to 2 million yuan" | http://scjgj.quanzhou.gov.cn/xxgk/ztzl/zxdt/202406/t20240603_3043353.htm | [official] [3] US Federal Trade Commission (FTC) | "Endorsement Guides: What People Are Asking" | https://www.ftc.gov/business-guidance/resources/ftcs-endorsement-guides-what-people-are-asking | [official] [4] National Laws and Regulations Database (Chinese: 国家法律法规数据库) | statutory text search entry point | https://flk.npc.gov.cn/ | [official] [5] Monzo | Annual Report 2026 (about 79% of new customers from word of mouth) | https://monzo.com/annual-report/2026/Monzo-Annual-Report-2026.pdf | [company annual report] [6] Robinhood | S-1 prospectus (referral program accounts for 80%+ of new customers) | https://www.sec.gov/Archives/edgar/data/1783879/000162828021013318/robinhoods-1.htm | [company public filing] [7] HONOR (Chinese: 荣耀) | "Invite for Gifts" campaign rules

(last-click attribution, required fields) | https://www.honor.com/cn/msale/yqylhdgz.html | [company official rules] [8] Amazon Associates | Tracking IDs (one ID per channel) | https://affiliate-program.amazon.com/help/node/topic/GS4DGSRUYG5BX8E3 | [official help center] [9] Jiemian News (Chinese: 界面新闻) | Pinduoduo's share-based purchase chain ("invite friends to join a group" embedded in the purchase-completion page; per its prospectus, 2017 customer acquisition cost about 11 yuan per person) | https://www.jiemian.com/article/2317958.html | [third-party verified: quoting the prospectus] [10] HBR podcast | "How Duolingo aims to diversify beyond language learning" (about 80% of acquisition from word of mouth and viral distribution) | https://hbr.org/podcast/2025/04/how-duolingo-aims-to-diversify-beyond-language-learning | [first-hand interview: HBR podcast]

# Chapter 20 How to Know If It's Working: Three Sets of Books and a Holdout Test

It used to be enough to look at one attribution report to judge whether this money was well spent—who brought the click, who brought the conversion, credit goes to them, and the budget follows the number. Now the customer's comparison process on the AI side carries no link, no source, and never enters your backend; the report both misses the real effect and credits some channel with credit that doesn't exist. What you need to do is change the question: stop asking "which conversion should this be credited to" and start asking "what would happen if this money weren't spent," then write that answer as a range you can show your boss and your finance team.

Every method above costs money, and the effect isn't visible on the customer's side—all you have is the books. So this chapter gives you 5 methods to answer to your boss and your finance team with a range.

---

## Method 1: Split one attribution report into four sets of books, each answering only one question

**What to do**: Drop the habit of "using one report to answer every question," split the numbers a decision needs into four sets of books, and state what each set of books is allowed to be used for.

**How**

1. The first set is the attribution report (multi-touch attribution, platform backend). It answers "how do already-happened conversions get allocated"; the data is observational and is only allowed for tactical optimization inside a channel: creative, audience, bids.
2. The second set is self-reported channel research. It answers "how buyers remember they found you"; the data is self-reported and carries systematic bias, so it is only used to see whether the "invisible part's share" is growing or shrinking.

3. The third set is the AI visibility baseline. It answers "do you hold a place at the answer layer"; the data is what you measure yourself, and it is used to judge asset health and decide which sources to fight for first.
4. The fourth set is the incremental experiment. It answers "what would you lose if this money weren't spent"; it is the only one of the four with causal evidence, and only it can decide whether a budget stays or goes.
5. Write down next to the report why attribution fails. First, decisions at the answer layer are digested inside the conversation itself—the user doesn't need to reach your page to get the conclusion. Second, sessions that jump over from AI clients usually carry no usable source information, and land on the page mostly as "direct / no source." Third, conversational research happens entirely outside your analytics system, with no tracking at all. Worse: that small fraction that does carry source information is itself a biased sample—the people willing to click through and the people who take the answer and leave are not the same group. The part that can be measured is often exactly the most biased subset [our judgment].
6. Before every budget meeting, state which set of books this decision uses. The vast majority of misjudgments happen when "the first set of books answers the fourth set's question." Also, label the AI channel's attribution number in the report as "an attribution floor, not the size of the channel," and anyone citing it must also report the self-reported figure.

**Example**: A public analysis from a video infrastructure company (relayed by Derivatex Agency in May 2026) claimed that in one month 27% of new sign-ups self-reported discovering the company through ChatGPT, Perplexity, and Claude, while Google Analytics credited those AI channels with only 0.5% in the same period—a 54-fold gap [third-party verified]. This figure has only been relayed, never traced back to the original analysis, and neither the sample size nor the question wording is known, so read it only as an order of magnitude. On the other side: external media that AI frequently cites get less than 1% of their referral traffic from AI platforms [third-party verified]. Both numbers are right; they aren't answering the same question—27% is "the path the buyer remembers," and 0.5% and 1% are "the path the system can trace." Addressability is shrinking too: an analysis of sell-side data curation and identity graphs says only about 30% of web traffic is addressable [third-party verified] — cut the denominator by seven tenths and any cross-channel causal claim gets thinner. Platforms also rewrite the calculation itself, and what they change is precisely

an account's comparability across time: from January 1, 2026, Amazon Ads applies a "shopping-signal-enhanced last touch" attribution model to on-site Sponsored Brands (vCPM), Sponsored Display and on-site Amazon DSP campaigns, crediting outcomes to the early discovery moments — exploratory browsing, general category search — and shortening the attribution window, so purchases, sales and ROAS in standard reports change; click-based attribution is unchanged, and the 14-day "all views" metric stays available as a parallel reading [official: platform notice]. A year-over-year comparison across January 1, 2026 is reading a change in definition plus a change in reality.

Cross-channel comparison needs one unit, and that unit has a definition: incremental ROAS (iROAS) is revenue that happened because of the media — excluding conversions that would have occurred anyway — divided by spend; the companion “adjustment factor” is the share of a platform's reported conversions that is genuinely incremental, used to convert platform numbers into decision numbers; platform figures cannot be compared across channels while incremental figures can, but incremental figures depend on experimental design and statistical power, and an “incremental result” from an underpowered sample is just noise [third-party verified: measurement vendor]. Its role is not “another ROI project” but everyday infrastructure that puts every channel on one coordinate system — without it, moving budget between channels has no common basis.

**Pitfalls**: Reading GA's 0.5% attribution directly as "the AI channel is worth roughly zero," then cutting it in the budget meeting. The conclusion looks data-backed, so no one has to be accountable for it.

## Method 2: Fix the wording of self-reported research so that bias becomes a trend

**What to do**: Run a buyer self-reported channel survey once a month, using fixed wording and a fixed position to lock the bias down, so different months can be compared.

**How**

1. Change the question. Don't ask "how did you hear about us"; ask "before contacting us, which of these things did you do." Self-report answers memory and belonging, not causation—the most recent and most face-saving touchpoint is easier to recall, email and offline events are easily forgotten, the more

"insider" answer is easier to give, and what position the question sits in the form, single- or multi-select, and whether sales fills it in for them all directly change the answer distribution. Calibration isn't eliminating the bias, it's fixing it in place.

2. Options should be multi-select, open, and non-leading, and must include "asked an AI assistant," split by platform: ChatGPT, Doubao, DeepSeek, Kimi, Yuanbao, Perplexity, WeChat Search.
3. There is a measured reason to split platforms: only 11% of domains are cited by both ChatGPT and Perplexity [third-party verified]. Counting "AI search" as one channel amounts to adding two completely different behaviors together.
4. Account for the three kinds of touchpoints separately and don't force single-touch attribution: the same buyer can be recorded both as "asked AI" and as "read a third-party review"; when analyzing, produce numbers under both a first-touch and a last-touch convention, each labeled.
5. Weight by strata. Survey response rates differ by channel; large accounts almost never fill in forms, so weight by response rate, or at least report numbers by customer stratum rather than merging them into one total percentage.
6. Set two calibration anchors, with fixed frequency and position: the change in branded terms and direct site visits, and the relative change in competitive-term bid cost. If the self-reported AI share is rising while branded direct visits stay flat, it is mostly recall bias; only when the two move in the same direction can it be treated as a trend. Same questions, same position, same order every month; a single month's number does not enter a budget decision.

**Example**: The payments company Fiska built its own attribution with "form self-reported channel + UTM"; reporting says its leads from large models rose from 3 a month to 14, more than 50% of total leads [vendor claim]. The value of this case isn't in the numbers but in the act of "building your own ruler." The quality of AI-referred traffic already has industry comparison data. Adobe figures cited in a 2026 retail AI report put AI-referred traffic to US retail sites up 62% year over year in July 2026, with those visitors spending 59% more time on site, viewing 22% more pages per visit, converting 60% better and generating 53% more revenue per visit; in the same report's survey of more than 3,000 US consumers, 68% had used a general AI chatbot for shopping in the past three months, but only 10% of purchasers completed the purchase inside AI, 60% of AI-assisted shoppers prefer to buy through the website or

app where the product is sold, and 93% still sought proof outside AI before deciding [third-party verified]. So the AI-referral number cannot be read as a volume figure alone; it has to be counted together with the conversion and retention it produces.

**Pitfalls**: Making the survey a single-select dropdown at the end of the form, collapsed by default, with options only "search engine / social media / friend referral / other"—"asked AI" gets filed under "other," the self-reported figure can come out even lower than GA's, and the team concludes from this that "customers themselves don't even remember the AI channel." Survey design alone can create error of more than tenfold, and it never shows up in any data-quality report.

## Method 3: Build your own visibility baseline table and test the core questions monthly

**What to do**: Using a fixed set of buyer questions and fixed entry points, measure for yourself "what AI says about us now"—build the table once, reuse it long term.

**How**

1. First define the question set: the 30–60 questions buyers ask most, in four categories—awareness, comparison, risk, and selection—with 5 fixed "must be mentioned" hard questions.
2. Then define the entry points, setting priority by where your customers are, not chasing full coverage: overseas, ChatGPT, Google AI Overviews, Perplexity; in Chinese, Doubao, DeepSeek, Kimi, Yuanbao.
3. Each record must keep at least these fields: test date, whether the brand was mentioned, what position it ranked, in what context it was mentioned (recommendation / comparison / risk warning), which third-party domains were cited, the share of parameters and credentials described incorrectly, which direct competitors also appeared in the same question, and whether running the same question 3 times produced consistent results.
4. Fix the frequency: test the 10 core questions every month and all the questions every quarter.
5. Fix the test discipline: open a fresh environment with no history, fix the time window and region, copy the questions verbatim without "optimizing" them, and run each one 3 times.

6. Use a set of current-state data to calibrate expectations, only to judge what counts as normal, not as a target: one in-house measurement, an analysis report covering 14,750 buyer questions across four platforms, found an overall brand citation rate of 14.0%, with 53% of brands completely absent on questions in their own category; by industry, finance and insurance 2.1%, legal and professional services 86% never cited [vendor claim]. The gaps between platforms are even more extreme: ChatGPT cites a specific brand in only 0.59% of answers, Perplexity 13.05%, and Grok 27% [third-party verified].

**Example**: Our own measurement used 6 real procurement questions, asked twice through each of 5 model channels (A–E), yielding 51 valid answers: asking the same model twice in a row, the overlap of the lists was only 22%–56%; switching to another model, only 4%–21% [our own measurement]. Another industry analysis notes that 40%–60% of cited sources get swapped out every month [third-party verified]—so a single sample isn't a conclusion; "we got recommended this month" is not an asset, "consistently cited for six straight months" is.

**Pitfalls**: Drawing a rising "we got recommended by ChatGPT" curve in the weekly report. Single platform, single run, single question, paired with a monthly 40–60% turnover, amounts to drawing noise as a trend. Worse is tying visibility to bonuses—a manipulable metric will be manipulated; buying mentions and fake reviews can push the number up short term while pushing the trust asset down, and the part that falls doesn't show up in any report.

## Method 4: Run one holdout test, and first work out how big a difference you can detect

**What to do**: Artificially create a control group where "this money isn't spent," and use the difference to answer the causal question—but work out the resolution first, then decide the design.

**How**

1. Pick a design by your market structure. Geographic control: split markets into test and control groups and pause or halve the spend in the test group, provided the markets can be operated independently with no cross-region spillover. Time-based on/off: instead of splitting geographically, rotate on/off weekly nationwide and compare each market's on-weeks against its off-weeks, which naturally

absorbs regional differences but costs more weeks. Dose ladder: don't go to zero; run three levels of −50% / flat / +50%, which is more noise-resistant than "on/off" and easier for the business side to accept.

2. Work out the resolution first—this step decides whether the experiment succeeds. Suppose there are 20 markets with 500 qualified leads per market per month, estimate the standard deviation of month-to-month variation across markets at three levels—75 / 150 / 250 leads—and with 10 test markets against 10 control markets, at 95% confidence and 80% power, the smallest detectable difference is about 94 / 188 / 313 leads per market per month, i.e. 19% / 38% / 63% relative. The implication is direct: what a geographic control can reliably answer is "is this channel actually worth nothing," not "did a 5% lift do anything."
3. Work the observation window the other way. Suppose 10,000 qualified leads a month and you want to verify a real increment of about 5% for an AI channel; take half the markets as test: detecting a 20% relative change needs about 6,700 leads per group, roughly 1.3 months; detecting 10% needs about 28,400 per group, roughly 5.7 months; detecting 50% needs only about 900 per group, roughly 0.2 months. In one sentence: small effects need long windows; what can be confirmed quickly is a large effect.
4. So set the first round's goal as "ruling out catastrophic misjudgment"—is this channel completely ineffective, and will things collapse if the budget is moved away. To measure a small effect there are only two paths: switch the unit of observation to revenue or contract value (a larger single value), or accept a two-quarter window.
5. Get five things right, or seasonality and natural variation will contaminate the conclusion: the test and control groups must be compared in the same week and under the same promotion cadence, never against the same period last year; also look at the same group's data for the 4–8 weeks before the experiment as a before-and-after baseline, to offset the gap the two groups already had; the smallest observation unit is a whole week, and any window under 7 days doesn't count; before the experiment, write down the hypothesis, the primary metric, the observation window, and the stopping rule, and don't peek mid-way; when reporting, give the effect size, not just the direction.

6. Set the stopping rule before you start. Looking at the results halfway and deciding when to stop based on them will lift a nominal 5% false-positive rate to over 30%; you also have to handle multiple comparisons, the novelty effect, and proxy metrics failing—you optimized form submissions but the deals didn't move.

**Example**: Experimentation itself can be built up—Booking.com runs more than 25,000 experiments a year, and its head of experimentation says there were as many as 1,000 online at once [third-party verified: citing an official statement]. A holdout test is not a special case; it is the crudest one in that system. The motive for doing this sits in the books: a 2026 marketing decision confidence index shows 78% of U.S. senior decision-makers believe at least 10% of their marketing spend is wasted because of insufficient measurement, with 7% putting the waste at 30% or more [third-party verified]. The methodology has also converged: a holdout test deliberately withholds spend from a random sample closely matched to the exposed population while everything else runs as usual, and reads the difference between the two groups as the incremental effect — based only on the brand's own transaction data rather than platform-reported attribution; note that in this vocabulary the exposed group is the control and the withheld holdout group is the experimental arm [third-party verified]. Where randomization is impossible — budgets cut, a channel launched, a program paused — a substitute now exists: detect the day the intervention happened automatically, then build a counterfactual for causal impact analysis, which on Meta's public GeoLift benchmarks estimates a pull-back experiment's effect within 6.2% of the randomized experiment's result [third-party verified: peer preprint]. The cost of not testing can be priced: a practitioner account of Google Ads incrementality testing reports a Performance Max campaign including brand search looking strong on the dashboard while shutting it off cut total sales by only 30% of the Google-attributed value (the other 70% would have happened through other channels or organic search), and that Google's own reporting overstated PMax's impact by about 33% [third-party verified]; an earlier Meta internal study across 15 advertisers concluded only 34% of "attributed conversions" were truly incremental [third-party verified]. Practitioner attitudes are converging: in a January 2026 survey of 500 senior U.S. decision-makers, 60% trusted independent incrementality testing most, 40% media mix modelling and only 37% in-platform reporting, while 78% said poor measurement wasted at least 10% of budget [third-party verified]. On the other side, AI has scrambled the experimentation system itself: when idea generation and variant

building cost almost nothing, "how many more tests we ran than last quarter" proves nothing, and leading teams measure maturity by learning speed and retained knowledge instead [vendor claim].

The choice of attribution window can be treated as an experimental variable: in one self-run geo-holdout, a difference-in-differences estimate put incremental revenue at about $9,711, while a 14-day attribution window reported $9,831, a 7-day window $9,947 and a 30-day window $9,493 — all in the same range, a $120 spread; the 1-day window reported $7,333 and lifetime came in clearly low at $6,650 [third-party verified: single-account self-test]. The author's own conclusion stays restrained: this does not show that attribution can replace incrementality testing, since the two measure different things and here they merely happened to agree [third-party verified]. The reusable move is worth more than the conclusion: run attribution and difference-in-differences on the same batch of data and you can see how much the window choice distorts your own account, which is usually cheaper than switching attribution tools.

**Pitfalls**: Treating "it didn't collapse when I paused it" as "the channel is useless." A holdout test is good at answering whether a channel has failed wholesale; it cannot answer whether a small increment exists. When the interval crosses 0, the right move is to extend the window, not to pick a prettier number.

## Method 5: Split the budget into three buckets, hard-code the keep-or-cut rules, and report ranges not point estimates

**What to do**: Split the budget into three buckets and state the thresholds, hard-code the keep-or-cut conditions; when reporting, give both the attribution number and the experiment number, plus an interval.

**How**

1. Split the buckets and state the thresholds: the attributable bucket, recommended at no more than 50% of the total budget, used for tactical optimization within the bucket; the proven-increment bucket, whose keep-or-cut is decided by the interval from the incremental experiment; the under-validation bucket, recommended at 5%–10%, treated as experiment cost, carrying no KPI except "finishing the design."

2. Hard-code three keep-or-cut rules. Add budget only after ruling out the "increment = 0" hypothesis and only when the incremental ROAS is higher than the marginal ROAS of the next-best channel; when the interval crosses 0, hold and extend the window; asset investments don't take part in this round of trimming—they serve search, AI entry points, sales, and existing-customer referrals all at once, and cutting them is cutting the pass to the whole chain.
3. For strategic judgment, look at the "water level," not the "share." AI arms every seller at once, so an unchanged industry-wide share is the norm and isn't itself evidence of "useless"; the standard changes from "how much did this channel grow" to "how long until this moat expires."
4. Write decision outputs always as intervals. Reference format: a company has an annual paid budget of 1 million yuan; GA credits the AI channel with 0.5% attribution, i.e. 5,000 yuan; self-reported research shows 27% of new customers remember asking an AI. Run one 8-week on/off experiment, and after 4 weeks on and 4 weeks off the output looks like this—"pause investment in this channel; the incremental qualified leads point estimate is −5%, interval −1% to −9%; corresponding annualized revenue impact X to Y ten-thousand yuan."
5. Tie the decision to that output: if the interval is entirely negative and the lower bound exceeds the next channel's marginal ROAS, add budget; if it crosses 0, extend the window to 12 weeks; if the point estimate is near 0 and the upper bound is small, taper per the under-validation bucket.
6. Two reporting disciplines. Report the attribution number and the experiment number to management at the same time, and label the model source—merging the two into one "blended attribution number" is the fastest way for a marketing department to lose credibility [third-party verified]; every budget change must carry an evidence grade—changes backed by an incremental experiment can be executed directly, while changes backed only by an attribution report may only be adjusted within the bucket, not moved across buckets.

**Example**: One industry measurement guide notes that multi-touch attribution models over-attribute digital channels by more than 30% [third-party verified]; another attribution guide cites a CMO survey saying 62% of CMOs listed "proving ROI to finance" as their biggest challenge [third-party verified]. There are ready tools to lower the bar too: Google open-sourced the aggregated-data modeling framework Meridian in March 2024 and opened it to all marketing teams starting January 2025 [official].

**Pitfalls**: Having a direction for budget migration but no rule. The typical move is to cut 20% of the paid-media budget and add it to content and GEO, on the reasoning that "paid-media ROAS is falling and AI is the trend"—the part that was cut was never incrementally validated, the part added in has no observation window either, and a year later you still can't say what that money brought; you have just switched to a different story for the report.

---

## Notes for this chapter

- Attribution numbers and experiment numbers cannot substitute for each other, and cannot be merged into one number. The former answers "how do things that already happened get allocated"; the latter answers "what would be lost if this money weren't spent"; a holdout with no control group cannot answer a causal question [our judgment].
- Reporting an interval means showing the decision maker uncertainty itself, not offering a more conservative point estimate; when the interval crosses 0, don't force a direction—that just shifts the cost of judgment onto the person who has to decide [our judgment].
- Several key numbers in this chapter currently exist only as relays and have not been traced back to their original source: the 54-fold gap between 27% and 0.5%, Boring Marketing's visibility status quo, the 40%–60% monthly turnover, and Leapd's 680-million-citation analysis. Use them only to judge orders of magnitude; don't put them in a budget table, and don't write them into external materials [our judgment].
- The resolution and sample-size formulas are demonstrations: the month-to-month variation uses three assumed levels—75 / 150 / 250—and doesn't include design effects or the variance inflation of stratified sampling; before going live, first estimate real variation from your own market-level data for the last 12 months, then work the window backward [our judgment].
- Some situations aren't suited to a straight holdout: regulated industries should switch to a dose ladder and document beforehand that this is an experiment, not a reduction of service; for consumer businesses with small ticket sizes and short cycles, switch the primary metric to a more sensitive proxy such as branded search volume and direct visits, and move the budget in small 10%–20% steps [our judgment].

## Sources for this chapter

[1] Pew Research Center | "Google users are less likely to click on links when an AI summary appears in the results" (with an AI summary, traditional results get 8% clicks and without one 15%; links inside the summary get 1%; 18% of searches contain an AI summary) | 2025-07-22 | https://www.pewresearch.org/short-reads/2025/07/22/google-users-are-less-likely-to-click-on-links-when-an-ai-summary-appears-in-the-results/ | [third-party verified] [2] SparkToro | "In 2026, Less than One Third of Google Searches Still Send a Click" (US zero-click rate 68.01% in the first four months of 2026; 60.45% in 2024, data source Datos; about 45% ten years ago, data source Jumpshot) | 2026-06-08 | https://sparktoro.com/blog/in-2026-less-than-one-third-of-google-searches-still-send-a-click/ | [third-party verified] [3] Google official blog | "Everything announced at Google I/O 2026" (AI Mode tops 1 billion monthly actives; Gemini 3.5 Flash is the default model) | 2026-05-20 | https://blog.google/innovation-and-ai/technology/ai/google-io-2026-all-our-announcements/ | [official] [4] Gmail official help page | Bulk senders: spam complaint rate target below 0.1%, must never reach 0.3% | https://support.google.com/mail/answer/14229414 | [official] [5] Blake, Nosko & Tadelis (eBay) | NBER Working Paper 20171 (large-scale field experiment: for existing customers of an already-known brand, brand keyword ads have very little incremental effect) | https://www.nber.org/papers/w20171 | [third-party verified] [6] Kohavi, Tang & Xu | "Trustworthy Online Controlled Experiments" | 2019, Cambridge University Press (standards for pre-registration, observation windows, and multiple comparisons) | [third-party verified] [7] Google Meridian | open-source MMM framework (open-sourced 2024-03, opened to all marketing teams from 2025-01) | https://developers.google.com/meridian | [official] [8] Meta Robyn | open-source MMM tool | https://github.com/facebookexperimental/Robyn | [official] [9] Our own measurement 1 | candidate-set multi-model measurement (2026-09-13: sampling and stability of the visibility baseline; 22%–56% overlap across two runs of the same question, 4%–21% across models) | repository `实测/实测1-候选集多模型/` | [our own measurement] [10] Video infrastructure company Gumlet | the 54-fold gap between the self-reported channel and GA attribution (relayed by Derivatex Agency 2026-05, not traced to the original analysis) | [third-party verified] [11] Similarweb | 2026 generative AI brand visibility index (external media frequently cited by AI get less than 1% of their referral traffic from AI platforms, relayed through a third-party summary) | [third-party verified] [12] Leapd | 680-million-citation analysis (only 11% of domains are cited by both ChatGPT and Perplexity) | [third-party verified] [13] Fiska | form self-

reported channel + UTM self-built attribution (leads from large models rose from 3 to 14 a month, over 50% of total leads) | [vendor claim] [14] Boring Marketing | in-house visibility study (14,750 buyer questions × four platforms; overall brand citation rate 14.0%; 53% of brands wiped out in their own category; finance and insurance 2.1%; legal and professional services 86% never cited) | [vendor claim] [15] Brand citation rate comparison (34,234 AI answers: ChatGPT 0.59% / Perplexity 13.05% / Grok 27%, relayed through a third-party summary) | [third-party verified] [16] eMarketer | cited sources turn over 40%–60% monthly | 2026-04 | [third-party verified] [17] CXL / Peep Laja | 12 A/B testing mistakes (the same test showed a 10%–20% lift in November and was completely ineffective in April) | [third-party verified] [18] House of MarTech | 2026 measurement guide (multi-touch attribution over-attributes digital channels by more than 30%; where measurement programs fail points to organization, not technology) | [third-party verified] [19] TapClicks | 2026 attribution guide (citing NIQ 2026 CMO Outlook: 62% of CMOs listed "proving ROI to finance" as their biggest challenge) | [third-party verified] [20] Observer | "Why experimentation drives high-performing companies" (citing Booking.com head of experimentation Lukas Vermeer: 25,000+ experiments a year, 1,000+ concurrent) | https://observer.com/2026/01/why-experimentation-drives-high-performing-companies/ | [third-party verified: citing an official statement]

[21] AI Digital (citing Haus, 2026 Marketing Decision Confidence Index) | 78% of U.S. senior decision-makers believe at least 10% of marketing spend is wasted through insufficient measurement, 7% put it at 30% or more | https://www.aidigital.com/blog/what-is-incrementality-testing-in-marketing | [third-party verified] [22] Measured | What is a holdout test (holdout mechanics; based only on first-party transaction data; the exposed group is the control) | https://www.measured.com/faq/holdout-test | [third-party verified] [23] arXiv | Pseudo-incrementality testing with naturally occurring interventions (on GeoLift benchmarks, a pull-back experiment estimated within 6.2% of the randomized experiment) | arXiv: 2609.18257 | https://arxiv.org/abs/2609.18257 | [third-party verified: peer preprint] [24] AdExchanger | Sell-side curation and the identity graph (only about 30% of web traffic is addressable) | https://www.adexchanger.com/data-driven-thinking/sell-side-curation-and-the-identity-graph/ | [third-party verified] [25] Haus | Google Ads incrementality testing (switching off a PMax campaign including brand search cut total sales by only 30% of Google-attributed value; platform reporting overstated PMax's impact by about 33%) | https://www.haus.io/article/google-ads-incrementality-testing |

[third-party verified] [26] AdLibrary | Incrementality (Meta internal study across 15 advertisers: only 34% of attributed conversions were truly incremental) | https://adlibrary.com/posts/incrementality | [third-party verified] [27] Digital Applied | Marketing attribution models 2026 (500 senior U.S. decision-makers: 60% trust independent incrementality testing most, 40% MMM, 37% in-platform reporting; 78% say poor measurement wastes ≥10% of budget) | https://www.digitalapplied.com/blog/marketing-attribution-models-2026-what-still-works | [third-party verified] [28] Amplitude | AI broke experimentation (with idea generation free, test volume no longer proves health; maturity measured by learning speed and retained knowledge) | 2026-06-01 | https://amplitude.com/blog/ai-broke-experimentation | [vendor claim]

[29] Attentive (citing Adobe's AI traffic trends report) | The 2026 State of AI in Retail (3,054 US consumers; Adobe: July 2026 US retail AI-referred traffic +62%, time on site +59%, pages per visit +22%, conversion rate +60%, revenue per visit +53%; 68% used AI to shop, 10% completed the purchase inside AI, 60% prefer to buy on the seller's site or app, 93% still sought proof outside AI) | 2026 | https://www.attentive.com/state-of-ai-in-retail-2026 | [third-party verified]

[30] Amazon Ads | Update to view-through attribution for Amazon on-site ads (from 2026-01-01, Sponsored Brands vCPM, Sponsored Display and on-site DSP adopt "shopping-signal-enhanced last touch" attribution, weighting early discovery moments such as exploratory browsing and general category searches and shortening the attribution window; purchases, sales and ROAS change in standard reports; click attribution is unchanged; a separate 14-day "all views" metric remains as a comparison) | 2026-01-01 | https://advertising.amazon.com/zh-cn/resources/whats-new/view-attribution-updates-for-amazon-store-ads | [official: platform announcement] [31] Measured | Incrementality Benchmarks by Channel (iROAS = revenue caused by the media, excluding conversions that would have occurred anyway, divided by spend; the adjustment factor is the genuinely incremental share of platform-reported conversions) | 2026 | https://www.measured.com/faq/incrementality-benchmarks-by-channel/ | [third-party verified: measurement vendor] [32] Lebesgue | Meta Ads Incrementality Geo-Holdout Test (difference-in-differences estimate about $9,711; 1/7/14/30-day windows at $7,333 / $9,947 / $9,831 / $9,493; lifetime $6,650) | 2026-10 | https://lebesgue.io/facebook-ads/meta-ads-incrementality-geo-holdout-test | [third-party verified: single-account self-test]

# Chapter 21 Don't Bet on One Gateway: Dependency Check and Hedging

In the past, customer acquisition meant picking the most effective gateway and stacking budget on it—the higher the position, the more customers, and that lesson held for twenty years. Now buyers have moved to a different place to pick vendors—AI answers, platform recommendations, other people's reviews, buyers' procurement agents—and at every gateway the rules, the price, and the survival are decided unilaterally by someone else; the line that worked best for you yesterday may not count today. What you need to do is this: before any gateway gets rewritten, measure once how dependent you are on it, then give that dependence a few fallbacks.

The gateways where customers live don't belong to you, and the day the platform changes its rules you still have to be alive. So this chapter gives you 5 methods, starting with a dependency check.

---

## Method 1: Answer these twelve questions to measure your own dependency

**What to do**: Use a fixed questionnaire to turn "how dependent am I on a given gateway" into a score you can record and track.

**How**

1. Answer once a quarter, 12 questions, each worth 0–3 points, 36 points total. The person who fills it in is the person being asked—don't hand it to an assistant.

| Group | Question | 0 points | 1 point | 2 points | 3 points |
|---|---|---|---|---|---|
| Gateway structure | What share of first deals does your single largest acquisition gateway contribute (by the | >70% | 50–70% | 30–50% | <30% |

| Group | Question | 0 points | 1 point | 2 points | 3 points |
|---|---|---|---|---|---|
| | signed / payment-received basis, not lead count) | | | | |
| Gateway structure | If this gateway's rules changed, could you find an equivalent substitute within 30 days | no substitute | substitute exists but the cost is unbearable | substitute exists, cost is bearable | multiple gateways already running in parallel |
| Gateway structure | Does your acquisition chain depend on one platform's sharing / redirect capability | fully dependent | mostly dependent | partly dependent | not dependent at all |
| Gateway structure | If the buyer goes through no platform at all, can they still find you | no | only by remembering the brand name | has its own channel | has its own touchpoint with sustainable return visits |
| Asset ownership | In whose database is your customer relationship record kept | all on the platform | mostly on the platform | mostly your own first-party data | first-party data complete and directly reachable |
| Asset ownership | If the gateway disappears, can your content / data assets still be found by buyers | no | only scattered remnants | yes, but with no third-party support | yes, and with cross-ecosystem third-party source support |
| Asset ownership | | entirely concentrated | mostly concentrated | somewhat dispersed | dispersed across ecosystems |

| Group | Question | 0 points | 1 point | 2 points | 3 points |
|---|---|---|---|---|---|
| | Are your third-party sources concentrated within a single platform ecosystem | | | | |
| Asset ownership | Does your attribution and measurement depend on a single platform's self-reported data | fully dependent | mostly dependent | has a partly self-built measure | has a self-built incremental-experiment measure |
| Protocols and interfaces | For agent / buyer-system integration, have you implemented only a single protocol or a single vendor's field spec | yes, with no alternative | yes, but with a migration plan | two paths already supported | protocol-neutral structured assets |
| Protocols and interfaces | Do price, inventory, lead time, credentials, and return / exchange terms have an authoritative, verifiable, machine-readable single source of truth | no | yes, but scattered across multiple systems | unified source, but not machine-readable | yes, with version history and a change audit trail |
| Governance | In the last four quarters, have you made an explicit decision about "gateway dependency" (rather than continuing to invest by default) | never discussed | discussed, no conclusion | conclusion reached, not recorded | every quarter has a resolution and a record |

| Group | Question | 0 points | 1 point | 2 points | 3 points |
|---|---|---|---|---|---|
| Governance | In budget / assessment, is there an item that explicitly rewards "reducing the share of any single gateway" | no, only lead counts are assessed | raised but can't be executed | there is a metric, with very low weight | a metric on par with lead counts |

1. Work out the total first, then pick the three lowest-scoring questions—those are the only direction you touch this quarter. 0–12 is high risk, 13–24 moderate, 25–36 relatively healthy.
2. Put this table into an existing quarterly meeting, no more than one page, answering only three things: how much the largest gateway's share changed (using the first-deal basis, not the lead basis), what the rule-change log updated (mark each one "affected / not affected / watching"), and what the exit plan updated.
3. Carve out a fixed proportion of the total acquisition budget, allowed to go to only two places: non-largest gateways, and building your own assets. This proportion has no universal number—it is determined by your cost structure—and cannot be raided mid-quarter.
4. List "share of any single gateway's contribution" and "own-asset increment" (net increase in the first-party list, net increase in case-library entries, third-party source coverage) as metrics on a par with lead counts. Set no absolute target, only a direction.
5. Write a one-page plan: if this gateway disappears tomorrow, what do I do in week one. Put it somewhere others can find it.

**Example**: In its FY2018 10-K, Zynga stated itself: 51% of revenue and bookings came from Apple, 38% from Google, and 9% from Facebook, the three totaling 98%; in 2017 it was 51%/33%/12%, barely moving over two years. This is a sample of a company writing "the gateway doesn't belong to me" into a regulatory filing in its own words—a dependency structure doesn't improve just because you know it exists [official: verbatim from SEC filing]. What revocation looks like is concrete too: a risk-control misjudgement at Kimi (Moonshot AI) wrongly banned accounts at scale, with the company ignoring the third-party developer ecosystem until developers protested

by halting work; the risk rules were then adjusted, the Code Agent developers contacted to refine User Agent detection, and every restricted account restored [third-party verified]. For a product living on someone else's interface, a ban is not an availability problem but an existential one — how fast you recover depends on whether the platform staffs developer relations.

**Pitfalls**: Handing the 12 questions down as a KPI. The team will prop up the most useless channels to score points, putting a little spend on three platforms with almost no customers, turning "multiple gateways" into a formality. Broad but same-source is still the same risk.

## Method 2: Turn rented position into assets you can take with you

**What to do**: Build a first-party identity so buyers can reach you directly, with no platform in between.

**How**

1. In the Chinese ecosystem this has a mature form: WeCom contacts, communities, official accounts / service accounts, mini programs, and SMS form a direct channel that doesn't depend on any single platform's algorithm. In the overseas ecosystem the counterpart is the email list, communities, and your own app.
2. Fit this into the actions you were going to take anyway, without adding a project: after each delivery, casually ask for a publishable customer testimonial; at each trade show, event, or support conversation, casually collect first-party contact details; give every content asset its own "can be contacted" entry point, and don't hide contact details behind a form. The acceptance standard changes from "the immediate effect of this touch" to "what this touch left behind."
3. The form of an asset can be a relationship, or it can be data. Turning content assets into data licensing is another path.
4. There is only one test: on the day the gateway disappears, can this asset call the customer back? If it can't, it doesn't count as an asset.
5. State the boundary clearly. A platform's compliance promise is not your asset: on 2025-04-22 Privacy Sandbox announced it would keep Chrome's existing third-party-cookie user-choice scheme and would not launch a new standalone

third-party-cookie consent prompt, and a batch of companies that had already rebuilt their attribution systems around "the death of the cookie" took the sunk cost [official]. First-party data doesn't depend on whether the platform keeps its word, which is why it is the asset.

**Example**: Watsons (Chinese: 屈臣氏) has 4,100+ stores and 65 million members across more than 490 cities in mainland China; in 2020 it launched full-chain private-domain operations, using "store staff adding contacts + WeCom / official account / mini program / community / livestream" to run layered member operations [third-party reported]; Luckin Coffee (Chinese: 瑞幸) built a super CDP system with over 20 million members and private-domain users and 35,000+ communities [third-party reported]. Stack Overflow took another path: licensing its community Q&A data to AI companies and pivoting to enterprise internal-knowledge solutions [third-party reported]. Same mechanism, two forms. The counter-example is Mailchimp: on its May 2026 earnings call Intuit gave growth twice — Global Business Solutions up 15% (17% excluding Mailchimp) and Online Ecosystem up 19% (22% excluding it) — while confirming Mailchimp revenue declined year over year and announcing a 17% workforce reduction of roughly 3,100 roles, with the wording naming "rightsizing in Mailchimp" directly [third-party verified]. The channel-level evidence speaks even more directly to this chapter: where new applications get built — AI agent ecosystems and app directories — Mailchimp appears in none of the lists, neither as an option nor as a connector or skill [third-party verified]. Its traffic has not disappeared; it is simply absent from the new gateways' lists. The same kind of decay follows a pattern: Chartbeat data shared with media shows publisher traffic from Google Search down 34% over the past year, and over two years small publishers lost 60% of search referrals, medium publishers 47% and large publishers 22% [third-party verified] — channel contraction is not an even haircut; it kills the weakest negotiators first. The other end of platform-imposed rule changes now has regulators in the room. Cloudflare filed a position with the UK Competition and Markets Authority on its proposed conduct requirements for Google, arguing that the only effective remedy is to force Googlebot to be split into separate crawlers by purpose so that publishers can allow search indexing while refusing generative AI use of their content, and noting that Daily Mail Group, The Guardian and the News Media Association have publicly backed that approach; its premise is entry dependency — publishers cannot afford to block Googlebot because search traffic is what keeps them alive, so the same content feeds both the index and, for free, generative AI [official: corporate position]. Data

sources are not assets either: on 30 September 2026 Reddit announced a phased shutdown of its public Data API — new public API requests stop on October 31, RSS for all subreddits stops on November 13, existing apps must be registered by January 12, 2027, and the free public API closes in March 2027; listening suites with commercial agreements keep their feeds, while what gets cut off is automation scripts, RAG pipelines that pull community posts into an internal assistant, and small tools that never made their data source clear [third-party verified]. Teams that use community content for topic selection and reputation monitoring either buy a commercial agreement or convert community insight into own-channel inputs — interviews, their own community, support records.

**Pitfalls**: Treating "how much private domain can grow" as the conclusion. These samples show the structural fact that "these kinds of companies treat first-party members as a hedging asset"; the multiples in them cannot be taken as an effectiveness promise.

## Method 3: Look at the failure correlation between gateways before deciding how many to spread across

**What to do**: Change "multiple gateways" from broad coverage into low failure correlation.

**How**

1. First build a counterexample on yourself: distribute the same piece of content to ten platforms, and if those ten platforms all depend on the same "content to click" logic, the same algorithm family, and the same regulatory jurisdiction, then it is still the same risk exposure across ten accounts—one classifier adjustment and ten places drop at once.
2. Compare on at least three dimensions: regulatory jurisdiction (are these gateways governed by the same law and the same regulator), medium (are they the same kind of distribution logic—search / recommendation / social / IM / offline), and buyer path (buyers researching on their own, peer referrals, or agent procurement).
3. Log each channel's "share of organic traffic" and "sensitivity to platform rule changes" in the same table, compute a concentration index, and look at its direction every quarter. The trend matters more than the absolute value.

4. Prepare for only the most likely way you would die. If your category is one AI can answer directly (question banks, news, tools, common Q&A), you are more likely to die from being replaced, so prioritize relationships and subscriptions; if distribution must pass through one big platform's sharing / payment / account system, you are more likely to die from being revoked, so prioritize multiple channels and your own touchpoints. Doing one thing well for the most likely death beats doing 10% on each of four lines.

**Example**: The daily-deal pair wrote "the correlation of gateway failure" most clearly: LivingSocial's valuation once reached $6 billion and it was ultimately sold for $0; Groupon's share price fell from $28 to around $4 [third-party verified: academic]. Two lives that looked like two gateways actually shared the same set of premises. Cross-market visibility also proves that a gateway is not one gateway: an AI-search index report built on more than 1.9 billion real user conversations says that in 24 of 30 multi-region industries at least one European market crowned a No. 1 brand that never cracked the U.S. top five, that only seven brands ranked first in every available market, that in markets where a net-new local brand led, 46%–59% of citations went to country-specific domains — 12–17 points higher than where no local brand led, and that after a June 2026 ChatGPT update citations per answer fell 10.4% and unique cited domains dropped in 56 of 58 industries while most leading brands held their position [vendor claim] — the same model update reshuffles who backs you, not what rank you hold.

The decline of the first gateway now has comparable longitudinal data: a randomised experiment with more than 1,000 Google users found outbound organic clicks down 39.8% when an AI Overview appeared and zero-click searches up 34.5%; Pew's observation of nearly 69,000 searches by 900 U.S. adults found 8% clicked a traditional result when a summary was present against 15% without one, while clicks on links inside the summary were 1%; on the traffic side, Chartbeat data across 2,576 publishers recorded Google Search referrals down 33% globally and 38% in the U.S. between November 2024 and November 2025, while Google Search still accounted for 7.3% of publisher referrals and ChatGPT for 0.02% — Google's referral volume is roughly 500 times ChatGPT's; generative AI referrals grew 86% year on year yet made up only about 0.7% of all web traffic [third-party verified]. What matters here is not how much AI took but how differently the two curves slope: the gateway's decline is gradual while the replacement is still small, and that gap is exactly when a second gateway has to carry weight.

**Pitfalls**: Computing concentration as a number you have to hit. This book gives no threshold for "a single channel should not exceed X%"—industries differ too much; use your own four consecutive quarters of data as the baseline and look only at the direction of the trend.

## Method 4: Turn the buyer's session into a relationship you can take with you

**What to do**: Beyond the session logs held by the buyer and the model, hold a relationship with the buyer yourself.

**How**

1. Be directly contactable: a clear brand name, plus a way to reach you that doesn't depend on a platform redirect. The brand name matters especially—after agents collapse the source away, a name the buyer can say out loud and search is the only thing that can still pass through the collapse. In the Chinese ecosystem this step is often just adding a contact; overseas it is subscribing to a newsletter.
2. Be directly returnable to: your own site, your own app, your own community, offline—the buyer can come back without going through any intermediary layer.
3. Evidence that can be remembered: the case library and third-party sources are the external form of "relationship" on the buyer's side, letting the buyer confirm you even when your salesperson isn't present.
4. There is only one acceptance standard: in the week the gateway disappears, how many customers can find you on their own. "How many contacts added, how many subscribers" is a vanity metric; this one is a structural metric.
5. Minimum-threshold action: every time before sending an outreach email, ask one question—"does this letter make the brand name clear?"

**Example**: Chegg's CEO said on the 2023-05-01 earnings call that students' interest in ChatGPT had surged since March and was affecting the new-customer growth rate; the next day the stock closed down 48.41% at $9.08 [third-party reported]. Stack Overflow's CEO, looking back, said ChatGPT put the company into an "existential moment," that he triggered an internal "Code Red," and that he reassigned about 10% of staff to tackle it [third-party reported]. Neither of these two was penalized by a platform; the buyers changed gateways themselves.

**Pitfalls**: Understanding "relationship" as having added a contact or pulled people into a group. An agent won't maintain the relationship for you—its goal is to maximize buyer utility and it has no need to remember whether you are reliable, so the relationship must be held by you.

## Method 5: Don't bet on a single agent protocol; make facts into protocol-neutral assets

**What to do**: Make verifiable facts into structured assets that aren't tied to any one protocol.

**How**

1. Price, inventory, lead time, credentials, return / exchange terms, and liability boundaries—these facts need an authoritative, verifiable, machine-readable single source of truth, with version history and a change audit trail. Such an asset can be verified under any protocol.
2. Treat integration implementations as replaceable parts, not long-term investments. Customize only for one vendor's field spec and you will have to redo it at the next protocol turnover, and with shorter notice.
3. Remember the division of labor: structuring and verifiability are protocol-neutral; integration implementation is protocol-bound. Put resources into the former.
4. Keep two agent-facing integration routes; don't pick a side.
5. Treat the scope of authorization as a variable that changes. Today the buyer lets an agent "look only, don't buy"; tomorrow it lets it "close deals automatically within budget"—when the scope of authorization changes, your acquisition strategy has to change with it, and you won't get advance notice.

**Example**: On 2025-09-17, Google released the Agent Payments Protocol (AP2), advanced jointly with more than 60 payments and technology companies. Its definition of the problem is blunt: existing payment systems assume "a human clicks buy on a trusted interface," and autonomous agents initiating payments break that assumption, so it has to solve three things—authorization, authenticity, and accountability—through Mandate, a tamper-proof encrypted contract the user signs with a verifiable credential signature, forming a non-repudiable audit chain; the text explicitly points to the B2B scenario, where enterprises can use it to "autonomously procure partner solutions through Google Cloud Marketplace" [official]. AP2 and

OpenAI's and Stripe's ACP are two parallel routes [official]. Platforms are widening the grounds for taking access away: on January 15, 2026 X announced a developer API policy change banning apps that "reward users for posting," on the grounds that they flood the platform with AI spam and spam replies, removing API access from InfoFi-style apps [third-party verified]. The implication for growth teams is direct: when the base of your acquisition or monetization model is somebody else's API and incentive rules, a single policy change can take away both distribution and revenue at once — and the stated reason is usually "content quality." The same mechanism has a second summary in the developer ecosystem: platforms are moving from open sharing to closed monetisation, and the API has gone from an ecosystem-building tool to an instrument of revenue growth and data defence, leaving developers exposed to extreme platform risk — a business built over years can go to zero on one policy change [third-party verified]. The agent-store slot has been falsified once as well: in July 2026 Doubao, Qwen and Yuanbao all took down their user-built agent marketplaces, with Doubao's marketplace holding about 8 million user-built agents against roughly 200 million daily actives yet costing RMB 130–240 million a day and earning under RMB 1 million a day; of roughly 3 million GPTs created only about 159,000 are public and active, with the median creator earning USD 0–100 a month [third-party verified]. Easy to list, zero distribution, negative unit economics for the developer — when a gateway's economics do not hold, spreading across more gateways will not save you.

**Pitfalls**: Treating protocol integration as "get on early for first-mover advantage." Protocols will turn over, and the notice period at turnover only gets shorter.

---

## Notes for this chapter

- The two ways a single gateway dies are not decided by the quality of your operation. Being replaced is "the gateway is still there, you're not" (Chegg, Stack Overflow, Stereogum and the like); being revoked is "you're still there, the gateway is gone" (ATT, WeChat's external-link blocking, the TikTok divestiture law). There is only one way to diagnose replacement: see whether exposure-type metrics diverge from real acquisition metrics—in a case PPA submitted to the UK Competition and Markets Authority, an automotive content publisher's top-ranked article saw traffic fall 25%, while over the same period search

visibility rose 7% and CTR fell from 2.75% to 1.71% [third-party reported]. If the divergence is confirmed for two consecutive quarters, it is not an optimization problem, it is a gateway problem.

- The "predictability" of platform policy is itself a source of risk. Privacy Sandbox, from launch to U-turn, is the best evidence: you can't make long-term investment decisions on "my judgment is this policy will land." That is why the hedging allocation in Method 1 can't follow the current-period ROI ranking—the pressure to raid it always shows up when the largest gateway is performing best, which is exactly when it should least be raided.
- Ownership and survival of the gateway are not in your hands. On 2025-09-02, the US Department of Justice won remedies in the Google search monopoly case; the court banned exclusive contracts and required opening part of the search index and user-interaction data, on the factual basis that Google had long held about 90% of all US search queries and the court had earlier issued a 277-page opinion finding monopoly [official]; on 2025-01-17, the US Supreme Court upheld the constitutionality of the TikTok divestiture law, and the distribution channel for 170 million US users faced "divest or shut down" [official: verbatim from the opinion]. Both can only be covered by multiple channels and your own assets; operational optimization can't cover them.
- The scoring tiers of the twelve questions and the four-part split of the hedge kit (assets / channels / relationships / protocols) are both [our judgment]; this is a work checklist, not a validity-tested measurement instrument. The 0/1/2/3 cut points are set by mechanism intuition; recalibrate them to your own industry structure and compare only against yourself.
- The figures in the Chinese-side private-domain samples (Watsons, Luckin, Perfect Diary, LYFEN, NetEase Zhiqi) are all media, vendor, or agency figures, and mostly self-reported single cases. This book uses them to illustrate the structural fact that "these kinds of companies treat first-party members as a hedging asset," and does not use the multiples in them as an effectiveness promise. Meta's own estimate of the ATT hit, about $10 billion, is used only for magnitude—management said on the same call that "we can't be precise, this is an estimate" [management statement, verified through third-party reporting].

## Sources for this chapter

[1] Google | "Introducing AI Overviews" | 2024-05-14 | https://blog.google/products/search/generative-ai-google-search-may-2024/ | [official] [2] Google | "Everything announced at Google I/O 2026" (AI Mode tops 1 billion monthly actives) | 2026-05-20 | https://blog.google/innovation-and-ai/technology/ai/google-io-2026-all-our-announcements/ | [official] [3] CNNIC (China Internet Network Information Center) | the 57th "Statistical Report on China's Internet Development" (602 million generative AI users, 42.8% penetration) | published 2026-02-05 | https://www.cnnic.com.cn/IDR/ReportDownloads/ | [official] (consistent across authoritative reprints such as Xinhua) [4] U.S. DOJ | "Department of Justice Wins Significant Remedies Against Google" | 2025-09-02 | https://www.justice.gov/opa/pr/department-justice-wins-significant-remedies-against-google | [official] [5] Google Cloud | "Powering AI commerce with the new Agent Payments Protocol (AP2)" | 2025-09-17 | https://cloud.google.com/blog/products/ai-machine-learning/announcing-agents-to-payments-ap2-protocol | [official] [6] American Bar Association (Business Law Today) | "BC Tribunal Confirms Companies Remain Liable for Information Provided by AI Chatbot (Moffatt v. Air Canada)" | 2024-02-14 | https://www.americanbar.org/groups/business_law/resources/business-law-today/2024-february/bc-tribunal-confirms-companies-remain-liable-information-provided-ai-chatbot/ | [third-party reported] [7] Ars Technica | "Company apologizes after AI support agent invents policy that causes user uproar" (Cursor) | 2025-04-17 | https://arstechnica.com/ai/2025/04/cursor-ai-support-bot-invents-fake-policy-and-triggers-user-uproar/ | [third-party reported] [8] FTC | "Trade Regulation Rule on Unfair or Deceptive Fees" | 2024-12-17 | https://www.ftc.gov/news-events/news/press-releases/2024/12/federal-trade-commission-announces-bipartisan-rule-banning-junk-ticket-hotel-fees | [official] (effective 2025-05-12) [9] arXiv:2507.05301 (Kai-Cheng Yang) | citations in AI search answers are highly concentrated | 2025-07-07 | https://arxiv.org/html/2507.05301v1 | [academic preprint, before peer review] [10] Our own measurement 1 · candidate-set multi-model measurement | 2026-09-13 | method and raw data at `实测/实测1-候选集多模型/` | [our own measurement] [11] CNBC (Lauren Feiner) | "Facebook says Apple iOS privacy change will result in $10 billion revenue hit this year" | 2022-02-02 | https://www.cnbc.com/2022/02/02/facebook-says-apple-ios-privacy-change-will-cost-10-billion-this-year.html | [management statement, verified through third-party reporting] [12] Google Privacy Sandbox Blog (Anthony Chavez) | "Next steps for Privacy Sandbox and tracking protections in Chrome" | 2025-04-22 | https://

privacysandbox.google.com/blog/privacy-sandbox-next-steps | [official] [13] Digiday (Ivy Liu) | "Google AI Overviews linked to 25% drop in publisher referral traffic" | 2025-08-15 | https://digiday.com/media/google-ai-overviews-linked-to-25-drop-in-publisher-referral-traffic-new-data-shows/ | [third-party reported] [14] AdExchanger (Anthony Vargas) | "The AI Search Reckoning Is Dismantling Open Web Traffic" | 2026-01-06 | https://www.adexchanger.com/publishers/the-ai-search-reckoning-is-dismantling-open-web-traffic-and-publishers-may-never-recover/ | [third-party reported] [15] Stereogum (Scott Lapatine) | "Getting Killed By AI" | 2025-11-10 | https://stereogum.com/2478838/stereogum-relaunch/news | [official: founder's own account] [16] Zynga Inc. | Form 10-K (FY2018), platform concentration disclosure (Apple 51% + Google 38% + Facebook 9% = 98%) | filed 2019-02-27 | https://www.sec.gov/Archives/edgar/data/1439404/000156459019005098/znga-10k_20181231.htm | [official: verbatim from SEC filing] [17] CNBC (Sarah Min) | "Chegg shares drop more than 40% after company says ChatGPT is killing its business" | 2023-05-02 | https://www.cnbc.com/2023/05/02/chegg-drops-more-than-40percent-after-saying-chatgpt-is-killing-its-business.html | [third-party reported] [18] The Verge《Decoder》 | "Stack Overflow users don't trust AI. They're using it anyway" | 2025-12-15 | https://www.theverge.com/podcast/844073/stack-overflow-ceo-ai-coding-chatgpt-code-red-interview | [third-party reported] [19] China News Service (Chinese: 中国新闻网) | "MIIT: Blocking website links is one of the key problems in the special rectification of the internet industry" (State Council Information Office press conference) | 2021-09-13 | https://www.chinanews.com/cj/2021/09-13/9564236.shtml | [third-party reported] [20] European Commission | "Commission provides guidance under Digital Markets Act to facilitate development of innovative products on Apple's platforms" | 2025-03-19 | https://digital-markets-act.ec.europa.eu/commission-provides-guidance-under-digital-markets-act-facilitate-development-innovative-products-2025-03-19_en | [official] [21] Supreme Court of the United States | "TikTok Inc. v. Garland, 604 U.S. ___ (2025)" | 2025-01-17 | https://www.supremecourt.gov/opinions/24pdf/24-656_ca7d.pdf | [official: verbatim from the opinion] [22] FTC | "FTC Announces Final Rule Banning Fake Reviews and Testimonials" (16 CFR Part 465) | 2024-08-14 | https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-trade-commission-announces-final-rule-banning-fake-reviews-testimonials | [official] (effective 2024-10-21) [23] FTC | "FTC Takes Action Against TruHeight" (thousands of employee-written five-star reviews + bot accounts; $4 million judgment) | 2026-04-13 | https://www.ftc.gov/news-events/news/press-releases/2026/04/ftc-takes-action-against-truheight-deceptive-unsubstantiated-

advertising-supposed-height-enhancing | [official] [24] Google | "New Gmail protections for a safer, less spammy inbox" | 2023-10-03 | https://blog.google/products-and-platforms/products/gmail/gmail-security-authentication-spam-protection/ | [official] [25] Everyone Is a Product Manager (Chinese: 人人都是产品经理, Yan Tao Sanshou column) | "A full teardown of Watsons' private domain" (65 million members, 4,100+ stores, 5,000+ communities) | 2022 | https://www.woshipm.com/operate/5452718.html | [third-party reported] [26] Jazzyear (Chinese: 甲子光年, interview with Luckin senior technical director Zhao Xu) | "Luckin reshapes the coffee business with AI" (super CDP, over 20 million members / private-domain users, 35,000+ communities) | 2025-05 | https://www.jazzyear.com/article_info.html?id=1496 | [third-party reported] [27] Aike SCRM (Chinese: 艾客 SCRM) | "A teardown of Perfect Diary's private-domain community operations" | 2023-01 | http://www.ikscrm.com/school/syll/1493.html | [third-party reported: vendor perspective] [28] Southern Metropolis Daily (Chinese: 南方都市报; republished by NetEase News) | "As the traffic dividend peaks, how to do marketing and customer service" (Tencent Qidian + LYFEN: 3,700+ stores, 45 million+ user records) | 2024-01 | https://www.163.com/dy/article/IOTLLO6Q05129QAF.html | [third-party reported] [29] NetEase Zhiqi (Chinese: 网易智企; reposted by Weiyangx) | "AI private-domain operations solution" (four-layer stratification; insurance customer case) | 2025 | https://www.weiyangx.com/441234.html | [vendor claim] [30] Wharton Knowledge | "The Death of the Daily Deal" (LivingSocial's $6 billion valuation → sold for $0; Groupon's share price $28 → $4) | https://knowledge.wharton.upenn.edu/article/death-daily-deal/ | [third-party verified: academic]

[31] Axios (citing Chartbeat) | Google Search traffic down 34% in a year; over two years small publishers −60%, medium −47%, large −22% | 2026-07-31 | https://www.axios.com/2026/07/31/google-search-publishers-seo-geo-llms-ai | [third-party verified] [32] SaaStr | How Mailchimp Went From $1B+ ARR to Shrinking Inside Intuit (15%→17% and 19%→22% excluding Mailchimp; 17% workforce reduction, about 3,100 roles; Mailchimp absent from agent lists) | https://www.saastr.com/how-mailchimp-went-from-1b-arr-and-20-growth-to-shrinking-inside-intuit/ | [third-party verified] [33] IT之家 (ITHome) | X removes API access from InfoFi-style apps (2026-01-15, on the grounds of generating AI spam) | https://www.ithome.com/0/914/499.htm | [third-party verified] [34] Sina Finance (Chinese: 新浪财经) | The Kimi wrongful-ban incident: a risk-control misjudgement hits the third-party developer ecosystem (all restricted accounts restored; User Agent detection refined) |

2026-05-27 | https://finance.sina.com.cn/wm/2026-05-27/doc-inhzivfi1630996.shtml | [third-party verified] [35] TechNews (Chinese: TechNews 科技新报) | Why platforms pulling API access is a warning for the ecosystem (the API turns from ecosystem building into revenue growth and data defence; developers carry platform risk) | https://cdn.technews.tw/ai-agent/anthropic-kill-lots-of-companies/322/ | [third-party verified] [36] Yage Daily Notes (Chinese: 鸭哥每日手记) | Why agent stores did not work (Doubao's roughly 8 million user-built agents earning under RMB 1 million a day; of about 3 million GPTs only about 159,000 public and active) | 2026-07-04 | https://yage.ai/share/uga-sunset-20260704.html | [third-party verified] [37] Profound | Profound Index Report: Summer 2026 (built on more than 1.9 billion real user conversations; in 24 of 30 multi-region industries a local No. 1 that never cracked the U.S. top five; only seven brands first in every available market; 46%–59% of citations go to country-specific domains where a local brand leads, 12–17 points above the control; after the June 2026 ChatGPT update citations per answer −10.4% and unique cited domains down in 56 of 58 industries, leading brands mostly holding position) | 2026 | https://www.tryprofound.com/reports-guides/profound-index-report-summer-2026 | [vendor claim]

[38] Cloudflare Blog | Google's AI advantage: why crawler separation is the only path to a fair Internet (position filed with the CMA on its proposed conduct requirements for Google; argues for mandatory crawler separation; publicly supported by Daily Mail Group, The Guardian and the News Media Association) | 2026-01-30 | https://blog.cloudflare.com/uk-google-ai-crawler-policy/ | [official: corporate position]

[39] Beri.net (Rajesh Beri) | Reddit's API Shutdown Spares Sprinklr, Not Your Python Scripts (announced on r/redditdev on 2026-09-30: no new public API applications accepted after 2026-10-31; all subreddit RSS feeds stop on 2026-11-13; existing apps must complete registration by 2027-01-12; the free public API closes in 2027-03; listening suites with signed commercial agreements keep their data feeds) | 2026-10-05 | https://www.beri.net/article/reddit-public-api-shutdown-march-2027-social-listening-ai-data-license-deadlines | [third-party verified] [40] Digital Content Next | The Evidence Is Clear: AI Search Is Costing Publishers (outbound organic clicks -39.8% and zero-click searches +34.5% when an AI Overview appears; Pew: 8% clicked a traditional result with a summary versus 15% without, links inside the summary 1%; Chartbeat across 2,576 publishers: Google referrals -33% globally and -38% in the U.S., with Google Search still at 7.3% of referrals and ChatGPT at 0.02%;

generative AI referrals +86% year on year but about 0.7% of all web traffic) | 2026-10-05 | https://digitalcontentnext.org/blog/2026/10/05/the-evidence-is-clear-ai-search-is-costing-publishers/ | [third-party verified]

# Chapter 22 Build the Foundation: First-Party Data, Compliance, and Omnichannel Orchestration

Every move in the previous two parts happens where the customer can see it — in the AI's answers, in the platform's feed, in the message you sent; the subject of this chapter never shows up on the customer's side at all.

This step does not happen on the customer's side, but it decides whether everything you did before can be logged, reviewed, and continued. So this chapter gives you 5 methods to lay the foundation first.

---

## Method 1: Get the customer list back into your own hands

**What to do**: Collect the customer contact details scattered across platform back-ends, sales reps' phones, Excel, and chat logs into one list you can export, contact, and delete person by person.

**How**

1. Settle on four fields first, and do not add a fifth: one reachable identifier (phone number or email), source (which entry point it came from), consent status (when and how it was obtained), and the time of the last touch.
2. Round up all the entry points: website forms, official accounts and WeCom, stores and trade shows, customer-service conversations, trial sign-ups, and closed customers. Assign one person per entry point to enter records as they come in; don't count on catching up at month-end.
3. Consent status must be tied to the individual and must be deletable person by person. Quarantine the segment whose source you cannot explain; it takes part in no marketing action.
4. Clean the list once a quarter: move those who withdrew consent or clearly said they do not accept marketing from the reachable zone to a quarantine zone, keeping one record for the audit trail.
5. Keep it in a system you can export from. A list stuck in a platform account back-end that you cannot export stops being yours the moment the platform changes a rule.

**Example**: The vendor's definition of a CDP (customer data platform) is this: integrate customer data from multiple sources into a unified customer profile, with specific capabilities that include identity resolution — matching and merging records of the same customer across devices and channels — and data activation — sending customer data to systems such as email platforms and ad networks [vendor claim]. Treat this passage only as a definitional reference for “how far you need to consolidate”. It promises no results, and it is not product-selection advice. The red line on lists has criminal samples: among the Supreme People's Court's typical cases on infringing citizens' personal information published in May 2026, an outsourcing firm developing and maintaining a hospital's online appointment system illegally obtained appointment users' personal information from the back end and imported it into its own database, reaching 2,878,070 records after de-duplication; the unit was fined RMB 300,000 and legal representative He and others received prison terms of five years six months to one year six months, while in another case a railway station conductor used his position to illegally query and sell passengers' travel information [official] — crossing the line on trading data for acquisition can bring a criminal judgment.

**Pitfalls**: Mixing bought or scraped lists into your own. The batch whose source you cannot explain can take down the credibility of the whole list with a single complaint or a single leak — and it never had the consent you claim for it in the first place.

## Method 2: Build “notice, consent, refusal” into the moment of collection

**What to do**: Let customers see how you will use their information before they leave contact details, and let them ask you to stop marketing to them after they do.

**How**

1. Give notice at the moment of collection, covering four things: the name and contact details of the processor; the purpose and method of processing, the types of information, and the retention period; the ways and procedures by which the individual can exercise their rights; and any other matters the law requires you to disclose. State them in a prominent way, in clear and easy-to-understand language, truthfully, accurately, and completely [official]. If these matters change later, tell the individual about the parts that changed.

2. Consent must be voluntary and explicit, and you may not trade it for access — except where the personal information is necessary to provide the product or service [official]. In practice it comes down to one line: make “I agree to receive marketing messages” a separate checkbox, and let the product work normally when it is left unchecked.
3. Provide two ways out: one for withdrawing consent, one for refusing personalized pushes and marketing. The provision on automated decision-making reads, verbatim, that you should also provide an option not targeting the individual's personal characteristics, or provide the individual with a convenient way to refuse [official].
4. Before giving personal information to another processor (a push platform, an ad platform, a data service provider), first disclose the recipient's name, contact details, processing purpose, processing method, and the types of personal information, and obtain separate consent [official].
5. Correction, deletion, and account closure must each have an entry point a user can actually click. These three were named one by one in the four-ministry special campaign of 2025: failing to provide effective functions for correcting and deleting personal information and closing accounts, and providing personalized information pushes without offering a convenient way to refuse [official].

**Example**: The U.S. FTC's “click-to-cancel” rule turned that sentence into a hard requirement — unsubscribing must be as easy as subscribing; the rulemaking drew more than 16,000 public comments, while consumer complaints rose from an average of 42 a day in 2021 to 70 in 2024 [official]. Consent and refusal must cost the same. Enforceable boundaries on the Chinese side have also been named: in enforcement cases announced by the Cyberspace Administration of China in September 2026, a Chongqing real-estate company installed more than 10 image-capture devices in its offices, collected and stored over 5,000 records of customers' facial information, and used them for customer identification and matching marketing with intermediary agencies; because it failed to truthfully, accurately and completely inform customers of the purpose, method and use of the collection and obtained no separate consent, it was ordered to correct the practice and given a warning and a fine [official]. Wiring biometric data into marketing matching is a high-risk zone, and separate consent plus data minimization are the items an inspection will check. Advertising-chain scrutiny is systematising: a law firm's reading of the 2026 special campaign on personal

information protection lists five priorities for internet advertising (excessive collection, inadequate notice and opaque third-party sharing, poor channels for exercising user rights, non-compliant opt-outs from personalised recommendations and so on), and requires companies to build security assessment and review for SDK onboarding, obtain data-security compliance commitments from SDK providers, and audit and supervise partners on an ongoing basis [third-party verified: law firm analysis]. These requirements are not on paper: on June 11, 2026 the Cyberspace Administration of China and two other departments notified 30 apps of problems in personal-information collection and use — 7 failed to publish their collection and use rules, 4 frequently demanded non-essential permissions, 5 failed to list SDK collection and use completely and accurately, and 14 provided no effective account-deletion function — and required the operators to complete rectification within 15 working days [official]. The 2026 enforcement priorities have been published: a joint announcement on 2026-04-02 by the Cyberspace Administration of China, the Ministry of Industry and Information Technology and the Ministry of Public Security lists illegal collection and use of personal information in the internet advertising sector as one of seven special campaigns, covering ad intermediaries and media platforms, with focus areas including collection beyond necessity, failure to state in privacy rules that data is used for advertising and profiling, failure to list the categories, purposes and recipients of data shared with third parties, no convenient route to correct or delete data or to switch off personalised recommendations, and continued collection after a user switches recommendations off; a set of policy Q&As published by the same administration in April 2026 also settles two counting questions — thresholds such as "processing the personal information of more than 10 million people" include the figure itself and count the natural persons currently being processed, excluding deleted records, and processing more than 10 million people requires a compliance audit at least every two years, more than 1 million but no more than 10 million every three to four years, and no more than 1 million every five years [official].

**Pitfalls**: Turning “consent” into a checkbox nobody has read, then interpreting it internally as a marketing authorization covering every channel. Consent is per-purpose — marketing, handing information to others, and using it abroad each stand on their own, and each must be withdrawable.

## Method 3: Do the threshold math before data leaves the country

**What to do**: Before customer data flows abroad (an overseas marketing cloud, an overseas ad platform, an overseas customer-service system, reports read by an overseas subsidiary), work out whether you need to file for a security assessment.

**How**

1. First list every segment where data leaves: overseas CRM and marketing clouds, overseas email and ad platforms, overseas customer-service ticketing systems, and reports read by overseas teams or agents.
2. Filing is required in four situations — providing important data abroad; operators of critical information infrastructure and processors of personal information on more than 1 million people providing personal information abroad; providing personal information on 100,000 people or sensitive personal information on 10,000 people abroad on a cumulative basis since January 1 of the previous year; and other situations requiring filing as prescribed by the national cyberspace administration [official].
3. Carry out a self-assessment of risk before filing: assess the legality, legitimacy, and necessity of the purpose, scope, and method of the outbound transfer; the scale, scope, types, and sensitivity of the data transferred; the protection capacity of the overseas recipient; and whether the contract you plan to sign spells out the security responsibilities in full [official].
4. Build the timing into your schedule: the assessment is completed within 45 working days of acceptance, extendable when the case is complex or supplementary materials are needed; a passed assessment is valid for 2 years, and if you still need to transfer data abroad when it expires, file again 60 working days before expiry [official].
5. The legal documents must spell out at least six things: the purpose, method, and data scope of the outbound transfer, and the use and method of the overseas recipient; the overseas storage location, the retention period, and the measures for handling the data after the period ends; the binding requirements for further transfers to other organizations or individuals; the security measures to be taken when actual control or the local legal environment changes;

remedies for breach, liability for breach, and the dispute-resolution method; and the requirements for emergency handling and the channels for protecting rights when data is tampered with, leaked, or illegally obtained [official].

**Example**: The most expensive lesson on the cross-border line comes with a price tag: the Irish Data Protection Commission fined Meta €1.2 billion for continuing to transfer EU user data back to the United States, finding it in breach of Article 46(1) of the GDPR [official]. China's thresholds have been quantified into two lines: official guidance states that a non-critical information infrastructure operator whose cumulative outbound transfers since 1 January of the year reach 100,000 or more but fewer than 1 million individuals' personal information (excluding sensitive personal information), or fewer than 10,000 individuals' sensitive personal information without important data, may meet its obligations through personal information export certification; splitting volumes to move what should be a security assessment onto the certification route is prohibited; and once cumulative transfers reach 1 million individuals' personal information or 10,000 individuals' sensitive personal information, a data export security assessment must be filed [official]. For cross-border acquisition the implication is direct: transferring CRM, marketing automation and overseas ad platform user data is quantified into scale thresholds, and "how many leads" decides which compliance route you take.

**Pitfalls**: Assuming that “we use overseas SaaS, so it counts as a transfer abroad”. Look at where the data actually lands, not where the vendor is registered or which party signed the contract [our judgment]. And treating the types and scale of personal information as a vague impression ends with the filing threshold miscalculated.

## Method 4: Label AI-generated content before you publish it

**What to do**: For text, images, audio, video, and virtual scenes generated or synthesized with AI, declare identity as required at every link of publishing and distribution, and make sure the label is still there after the file is exported.

**How**

1. First tell the two kinds of label apart: an explicit label is added to generated or synthesized content or to the interaction interface and is clearly perceivable by users, in the form of text, sound, or graphics; an implicit label is written into the file data by technical means and is not easily perceived by users [official].

2. Apply it per medium: for images, add a prominent notice in an appropriate position; for video, add it in the opening frames and around the playback area, and it may also be added at the end and in the middle; for virtual scenes, add it in the opening frames; for text, add a text or symbol notice in an appropriate position at the start, end, or middle, or add a prominent notice in the interface or around the text [official].
3. Write the implicit label into the file metadata, covering at least the attribute information of the generated or synthesized content, the service provider's name or code, the content number, and other production element information [official].
4. Where download, copy, or export functions exist, make sure the file contains an explicit label that meets the requirements — a label that only sits in the back-end and drops off on export does not count as done [official].
5. If you distribute through an app store, the labeling materials will be checked at listing or launch review; when publishing generated or synthesized content, declare it proactively and use the labeling function the platform provides; no organization or individual may maliciously delete, tamper with, forge, or conceal these labels [official].
6. Remember the date: the measures took effect on September 1, 2025 [official].

**Example**: Article 9 of the Measures for Labeling AI-Generated and Synthesized Content gives teams making ad creative the most practical reading: when a user requests generated or synthesized content without an explicit label, the service provider may supply it after making the user's labeling obligations and usage responsibilities clear through the user agreement, and must keep logs of the recipient and related information for no less than six months in accordance with the law [official]. The moment you get “clean” material, the labeling obligation and usage responsibility shift to your side — and the logs still have to be kept. Labeling obligations are becoming an exercise across jurisdictions: a law firm's comparative analysis notes that China's Measures regulate the whole lifecycle of AI-generated synthetic content — production, distribution, dissemination and use — with duties covering generative service providers, app distribution platforms, application service providers, online content dissemination providers and users; the EU relies mainly on Article 50 of the AI Act, whose labeling requirements take effect on August 2, 2026; the U.S. has no dedicated federal legislation, but California's AI Transparency Act takes effect on January 1, 2026 and explicitly separates "explicit disclosure" from "implicit

disclosure," the latter required to convey the provider's name, system name and version, the time of creation and a unique identifier [third-party verified: law-firm analysis]. Enforcement has landed as well: on April 28, 2026, cyberspace authorities investigated Jianying (Chinese: 剪映) and other platforms for violating the rules on labeling generated synthetic content [official]. Short-video platforms fell in line just as promptly: on June 14, 2026 major short-video platforms fully implemented the special management requirements for labeling AI-generated content, launching a platform-wide mandatory tag and requiring the relevant videos to carry a conspicuous AI-generated label within the first 3 seconds, with the label no smaller than 5% of the frame width; content without compliant labeling is not distributed, and deliberate concealment can bring takedowns, throttling or bans [third-party verified: state-media commentary].

**Pitfalls**: Assuming only face-swaps and digital humans need labeling. The scope is “text, images, audio, video, and virtual scenes generated or synthesized using AI technology” [official] — mass-generated product images, AI-voiced talking-head videos, and AI-written scripts and copy all fall inside it.

## Method 5: Link your touchpoints into a chain that reports back

**What to do**: Connect “which channel a list segment is pushed to, what that channel brings, and how the result is read back” into a single line, one line one set of definitions, reviewed channel by channel.

**How**

1. One push line per channel: cut a segment from the list, push it to that channel, and give the line its own name so it does not blend into organic traffic.
2. Fix the report-back fields, these six and no more: spend, impressions, cost per thousand impressions, clicks, click-through rate, and cost per click [official]. Read only these for every line; do not swap definitions on the fly.
3. When the same group is pushed to two channels, look at the two sides separately. Different channels have different cost structures, and blending them into one average will wrongly convict one and wrongly acquit the other at the same time.

4. The read-back must land on “the group you pushed out”. The vendor's own documentation locks the definition down: it analyzes only the advertising and marketing performance of ads served to the audience pushed from its own platform to that ad platform [official]. Your table must lock down the same limitation.
5. Reconcile the six fields with closed deals in the CRM once a quarter. If the channel report claims conversions and you cannot find the corresponding people in the CRM, fix the definitions first, then talk about adding budget.
6. Omnichannel orchestration has no authoritative official definition; just follow the practice [our judgment]. The deliverable you hand in is one table: channel × pushed list × six fields × quarterly trend.

**Example**: Alibaba Cloud's (Chinese: 阿里云) help documentation gives a very concrete path: an audience-push task pushes an audience to Ocean Engine (Chinese: 巨量引擎), then on the performance analysis page you select the time range, the Ocean Engine account, the ad audience, and the ad order to look back at spend, impressions, cost per thousand impressions, clicks, click-through rate, and cost per click, and the bottom of the page lists key metrics order by order [official]. The same document limits it: it analyzes only ads served to the audience pushed from it to that platform [official].

**Pitfalls**: Pushing without reading back — or reading back only against your own last week. Without a control, you cannot tell whether the list did the work or the channel would have produced volume anyway.

---

When the foundation is weak, every method above gets discounted. If the list cannot be gathered, the results of all your visibility, content, and outreach work never settle; if consent status is unclear, leads coming back from channels to the CRM dare not be touched again; if the outbound threshold is uncalculated, the overseas set of moves can be halted at any time; if labels are missing, shipping creative is accumulating risk; if read-back definitions are not unified, not one set of books will hold. None of these five things happen on the customer's side; they do one thing only — let what you did before be logged, reviewed, and continued.

## Notes for this chapter

- This chapter describes the “notice–consent” requirement itself and gives no article numbers. Using customer data “for marketing” is not something a single clause of the Personal Information Protection Law covers; it is assembled from the requirements on notice, consent, separate consent, and convenient refusal, and pinning one article number on it would compress several provisions into one. If you need an article number, go back to the original text, check it clause by clause, and only then cite it.
- For automated decision-making, this chapter quotes the original wording — where it is used for information pushes and commercial marketing, an option not targeting the individual's personal characteristics should also be provided, or the individual should be given a convenient way to refuse; the original text also states that unreasonable differential treatment of individuals in transaction prices and other transaction terms is prohibited [official]. This chapter does not interpret the enforcement intensity of this provision, nor does it expand on the definitions for algorithm filing and algorithm recommendation management.
- “Omnichannel orchestration” has no authoritative official definition. It is neither a legal concept nor was a unified definition ever given by a regulator or standards body. This chapter writes it as a practice — cut a list segment and push it out, report back six fields, review channel by channel, reconcile with the CRM once a quarter — naming no products and not presenting it as an industry standard [our judgment].
- The first-party data section cites Microsoft's definition of a CDP; that is a vendor's reading, to be used only as a definitional reference for “how far you need to consolidate”, not as a product capability or a promise of results [vendor claim]. The Alibaba Cloud document cited in Method 5 is an operational definition that explains how the fields are read; it is likewise not a promise of results [official].
- The numbers in the outbound threshold (1 million people, 100,000 people, 10,000 people, 45 working days, 2 years, 60 working days) all come from the original text of the Measures for the Security Assessment of Data Cross-Border Transfer, which took effect on September 1, 2022 [official]. Before actually filing, go by the currently effective text and the reading of the cyberspace administration authority where you are located.

## Sources for this chapter

[1] Personal Information Protection Law of the People's Republic of China (adopted August 20, 2021; effective November 1, 2021) | full text on the Cyberspace Administration of China website | https://www.cac.gov.cn/2021-08/20/c_1631050028355286.htm | [official] [2] Cyberspace Administration of China Order No. 11, Measures for the Security Assessment of Data Cross-Border Transfer (effective September 1, 2022) | full text on the Cyberspace Administration of China website | https://www.cac.gov.cn/2022-07/07/c_1658811536396503.htm | [official] [3] Measures for Labeling AI-Generated and Synthesized Content (issued March 7, 2025; effective September 1, 2025) | full text on the Cyberspace Administration of China website | https://www.cac.gov.cn/2025-03/14/c_1743654684782215.htm | [official] [4] Office of the Central Cyberspace Affairs Commission, Ministry of Industry and Information Technology, Ministry of Public Security, State Administration for Market Regulation | Announcement on Carrying Out the 2025 Special Campaign on Personal Information Protection (March 28, 2025) | full text on the Cyberspace Administration of China website | https://www.cac.gov.cn/2025-03/28/c_1744867353112759.htm | [official] [5] Alibaba Cloud | Ocean Engine performance analysis (the fields and limiting definition for reading back results after a Quick Audience audience push) | https://help.aliyun.com/zh/document_detail/170392.html | [official] [6] Microsoft | What is a CDP? (the customer data platform definition, identity resolution, and data activation capabilities) | https://www.microsoft.com/zh-cn/dynamics-365/resources/what-is-a-cdp | [vendor claim] [7] U.S. FTC | Final Click-to-Cancel Rule (unsubscribing must be as easy as subscribing; 16,000+ public comments) | 2024-10 | https://www.ftc.gov/news-events/news/press-releases/2024/10/federal-trade-commission-announces-final-click-cancel-rule-making-it-easier-consumers-end-recurring | [official] [8] Irish Data Protection Commission (DPC) | decision against Meta Ireland: €1.2 billion (breach of Article 46(1) of the GDPR) | http://www.dataprotection.ie/en/news-media/press-releases/Data-Protection-Commission-announces-conclusion-of-inquiry-into-Meta-Ireland | [official]

[9] Cyberspace Administration of China | Recent enforcement cases on cybersecurity, data security and personal information protection (a Chongqing real-estate company using over 5,000 facial records for matching marketing without separate consent; warning and fine) | 2026-09-15 | https://www.cac.gov.cn/2026-09/15/c_1790876152357946.htm | [official] [10] Haiwen & Partners | A three-jurisdiction comparison of AIGC labeling (EU AI Act Article 50 effective 2026-08-02;

California's CAITA effective 2026-01-01; China's labeling measures covering the whole chain of duty bearers) | http://www.haiwen-law.com/35/1700 | [third-party verified: law-firm analysis] [11] People's Daily Online (relaying Xinhua) | Cyberspace authorities investigate Jianying and other platforms over violations of the labeling rules for generated synthetic content (2026-04-28) | http://society.people.com.cn/n1/2026/0428/c1008-40710527.html | [official] [12] AllBright Law Offices (Chinese: 上海市锦天城律师事务所) | 2026 special campaign on personal information protection: five priorities for internet advertising (security assessment and review for SDK onboarding; ongoing partner audits and supervision) | https://www.allbrightlaw.com/CN/10475/121a2e00ab688315.aspx | [third-party verified: law firm analysis] [13] Ministry of Commerce / Cyberspace Administration of China | Q&A on data export security management policy and regulations (2026-09) (the 100,000 / 1 million lines; no splitting volumes; security assessment required at 1 million) | https://fec.mofcom.gov.cn/article/ggfw/crjfw/crjfwxx/202609/8055.html | [official] [14] JunHe Law Offices (Chinese: 君合律师事务所) | Ten questions and trends worth watching in China's data field in 2026 | https://www.junhe.com/legal-updates/2650 | [third-party verified: law firm analysis] [15] Supreme People's Court | Typical cases of crimes infringing citizens' personal information and related crimes (an outsourcing firm maintaining a hospital appointment system illegally obtained appointment users' information, 2,878,070 records after de-duplication; RMB 300,000 fine; He and others sentenced to five years six months to one year six months) | 2026-05-08 | https://www.court.gov.cn/zixun/xiangqing/499271.html | [official] [16] Xinhua (Cyberspace Administration of China General Office) | CAC notifies 30 apps over personal-information collection and use problems (7 failed to publish collection-use rules, 4 frequently demanded non-essential permissions, 5 failed to list SDKs, 14 had no effective account-deletion function; rectify within 15 working days) | 2026-06-11 | https://app.xinhuanet.com/news/article.html?articleId=20260611ac2e1e2330bc4bd2983994c3c18543f7 | [official] [17] China.com.cn | Mandatory AI-content labeling draws a red line for innovation (major short-video platforms fully implement the labeling requirement from 2026-06-14; mandatory tags; a conspicuous AI-generated label within the first 3 seconds; label no smaller than 5% of frame width; non-compliant content not distributed, deliberate concealment brings takedowns, throttling or bans) | 2026-06-23 | https://www.china.com.cn/opinion2020/2026-06/23/content_118559677.shtml | [third-party verified: state-media commentary]

[18] Cyberspace Administration of China, MIIT and MPS | Announcement on the 2026 special campaigns for personal information protection (internet advertising named as a priority sector, with five focus areas) | 2026-04-02 | https://www.cac.gov.cn/2026-04/02/c_1776867645836849.htm | [official] [19] Cyberspace Administration of China | Personal information protection policy Q&A (April 2026) (the "more than 10 million people" threshold includes the figure itself and counts current records excluding deleted ones; compliance audit frequency 2 years / 3–4 years / 5 years) | 2026-04-29 | https://www.cac.gov.cn/2026-04/29/c_1779200509387274.htm | [official]

# Chapter 23 Who Does It and What It Costs: Three Minimum Configurations

When AI projects stalled in the past, people tended to blame the model: not a new enough version, not big enough parameters, prompts not written well. Now the tools are on the table and the capability is sufficient; what blocks AI customer acquisition is almost always the organization — who does it, whose money pays for it, what goes to machines, and what must stay with people. What you have to do is: sort the work into buckets first, then staff it and track every dollar by company size.

Every method above ultimately comes down to “who is doing it”. So this chapter gives you 5 methods to allocate people and money by company size.

---

## Method 1: Decide up front which work machines may not touch

**What to do**: Use three criteria to sort the process into buckets, then write the “things machines may not touch” into a one-page boundary.

**How**

1. Hand work to a machine only if all three criteria are met: repetitive (how many times a week or a month it happens, and whether the pattern is stable); information-dense (how much context must be read and how many systems must be checked to do it correctly); and reversible (can an error be caught, can it be undone, is the cost controllable). The third is the core one, and it is a business property, not a technical one: an outbound email cannot be recalled once sent, and a quote in the buyer's hands constitutes a commitment.
2. By the criteria, drop every step of the process into one of three buckets: steps expressible as if-then rules get ordinary code — if plain code can solve it, don't use a model; steps that have thousands of human judgment records with outcomes and low-risk output forms (yes/no, routing, labeling) go to a large model with full measurement; steps that are high-risk or short on samples go to

a human in the loop, with an agent assembling the evidence and a person making the decision in 30 seconds. After bucketing, people move from execution down to review — no fewer people, a different position.

3. Put the escalation threshold on “emotion and failure-risk signals”, not on “problem type”. Routing by type (pre-sales to AI, after-sales to humans) will miss the most expensive class of customers: the one who is already unhappy but has not yet decided whether to leave.
4. Write a one-page usage boundary, with six rules locked down: externally facing facts may come only from the fact list, and anything with no entry cannot be said by a machine; questions of price, qualifications, and commitments must be signed off by a person; emotion and failure-risk signals trigger escalation to a human; every automated action is reversible, leaves a trace, and has a circuit breaker, with the stop conditions written at procurement time; any statement involving “research shows” or “data shows” must be able to name the original publisher, and anything you cannot verify gets deleted; if employees are using unapproved tools, the approved tools have a gap, and the right move is to speed up delivery of a governable platform.
5. Delete “equivalent to N people” from your reports and replace it with three layered metrics: automatic resolution rate for routine questions, escalation rate for emotion and failure-risk signals, and fact-statement rollback rate (the number of times a machine got it wrong).

**Example**: Klarna's official press release said its AI assistant handled 2.3 million conversations in its first month, work equivalent to 700 full-time agents [vendor claim]; in 2025 the CEO acknowledged that an organization built around cost first “ends up delivering lower quality” and announced it was hiring human agents again [third-party reported]; by 2026 the workload carried by AI had grown to roughly the equivalent of 850 full-time agents, alongside a rebuilt in-house team of about 100 highly skilled staff [third-party reported]. Stack the three together and you get a certainty-and-agent bucket expanding while the human-in-the-loop bucket shrinks but refines. The enterprise adoption curve has an official figure: Salesforce's Agentic Enterprise Index says the number of agents enterprises have activated grew nearly 3× [vendor claim].

**Pitfalls**: Treating “AI and frustrated customers don't get along” as the model not being strong enough. A field experiment by Dartmouth's Tuck School of Business based on Alibaba's customer service found that AI shortened routine inquiries but struggled in sessions with high emotion and high service-failure risk, even with human involvement [third-party research reported] — this is a threshold-design problem. Nor should you put “equivalent to N people” into your own reports: it blends volume handled, complexity distribution, the quality definition, and cost shifting into a single number.

## Method 2: Turn the acquisition budget into asset purchases

**What to do**: Turn the budget table from three lines (media, agency fees, labor) into four asset line items, then add the two lines of cost most often left out.

**How**

1. Give each of the four asset types its own line item: the fact foundation (every fact has a version, an owner, and a change flow; pages are structured, documents are machine-readable, and the customer-service and sales AI knowledge bases hang off the same fact source); first-hand evidence (data you measured yourself, retellable customer cases, verifiable controlled experiments — one investment, many rounds of reuse); third-party sources (what you buy is “the name of the source when it is cited”); and reversible automation and measurement (the steps that land in the agent bucket, plus the cost of building the three layers of evidence).
2. Add the two overlooked cost lines: compute and tokens — this is sliding from the IT line to the acquisition line, and mass generation plus mass outreach make token cost rise along with lead volume; and the hidden labor of human-machine collaboration — teams that adopt digital employees still spend time feeding context, checking output, debugging, rerunning prompts, and cleaning up “confident but wrong” answers, about 6.4 hours a week, and those 6.4 hours are not an efficiency dividend but a new operating cost [third-party reported] (the sample size of that survey could not be traced to the original report page this round; this chapter quotes only the duration — see the notes for this chapter).
3. Do not set up an “AI acquisition project”. The Salesforce CEO has made a public judgment: if there is a standalone “AI budget”, you are doing AI the wrong way [third-party reported]. The typical ending of a project budget is that the year

ends, the project closes, and the assets go to zero; to judge whether a budget has migrated, ask one question only — what share of the money will leave something behind when this quarter ends [our judgment].

4. Testing ammunition must be a fixed line in the budget table, not what is left at year-end.
5. Tie tool budget to process change: for every tool expense, write down “which step of the process changes, who changes it, and what number you watch after the change”; if you cannot write it down, don't buy yet. Prioritize the workflow with the most handoffs, not the one with the largest volume — time hides in queues and handoffs.
6. Treat shadow AI as a diagnostic signal, not a violation: 70%–80% of employees use unapproved tools, and inside one chip company 1,500–1,600 tools were in use at the same time [third-party reported]. When employees use unapproved tools, the approved tools have a product gap.

**Example**: A 2026 marketing budget allocation guide offers this structural reference: digital channels 40%–50%, content and brand 20%–30%, events and sponsorship 10%–15%, research and contingency 10%–20%, with a fixed 8%–12% set aside for testing new channels, new audiences, and new tactics [vendor claim]. The counterexample comes from a Chinese-side first-person account: the company gave every department AI tools and ran dozens of training sessions, and the year-end review showed actual AI usage under 15% [third-party reported]; overseas, the same phenomenon has the structure of 5%–10% heavy users, 20% occasional users who use it crudely, and 70% who barely use it at all [third-party reported]. A harder set of numbers comes from Gartner's 2026 CMO Spend Survey (401 CMOs and marketing leaders surveyed between January and March 2026): marketing budgets edged up to 7.8% of company revenue, CMOs allocate an average of 15.3% of the marketing budget to AI initiatives, and while 70% say becoming an AI leader is a critical goal for 2026, only 30% report mature or fully developed AI readiness capabilities [official: institutional survey]. The money has moved and the capability has not caught up — that gap is what the "buy assets or buy tools" question is scored on. The external environment is tightening at the same time: the 35th edition of The CMO Survey, run by Duke Fuqua, gives three findings that constrain each other — marketers' optimism about the economy is the lowest since 2020; under pressure to show returns, activity contracts, spending falls and existing customers take priority over expansion; and firms expect AI to carry more than half of marketing activity within three years even

though technology is outpacing organisational readiness, with the survey stating that no marketing technology activity is currently delivering its full potential [third-party verified]. Gartner's 2026 CMO survey calls this position a "trilemma", with flat budgets and more tasks [third-party verified]. The martech line item is shrinking too: in the same Gartner 2026 CMO spend survey (401 CMOs), martech fell to 19.4% of marketing budgets, a five-year low against 26.6% in 2021, while marketers use only 49% of their stack's capabilities [third-party verified: citing Gartner]. The Chinese total is rising: advertiser research puts 2026 marketing investment up about 10% on average, with 44% of advertisers who already fund content marketing saying they will spend more on it [third-party verified] — the money has not disappeared; it is moving from buying tools to buying content and data that leave something behind.

Where the money goes is being repriced as well: the Aleph × Benchmarkit 2026 B2B SaaS and AI-native benchmarks (published June 2026, covering full-year 2025, with 198 of 342 companies disclosing) put median CAC payback at 16 months, an 11% improvement on CY2024's 18 months and the largest single-year gain in four years of the dataset; two easily missed methodological points are that the denominator must use gross-margin-adjusted revenue (at 80% margin, “12-month payback” is really 15) and that it must be computed per channel and per cohort — a blended CAC hides a channel at 28 months coexisting with one at 4 [third-party verified]. AI's effect on budget is not one-directional compression either: a Semrush survey of 1,008 U.S. marketing budget decision-makers found 42% increased budgets because of AI against only about 16% cutting, and roughly a third reporting a significant rise in headcount as well [third-party verified]. Read together, “AI saves headcount” is a weak case to bring into a second-year budget review; cost per unit of output and the number of executable experiments added are sturdier.

**Pitfalls**: Evaluating tools by satisfaction and the number of users. The UK Department for Business and Trade's pilot of M365 Copilot (1,000 seats, 3 months) showed only 1.14 actions per person per day, PowerPoint 7 minutes faster but half the quality, and Excel slower and worse, with the self-assessment concluding that it “found no robust evidence that time savings translated into productivity gains” — yet 72% of users reported satisfaction [third-party reported]. Satisfaction decoupled from output is exactly why shelfware can stay hidden for three years.

## Method 3: Staff and fund by company size — three tiers listing only what fails without it

**What to do**: Find your own seat and work out the minimum you cannot do without at your tier.

**How**

1. **1–2-person companies** (micro team, single category): the only thing to do at this tier is start accumulating assets; do not touch automation. Fix a question set of 20–30 real buyer questions, run an entry-point audit monthly (fixed question set → fixed entry points and frequency → record only four fields each time → watch the trend, not a single run); build a fact list v0, starting with 20 entries, each with one owner, one version number, and one effective date, and the owner is the boss; no agents, no outbound automation, no mass content. Spend the money on running the question set and maintaining the fact list, not on tools. The test: three months later, can you state “what our stability rate is, and which facts were stated wrongly”.
2. **5–15 people** (small to midsize team, single market): this tier has to put human-machine division of labor in place and start having dedicated roles. One full-stack growth engineer (can be grown internally), fully accountable for “whether the buyer is described correctly”; one to two content and source operators, whose assessment switches to evidence KPIs (number of first-hand materials filed, number of cases with all four required elements, consistency rate of third-party statements); one AI visibility owner (concurrent role), tied to the fact-list owner; use the three criteria from Method 1 to bucket the first three steps of the acquisition process, landing no more than 2–3 agents plus 2 human checkpoints. Spend the money on people and process change, with tool spending tied to the process. The test: at least one process achieves “people demoted from execution to review” without the failure rate rising.
3. **50+ people** (multiple categories, multiple markets): the core at this tier is cross-functional authority. Form a cross-functional squad — full-stack growth engineer, content and sources, data, legal/compliance liaison — reporting to someone who can command marketing, sales, and product at the same time, otherwise the process redesign will not move; one question set per language per market; visibility written into the assessment of one business-line owner, never as “jointly owned across departments”; appoint an STO (single-threaded owner),

one unique owner per important thing. Spend the money on setting up the cross-functional squad's authority and on the bilingual question sets and fact list. The test: if there is a single instance of “nobody can say who owns this”, this tier is not yet standing.

**Example**: Someone has priced a “fully staffed” growth team: a few engineers plus 1 PM and 1 designer, about $1 million a year; the rule of thumb for carrying that cost is an ARR of at least $5 million [third-party verified: first-hand from a practitioner].

The tool stack's ledger can be laid out first: a benchmark built on a nationally representative sample of 2,400 marketers plus a full audit of 240 client accounts puts the average stack at 121 tools (91 in 2022, 24 in 2014), with almost all of the increase coming from AI-native tools (14% of the stack); only 34.1% of tools are actually used in any given month (42% in 2022), which means paying monthly for about 80 tools nobody used; stack spend is 22.7% of the marketing budget (29.5% in 2022), AI tool adoption is 84.1%, and average tool tenure has compressed from 27.1 months to 18.4 [third-party verified]. The dormancy breaks down further: dormant licences 32.4%, duplicate functionality 18.7%, stalled pilots 8.4%, leftovers from departed staff 6.4% [third-party verified]. For staffing and budgeting, the implication is that the bottleneck is rarely a missing tool but a gap at the identity, data and context layer — and the thicker the stack, the better it covers the problem underneath.

**Pitfalls**: Treating the three tiers as a ladder and assuming that with more people you should naturally move up. The three tiers list only the moves that fail without; everything beyond them can be postponed. The way the vast majority of companies stumble at this step is remarkably consistent: the tools are in place, the process is unchanged; the process changes, the assessment is unchanged; the assessment changes, the people are not.

## Method 4: Assign three things to one named person each

**What to do**: Assign one owner each to three new roles and to the question of “how AI retells us”.

**How**

1. The full-stack growth engineer. The old division of labor was an assembly line — the content team makes content, the SEO team does rankings, the paid team buys traffic, SDRs make calls — and each role owns one segment, with nobody

accountable for “whether the buyer was described wrongly”. This role has to touch pages, data, structured markup, and experiments at the same time; its value is making one person fully accountable for “whether the buyer is described correctly”. The pricing signal from the market side is already out: FDE (forward-deployed engineer) hiring grew more than 1,000% year over year, with total compensation of $300,000–$550,000 for mid-to-senior levels, and the most valued skills have shifted to customer discovery, problem decomposition, and AI product judgment [third-party reported]. The industry's shared conclusion about this kind of role is an anti-hiring line: “You can't hire an FDE, you can only grow one.”

2. The content and source operator. The object of work shifts from “content output volume” to: digging up and filing first-hand materials, whether the four required elements of the case library are complete, whether statements in third-party sources match reality, and whether community answers conflict with the fact list. The tightest thing for this role is consistency of outward statements — when the fact list changes, what must change with it is usually not just the website but also structured markup, third-party platform profiles, sales scripts and quote templates, the customer-service and sales AI knowledge bases, and archived historical content. There is no technical substitute for this; it needs one clear owner.
3. Ownership of AI visibility. Don't argue about “which department”; give it to whoever “owns the outward facts” — the position where the fact-list owner sits. Visibility's downstream is citation and retelling, its upstream is facts and structure, and the middle is third-party sources; the true owner is the common owner of these three things. The owner must be a person, not a department — writing “marketing department” actually means nobody.
4. Every volume target must carry a quality guardrail metric; every important thing gets a unique owner with enough authority and decision rights; each business pillar names 1–2 full-time champions who spend 50%–100% of their time on process change. The champion must be dedicated — a part-time champion is no champion at all.

**Example**: In the Chinese ecosystem the same evolution is stated more bluntly: Anker Innovations (Chinese: 安克创新) proposed that by 2027 it will essentially eliminate segmented roles such as front-end and QA, keeping only three roles — business engineer (full stack), AI engineer, and product manager [third-party

reported]; another organization treats the “end-to-end owner” as the smallest unit of reform [third-party reported]. The AI-native organization playbook from Meta's ads team gives two handles, the STO and the Agent Champion [third-party reported]; a RevOps playbook requires every volume target to carry a quality guardrail metric, on the grounds that benchmarking showed a high rate of reward hacking [vendor claim]. Chinese samples show the sales function changing position: Zhipu (Chinese: 智谱) reported H1 2026 revenue of RMB 954 million, up 399.7% year on year, while sales and marketing spend fell 14.8% to RMB 178 million; MiniMax revenue reached RMB 786 million, up 283.1%, with sales and distribution costs down 17.9% — because the transaction changed shape, with Zhipu's open-platform and API revenue jumping from RMB 29.1 million to RMB 825 million, 86.5% of total revenue, as customers self-serve by topping up tokens. The report's summary is that sales shifted from pre-sale acquisition to in-life and post-sale, and sales OKRs moved from user counts to retention [third-party reported]. Frontline sellers put it more plainly: whether a customer buys an AI office suite does not depend on model capability but only on how much money it saves [third-party reported].

**Pitfalls**: Reading “full stack” as “cheap labor that does everything”. Once KPIs slice it into “how many posts you published” and “how many pages you built”, the role degenerates into the previous era's generalist, and all you save is headcount.

## Method 5: Change the hiring filter and the review yardstick — move one lever per quarter

**What to do**: Settle three things in one pass: whom to hire, what to assess, and at what pace to change.

**How**

1. On the hiring side, change the filter. The old filter was a skills checklist; the new one is learnability and verification ability. Two checkable signals: can the candidate design an evaluation loop (eval, regarded in the industry as the strongest signal of real large-model experience); and does the candidate have “fabrication literacy”, able to state clearly within 30 seconds which model they are using and what it costs [third-party reported].

2. Guard against one side effect: a hollowed-out junior pipeline. Redefine entry-level roles as “verification roles” (fact spot-checks, hallucination checks, spec writing) rather than simply freezing headcount — otherwise in three to five years there will be a gap in senior acquisition people who understand the business.
3. On the assessment side, change the yardstick. Content output volume becomes the number of first-hand materials filed and the cleanliness of the fact list; ranking position becomes stability rate; lead count becomes the count of valid leads after five-tier segmentation and the lead lift verified by incremental experiments; reach becomes reply rate plus quality guardrail metrics (unsubscribe rate, complaint rate).
4. Move only one lever per quarter, in this recommended order: Q1 the fact foundation (fact list v0 plus structured-content work, no process change, no tools bought); Q2 human-machine division of labor (bucket by the three criteria, land 2–3 agents plus 2 human checkpoints); Q3 sources and external evidence; Q4 measurement and assessment. Moving one at a time is non-negotiable.
5. Set the budget for continuous reinforcement, not for a one-time push. The shape of AI tools changes every month, and the finish line keeps moving.

**Example**: Indeed data shows overall hiring is weak (postings are only about 6% above the February 2020 baseline), but postings mentioning AI are more than 130% above baseline and make up 4.2% of U.S. postings, with the AI mention rate in marketing roles jumping from 8.4% to 14.9% [third-party reported]. On the other side is the failure rate of change: a Gartner survey shows that only about 32% of mid- and senior-level business leaders believe their most recent change achieved “healthy change adoption”, and 79% of employees have low trust in change, while “routinizing” change raises the probability of healthy adoption on time by about 3 times [official claim]. The difficulty of proving org-wide return has numbers too: Atlassian's 2026 State of Teams survey shows 94% of marketers use AI at work, yet only 3% of CMOs are certain they can see org-wide AI ROI, and the report attributes the shortfall to a "fragmentation tax" from weak cross-role collaboration [third-party verified]. Bain's survey points the same way (1,125 sales and marketing leaders, 18 industries, 40 countries): more than 90% of commercial organizations are experimenting with AI and most are not yet seeing meaningful productivity gains, while companies in the top revenue-growth quartile capture twice the AI-attributed revenue impact and 30% more productivity by rewiring sales and marketing workflows [third-party verified]. Same budget, different outcome — the difference is not the number of pilots but

whether the workflow was rewritten. Where the money goes also has a sample: an annual report based on a survey of 2,500 marketing professionals says 87% of enterprise marketing teams now use AI tools and the average team allocates 18% of its total budget to AI tools and platforms, up from 11% a year earlier; it lists quality control (62%), team training (54%) and integration complexity (47%) as the top challenges and reports that teams addressing all three systematically see 2.3x higher ROI [third-party verified: self-reported].

**Pitfalls**: Talking about AI transformation without changing assessment is forcing employees to pretend to use AI. And don't turn assessment into another kind of pretending — the output of someone who truly knows how to use AI in 3 hours a day can exceed 10 hours from someone who does not; assess by hours worked and it is the productive one who loses.

---

## Notes for this chapter

- Attributing AI implementation failure to “the model not being strong enough” is rejected by all three research efforts: BCG surveyed 59 countries and 1,000 CxOs and found that only 26% of companies got value from AI beyond proof of concept, with about 70% of obstacles coming from people and process and only 10% related to the algorithm itself [official claim]; a Stanford study covering 51 real deployments concluded that in 77% of cases the hardest challenges were invisible (change management, data quality, and process reengineering) [third-party reported]; and MIT NANDA's report is widely retold as 95% of enterprise AI pilots producing no measurable return [third-party reported]. All three are self-reported or survey findings; agreement in direction is enough.
- Three items to be verified serve only as directional evidence: MIT NANDA's “95%/5%” could not be traced to the original report page or sample definition, and it measures “whether an AI pilot produced a fast P&L impact within 6 months”, which is often misread as a project failure rate; Glean's “6.4 hours a week of botsitting” comes from a research institute the vendor set up; the sample window and de-duplication definition behind the FDE hiring figure “+1,000% year over year” are not public. The Uber CTO's “burned through a whole year's token budget” comes through a third-party retelling and is used only as an order of magnitude.

- The three minimum configurations and the assessment-metric switchover table are [our judgment]; they are design values, not benchmarks. We suggest trial-running them for one quarter at one or two companies, then backfilling the thresholds once you have actual numbers for action volume and stability rate.
- On the Chinese side there are public disclosures of the same kind: a bank claims AI efficiency saved 15.56 million labor hours, equivalent to the workload of 8,000 employees, and a logistics company uses 28 “digital employees” to take on standardized work (both unnamed internal readings) [third-party reported]. Read numbers like these with the same question list: by what definition is the equivalent headcount converted? Is the equivalence in volume handled or in outcomes handled? Where did the released labor go?
- The ordering under “move only one lever per quarter” is our recommendation; it is the discipline of “one at a time” itself that this chapter is most willing to stand behind.

## Sources for this chapter

[1] BCG | “Where's the Value in AI?” | 2024-10-24 | https://www.bcg.com/press/24october2024-ai-adoption-in-2024-74-of-companies-struggle-to-achieve-and-scale-value | [official claim] [2] Stanford Digital Economy Lab | “The Enterprise AI Playbook: Lessons from 51 Successful Deployments” (Pereira / Graylin / Brynjolfsson, 116 pages) | 2026-04 | https://digitaleconomy.stanford.edu/publication/enterprise-ai-playbook/ | [third-party reported] [3] MIT NANDA | “GenAI Divide” (95% of pilots no return / 5% scaled) | [to verify: original report page and sample definition] | [third-party reported] [4] Gartner | press release: only 32% of leaders achieve “healthy change adoption”, 79% of employees have low trust in change | 2025-07-08 | https://www.gartner.com/en/newsroom/press-releases/2025-07-08-gartner-hr-research-finds-just-32-percent-of-business-leaders-report-achieving-healthy-change-adoption-by-employees | [official claim] [5] Digital Applied | “AI adoption change management 2026 playbook” | 2026-06 | https://www.digitalapplied.com/blog/change-management-ai-adoption-2026-overcoming-resistance-playbook | [third-party reported] [6] Klarna | official press release: AI assistant handled 2.3 million conversations in its first month, equivalent to 700 full-time agents | 2024-02-27 | https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats-in-its-first-month/ | [vendor claim] [7] Maginative (relaying a Bloomberg interview) | Klarna admits it “went too far” and is hiring human agents again | 2025 | https://www.maginative.com/article/

klarna-dials-back-its-ai-customer-service-strategy-now-its-hiring-humans-again/ | [third-party reported] [8] Forbes (Bernard Marr) | How Klarna's AI Agent Strategy Backfired But Became a Useful Lesson | 2026-07-16 | https://www.forbes.com/sites/bernardmarr/2026/07/16/how-klarnas-ai-agent-strategy-backfired-but-became-a-useful-lesson/ | [third-party reported] [9] Dartmouth Tuck School of Business | Agentic AI and Human-in-the-Loop Interventions (Lu, field experiment on Alibaba's customer service) | 2026-07 | https://tuck.dartmouth.edu/news/articles/even-with-humans-in-the-loop-agentic-ai-systems-struggle | [third-party research reported] [10] Varick Agents (Vasuman Moza) | “AI Adoption is a Myth” (three-bucket method, UK DBT pilot, barbell distribution) | 2026 | [vendor first-hand methodology: the author sells this service] [11] Matt Beane | AI-Native Organizational Transformation (Uber CTO token budget, network constraint theory) | 2026-09-01 | https://www.wildworldofwork.org/p/ai-native-organizational-transformation | [third-party reported] [12] Glean Work AI Institute | Work AI Index (6,000 digital employees; 6.4 hours a week of botsitting) | 2026-08 | [to verify: original report page and sample definition] | [third-party reported] [13] ByteByteGo (Shah Rahman, Meta) | AI-Native Leaders: The Organizational Playbook (STO / Agent Champion) | 2026-06-22 | https://blog.bytebytego.com/p/ai-native-leaders-the-organizational | [third-party reported] [14] 36Kr (Chinese: 36氪) authorized release | “What Is an AI-Native Organization?” (AI usage under 15%) | 2026-05-21 | https://eu.36kr.com/zh/p/3818273462387592 | [third-party reported] [15] Tencent (Chinese: 腾讯) “AI-Native Ways of Working” survey (Anker Innovations role convergence) | 2026-08 | [to verify: original report source and sample] | [third-party reported] [16] Improvado | 2026 marketing budget allocation guide (digital channels 40–50% / content and brand 20–30% / testing 8–12%) | 2026 | https://improvado.io/blog/marketing-budget-allocation | [vendor claim] [17] SaaStr | Marc Benioff: a standalone “AI budget” is wrong | 2026 | https://www.saastr.com/marc-benioff-ceo-salesforce-on | [third-party reported] [18] RevSure | Who Owns the Agents (RevOps as agent-fleet management) | 2026 | https://www.revsure.ai/research-playbooks/who-owns-the-agents | [vendor claim] [19] Indeed Hiring Lab | AI Jobs and Labor Market (AI mention rate: marketing 8.4% → 14.9%) | 2026-07 | [to verify: original publication page] | [third-party reported] [20] BCG Henderson Institute | AI Will Reshape More Jobs Than It Replaces (50–55% reshaped / 10–15% eliminated) | 2026-04 | [third-party reported] [21] The Pragmatic Engineer | What is Growth Engineering (a fully staffed growth team costs about $1 million a year; $5 million ARR threshold) | https://newsletter.pragmaticengineer.com/p/what-is-growth-engineering | [third-party verified: first-hand from a practitioner]

[22] Gartner | 2026 CMO Spend Survey (AI at 15.3% of marketing budgets; marketing budget at 7.8% of company revenue; 70% call becoming an AI leader critical, only 30% report mature readiness) | 2026-05-11 | https://www.gartner.com/en/newsroom/press-releases/2026-05-11-gartner-2026-cmo-spend-survey-finds-cmos-allocate-15-point-3-percent-of-marketing-budgets-to-ai-but-only-30-percent-are-ready-to-scale-ai-capabilities | [official: institutional survey] [23] Atlassian | State of Teams 2026: CMO Brief (94% of marketers use AI, only 3% of CMOs are certain of org-wide AI ROI; the fragmentation tax) | 2026-06-01 | https://www.atlassian.com/blog/guides-research/the-state-of-teams-2026-cmo-brief | [third-party verified] [24] Bain | AI Won't Transform Your Go-to-Market Until You Transform the Work (1,125 sales and marketing leaders, 18 industries, 40 countries; 90%+ experimenting; winners with 2x revenue impact and 30% more productivity) | https://www.bain.com/insights/ai-wont-transform-your-go-to-market-until-you-transform-the-work-b2b-growth-agenda-2026/ | [third-party verified] [25] Salesforce | Agentic Enterprise Index 2026 (agents activated by enterprises grew nearly 3×) | https://www.salesforce.com/news/stories/agentic-enterprise-index-2026/ | [vendor claim] [26] Duke Fuqua / The CMO Survey (35th edition) | CMOs Face Headwinds (optimism lowest since 2020; activity contracting, existing customers prioritised; AI expected to carry over half of marketing activity within three years) | https://www.fuqua.duke.edu/duke-fuqua-insights/CMOs-Face-Headwinds-Even-as-Marketing-Value-and-AI-impact-grow | [third-party verified] [27] The Sword and the Script | Gartner CMO trilemma 2026 (flat budgets, more tasks) | https://www.swordandthescript.com/2026/07/gartner-cmo-trilemma-2026/ | [third-party verified] [28] TMTPost (Chinese: 钛媒体) | Zhipu and MiniMax double revenue while sales spend falls (Zhipu revenue RMB 954 million, +399.7%, sales and marketing RMB 178 million, −14.8%, API revenue RMB 825 million or 86.5%; MiniMax revenue RMB 786 million, +283.1%, sales and distribution −17.9%) | 2026-09-16 | https://www.tmtpost.com/8141022.html | [third-party reported] [29] Growth Hacker AI Weekly (compilation) | The AI office-suite war as frontline sales sees it (buyers decide on money saved, not model capability) | https://mp.weixin.qq.com/s?__biz=MzA3MDQyNDUzMA==&mid=2649928659&idx=1&sn=48f40bdd248d45a838c1684832ae3c54 | [third-party reported]

[30] AI-CMO | State of AI Marketing 2026 (survey of 2,500 marketing professionals: 87% of enterprise teams use AI; AI's share of marketing budget 11% → 18%; top challenges quality control 62%, training 54%, integration 47%) | 2026-01-15 | https://ai-cmo.net/assets/reports/state-of-ai-marketing-2026.pdf | [third-party verified: self-reported]

[31] CMSWire (citing Gartner's 2026 CMO Spend Survey, 401 CMOs) | Why are marketers still paying for martech they never use? (martech at 19.4% of the marketing budget, a five-year low, down from 26.6% in 2021; marketers use on average 49% of their stack's capability; 65.7% of organisations rank data integration as the top stack-management challenge) | 2026 | https://www.cmswire.com/digital-marketing/why-are-marketers-still-paying-for-martech-they-never-use/ | [third-party verified: citing Gartner] [32] China Daily | Advertiser survey: China's 2026 marketing investment expected to grow 10% on average (among advertisers that invest in content marketing, 44% expect content-marketing spend to rise in 2026; sports events lead with a 31% increase) | 2026-01-21 | https://cn.chinadaily.com.cn/a/202601/21/WS69704bf0a310942cc499c24c.html | [third-party verified] [33] SaaSHero (citing the Aleph x Benchmarkit 2026 benchmarks) | How to Improve CAC Payback Period (median 16 months, previously 18; 198 of 342 companies disclosing; denominator must use gross-margin-adjusted revenue; compute per channel and per cohort) | 2026-10-21 | https://www.saashero.net/strategy/improve-cac-payback-period-saas | [third-party verified] [34] MarTech (citing a Semrush survey) | The Real Impact of AI on Budgets, Stacks and Teams (42% increased budgets because of AI, about 16% cut; roughly a third report significant headcount growth; sample of 1,008 U.S. marketing budget decision-makers) | 2026-04-01 | https://martech.org/the-real-impact-of-ai-on-budgets-stacks-and-teams/ | [third-party verified] [35] Visionary Marketing | MarTech Stack Statistics 2026 (2,400 marketers plus a full audit of 240 accounts: 121 tools on average; 34.1% monthly active utilisation; 84.1% AI tool adoption; stack at 22.7% of budget; 18.4 months average tool tenure; dormant licences 32.4%, duplicate functionality 18.7%) | 2026-06-03 | https://visionary-marketing.co.uk/blog/martech-stack-statistics-2026 | [third-party verified]

# Chapter 24 Don't Cross These Lines: Six Failure Files

Every earlier chapter was about how to do things; this one is the reverse: it collects only the ways to die that have a name, a date, and a number. Someone has crossed the line and paid the price at every step above. So this chapter gives you 7 methods — one for each of the six files, plus a closing one — starting with how others went down.

A risk checklist gets read, nodded at, and put in a drawer; a failure file does not — someone really paid for each of them.

---

## Method 1: Don't amplify a claim you can't back

**What to do**: Cross “a position in AI answers” off the purchase order, and don't use fake reviews to amplify an unsupported claim.

**How**

1. Name every customer logo on the wall, every award list, every five-star rating, and every review piece, one by one, and each must point to a real contract or authorization; take down first anything that cannot.
2. Get written authorization for names and logos before they go live, spelling out the use and how it can be revoked; a trial customer cannot be displayed as a customer.
3. For every performance claim, find the internal source first, and delete it if there is none; write “who signed off” into the creative ledger.
4. Don't buy services that fake orders, fake reviews, or buy award lists, and don't sell them either: both directions are on the enforcement list.

**Example**: The “customers” on 11x's website logo wall, including ZoomInfo and Airtable, were found by a media investigation not to be real customers (only about a one-month trial); ZoomInfo sent a letter alleging false advertising and trademark infringement, and the report said total customer retention was below 50% [third-party verified]; Artisan was banned by LinkedIn at the end of 2025 and reinstated about two

weeks later in January 2026, with its lead pipeline cut off outright during the suspension [third-party verified]; the U.S. FTC passed the Consumer Reviews and Testimonials Rule (16 CFR Part 465) 5-0 on August 14, 2024, and the April 2026 TruHeight case settled for $4 million, with each violation of a consent order carrying a civil penalty of up to $53,088 [official]. AI selling points from the same period have also been priced: the SEC's first “AI washing” enforcement in March 2024 brought a $400,000 fine, and DoNotPay settled for $193,000 in one of the five FTC Operation AI Comply cases [official].

**Pitfalls**: Don't use “everyone is faking it” as a reason — the first round of pushback hurts the honest sellers in the trade. A single fabricated review can influence a model output within 72 hours [to verify: 1].

## Method 2: Don't mistake “can do” for “should do”

**What to do**: Take personalization only as far as the granularity of consent, and don't use the anxiety you found to push up the ticket size.

**How**

1. Write the purpose of the data as one sentence that can be read aloud: data collected for purpose A is not used for purpose B, and a change of purpose requires fresh consent.
2. Record consent by “for what, at what frequency, on what channel”; carry out a personal information protection impact assessment before automated decision-making, and provide a convenient way to refuse.
3. Self-test before reaching out: if this message were screenshotted and posted in a public community, would you still be okay with it?
4. Every touch by AI outbound calling and AI customer service must be traceable to a consent chain; anything that cannot be traced, stop first.

**Example**: Qualtrics's October 2025 survey (14 countries, 20,000+ consumers): 53% of consumers worry about data misuse in automated AI interactions (up 8 percentage points year over year), and after a bad experience only 29% give feedback directly while 47% cut spending [vendor claim]; Usercentrics's “State of Digital Trust 2026” (7 countries, 11,000 people): 52% of consumers trust AI with personal data less, and in Sweden 77% find AI personalization intrusive [official claim]. On the Chinese side: Shenzhen Zhenai (Chinese: 珍爱网) used “virtual humans” and age

anxiety to push the ticket size to tens of thousands of yuan, and got a 1.7 million yuan fine in return [third-party verified]; a central-media investigative report described AI outbound calling at 800–1,500 calls per machine per day and costs as low as a little over one jiao (0.1 yuan) per minute, grounded in Article 24 of the Regulations for the Implementation of the Law on the Protection of Consumer Rights and Interests, which bans commercial calls without consent [official]. A scaled sample of AI impersonating real people has appeared as well: one Chinese app studio running 20-plus dating apps abroad put large language models into the matchmaking pool to pose as female users — over a two-week investigation the system generated more than 4,700 distinct AI personas and handled 2.36 million messages, with a single model talking to about 25,000 people at once and an AI-to-human match ratio of about 4:1, humans only appearing on camera and in live video calls; individual transfers were just RMB 20–50 under an industry rule of "chat three times, ask once" [third-party reported: media investigation]. Turning one-to-one emotional labour into a batch-instantiated persona pool is a direct reason why risk and compliance detection costs are rising.

**Pitfalls**: The currency of personalization is trust, and the credit limit is held by the user — technology decides how accurately you can speak, tolerance decides how close you can get.

## Method 3: Don't let AI promise where you have no facts

**What to do**: Write the cost of “AI saying the wrong thing” into the process, not into an apology letter: for price, terms, delivery time, payment terms, and qualifications, AI produces drafts only.

**How**

1. List what AI may not state as settled: price, terms, delivery time, payment terms, qualifications, and compensation.
2. Add a “who signs off on this sentence” field to every customer-facing AI channel; with no signer, it does not go out.
3. The first action after a hallucination is to clarify proactively and trace back the people affected.
4. Flag AI's limits on the welcome page and in prominent places in the interface, and route key answers through retrieval-augmented generation (look up the material first, then answer).

**Example**: In April 2025, Cursor's support bot invented a new policy that “one subscription only covers one device”, Reddit and Hacker News took it as an official change, and users publicly announced they were cancelling their subscriptions; co-founder Michael Truell apologized afterward, attributing it to hallucination and “non-deterministic output” [third-party verified]. The Air Canada chatbot promised Jake Moffatt he could buy a ticket at the original price and apply for a bereavement discount refund within 90 days; after the company refused the refund it argued that “the chatbot is a separate legal entity responsible for its own actions”, was rebuked by the British Columbia Civil Resolution Tribunal as “remarkable” and dismissed, and was ordered to pay about 650 Canadian dollars [third-party verified] [to verify: 2]. In December 2025, the Hangzhou Internet Court issued a first-instance ruling in the country's first generative AI “hallucination” tort case: AI does not have the standing of a civil subject, and a commitment it generates on its own does not constitute a declaration of intent by the service provider; because the defendant had prominently flagged the limitations and used retrieval-augmented generation, the court found no fault and dismissed the claims [third-party verified]. The documented cases already form a list. A compilation of nine enterprise-grade hallucination incidents leads with Air Canada's chatbot inventing a policy that a full-price ticket had to be bought before claiming a bereavement discount, which a tribunal made the airline honour, alongside lawyers sanctioned in Mata v. Avianca for citing fabricated cases and Deloitte Australia partially refunding a government-commissioned report that carried invented sources; its judgment is worth keeping — the most expensive hallucinations are not funny screenshots but fluent, specific and wrong output, which is why they survive review and reach customers, regulators and courts, with reasoning models generating more claims, multi-step agent workflows compounding errors and retrieval itself not guaranteeing grounding, so the fix is architectural: per-output source traceability, influence scoring, full agent execution logs and human contestability [third-party verified].

**Pitfalls**: Hangzhou's case won an exemption from liability, and the price was that a customer's mistaken belief had already occurred and public credibility had been tested in open litigation: an exemption saves money, not trust.

## Method 4: Don't buy an amplifier with no signal

**What to do**: Justify the project by “how many hours a week a job eats”, not by “what this AI can do”.

**How**

1. Run a time audit first: have the people who will use it report the hours a week they spend on preparatory work, and sort by hours.
2. Before buying, ask “who will change which habit because of it”; if you cannot name the person and the habit, don't spend the budget yet.
3. After launch, measure adoption speed, not capability breadth — a simple agent at 90% adoption beats a complex agent at 10%.
4. Design for deletion: with every tool you add, kill one manual task at the same time.

**Example**: Boomi's Luke Hagstrand has a post-mortem that includes negative data: the team built a multi-agent chatbot for the customer success department that “could answer any CSM question”, with a beautiful demo and a high-profile launch, adoption under 20%, and they finally decided to kill it: it had become “one more thing to stuff into the daily workflow” [vendor claim]. After deconstructing it, they rebuilt it as a set of small, specialized, nearly invisible agents that automatically aggregate and push signals from Salesforce, Gong, and Gainsight every week; 30 days after 2.0 launched, adoption hit 85%, saving each CSM 4–6 hours a week [vendor claim]. The post-mortem also relays two numbers: MIT says 95% of organizations get zero return on generative AI, and the industry line is that “73% of companies are stuck in PoC purgatory” [vendor claim] [to verify: 3]. Built-in traffic is not demand either: Lionel Messi's sports drink Mas+ shut down less than two years after launch, and Alex Cooper's functional drink Unwell announced it would stop after that year's Halloween flavours, with reporting attributing both to market maturity, consumer fatigue and a gap between mass-market demand and a creator's audience — follower scale is not the same as sustainable repeat demand [third-party verified].

**Pitfalls**: Adoption rate measures behavior change and has nothing to do with how much the system can do. A purchase that cannot report an adoption rate bought a subscription invoice.

## Method 5: Don't stare at a number that doesn't mean a deal

**What to do**: Write the “definition contract” first, then turn on AI scoring: what counts as a qualified lead, written in buyer language as a definition a system can read.

**How**

1. The definition contract must answer at least three things: what counts as a qualified lead; what counts as a mismatch and who adjudicates it; and who signs off on the conversion rate from one tier to the next.
2. Run intent data as a “candidate set”: what you buy is a ranking, and handing it straight to sales means letting AI replicate an uncalibrated assumption many times over.
3. Once a quarter, pull lead counts, MQL counts, reply rate, and meeting counts back against the closed-deal denominator; if the metrics are up and deals are flat, check the definitions first.
4. Write a one-sentence definition for each of the five labels, with the signer and the review date noted.

**Example**: A B2B marketing agency's benchmark library (including an operations audit of 47 ABM deployments): intent-prioritized accounts converted to opportunities at 21.3% versus 8.4% for non-prioritized ones, and median sales cycle shortened by 28 days; the same document's “quality warning”: 62% of buyers reported that fewer than 70% of flagged accounts showed CRM-documented activity within 30 days, 47% of intent records expired within 14 days, and 29% still had no two-way CRM integration 6 months after deployment [third-party verified]. A cohort study of 14 B2B SaaS sales organizations: AI-SDR outreach reply rates decayed by more than 60% within 18 months, from 11.2% at launch → 6.1% at month 12 → 4.4% at month 18 [third-party verified]. On the Chinese side, Hangzhou Shantao Technology (Chinese: 杭州杉淘科技) faked 589,000 orders on a “shoot A, ship B” scheme and was fined 725,000 yuan [third-party verified]. The field itself is being harvested by channel sellers first: one of GEO's earliest working businesses is recruiting city agents — the pitch is usually "no franchise fee", yet first stock, prepaid compute, system fees, deposits and backend packages always find a name to collect the money under (the listed prices include a city-exclusive agency starting at RMB 50,000 of first stock, 50%–70% commission and self-reported payback claims of two to three months); the structural mismatch is that the companies actually suited to long-term GEO work — those with brand equity, users who genuinely ask AI, and the ability to keep producing credible content — are not the restaurants, renovators, training schools and shops

that regional agents can most easily reach — "even if you hit 100% mention rate on the ten questions the vendor designed, that does not mean any consumers actually asked them" [third-party reported: investigative reporting].

**Pitfalls**: Early decay is often misattributed to seasonality, ICP drift, or product changes, and most teams only notice after month 12 — the very period when the metrics were rising is the period when mismatch was accumulating.

## Method 6: Don't leave compliance until after launch

**What to do**: Treat compliance as the title registration of your acquisition assets, done before launch — content states clearly who generated it, and data states clearly where it came from.

**How**

1. Do explicit and implicit labeling together: explicit is the notice users can see (text at the start, end, or middle; a spoken notice added to audio; prominent placement for images and video); implicit is written into the metadata, covering generation attributes, the service provider's name or code, and the content number, with digital watermarks encouraged.
2. The label must still be there after material is downloaded, copied, or exported; when a user requests content without an explicit label, logs must be kept for no less than six months in accordance with the law.
3. Lock down one internal rule: no maliciously deleting, tampering with, forging, or concealing labels, and no providing such tools or services to outsiders.
4. For EU customers, four categories — AI interacting directly with people, generating synthesized content, emotion recognition or biometric categorization, and deepfakes — must be disclosed to a “clearly discernible” standard; fine print in the footer and a label that flashes by do not qualify.

**Example**: The four-ministry Measures for Labeling AI-Generated and Synthesized Content (CAC General Office Document [2025] No. 2) took effect on September 1, 2025, and platforms must verify implicit labels [official]; the transparency obligations in Article 50 of the EU AI Act (Regulation (EU) 2024/1689, published in the Official Journal on July 12, 2024) apply from August 2, 2026 [official] [to verify: 4]; on the data side, the practical reading of the Personal Information Protection Law requires a personal information protection impact assessment before profiling-based automated

decision-making [third-party verified] [to verify: 5]. The other half of the obligation — the conversational and emotionally interactive scenarios — is being written into law: on January 8, 2026 the Attorney General of Kentucky sued Character Technologies and its product Character.AI, becoming the first U.S. state to sue an AI chatbot company, alleging it put profit above children's safety — the platform has over 20 million monthly users, has been blamed for at least two adolescent suicides (a 14-year-old Florida boy in 2024 and a 13-year-old Colorado girl in 2025), and is used by tens of thousands of Kentuckians including thousands of minors with no age verification, in violation of the Kentucky Consumer Protection Act and the Kentucky Consumer Data Protection Act [official]. It is not confined to one state: a law firm's tracker says 78 chatbot-related bills were filed across 27 states in the first weeks of 2026, at least six more states advanced legislation after California's companion law took effect on January 1, 2026, and an analysis of 284 deployer-facing AI litigation matters found chatbot wiretap suits (brought under the Electronic Communications Privacy Act and state wiretap statutes) grew from 2 in 2021 to 30 in 2025, the fastest-growing category — and they govern what the bot collects, not what it says, so satisfying disclosure rules does not satisfy the prior consent a wiretap statute requires [third-party verified: law firm analysis].

**Pitfalls**: Material without a label cannot be distributed compliantly, and what cannot be distributed is worth zero; “deleting the label” itself turns a violation from “forgot” into “on purpose”.

## Method 7: Don't automate a funnel that doesn't hold

**What to do**: Validate the funnel before doing any AI transformation; if any one of the three questions cannot be answered, work may not start.

**How**

1. Ask: does this funnel already hold in the segment where “the buyer has not contacted you yet”? Or is it waiting for a customer who is no longer present to walk in?
2. Ask: can every tier be traced to a closed deal? For a tier that cannot, what justifies keeping it in the process?
3. Ask: with all the AI switched off, does this funnel still turn? Run it manually on a small sample; if it does not run, it does not hold.
4. On the project charter, write “is the entry point still valid” as the first checkpoint.

**Example**: At Chegg's earnings call on the evening of May 1, 2023, CEO Dan Rosensweig said that student interest in ChatGPT had surged since March and was affecting new customer growth; the stock closed down 48.41% at $9.08 the next day, quarterly revenue guidance of $175–178 million came in far below analysts' expected $193.6 million, and yet Q1 revenue and earnings per share both beat expectations [third-party verified]. On the content side: Charleston Crafted (a home-improvement blog) lost 70% of its traffic and 65% of its ad revenue from March to May 2024; The Planet D (a travel blog founded in 2008) first lost half its traffic, then fell another 90% after layoffs, and closed in 2025; Stereogum's founder said in his own account that ad revenue was cut by 70%; Business Insider's organic search traffic fell 55% from April 2022 to April 2025 and it cut 21% of staff in May 2025 [third-party verified]. Pew, July 2025: when an AI summary appears, only 1% click the cited link inside the summary and 8% click the organic results below (15% when there is no summary) [third-party reported]. Stack Overflow's CEO in post-mortem: ChatGPT put the company through an “existential moment” and triggered a “Code Red”, and three years later it transformed into a SaaS business aimed mainly at enterprises; more than 80% of users want to use or already use AI for code, but only 29% truly trust AI with valuable work [third-party verified]. On the buyer side, tender records published by the City of Bellingham in January 2026 show that an employee asked ChatGPT to write tender requirements “more favorable to a certain vendor”, and at least 16 passages of the model's wording went verbatim into the final 350 requirements, with a five-year, $2.7 million contract awarded to the higher bidder (about $2.5 million versus about $1.45 million for the runner-up) [third-party verified].

**Pitfalls**: “The product didn't break, the execution wasn't wrong, the buyer changed entry points” — that line is the mechanism itself.

---

## Notes for this chapter

- **The background, amounts, dates, and regulatory document numbers of the six files follow the original manuscript, unchanged, with no new facts or cases added.** Phenomena for which named, checkable evidence was not obtained are left blank by preference.

- **Vendor self-reported data is always read at a discount**: the performance numbers from Qualtrics, Usercentrics, Boomi, The Starr Conspiracy, and Digital Applied are vendor readings and have not been reproduced by an independent third party; the denominator behind “a X-fold / X% improvement” is defined by the publisher.
- **The 12 questions at the end of the chapter, the three principles, and the three questions in Method 7 are all [our judgment]**; they are a question order, not a scoring standard; each company should first run a quarter under its own definitions before setting thresholds.
- **Items still to be verified** (original sources not traced; do not cite them alone): ① the 2026 CCTV 3·15 “AI poisoning” item is relayed only from a vendor press release, and “generating a dozen or more articles in minutes” and “making a nonexistent product a top recommendation in 72 hours” were not obtained from CCTV's original program or CCTV.com's original reporting [to verify: 1 / 6]; ② the Air Canada award has two readings (about 650 Canadian dollars / 812 Canadian dollars), and the original decision text was not checked [to verify: 2 / 8]; ③ “MIT says 95% zero return” and “73% of companies stuck in PoC purgatory” in the Boomi post-mortem are secondhand relays [to verify: 3 / 7]; ④ “applies from 2026-08-02” for Article 50 of the EU AI Act was not checked sentence by sentence against the original text of Article 113 [to verify: 4 / 9]; ⑤ the PIPL article numbers and 11x's “retention below 50%” (about 24 sources; whether it is logo, contract, or revenue retention is unclear) have only a single secondary source [to verify: 5 / 10].

## Appendix: 12-question self-check

Use it as a red-light checklist. If any single light is on, do not scale up yet — because what scales is the mistake.

| # | Question (the red-light signal is in parentheses) |
|---|---|
| 1 | Have we run this funnel manually on a small sample **without AI**? (never ran it / ran it and didn't look at the data) |
| 2 | Can we make mainstream models **accurately restate** who we are, what we do, and **who we are not for**? (cannot produce an accuracy baseline) |
| 3 | |

| # | Question (the red-light signal is in parentheses) |
|---|---|
| | How many third-party sources are **not under our control** (editorial media, review sites, communities, podcasts)? (our own sites are more than half) |
| 4 | Can we explain clearly why “this lead is qualified” — in **buyer language** rather than a score range? (all we can give is a score and a model name) |
| 5 | For the five tiers — lead / prospect / opportunity / win / payment — the **definition contract**: who signs, how often is it reviewed? (nobody signs / the definition is stuck in an old deck) |
| 6 | For our AI outreach, can we state for each person **at what granularity they consented to which purpose**? (using “signing up means consent” to cover everything) |
| 7 | For every AI-generated piece aimed at buyers, **who signs off, who can pull it, how often is it reviewed**? (there is a publishing process but no retraction process) |
| 8 | Do our AI-generated materials have both **explicit and implicit labels**? Whose service-provider code is in the metadata? (explicit only, with metadata set by the outsourcing vendor) |
| 9 | Do our customer-facing chatbots and voice agents meet disclosure obligations **where the user is located** (including frequency and protection of minors)? (compliant only by place of registration) |
| 10 | In our most recent “**turn AI off for a week**” controlled experiment, which metrics did we look at? (never done, or only activity volume) |
| 11 | What is the **adoption rate** of the AI tools and agents we have bought? Why don't the non-adopters use them? (cannot give an adoption rate, only renewal records) |
| 12 | If a platform changed its rules or a model changed version tomorrow, which of our acquisition assets **would go to zero overnight**? (cannot list them, let alone hedge) |

Questions 8 and 9 are the ones Chinese companies most often miss: **labeling is about distribution eligibility, not aesthetics**; question 12 maps to the book's hedging logic — **dependency cannot be eliminated, but it can be inventoried, priced, and spread**.

## Three principles

Six files, one closing method, twelve red lights, distilled into three sentences.

### Principle One: Pushback always arrives

Wherever **a signal that can be faked at scale** is used for acquisition, pushback will certainly arrive — only the timing and the form are unknown. This book has seen four ways it shows up: platform rules (the LinkedIn ban), regulatory process (FTC fines, typical cases from the State Administration for Market Regulation), buyer detection and filtering (recipients spotting AI templates, buyers using AI to rewrite tender requirements), and rivals' retaliatory imitation (once the technique is copied, the signal's value goes to zero).

The operational meaning is not “don't use fakeable signals”, but: **while the signal still has value, move part of the budget to producing the unfakeable part** — first-hand data, real customer testimonials, a verifiable operating process, professional judgment someone is willing to put their name to. What they share is this: **the cost of faking them is higher than the faker is willing to pay.**

### Principle Two: An amplifier has no direction

AI does not produce direction; it multiplies the slope of an existing direction by a larger number. **A negative number times a large number gives a more negative number.** This explains four of the six files in this chapter: when the funnel does not hold, automation makes it fail to hold faster; when the lead definition is fuzzy, scoring makes it fuzzy many times over; when the script has no factual foundation, generation makes it fabricate more loudly; when compliance is not done, scale makes it break the rules more expensively.

The order is therefore fixed: **establish that the slope is positive first, then talk about amplification.** The way to judge the sign of the slope is crude, and that is exactly why so few people do it — run it manually on a small sample and see whether it actually closes.

### Principle Three: Closing the loop comes before scale

Scale is a multiplier on the loop, not a substitute for it. A minimum viable loop must answer four questions and leave checkable evidence: **who was reached → why him → who closed → why he closed.** When you cannot answer, the larger the scale the further you are from the answer — the first effect of scale is to worsen the signal-to-noise ratio.

One sentence you can put into a quarterly goal: **move only one thing per quarter, and it must be the broken link inside the loop, not a new tool outside it.**

---

## Sources for this chapter

**Regulatory documents and statutes ([official])** - U.S. FTC | Final Rule Banning Fake Reviews and Testimonials (16 CFR Part 465) | 2024-08-14 | https://www.ftc.gov/news-events/news/press-releases/2024/08/federal-trade-commission-announces-final-rule-banning-fake-reviews-testimonials; FTC v. TruHeight ($4 million judgment) | 2026-04-13 | https://www.ftc.gov/news-events/news/press-releases/2026/04/ftc-takes-action-against-truheight-deceptive-unsubstantiated-advertising-supposed-height-enhancing - U.S. FTC | Operation AI Comply (including Rytr, DoNotPay) | 2024-09-25 | https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes | see also https://www.ftc.gov/legal-library/browse/cases-proceedings/donotpay; U.S. SEC | AI washing enforcement ($400,000) | 2024-03-18 | https://www.sec.gov/newsroom/press-releases/2024-36 - Cyberspace Administration of China and three other departments | Measures for Labeling AI-Generated and Synthesized Content (CAC General Office Document [2025] No. 2, effective 2025-09-01) | https://www.cac.gov.cn/2025-03/14/c_1743654684782215.htm - European Union | Regulation (EU) 2024/1689 (AI Act), OJ L, 2024/1689, 12.7.2024 | https://eur-lex.europa.eu/eli/reg/2024/1689/oj - State Council | Article 24 of the Regulations for the Implementation of the Law on the Protection of Consumer Rights and Interests | effective 2024-07-01 | https://www.gov.cn/zhengce/content/202403/content_6940158.htm - Standing Committee of the National People's Congress | Personal Information Protection Law of the People's Republic of China | effective 2021-11-01 | https://www.cac.gov.cn/2021-08/20/c_1631050028355286.htm

**Court judgments and regulatory cases ([third-party verified])** - Moffatt v. Air Canada, 2024 BCCRT 149 | American Bar Association analysis | 2024-02-14 | https://www.americanbar.org/groups/business_law/resources/business-law-today/2024-february/bc-tribunal-confirms-companies-remain-liable-information-provided-ai-chatbot/ - Hangzhou Internet Court | first-instance judgment in the country's first generative AI “hallucination” tort case | People.cn (People's Court Daily) | 2026-01-26 | http://society.people.com.cn/n1/2026/0126/c1008-40652440.html - State Administration for Market Regulation | seven typical cases of online unfair competition | Xinhua Net | 2024-12-31 | http://www.news.cn/food/20241231/ceec0fa0841345fb872ed48b07dd3748/c.html

**Media, companies, and third-party research** [third-party verified] unless otherwise noted - CNBC | Chegg stock falls 48.41% | 2023-05-02 | https://www.cnbc.com/2023/05/02/chegg-drops-more-than-40percent-after-saying-chatgpt-is-killing-its-business.html - The Verge | Decoder interview with Stack Overflow's CEO | 2025-12-15 | https://www.theverge.com/podcast/844073/stack-overflow-ceo-ai-coding-chatgpt-code-red-interview - AdExchanger | The AI search traffic collapse | 2026-01-06 | https://www.adexchanger.com/publishers/the-ai-search-reckoning-is-dismantling-open-web-traffic-and-publishers-may-never-recover/ - Stereogum (Scott Lapatine) | Getting Killed By AI | 2025-11-10 | https://stereogum.com/2478838/stereogum-relaunch/news | [official: founder's own account] - KNKX / Cascade PBS | City of Bellingham employee used ChatGPT to exclude bidders | 2026-01-05 | https://www.knkx.org/government/2026-01-05/city-of-bellingham-chatgpt-ai-contract-vendor - Ars Technica | Cursor's support bot invents a policy | 2025-04-17 | https://arstechnica.com/ai/2025/04/cursor-ai-support-bot-invents-fake-policy-and-triggers-user-uproar/ - TechCrunch | 11x claims customers it does not have | 2025-03-24 | https://techcrunch.com/2025/03/24/a16z-and-benchmark-backed-11x-has-been-claiming-customers-it-doesnt-have/ - TechCrunch | Artisan raises $25 million; CEO admits early hallucinations | 2025-04-09 | https://techcrunch.com/2025/04/09/artisan-the-stop-hiring-humans-ai-agent-startup-raises-25m-and-is-still-hiring-humans/ - TechCrunch | LinkedIn bans and reinstates Artisan | 2026-01-07 | https://techcrunch.com/2026/01/07/yes-linkedin-banned-ai-agent-startup-artisan-but-now-its-back/ - Boomi (Luke Hagstrand) | Built an AI agent nobody used | 2025-10-01 | https://boomi.com/blog/customer-success-ai-agent/ | [vendor claim: includes negative data] - Qualtrics | AI customer service fails at 4 times the rate of other uses (14 countries) | 2025-10 | https://www.qualtrics.com/news/ai-powered-customer-service-fails-at-four-

times-the-rate-of-other-tasks/ | [vendor claim] - Usercentrics (Sapio Research) | State of Digital Trust 2026 (7 countries, 11,000 people) | 2026-03 | https://usercentrics.com/resources/state-of-digital-trust-report-2026/ | [official claim] - The Starr Conspiracy | B2B intent data benchmarks 2025 (refreshed 2026-05) | https://www.thestarrconspiracy.com/insights/benchmarks/b2b-intent-data-benchmarks-2025 - Digital Applied | The case against AI SDRs 2026 (cohort study of 14 companies) | 2026-05 | https://www.digitalapplied.com/blog/case-against-ai-sdrs-contrarian-analysis-2026 - China.com.cn | AI outbound calls at a thousand a day (including a Shanghai Consumer Council survey) | 2026-07-28 | http://news.china.com.cn/2026-07/28/content_118621717.shtml - China Economic Net | AI-generated content must carry labels (over 490 large models filed) | 2025-09-01 | http://www.ce.cn/xwzx/gnsz/gdxw/202509/t20250901_2460291.shtml - King & Wood Mallesons (Chinese: 金杜律师事务所) | Retail marketing data compliance | 2023-06 | https://www.kingandwood.com/cn/zh/insights/latest-thinking/uni-marketing-data-compliance-in-retail-marketing-in-era-of-digitalization.html - Sohu (Chinese: 搜狐) (vendor press release, relaying the 2026 CCTV 3·15 “AI poisoning” item) | 2026-08-31 | https://www.sohu.com/a/1068667519_100119123 | [vendor press release: quoted secondhand] - Pew Research Center | Study of click behavior when an AI summary appears (1% / 8% / 15%) | 2025-07-22 | https://www.pewresearch.org/short-reads/2025/07/22/google-users-are-less-likely-to-click-on-links-when-an-ai-summary-appears-in-the-results/

- Growth Hacker AI Weekly (investigative reporting) | GEO franchising looks like a scam: the earliest working GEO business is collecting franchise fees from city agents (self-reported pricing such as RMB 50,000 of first stock for a city exclusive and 50%–70% commission; structural mismatch between who GEO suits and who agents can reach) | https://mp.weixin.qq.com/s?__biz=MzA4MTc3NzY0Mg==&mid=2652351009&idx=1&sn=b1ef2564fc89ae3083641db5096a2 | [third-party reported: investigative reporting]
- Growth Hacker AI Weekly (media investigation) | Dating apps filling the matchmaking pool with LLM personas posing as women (4,700+ AI personas and 2.36 million messages in two weeks; AI-to-human ratio about 4:1; the "chat three times, ask once" rule) | https://mp.weixin.qq.com/s?__biz=Mzg5ODU4OTU3MQ==&mid=2247502366&idx=1&sn=07d845ffa276a6eb68edd0a82cf0 | [third-party reported: media investigation]

- Modern Retail | Why more celebrity-led brands are shutting down (Messi's Mas+ closed; Alex Cooper's Unwell discontinued; the gap between mass-market demand and a creator's audience) | 2026-08-25 | https://www.modernretail.co/operations/why-more-celebrity-led-brands-are-shutting-down/ | [third-party verified]
- Kentucky Attorney General's Office | AG Coleman Sues AI Chatbot Company for Preying on Children (2026-01-08, the first state-level lawsuit in the U.S.; Character Technologies and Character.AI; over 20 million monthly users; blamed for the 2024 suicide of a 14-year-old Florida boy and the 2025 suicide of a 13-year-old Colorado girl; tens of thousands of Kentuckians including thousands of minors with no age verification; violations of the Kentucky Consumer Protection Act and the Kentucky Consumer Data Protection Act) | 2026-01-08 | https://kentucky.gov/Pages/Activity-stream.aspx?n=AttorneyGeneral&prId=1857 | [official]
- Baker Botts (Our Take) | AI Chatbot Regulation: 78 State Bills, 58 Lawsuits (78 chatbot-related bills in 27 states in the first weeks of 2026; at least six more states advanced legislation after California's law took effect on 2026-01-01; among 284 deployer-facing AI litigation matters, chatbot wiretap suits grew from 2 in 2021 to 30 in 2025, the fastest-growing category; disclosure compliance does not equal wiretap prior consent) | 2026-02-24 | https://ourtake.bakerbotts.com/post/102mipe/ai-chatbot-regulation-78-state-bills-58-lawsuits | [third-party verified: law firm analysis]
- Seekr | AI Hallucination Examples: 9 Real Enterprise Cases (fabricated case law in Mata v. Avianca, the Air Canada chatbot ruling, invented citations in a paid Deloitte report; the most expensive hallucinations are fluent and specific; recommends per-output traceability, influence scoring and full agent execution logs) | 2026-09-08 | https://www.seekr.com/resource/ai-hallucination-examples/ | [third-party verified]

# Appendix A · Case Index and Sources

This appendix is the book's tiered ledger: every source carries its **nature**, its **data point in time** and its **link test status**; every case carries whether it is **named and verifiable**. Any figure whose original source cannot be traced is not written into the book; any passage that uses a composite profile is explicitly flagged, both in the body and in this section.

## A.1 How source tiers are defined

| Source tier | Meaning | Rules for use |
|---|---|---|
| [official] | Official documents, official blogs, financial filings and official releases from an institution or platform | May be cited directly, with the publisher and date noted |
| [third-party verified] | Independent third-party research, academic papers, authoritative media reports | Must be checked against the original report/paper page |
| [vendor claim] | Vendor self-reported data, self-run benchmarks, advertorials and press releases | A lead only, **never to be treated as fact**; the vendor's identity and interest must be disclosed when citing |
| [our judgment] | The authors' inference | Kept in a separate column from fact; readers may challenge it |
| [our own measurement] | Evidence collected by this book itself | Method, sample, time and boundaries must be given; reproducible and falsifiable |

## A.2 Key-figure verification ledger (verified)

| No. | Figure | Original source | Nature |
|---|---|---|---|
| N-01 | | | |

| No. | Figure | Original source | Nature |
|---|---|---|---|
| | US Google zero-click rate 68.01% (Jan–Apr 2026); 60.45% in 2024; about 45% a decade ago | SparkToro, 2026-06-08 (data source Similarweb; Datos for 2024) | [third-party verified] |
| N-02 | AI Mode monthly actives above 1 billion (Gemini 3.5 Flash as the default model) | Google official blog, 2026-05-20 (announced at I/O 2026) | [official] |
| N-03 | With an AI summary present, traditional-result clicks are 8% versus 15% without one; links inside the summary get 1% clicks; 18% of searches contain an AI summary; 88% of summaries cite more than three sources | Pew Research Center, 2025-07-22 (sample: 900 US adults, March 2025 data) | [third-party verified] |
| N-04 | Google bulk-sending requirements: spam-complaint-rate target < 0.1%, never to reach 0.3%; SPF/DKIM/DMARC and one-click unsubscribe required | Gmail official help page | [official] |
| N-05 | FCC declaratory ruling: AI-generated voice counts as "artificial or prerecorded voice" under the TCPA (2024-02-08) — without prior consent it is unlawful | FCC official ruling document (FCC-24-17) | [official] |
| N-06 | FTC final rule on fake reviews and testimonials: announced 2024-08-14, effective 2024-10-21 (16 CFR Part 465) | FTC press release + US Federal Register | [official] |
| N-07 | China generative-AI users 602 million, up 141.7% from the end of 2024, penetration 42.8% (as of 2025-12) | CNNIC 57th Statistical Report on China's Internet Development (published 2026-02-05) | [official] |
| N-08 | 6sense 2025 Buyer Experience Report: 85%–95% of purchases come from the "day-one list"; 80% of won deals go to the vendor preferred before first contact; 61% of evaluation is completed before contact | 6sense Science of B2B | [vendor claim] |

| No. | Figure | Original source | Nature |
|---|---|---|---|
| N-09 | Gartner forecast: traditional search volume down 25% by 2026 (published 2024-02-19; a forecast) | Gartner press release | [official: forecast] |

## A.3 Named case index

| Case | Chapter | Point | Verifiability |
|---|---|---|---|
| Chegg admitted on its earnings call that ChatGPT was hitting new-customer growth; shares fell −48.41% the next day | Ch.2 / Ch.10 | The direct impact of gateway migration on an incumbent acquisition model | Named and verifiable (earnings filing and media reports) |
| Stack Overflow "Code Red": 10% of staff shifted to firefighting; three years later it moved to B2B SaaS | Ch.2 / Ch.10 | An organization-level response after a traffic gateway is drained | Named and verifiable (public interviews) |
| Stereogum ad revenue −70%, move to subscriptions (first-person account signed by the founder) | Ch.1 / Ch.10 | Media survival in the zero-click era | Named and verifiable (signed first-person account) |
| Cursor's support bot invented a "single-device policy"; users publicly cancelled and the founder apologized | Ch.7 / Ch.17 | The cost of AI promises that are not true | Named and verifiable (media reports) |
| Bellingham, US: a staffer used ChatGPT to draft slanted tender conditions; a $2.7 million contract went to a pricier preferred vendor; the city opened an investigation | Ch.7 / Ch.13 | Buy-side AI use equally needs authorization and an audit trail | Named and verifiable (media reports) |
| Moffatt v. Air Canada (2024 BCCRT 149): a company is liable for its chatbot's misstatements | Ch.7 / Ch.17 | "The AI said it" is no defence | Named and verifiable (ruling text) |

| Case | Chapter | Point | Verifiability |
|---|---|---|---|
| A Chevrolet dealership's AI support agent was induced to agree to "sell a car for $1" | Ch.13 / Ch.17 | A textbook case of missing price authority | Named and verifiable (media reports) |
| CNET used an in-house AI tool to mass-produce personal-finance articles, paused it after corrections | Ch.14 | Industrial-scale content burns through byline credibility | Named and verifiable (media investigation) |
| WeChat (Chinese: 微信) banned the third-party add-on WeTool (2019) | Ch.10 / Ch.11 | Automating personal WeChat is a red line | Named and verifiable (media reports) |

## A.4 Composite-profile list

Any passage in the body that uses a "composite profile" is built from widely observed industry patterns and is **not a named event**; the body flags this explicitly. Current list:

| Location | Profile subject | Label |
|---|---|---|
| Ch.1 | An industrial-equipment maker that treats a single AI-visibility test as a conclusion | Composite profile ｜ Built from widely observed industry patterns, not a named event |

## A.5 Our own measurement archive

| Measurement | Content | Method archive |
|---|---|---|
| Measurement 1 | Multi-model test of the candidate set: 6 questions (3 Chinese–English pairs) × 5 model channels × 2 runs; 51 valid answers | See the repository folder `实测/实测1-候选集多模型/`: test report, raw answers, extraction results, results table, re-run script |

## A.6 Figures not adopted (listed for reference)

The following figures appear in industry discussion but are **not adopted by this book**, either because the original source could not be traced or because the definition does not match. If a reader sees these claims in some article, they can come back to this table to check why the book did not write them: median industry conversion rates for Leads→MQL→SQL, monthly actives and search penetration for the various content platforms (QuestMobile / Xiaohongshu / Douyin official figures), G2 and Forrester surveys of buyers' AI usage, Ahrefs and Seer studies of CTR decline, McKinsey's day-one-list share, and GDPR and EU AI Act fine and timeline details.

## A.7 Industry figures read but not written into the body (listed for reference)

These sources appeared in early drafts as reference points for industry differences; after the formal draft deleted the whole "industry landing table" research device, the body no longer cites them. They are kept here so that readers and later versions can trace the book's choices:

- Bartoli Consulting Group | B2B Manufacturing SEO Case Study: Electronic Power Design | https://www.bartoliconsulting.com/case-studies/seo-case-studies-b2b-manufacturing | [vendor claim]
- Dayu Zhiyuan (Yixunpan) (Chinese: 大鱼致远（易询盘）) | B2B Foreign-Trade Social-Media Marketing Project Fit Guide | https://www.rocforever.com/cn/B2B-n6595.html | [vendor claim]
- CNBC | Amazon faces a dilemma: fight AI shopping agents or join them | 2025-12-24 | https://www.cnbc.com/2025/12/24/amazon-faces-a-dilemma-fight-ai-shopping-agents-or-join-them.html | [third-party reported]
- Xiaoman Technology OKKI (Chinese: 小满科技 OKKI) | Relayed from a Shuzhi Qianxian interview piece (AI overall adoption rate 75%; iResearch 2024 foreign-trade B2B SaaS tools share 14.7%) | https://mp.m.ofweek.com/ai/a456714443567 | [third-party reported]
- Laigu Pro (Chinese: 来鼓 Pro) | New Oriental Douyin (Chinese: 抖音) DM AI reception case | 2026-07 | https://laigu.com/blog/教育行业抖音获客用ai私信自动回复30天提升获线/ | [vendor claim]

- 21st Century Business Herald | Securities-firm digital customer acquisition: the SAC's review of 19 securities firms' wealth-management digitalization cases | 2025-08-08 | https://www.21jingji.com/article/20250808/herald/15e5a6d995898e427b2a6496ee4e820a.html | [third-party reported]
- eMarketer | FAQ on GEO and AEO: Where AI search and SEO overlap in 2026 | 2026-04 | https://www.emarketer.com/content/faq-on-geo-aeo--where-ai-search-seo-overlap-2026 | [third-party reported]

## A.8 Case pool (named and verifiable, including those not used in the body)

This table is the complete ledger of the case-gathering effort: every entry meets three conditions — **named** (company / product / institution), **URL tested reachable**, and **a sentence supporting the conclusion readable at the source**. Those marked in the "Used in" column have been written into the body; those marked "not used in the body" were collected this round but not used, left for later versions and for readers to check themselves.

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Amazon** : Customer Obsession + Working Ba | [official] (codes) + [third-party: former Amazon executive | https://www.amazon.jobs/content/en/our-workplace/leadership-principles ; https://workingbackward | not used in the body |
| Microsoft** : Customer Connection Program ( | [official] | https://techcommunity.microsoft.com/blog/microsoftintuneblog/announcing-the-microsoft-management | Ch.1, Method 4 |
| Intuit** : Design for Delight — Follow-Me- | [official] (Intuit Inc. copyright doc | https://static1.squarespace.com/static/62726b549e95bf2087fc2759/t/62ffb171b5c6f01e0ed15a5a/16609 | not used in the body |
| | | https://review.firstround.com/how-superhuman-built-an-engine-to-find-product-market-fit/ | |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Superhuman** : Product-Market Fit Engine (putting | [third-party: first-person account signed by the founder] | | Ch.1, Method 4 |
| Cloudflare** : block AI crawlers by default + Pay Per Crawl | [official] | https://www.cloudflare.com/press/press-releases/2025/cloudflare-just-changed-how-ai-crawlers-scr | not used in the body |
| Google × Reddit** content licensing (reported at about $60M/year) | [official] (scope of the partnership); the amount is [third-party reported | https://blog.google/company-news/inside-google/company-announcements/expanded-reddit-partnership | not used in the body |
| The Atlantic × OpenAI** content and product partnership | [official] | https://www.theatlantic.com/press-releases/archive/2024/05/atlantic-product-content-partnership- | not used in the body |
| News Corp × OpenAI** multi-year global partnership | [third-party verified: authoritative media] | https://www.cnbc.com/2024/05/22/newscorp-and-openai-strike-multi-year-partnership-deal-to-use-jo | not used in the body |
| OpenAI ChatGPT Instant Checkout** (buyer age | [third-party verified: authoritative media] | https://www.cnbc.com/2025/09/29/chatgpt-instant-checkout-etsy-shopify.html | not used in the body |
| Amazon Rufus** (in-site AI shopping assistant) | [vendor claim] (Amazon self-reported) | https://www.aboutamazon.com/news/retail/amazon-rufus-ai-assistant-personalized-shopping-features | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Estée Lauder × OpenAI** : 240+ custom GPT | [third-party reported] (official customer story at openai | https://www.ciodive.com/news/Estee-Lauder-ELC-generative-AI-OpenAI-partnership-ChatGPT/733005 | not used in the body |
| Duolingo** : 'AI-first' memo and the reputational backlash | [third-party verified: authoritative media] | https://fortune.com/2025/06/09/duolingo-ceo-surprised-backlash-ai-first-company-announcement | not used in the body |
| Shopify** : 'prove AI can't do it, then add people' | [third-party verified: authoritative media] (memo text on X | https://www.cnbc.com/2025/04/07/shopify-ceo-prove-ai-cant-do-jobs-before-asking-for-more-headcou | not used in the body |
| Shopify** : orders from AI search up 15× in a year (positive case) | [third-party compilation: relayed from an earnings call] (original source at | https://detailed.com/public | not used in the body |
| IAC / People Inc.** (formerly Dotdash Meredith) : | [third-party verified: authoritative media] | https://digiday.com/media/publishers-swap-traffic-angst-for-strategy-in-q3-earnings | not used in the body |
| Penske Media v. Google** : the first publisher suit over AI summaries | [third-party verified: authoritative media] | https://techcrunch.com/2025/09/14/rolling-stone-owner-penske-media-sues-google-over-ai-summaries | not used in the body |
| Encyclopaedia Britannica v. Perplexity** | [official] (plaintiff press release) | https://www.prnewswire.com/news-releases/britannica-files-copyright-and-trademark-infringement-l | |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| | | | not used in the body |
| Epic Games 'Fortnite' | [authoritative media: citing an Apple official statement] | https://www.cnbc.com/2020/08/13/apple-kicks-fortnite-out-of-app-store-for-challenging-payment-ru | Ch.10, Method 1 |
| Valve 'Steam Link' (iOS version) | [authoritative media: citing a Valve statement] | https://www.cnet.com/tech/mobile/valve-accuses-apple-of-rejecting-steam-link-mobile-streaming-ap | not used in the body |
| Apple (App Store, China region) | [authoritative media] | https://xinwen.bjd.com.cn/content/s5f2624fee4b0d7d5fee65d92.html | Ch.10, Method 1 |
| Apple (corroborating item in the same case · The Paper) | [authoritative media] | https://www.thepaper.cn/newsDetail_forward_8420791 | not used in the body |
| Patuoxun (Chinese: 帕拓逊) (main brand Mpow) | [authoritative media] | https://m.bjnews.com.cn/detail/162969795414530.html | Ch.10, Method 5 |
| Sunvalley (Chinese: 泽宝) (parent Xinghui Co.; brands RAVPower / Taotr | [authoritative media] | https://m.bjnews.com.cn/detail/162969795414530.html / https://m.caixin.com/m/2021-07-21/10174347 | Ch.10, Method 5 |
| | [authoritative media] | https://m.bjnews.com.cn/detail/162969795414530.html | |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Aukey (Chinese: 傲基) (a different company from Sunvalley above) | | | not used in the body |
| Youshu (Chinese: 有棵树) (a subsidiary of Tianze Information) | [authoritative media: citing a listed-company announcement] | http://www.nbd.com.cn/rss/toutiao/articles/1853164.html | not used in the body |
| Tongtuo Technology (Chinese: 通拓科技) (a subsidiary of Huading Co.) | [authoritative media: citing a listed-company announcement] | https://m.bjnews.com.cn/detail/162969795414530.html | not used in the body |
| Winona (Chinese: 薇诺娜) (domestic beauty brand) | [official: Tencent Marketing Academy (platform-reported results)] | https://eschool.qq.com/Solution/ListDetail/pd-7788 | Ch.10, Method 3 |
| Sephora | [official: Tencent Marketing Academy (platform-reported results)] | https://eschool.qq.com/Solution/ListDetail/pd-7788 | not used in the body |
| OMEGA | [official: Tencent Marketing Academy | https://eschool.qq.com/Solution/ListDetail/pd-7788 | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| | (platform-reported results)] | | |
| HP | [vendor claim: Amazon Ads official case library] | https://advertising.amazon.com/solutions/products/sponsored-brands | Ch.10, Method 4 |
| Taobao listings (Chinese: 淘宝宝贝) (Zhitongche / keyword bidding, Alimama official figures) | [official: Alimama] | https://www.alimama.com/news_detail.htm?contentId=430 | Ch.10, Method 4 |
| Zynga | [authoritative media: citing a Zynga SEC filing] | https://www.cnet.com/home/smart-home/zynga-shares-soar-on-facebook-ipo | Ch.10, Method 5 |
| Facebook (2018 News Feed algorithm overhaul) | [authoritative media: citing a Facebook official statement ( | https://www.theguardian.com/technology/2018/jan/11/facebook-news-feed-algorithm-overhaul-mark-zu | not used in the body |
| WeChat Open Docs 'Mini Program Search Optimization Guide' | [official] | https://developers.weixin.qq.com/miniprogram/dev/framework/search/seo.html | not used in the body |
| WeChat Open Docs 'Service Providers Receiving Mini Program Violation | [official] | https://developers.weixin.qq.com/doc/oplatform/Third-party_Platforms/2.0/product/illegalEvent.ht | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Penalties on Behalf of Developers' | | | |
| WeChat Open Class PRO 2023 (reprinted by Xinhuanet) | [official: platform-reported] | https://www3.xinhuanet.com/tech/20230111/b4286147d13f4bb6a444666b31f8a4d4/c.html | not used in the body |
| Amazon seller success-story library (official Amazon) | [vendor claim] | https://sell.amazon.de/informationen/verkaufer-geschichten?mons_sel_locale=zh_CN | not used in the body |
| Gary Vaynerchuk / VaynerMedia** (personal brand + team) | [first-person account] | https://garyvaynerchuk.com/the-garyvee-content-strategy-how-to-grow-and-distribute-your-brands-s | Ch.12, Method 1 |
| Moz** (Whiteboard Friday series) | [institutional methodology] | https://moz.com/blog/repurpose-content-with-ai-whiteboard-friday | not used in the body |
| East Buy (Chinese: 东方甄选)** (Dong Yuhui clip licensing) | [authoritative media: Sohu Tech] | https://www.bianews.com/news/details?id=165035 | Ch.12, Method 1 |
| WeChat Official Accounts Platform (Chinese: 微信公众平台) **first case of its 'content-laundering complaint panel' (the** | [official reprint: National Copyright Administration site] | https://www.ncac.gov.cn/xxfb/yjdt/201901/t20190107_49426.html | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| **account** Jingyingshuo** (Chinese: 精英说) accused of laundering from * | | | |
| Douyin** (copyright-infringement governance, first half of 2021) | [authoritative media: China Securities Journal] | https://cs.com.cn/sylm/jsbd/202106/t20210624_6178255.html | not used in the body |
| Douyin** 'Code of Conduct for Misbehavior by Matrix Accounts' | [authoritative media: Xinhuanet (relaying an official new rule)] | http://www.xinhuanet.com/tech/20231130/c932193c6db34a079588421d9706516c/c.html | Ch.12, Method 4 |
| YouTube** (channel monetization policies: | [official rules] | https://support.google.com/youtube/answer/1311392?hl=en | Ch.12, Method 4 |
| WeChat Official Accounts Platform** (new rule against 'non-human automated creation') | [authoritative media + official response] | https://www.yicai.com/news/103125654.html | not used in the body |
| Gong** (Gong Labs) | [vendor self-reported] | https://www.gong.io/blog/taking-the-guesswork-out-of-sales-call-effectiveness | Ch.12, Method 5 |
| Ahrefs** (study of its own logs + traffic data) | [vendor self-reported] | https://ahrefs.com/blog/llmstxt-study | |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| | | | not used in the body |
| Strava** (Year in Sport: Trend Report) | [vendor self-reported: official press release] | https://press.strava.com/articles/strava-releases-12th-annual-year-in-sport-trend-report-2025 | Ch.12, Method 5 |
| Spotify** (annual Wrapped campaign) | [official] | https://newsroom.spotify.com/2025-12-03/2025-wrapped-user-experience/ | not used in the body |
| Cyberspace Administration of China (Chinese: 中央网信办)** notice on cases of 'self-media' failing to label properly (named accounts: Douyin's 'Qingqing Guoji' and 'Jike Kepuguan | [official notice] | https://www.cac.gov.cn/2026-05/03/c_1779492291101867.htm | not used in the body |
| WeChat Channels (Chinese: 微信视频号)** 'Announcement on Strengthening the Governance of AI-Generated Content' | [state media reprint of an official announcement: Qingdao Daily (source: WeChat Channels | https://www.dailyqd.com/guanhai/438720_1.html | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Xiaohongshu** 'AI Governance Rules Announcement' | [authoritative media relay: ITHome / Sina Tech] | https://finance.sina.com.cn/tech/digi/2026-08-07/doc-inimnram0019917.shtml | not used in the body |
| CAC** first phase of the 'Qinglang · Cracking Down on AI Technology Abuse' campaign (involving Tencent, Weibo, Douyin, Ali | [official notice] | https://www.cac.gov.cn/2025-06/20/c_1752129980667315.htm | not used in the body |
| HubSpot Website Grader | [official] (HubSpot official press release) | https://www.hubspot.com/blog/bid/5539/website-grader-analyzes-over-2-million-sites | Ch.15, Method 1 |
| Shopify Free Business Name Generator | [official] | https://www.shopify.com/tools/business-name-generator | Ch.15, Method 5 |
| Ahrefs Free Backlink Checker | [official] | https://ahrefs.com/backlink-checker | Ch.15, Method 1 |
| Heroku free plan discontinued | [official] | https://help.heroku.com/RSBRUH58/removal-of-heroku-free-product-plans-faq | Ch.15, Method 4 |
| Railway free allowance: first shut down, then reopened and hard-coded | [vendor claim] (vendor self-reported, including self-reported unit-economics figures) | https://blog.railway.com/p/free-plan | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Figma pricing, seat and billing revamp | [official] | https://www.figma.com/blog/billing-experience-update-2025 | not used in the body |
| Supabase Free plan limits | [official] | https://supabase.com/pricing | not used in the body |
| Cloudflare Workers Free daily limits | [official] | https://developers.cloudflare.com/workers/platform/limits | not used in the body |
| Stripe MCP server | [official] | https://docs.stripe.com/mcp | not used in the body |
| Linear MCP server | [official] | https://linear.app/docs/mcp | not used in the body |
| GitHub MCP Server | [official] | https://github.com/github/github-mcp-server | not used in the body |
| MCP Registry (the official MCP registry) | [official] | https://blog.modelcontextprotocol.io/posts/2025-09-08-mcp-registry-preview | not used in the body |
| GitHub MCP Registry | [official] (GitHub official blog) | https://github.blog/ai-and-ml/github-copilot/meet-the-github-mcp-registry-the-fastest-way-to-dis | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Stripe Directory | [official] | https://docs.stripe.com/directory | not used in the body |
| CoSchedule Headline Analyzer | [vendor claim] ('#1' is self-assessed) | https://coschedule.com/headline-analyzer | Ch.12, Method 4 |
| Zapier MCP | [vendor claim] (9,000+ / SOC 2 are self- | https://github.com/zapier/zapier-mcp | not used in the body |
| Sentry MCP | [official] | https://github.com/getsentry/sentry-mcp | not used in the body |
| Cloudflare remote MCP server capability launch | [official] | https://blog.cloudflare.com/remote-model-context-protocol-servers-mcp | not used in the body |
| Cloudflare MCP Server Portals | [official] | https://blog.cloudflare.com/zero-trust-mcp-server-portals | not used in the body |
| Claude Connectors directory | [official] | https://claude.com/connectors | not used in the body |
| MCP Registry publishing guide (quickstart) | [official] | https://modelcontextprotocol.io/registry/quickstart | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Heroku new low-cost plans Eco / Mini | [official] | https://blog.heroku.com/new-low-cost-plans | not used in the body |
| Ahrefs Free SEO Tools (tools collection page) | [official] | https://ahrefs.com/free-seo-tools | Ch.15, Method 1 |
| Australia's Robodebt (automated debt collection) Royal Commission report | [authoritative media investigation] (Guardian relaying the royal | https://www.theguardian.com/australia-news/2023/jul/07/robodebt-royal-commission-final-report-re | Ch.17, Method 5 |
| Zillow Offers (automated home buying / automated offers) shutdown post-mortem | [company's public post-mortem] (earnings press release filed with the SEC) | https://www.prnewswire.com/news-releases/zillow-group-reports-third-quarter-2021-financial-resul | Ch.17, Method 5 |
| 11x.ai (AI SDR / automated outreach) metrics independently disproven | [authoritative media investigation] (TechCrunch, multiple | https://techcrunch.com/2025/03/24/a16z-and-benchmark-backed-11x-has-been-claiming-customers-it-d | Ch.17, Method 5 |
| Air Canada support bot misstated bereavement fares (Moffatt v. Air Ca | [institution/law-firm analysis] (law firm quoting the ruling text paragraph by paragraph; ru | https://fosterandcompany.com/air-canada-found-liable-for-negligent-misrepresentation-by-chatbot | Ch.17, Method 7 |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Netherlands' SyRI welfare-fraud risk system (NJCM et al. v. the Dutch State) | [regulator document] (official judgment publication by the Dutch judiciary, | https://uitspraken.rechtspraak.nl/details?id=ECLI%3ANL%3ARBDHA%3A2020%3A1878 | Ch.17, Method 7 |
| Budapest Bank emotion AI analyzing support calls (Hungarian DPA: NAIH-85-3/202 | [regulator document relay] (decision collected by GDPRhub | https://gdprhub.eu/NAIH_(Hungary)_-_NAIH-85-3/2022 | not used in the body |
| Dutch childcare-benefits scandal (Toeslagenaffaire) | [authoritative media investigation] (CNBC) | https://www.cnbc.com/2021/01/15/dutch-government-resigns-after-childcare-benefits-scandal-.html | not used in the body |
| iTutorGroup hiring software automatically rejected applicants aged 55/60+ (EEOC v. | [regulator document] (EEOC press release) | https://www.eeoc.gov/newsroom/itutorgroup-pay-365000-settle-eeoc-discriminatory-hiring-suit | Ch.17, Method 8 |
| Rite Aid store AI facial-recognition false alerts (FTC v. Rite Aid, 2 | [regulator document] (FTC case file + press release) | https://www.ftc.gov/legal-library/browse/cases-proceedings/2023190-rite-aid-corporation-ftc-v | not used in the body |
| RealPage rental algorithmic pricing (DOJ et al. v. RealPage, 2 | [institution/law-firm analysis] (Paul Weiss paragraph-by- | https://www.paulweiss.com/insights/client-memos/practical-takeaways-from-the-doj-s-algorithmic-p | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| DoNotPay 'robot lawyer' + FTC 'Operation AI Comply | [regulator document] (FTC press release + case file) | https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai | not used in the body |
| FCC penalty over AI-faked Biden voice robocalls (Steve Kramer; Lin | [authoritative media investigation] (NH Public Rad | https://www.nhpr.org/nh-news/2025-06-13/political-operative-fake-biden-robocalls-nh-primary-foun | not used in the body |
| Amazon experimental hiring algorithm discriminated against women | [authoritative media investigation] (Guardian relaying a Reuters investig | https://www.theguardian.com/technology/2018/oct/10/amazon-hiring-ai-gender-bias-recruiting-engin | not used in the body |
| Cursor / Anysphere support bot invented company policy | [institution database] (AI Incident Da | https://incidentdatabase.ai/cite/1039 | not used in the body |
| Klarna AI support 'equal to 700 full-time agents' | [vendor claim] (Klarna official press release) | https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-se | not used in the body |
| DPD support bot went off the rails and was shut down | [authoritative media investigation] (BBC) | https://www.bbc.com/news/technology-68025677 | Ch.17, Method 8 |
| | [institution database] (OECD.AI | https://oecd.ai/en/incidents/2024-03-29-3dce | |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| New York City's MyCity official business Q&A bot gave unlawful advice | incident database, index | | not used in the body |
| FTC 6(b) inquiry into eight 'surveillance pricing' companies | [regulator document] (FTC press release) | https://www.ftc.gov/news-events/news/press-releases/2024/07/ftc-issues-orders-eight-companies-se | not used in the body |
| Luckin Coffee (Chinese: 瑞幸咖啡) | [company filing: SEC prospectus] | https://www.sec.gov/Archives/edgar/data/1767582/000091205719000058/filename1.htm | not used in the body |
| Pinduoduo (Chinese: 拼多多) (PDD Holdings, formerly Pinduoduo) | [company filing: SEC annual report] | https://www.sec.gov/Archives/edgar/data/1737806/000110465922049283/pdd-20211231x20f.htm | not used in the body |
| Dropbox | [company website / official help center] | https://www.dropbox.com/refer | not used in the body |
| Airbnb (Chinese: 爱彼迎) | [company official help center] | https://www.airbnb.com/help/article/3622 | not used in the body |
| Robinhood | [company filing: SEC] | https://www.sec.gov/Archives/edgar/data/1783879/000162828021013318/robinhoods-1.htm | Ch.19, Method 1 |
| Monzo (UK digital bank) | [company annual report] | https://monzo.com/annual-report/2026/Monzo-Annual-Report-2026.pdf | |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| | | | Ch.19, Method 1 |
| HONOR (Chinese: 荣耀) | [company website campaign rules] | https://www.honor.com/cn/msale/yqylhdgz.html | Ch.19, Method 4 |
| Amazon Associates | [company official help center] | https://affiliate-program.amazon.com/help/node/topic/GS4DGSRUYG5BX8E3 | Ch.19, Method 4 |
| Taobaoke / Alimama (Chinese: 淘宝客 / 阿里妈妈) (Alibaba) | [company official developer documentation] | https://open.alitrip.com/docs/api.htm?apiId=62201 | Ch.10, Method 4 |
| Fashion Nova (US fast fashion) | [enforcement action: FTC press release] | https://www.ftc.gov/news-events/news/press-releases/2022/01/fashion-nova-will-pay-42-million-par | not used in the body |
| Sunday Riley Modern Skincare + Devumi | [enforcement action: FTC press release] | https://www.ftc.gov/news-events/news/press-releases/2019/10/devumi-owner-ceo-settle-ftc-charges- | not used in the body |
| The Bountiful Company (parent of Nature's Bounty | [enforcement action: FTC press release] | https://www.ftc.gov/news-events/news/press-releases/2023/02/ftc-charges-supplement-marketer-hija | not used in the body |
| Rytr (AI writing tool) | [enforcement action: FTC press release] | https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Sitejabber (GGL Projects, Inc.) | [enforcement action: FTC press release (final order)] | https://www.ftc.gov/news-events/news/press-releases/2025/01/ftc-approves-final-order-against-sit | not used in the body |
| Roomster and its operators + Jonathan Martinez (d/b/a | [enforcement action: FTC press release] | https://www.ftc.gov/news-events/news/press-releases/2022/08/ftc-states-sue-rental-listing-platfo | not used in the body |
| Amazon sues AppSally and Rebatest (two fake-review brokers) | [company official press release: platform enforcement] | https://press.aboutamazon.com/2022/2/amazon-sues-fake-review-brokers-who-attempt-to-profit-from- | not used in the body |
| Amazon + the US Better Business Bureau (BBB) sue Skitsolutionbd; Am | [company official press release: platform enforcement] | https://www.aboutamazon.com/news/policy-news-views/amazons-latest-actions-against-fake-review-br | not used in the body |
| Tripadvisor (Chinese: 猫途鹰) | [company official report / platform enforcement] | https://tripadvisor.mediaroom.com/2025-03-18-Tripadvisors-2025-Transparency-Report-reveals-stron | not used in the body |
| Beijing Haipu Advertising Co., Ltd. (Chinese: 北京嗨噗广告有限公司) | [enforcement action: SAMR official site] | https://www.samr.gov.cn/jjj/sjdt/gzdt/art/2023/art_885915ef69c94901bd4ce6b44d9029bc.html | not used in the body |
| Baoji Haisheng Modern Agriculture Co., | [enforcement action: Shaanxi Provincial Administration | https://snamr.shaanxi.gov.cn/sy/sjyw/202309/t20230919_2604145.html | not used in the body |

| Case (named) | Source tier | Source | Used in |
|---|---|---|---|
| Ltd. (Chinese: 宝鸡海升现代农业有限公司) | for Market Regulation official site] | | |
| Hangzhou Dongtian Technology Co., Ltd. (Chinese: 杭州东田科技有限公司) | [enforcement action: SAMR release, reprinted on a government site] | https://www.qinshui.gov.cn/ztzl_369/cxxcjy_25/bgsx_36/202309/t20230901_1847663.shtml | not used in the body |
| FTC 'Rule on the Use of Consumer Reviews and Testimonials' (16 CFR Part 465) | [regulatory rule text + FTC press release] | https://www.ecfr.gov/current/title-16/chapter-I/subchapter-D/part-465 | not used in the body |

## A.9 Second-round material-mining ledger (growth-hacking material library)

In the second round I swept the growth-hacking material library gathered earlier (`~/wiki/topics/growth-hacking/raw/articles/`, 339 pieces, of which 286 had not previously been cited by this book), section by section, yielding **90** candidates in total. The table below is the complete ledger: items marked "Ch.X, Method N" have been written into the body; those marked "not used in the body" were collected this round but not used, kept for later versions and for readers to check themselves.

Collection discipline: a material summary is a lead only; every item requires fetching the original page at `source_url` and pulling the sentence that supports the conclusion from that page; every URL is tested for reachability one by one; figures whose original publisher cannot be traced are not collected.

| Case (named) | Source tier | Source |
|---|---|---|
| | | https://www.polaranalytics.com/case-studies/how-sage-and-paige-their-digital |

| Case (named) | Source tier | Source |
|---|---|---|
| Sage and Paige **(US women's-wear DTC brand)** × Polar Ana | [vendor claim] (the vendor's own case page; the client is named and includes a client-side signed | |
| Apple ATT / Meta** (AdLibrary five-year retrospective) | [third-party verified: industry retrospective] (AdLibrary independent ana | https://adlibrary.com/posts/ios-14-att (**200**) |
| Focus Digital** (digital-marketing agency, 2026 CAC Trends R | [third-party verified: agency's own data report, methodology and definitions public | https://focus-digital.co/customer-acquisition-cost-trends/ (**200**) |
| Digital Applied** (consultancy, CAC Benchmarks 2026 | [third-party verified: compiles 8 sets of public benchmarks, each source labeled] ( | https://www.digitalapplied.com/blog/customer-acquisition-cost-ben industry ( |
| Jungle** (AI study app, founders Julian / David) × ** | | https://superwall.com/blog/this-ai-study-app-makes-usd100k-mo-h code (**200** |

| Case (named) | Source tier | Source |
|---|---|---|
| | [third-party reported: founder interview] (Superwall | |
| Recurly** (subscription-billing infrastructure vendor, 2025-12-03 win-back guide) | [vendor claim: including industry data] (Recurly is a billing product | https://recurly.com/blog/customer-winback-strategies-for-subscript |
| N26 **(German digital bank)** × Sacra** (company-research firm) | [third-party verified: firm estimate] (Sacra company research, some | https://sacra.com/research/n26 (**200**) |
| Antenna** (US subscription data firm, 2024 year-end roundup) | [third-party verified: data-firm estimate] (Antenna is a third | https://www.antenna.live/insights/antennas-2024-top-subscription-churn (**2 |
| Benchmarkit** (2026 B2B SaaS & AI-Native Me | [third-party verified: citing a Benchmarkit report] (ori | https://www.thesaascfo.com/saas-grr-benchmark-2026/ (**200**) |
| | [vendor claim: vendor network | https://recurly.com/research/churn-rate-benchmarks/ (**200**) |

| Case (named) | Source tier | Source |
|---|---|---|
| Recurly** (subscription-billing platform network data, 2026-07) | data] (Recurly is a billing | |
| ChartMogul** (SaaS retention report, 2,500+ SaaS companies) | [third-party verified] (based on ChartMogul's own platform | https://chartmogul.com/reports/saas-retention-the-new-normal/ (**2** |
| Booking.com **(experimentation culture)** × Airbnb **(photography experiment)** / Ob | [third-party reported: media] (Observer column, B | https://observer.com/2026/01/why-experimentation-drives-high-pe companies/ (**20 |
| Case name (named) | | |
| SaaS Capital** 2025 private-SaaS retention benchmarks | [third-party verified] (its own annual survey running over a decade, sample = ARR> | https://www.saas-capital.com/blog-posts/what-is-a-good-retention- private-saas |
| Lenny's Newsletter** GOOD/GREAT retention benchmarks | [third-party verified] (compiles public data with 20 practitioners) | https://www.lennysnewsletter.com/p/what-is-good-retention-issue- |
| | | https://blog.duolingo.com/growth-model-duolingo (200) |

| Case (named) | Source tier | Source |
|---|---|---|
| Duolingo** seven-state growth model | [official] (Duolingo official blog, by Er | |
| Remery** 17 pricing experiments (12 losses, 5 wins) | [third-party reported: the company's own account of its experimental data] | https://remery.ai/blog/saas-pricing-experiments-17-tests-results (2 |
| Tsinghua Business Review A/B testing** (including iQIYI / China Mobile) | [third-party verified: academic journal] (Tsinghua SEM's Tsinghua Business Review 2 | https://www.sem.tsinghua.edu.cn/info/1171/35965.htm (200) |
| MarketerHire** paid-acquisition framework | [third-party verified: industry content] | https://marketerhire.com/blog/paid-acquisition-strategy (200) |
| Zylo** 2026 SaaS pricing trends | [third-party verified] (survey of Zylo's customer network; Gartne | https://zylo.com/blog/saas-pricing-trends (200) |
| Polytraffic** CAC benchmarks by channel | | https://polytraffic.com/articles/customer-acquisition-cost-by-chann |

| Case (named) | Source tier | Source |
|---|---|---|
| | [third-party verified: compilation] (aggregating HubSpot/Prof | |
| Yueshengshi (Chinese: 粤省事)** (Guangdong provincial government / Digital Guangdong, WeChat mini program + App) | [official] (Digital Guangdong website) + [third-party verified] (Tencent Cloud do | https://www.secrss.com/articles/14012 (200); https://www.digitalgd yss/ywly_y |
| Shanghai "One-Stop Online Service" / "Suishenban" (Chinese: 上海"一网通办" / "随申办")** (Shanghai municipal government) | [third-party verified] (Jiemian News, citing the deputy director of the Shanghai Big Data Center at | https://m.jiemian.com/article/11828626.html (200) |
| Zheliban (Chinese: 浙里办)** (mobile end of the Zhejiang government service portal) | [official (relayed by official media)] (Xinhuanet relaying Zhejiang Daily, with offici | https://www.news.cn/government/ 20240625/74eca00cd3bb49129d0318b909d6b1f6/c.html (200) |
| | [third-party verified: citing the | https://pro.huxiu.com/article/4215577.html (200) |

| Case (named) | Source tier | Source |
|---|---|---|
| Lalamove (Chinese: 货拉拉)** (intra-city freight platform, prospectus figures) | prospectus] (Huxiu Miaotou based on a 2025-04 | |
| Full Truck Alliance / "Yunmanman" (Chinese: 满帮集团 / "运满满")** (digital freight platform, YMM.US) | [third-party verified: citing financial filings] (Securities Daily = relayed from filings; Sina Tech | http://www.zqrb.cn/gscy/qiyexinxi/2025-03-06/A1741232211313.h https://finance.s |
| Temu** (PDD, parent of Pinduoduo, going global) | [third-party verified] (Harvard Business School Working Know | https://www.library.hbs.edu/working-knowledge/how-shein-and-ten fast-fashion- |
| Airbnb** (acquiring customers via the Craigslist platform) | [third-party reported: single source, **second-hand relay**, used with caution] | https://growthexpertz.com/growth-hacking-case-studies (200) |
| Shopify partner ecosystem** (App Store + referral partners) | [third-party reported: single source, **second-** | https://partnerinsight.io/insights/what-does-it-take-to-build-a-thrivi marketplac |

| Case (named) | Source tier | Source |
|---|---|---|
| | **hand relay**, used with caution] | |
| Fastenal** (US industrial distributor/ fastener supplier) | [third-party verified: citing the company annual report] (Dorn Group/M | https://www.dorngroup.com/fastenal-investments-line-extensions-c (200; **requires with |
| Kuaidi100 "Baidiyun" × Langsheng Group (Chinese: 快递100「百递云」 × 朗生集团)** (logistics-SaaS customer case) | [vendor claim: platform self-reported] (a case from Kuaidi100's own official site | https://www.kuaidi100.com/enterprise/customer/langsheng.shtml ( |
| Sunnystep × SHOPLINE** (footwear brand, Singapore / overseas DTC) | [third-party reported: citing a vendor client, used with caution] (Xiaguangshe report | https://mp.ofweek.com/digitaleconomy/a256714077547 (200) |
| Hotmail** ("PS: I Love You" email signature) | [third-party verified: book excerpt] (TechCrunch book excerpt, fig | https://techcrunch.com/2009/10/18/ps-i-love-you-get-your-free-em (200) |

| Case (named) | Source tier | Source |
|---|---|---|
| Amazon third-party seller marketplace | [third-party reported: single source, second-hand relay] | https://zetetikos.substack.com/p/case-study-how-amazons-flywhee `.../2026- |
| "Qinwuyuan" (Chinese: "秦务员") (Shaanxi integrated government-services platform) | [official] (published by People's Daily Online; contributing unit = the platform's operations department) | http://politics.people.com.cn/BIG5/n1/2025/1223/c458474-406302 `.../2026-08 |
| Fan Deng Reading (Chinese: 樊登读书) | [third-party verified: self-media relay, figures are the company's own, used with caution] | https://www.sohu.com/a/684937054_121716213 (200) |
| Substack recommendation network | [third-party verified: Chinese second-hand relay (Runwise → Maimai), | https://maimai.cn/article/detail?fid=1797931979&efid=sZR2w6WLVHbUNk2W33BGVQ (200) |
| a16z 'Top 100 Gen AI Consumer Apps' 6th edition (Ca | [third-party verified] (a16z based on SimilarWeb | https://a16z.com/100-gen-ai-apps-6/ (200) |

| Case (named) | Source tier | Source |
|---|---|---|
| Guixingren Pro (Chinese: 硅星人 Pro) 'The 10,000 wild ways to growth-hack an AI product' (Wordware | [third-party verified: media case compilation, each case's data is a media relay, downgr | https://www.woshipm.com/ai/6299942.html (200) |
| Spotify · Discover Weekly | [third-party verified: product data citing Time/ Fast Co | https://www.markhub24.com/post/spotify-s-discover-weekly-how-p became-a-co |
| Netflix · recommendation engine (NRE) | [vendor claim: an e-commerce tool vendor's blog relaying Netflix | https://www.rebuyengine.com/blog/netflix (200) |
| 1ClickReport: the May 2026 core update and "the death of programmatic SEO" | [vendor claim: an SEO tool vendor's first-hand client observation, a single firm's 6 | https://www.1clickreport.com/blog/google-may-2026-core-update-p seo-dead (200 |
| Luckin Coffee (CGO Yang Fei) | [official: financial-filing data] + [first-person account: CGO interview] | https://www.centurium.com/press/…(Centurium Capital interview) |

| Case (named) | Source tier | Source |
|---|---|---|
| Intercom · Fin AI support | [vendor claim: Intercom engineering team's own account (via Ho | https://www.honeycomb.io/resources/case-studies/how-honeycomb intercom-observe-an |
| Reforge 'growth loops replace the funnel' · Pinterest content loop | [third-party verified] (Balfour / Winters | https://www.reforge.com/blog/growth-loops (200) |
| Notion (disclosed by the company to CNBC) | [third-party verified: disclosed by the company to CNBC] | https://www.cnbc.com/2025/09/18/notion-launches-ai-agent-as-it-c million-in-an |
| Zoom | [third-party verified] (revenue and customer counts traceable to Zoom's filings; | https://www.getmonetizely.com/articles/case-study-how-zoom-scal strategy-durin |
| Remery** (B2B workflow-automation SaaS) | [vendor claim] (the company's own account of its A/B experimental data; full pre- | https://remery.ai/blog/saas-pricing-experiments-17-tests-results |
| Shopify App Store (+ partner ecosystem) | | https://partnerinsight.io/insights/what-does-it-take-to-build-a-thrivi marketplac |

| Case (named) | Source tier | Source |
|---|---|---|
| | [third-party verified] (Partner Insight's | |
| Duolingo** (soft-wall signup page A/B) | [third-party reported] (NoGood post-mortem, citing Duolin | https://nogood.io/blog/duolingo-case-study/ |
| Duolingo** (economics of organic acquisition) | [official] (HBR podcast / case figures; the $2.5 is | https://hbr.org/podcast/2025/04/how-duolingo-aims-to-diversify-be learning |
| Hotmail | [third-party relay] (TechCrunch publishing 'Vira | https://techcrunch.com/2009/10/18/ps-i-love-you-get-your-free-em |
| Slack | [third-party reported] (First Round Review | https://review.firstround.com/from-0-to-1b-slacks-founder-shares-t strat |
| Decagon × Duolingo (Duolingo English Test) | [vendor claim] (Decagon official case page, the client is D | https://www.decagon.ai/case-studies/duolingo |
| Decagon × Chime | [vendor claim] (Decagon official case page; the | https://www.decagon.ai/case-studies/chime |

| Case (named) | Source tier | Source |
|---|---|---|
| | page notes its own figures | |
| UC Berkeley 'California Management Review' ( | [third-party verified] (UC Berkeley academic journal I | https://cmr.berkeley.edu/2026/04/chatbot-frustration-is-real-hidden pract |
| Braintrust 'Best AI customer service agents | [third-party reported] (Braintrust platform comparison report; | https://www.braintrust.dev/articles/best-ai-customer-service-agent |
| NetEase Zhigi · Yunshang (NetEase Yunshang) (Chinese: 网易智企·云商（网易云商）) AI support / AI private domain | [vendor claim] (NetEase Zhigi · Yunshang official content; case clients are anonymous, | https://grow.163.com/cms/ai-si-yu-wang-yi-zhi-qi-2026-YcKhIBbi.ht |
| Luckin Coffee** (co-founder and CGO Yang Fei interviewed, Centurium Capital · Founder Talk) | [vendor claim: a company executive's own interview account, single source] (the publisher is an investor | https://www.centurium.com/press/ %E7%91%9E%E5%B9%B8%E5%92%96%E5%95%A1%E6%9D%A8 |

| Case (named) | Source tier | Source |
| --- | --- | --- |
| Watsons (Chinese: 屈臣氏)** (a full teardown in the Woshipm private-domain operations column) | [third-party reported: column teardown, relaying the company's figures] | https://www.woshipm.com/operate/5298724.html(200) |
| MINISO (Chinese: 名创优品)** (teardown in the Woshipm private-domain operations column) | [third-party reported: column teardown, relaying the company's figures] | https://www.woshipm.com/operate/5288035.html(200) |
| Perfect Diary (Chinese: 完美日记)** (Growth Box OSINT teardown, reprinted by CBNData) | [third-party reported: OSINT estimate, single source, used with caution] | https://www.cbndata.com/information/40929(200) |
| Soochow Securities Digital Intelligence Branch (Chinese: 东吴证券数智分公司)** (Securities Daily report, reprinted by Securities Times online) | [third-party reported: authoritative industry media] (Securities Daily, citing the company | https://www.stcn.com/article/detail/3588663.html(200) |
| Runwise post-mortem ｜ SMS win-back frequency experiment** | [third-party reported: method post- | https://runwise.co/digital-growth/growth-hacking/5455/(200) |

| Case (named) | Source tier | Source |
|---|---|---|
| (a Runwise innovation-community method post-mortem) | mortem] (Runwise's own media | |
| Revolut** (2025 annual report + company press release published with it) | [official: company press release (published with the 2025 annual report)]* | https://www.revolut.com/en-US/news/revolut_reports_record_profit_of_2_3bn_for_2025_as_re |
| Airbnb** \| Referrals 2.0 rebuild (teardown by referral-marketing vendor Extole | [vendor claim: referral-marketing vendor relaying Airbnb's growth team | https://www.extole.com/blog/how-airbnbs-marketing-strategy-attra with-the- |
| Dropbox** \| two-sided referral reward (post-mortem on a sweepstakes/referral SaaS blog) | [vendor claim: referral SaaS blog, second-hand relay, single source] — | https://www.viraloo.org/blog/dropbox-referral-program-case-study- growth(200) |
| Pinduoduo** (in-depth teardown by Jiemian News / Huanfan Caijing) | [third-party reported: authoritative media] (Jiemian News / Huanfan Caijing; " | https://www.jiemian.com/article/2317958.html(200) |

| Case (named) | Source tier | Source |
|---|---|---|
| MINISO** \| referral acquisition (coupon referrals + invite-a-friend-for-cash) | [third-party reported: column teardown, relaying the company's figures]** + note | https://www.woshipm.com/operate/5288035.html(200) |
| Three Youzan (Chinese: 有赞) customer-referral cases**: Huanglaowu Foods / Liangjiangren / Chunke CENO (Youzan New Retail Encyclopedia, blog | [vendor claim] (promotional material from customers of Youzan's own toolchain, not third-party | https://www.cnblogs.com/youzan/articles/19099271(200) |
| Starbucks Rewards | [third-party reported: loyalty SaaS compilation, relaying the company's disclos | https://www.rivo.io/blog/best-rewards-programs-dtc-brands(200) |
| Sephora Beauty Insider | | |
| N26 (Contrary Research report) | [third-party reported: research firm] | https://research.contrary.com/company/n26(200) |
| | | https://knowledge.wharton.upenn.edu/article/death-daily-deal/ |

| Case (named) | Source tier | Source |
|---|---|---|
| Groupon / LivingSocial** (Wharton professor Eric C | [third-party reported] (Wharton official publication, citing the professor | |
| Casey Winters** (former head of growth at Pinterest / Grubhub | [third-party reported] (a signed article by a named practitioner) | https://www.caseyaccidental.com/p/when-growth-plateaus-how-an |
| Allbirds** (NASDAQ: BIRD, IPO 2021-11) | [official] (transaction amounts, recorded by the SEC) + [third-party reported | |
| FTC "click-to-cancel" rule** (Click-to-Cancel, amending the 1973 N | [official] (FTC press release text) | https://www.ftc.gov/news-events/news/press-releases/2024/10/fed commission-ann |
| Meta (Meta Platforms Ireland)** cross-border data transfer case | [official] (DPC official press release text) | http://www.dataprotection.ie/en/news-media/press-releases/Data-P Commission-ann |
| N26 / Germany's BaFin** growth cap and anti-money-laundering penalty | [third-party reported] (Banking Dive named repo | https://www.bankingdive.com/news/n26-bafin-growth-cap-aml-com tayenthal |
| Growth Engineering / Pragmatic Engineer × | [third-party reported] (in-depth | https://newsletter.pragmaticengineer.com/p/what-is-growth-engine |

| Case (named) | Source tier | Source |
|---|---|---|
| | interview with a named expert, author Gerg | |
| PostHog growth-engineer team | [vendor claim] (PostHog self-reported, and it is a product-analytics ven | https://posthog.com/blog/what-is-a-growth-engineer |
| Atlassian product growth engineer** (official blog) | [official] (Atlassian official blog; Atla | https://www.atlassian.com/blog/how-we-build/what-does-a-product engineer-work-on |
| Quibi** (Jeffrey Katzenberg + Meg Whitman) | [third-party reported] (Babson College official publication, with named | https://entrepreneurship.babson.edu/lessons-from-billion-dollar-fai |
| Dark Patterns** (ACM 'Communications o | [third-party reported] (peer-reviewed journal article; a textbook-level reference | https://cacm.acm.org/practice/dark-patterns/ |
| Robinhood / FINRA $26 million fine | | https://www.finra.org/media-center/newsreleases/2025/finra-order financial-pa |

| Case (named) | Source tier | Source |
|---|---|---|
| | [official] (FINRA press release text) | |
| Case/study (named) | | |
| Bain 'AI Is Transforming Productivity, but | [third-party reported] (Bain official research report) | https://www.bain.com/insights/ai-transforming-productivity-sales-r frontier-te |
| Salesforce 'State of Sales' 2026** (7th edition, 23 | [vendor claim] (Salesforce's own annual survey) | https://www.salesforce.com/news/stories/state-of-sales-report-announcement-2026/ |
| UC Berkeley 'California Management Review'* | [third-party reported] (academic journal Insights; of which | https://cmr.berkeley.edu/2025/02/balancing-personalized-marketin privacy-in-th |
| Osano '2025 Marketing Data Privacy Compliance Guide' | [vendor claim] (practical guide from a privacy-compliance SaaS vendor) | https://www.osano.com/articles/marketing-data-privacy-guide |

# Appendix C · Version History and Update Notes

This book is a "living book": continuous gathering → scheduled merging → version increments. This appendix records the substantive changes in each version; the detailed commit record is in the repository CHANGELOG (https://github.com/Jweokk/ai-customer-acquisition-book/blob/main/CHANGELOG.md).

## v0.3.0 — 2026-09-13 (English edition)

- **First complete English edition**: 31 files under `book-en/` — preface, 4 part openers, 24 chapters, Appendix A (case index and both ledgers) and Appendix C — mapped one-to-one onto the Chinese edition.
- Translation rules live in `book-en/README.md`; terminology and source-tier mapping live in `assets/glossary-en.md`.
- `llms.txt` now pairs every chapter with both its Chinese site page and its English file on GitHub.
- The English PDF is published at the repository root as `ai-customer-acquisition-en.pdf` (GitHub distribution only).
- The Chinese content is unchanged in this release; the version number moves with it.

## v0.1.0 — 2026-09-13

**First edition (in progress).**

- Published the preface and the skeleton and first draft of the four parts and seventeen chapters;
- The book's proposition and boundaries: customer acquisition = earning the qualification to enter the buyer's decision process; the finish line is drawn at the first completed transaction / signed contract;
- Dual-track perspective: the Chinese ecosystem and overseas ecosystems written side by side;

- Four sourcing disciplines (trace to the original publisher, label the source tier, present differences in official figures side by side, delete what cannot be verified);
- **Added the book's own measurements**: the multi-model candidate-set measurement (6 questions × 5 model channels × 2 runs, 51 valid answers), concluding that a candidate set is a distribution rather than a list, that Chinese and English gateways return two candidate sets that barely intersect, and that the bare model layer carries no citations;
- Established Appendix A's tier ledger (source tier, data date, verifiability).

## v0.2.0 — 2026-09-13 (structure edition)

- **Established the book's main line**: buyers have changed where they pick vendors. Part 3 was reordered around the buyer's four steps (Find → Shortlist → Close → Bring the next one), and the second sentence of every chapter opening now returns to the main line; the preface carries the main line in one sentence, the reader-positioning table and that formula.
- **Chapter list changed to 24 chapters**: expanded from 17 chapters to 24 (adding eight chapters — advertising, in-platform acquisition, letting partners sell for you, the content factory and matrix distribution, product-led acquisition, AI answering customers around the clock, word-of-mouth referral and virality, and building the foundation first; the dark funnel merged into Chapter 1); chapter names changed to plain language you can grasp at a glance.
- **Numbering migration**: all chapter-number references, navigation and indexes synced to the 24-chapter numbering; appendices moved to the 90–93 prefix.
- **Writing style settled as the method format**: the tactic name is the subsection name, each method has four blocks (what to do / how / example / pitfalls to avoid), and each chapter ends with two items (Notes for this chapter, Sources for this chapter);
- **Sourcing discipline unchanged**: numbers that cannot be traced to the original publisher do not go into the body; effect multiples reported by vendors are marked only [vendor claim] and are not treated as budget assumptions.

## v0.2.1 — 2026-09-13 (case round)

- 16 methods that previously could only say "no publicly verifiable case available" were replaced with named, verifiable cases; every case URL was tested one by one and confirmed reachable, and the verifiable original sentence is kept in the sources;
- Added A.8, the case pool, registering the cases gathered this round but not written into the body, for reference;
- Vendor-reported effect figures are still marked only [vendor claim] / [vendor self-reported], and are not treated as fact.

## v0.2.2 — 2026-09-13 (sourcing and title round)

- Subtitle changed to "Buyers are asking AI — are you in the answer?"; the part opener subsection renamed "Chapters in This Part";
- The self-test channels no longer write specific model vendors and version numbers, and instead use capability tiers A–E (reasoning / general-purpose / lightweight / open-weight / spot-check), with the raw data and scripts anonymized accordingly.

## v0.2.3 — 2026-09-13

- Subtitle fixed as "Use AI to Acquire Customers: Find, Shortlist, Close, Bring the Next One";
- Added Appendix E · References for the Whole Book: all the book's sources merged and de-duplicated into 299 entries, uniformly numbered with the citing chapters listed against each one; the body carries no numeric superscripts (the tier stays at the end of the sentence), and the end-of-chapter lists are kept.

## v0.2.5 — 2026-09-13

- One round of salvaging from the earlier growth-hacking material library: 90 candidates, 24 written into the body (Chapters 10, 11, 12, 13, 16, 18, 19, 20, 21, 22, 23), with a new A.9 ledger; Appendix E updated to 314 entries.

## Version Plan

- v0.2.x: chapter-by-chapter consolidation (unified terminology, sources merged into Appendix A and numbered, every chapter's unverified items brought to zero or deleted);
- v0.3.x: complete the in-house measurements for Chapter 8 and Chapter 20; add concrete cases to the industry difference matrix;
- v1.0.0: whole-book consolidation complete, Appendices A–D complete, PDF finalized.

## How Updates Work

New content is merged into existing chapters in the normal chapter structure, so readers cannot see "which paragraph was added in which version"; version notes appear only in this appendix and the CHANGELOG. The body carries no version markers.

# Part 1 · See the Battlefield

> First confirm one thing — where buyers are actually choosing right now, and whether you're present there.

This part has only one chapter, but it decides whether everything you read after it has anything to do with you.

Buyers often have the list already drawn up before they meet you: they ask an AI first, look at other people's reviews first, run a round of comparison themselves first. You can't see these steps, and you can't appeal them. So the first thing to do isn't to learn methods — it's to run a gateway audit and get clear on whether you're present: measure that you are, and you follow the map for the rest of the book; measure that you aren't, and you know which lesson you need to make up.

## Chapters in This Part


- Chapter 1 See the Battlefield: How Buyers Choose Vendors Now (5 methods)

# Part 2 · New Methods for the AI Era

> These are the roads buyers have just been given — you have to hold them first.

Buyers now ask AI outright "which one should I choose", let an agent do the filtering for them, look at what platforms push and at what other people say. These roads didn't exist before AI spread, or cost so much nobody walked them — so everything this part covers is a **newly added action**: not doing the old methods better, but building visibility from zero in a new place.

This part covers the first two steps of the main line: where the buyer chooses (Find), and on what grounds he keeps you on the list (the first half of Shortlist).

## Chapters in This Part


- Chapter 2 Get Mentioned in AI Answers: How to Do GEO/AEO (10 methods)
- Chapter 3 Get Onto the Buyer's First Shortlist (6 methods)
- Chapter 4 Let Third Parties Speak for You (6 methods)
- Chapter 5 Lead With Your Own Pitfalls (6 methods)
- Chapter 6 Be Readable by Machines (6 methods)
- Chapter 7 Win the Buyer's AI Agent (5 methods)
- Chapter 8 Three Extra Moves for Chinese Markets (5 methods)

# Part 3 · Upgrading the Old Methods with AI

> None of the old roads have disappeared; only the way you walk them has changed.

Advertising, outreach, platforms, partners, content, phone — all of these channels are still in use, and buyers still show up in these places. What changed is the process and the price: targeting has been handed over to the systems, creative capacity has become the competitive edge, and the receiving end and the regulators have tightened the gate at the same time.

This part is ordered by **which step the buyer has reached**, down one road to the end:

- **Step 1 · Find** (Chapters 9–13) — get the buyer to run into you for the first time: buy position, borrow platforms, borrow partners, produce your own content, reach out yourself
- **Step 2 · Shortlist** (Chapters 14–16) — he starts comparing, and you have to be kept: read first, try second, ask third
- **Step 3 · Close** (Chapters 17–18) — he's willing to talk: catch the leads and the conversation, and hold the red lines at the same time
- **Step 4 · Bring the next one** (Chapter 19) — he became a customer, and brings the next customer over

## Chapters in This Part

**Step 1 · Find**


- Chapter 9 Advertising: From "Picking an Audience" to "Setting a Goal" (6 methods)
- Chapter 10 Acquire on Other People's Platforms: App Stores, In-Platform E-commerce, Industry Directories (5 methods)
- Chapter 11 Let Partners Sell for You: Channels, Resellers, Ecosystem Directories (5 methods)

- Chapter 12 Turn One Piece of Content Into Thirty: The AI Content Factory and Matrix Distribution (5 methods)
- Chapter 13 Outreach: From "Sending More" to "Sending Precisely" (5 methods)

**Step 2 · Shortlist**

- Chapter 14 Content Others Cannot Copy: First-Party Data and Real Cases (5 methods)
- Chapter 15 Let the Product Acquire: Free Tools, Trials, Open APIs (5 methods)
- Chapter 16 Let AI Take Over Your Customer Inquiries 24 Hours a Day (5 methods)

**Step 3 · Close**

- Chapter 17 Use AI to Find Customers, Manage Leads and Run the Sale (8 methods)
- Chapter 18 Phone and Private Domains: What Still Works, What Blows Up on Contact (5 methods)

**Step 4 · Bring the next one**

- Chapter 19 Let Customers Bring Customers: Word of Mouth, Referrals and Virality (5 methods)

# Part 4 · Don't Waste Money or Build on Sand

> Whether the two parts above can run at all depends on this one.

This part is not on the buyer's journey. The buyer won't look at you once more because you did the things here — but every method in the previous 19 chapters rests on this part holding true.

It governs five things: **the books** (whether the money you spent actually did anything), **dependency** (none of the gateways the buyer sits in belong to you), **the foundation** (first-party data, compliance, omnichannel orchestration), **people** (who ends up doing all of the above), and **red lines** (the mistakes others have already paid the price for on your behalf).

Reading this part before you act is far cheaper than acting first and coming back to patch it up.

## Chapters in This Part


- Chapter 20 How to Know It's Working: Three Sets of Books and a Holdout Test (5 methods)
- Chapter 21 Don't Bet on One Gateway: Dependency Check and Hedging (5 methods)
- Chapter 22 Build the Foundation First: First-Party Data, Compliance, Omnichannel Orchestration (5 methods)
- Chapter 23 Who Does It and What It Costs: Three Minimum Configurations (5 methods)
- Chapter 24 Don't Cross These Lines: Six Failure Files (7 methods)